



# Computer Organization and Software Systems

Contact Session 1  
Introduction to Computer Systems

Dr. Lucy J. Gudino



# Team



## Instructors:

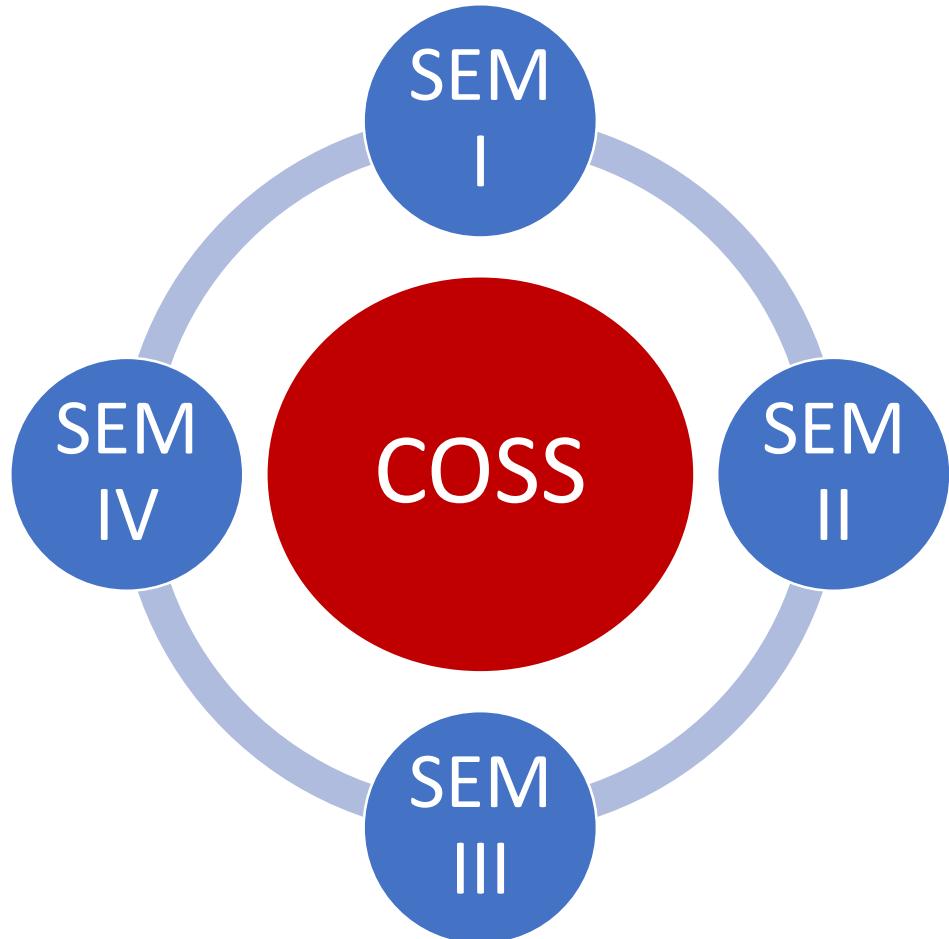
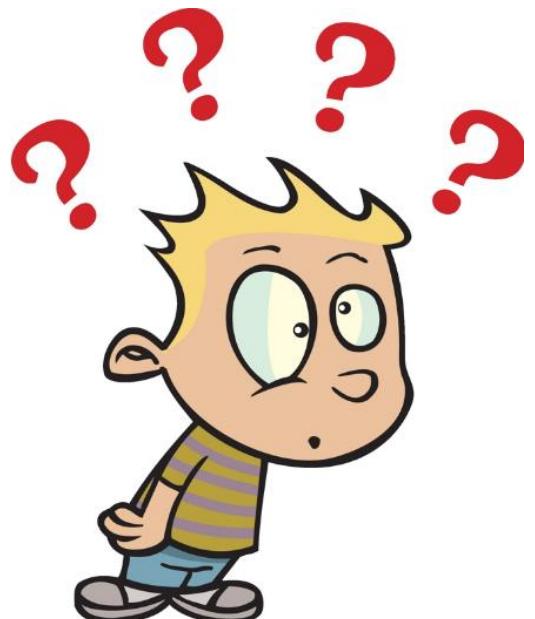
- Dr. Lucy J Gudino ( IC)
- Prof. Lakshmikanta
- Prof. Pradeep H K

## Learning Facilitators:

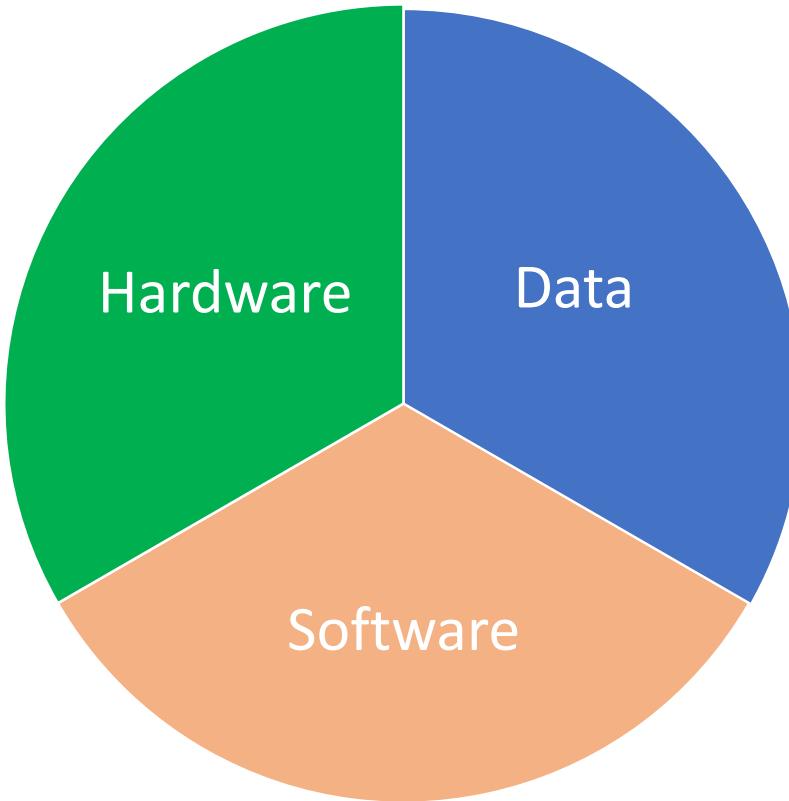
- Dr. VAIBHAV JAIN (Lead TA)
- Prof. Vivekananda M R
- Dr. Shamanth Nagraj
- Prof. Puneet Kumar

# Introduction

- Why Study COSS?



# Introduction



Data analytics: is the process of examining **data sets** in order to draw conclusions about the information they contain, increasingly with the aid of **specialized systems** and **software**.



# Three courses

- Computer Organization and Software Systems (Core Course)
- Systems for Data Analytics (Elective)
- Big Data Systems (Elective)

# Benefits

- Understanding System Architecture
- Efficient Code Implementation
- Performance Optimization
- Memory Management
- Parallel Computing
- System-level Troubleshooting
- Integration with Existing Systems
- Scalability Considerations
- Resource Utilization

```
int count;
```

```
for (count = 1; count <= 10; count++)  
    printf("Count = %d", count);
```



# Text Books and Reference Books

## Text Books:

- (T1) W. Stallings, Computer Organization & Architecture, PHI, 10th ed., 2010.
- (T2) A Silberschatz, Abraham and others, Operating Systems Concepts, Wiley Student Edition, 8th Edition

## Reference Books:

- (R1) Patterson, David A & J L Hennenssy, Computer Organization and Design - The Hardware/Software Interface, Elsevier, 5th Ed., 2014.
- (R2) Randal E. Bryant, David R. O'Hallaron, Computer Systems - A Programmer's Perspective, Pearson, 3rd Ed, 2016.
- (R3) Tanenbaum, Modern Operating Systems: Pearson New International Edition, Pearson Education, 2013 (Pearson Online)
- (R4) Stallings, Operating Systems: Internals and Design Principles : International Edition, Pearson Education, 2013 (Pearson Online)



# Evaluation Scheme

- 5 unit course.

SI No.	Evaluation Component	Weightage %	Nature of Component
1	Mid Sem Exam	30%	CB
2	Comprehensive Examination	40%	OB
3	Quiz	5% (Two quizzes - Best of two)	OB
4	Assignments	25%	OB

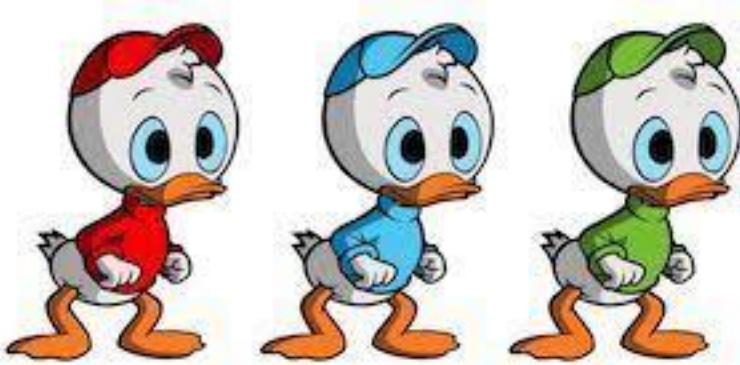


# Assignments

- Two assignments:
  - One pre-midsem exam : 10%
  - One post-midsem : 15%
- Lab based
- Simulator to be used : CPU-OS simulator
  - Nuvepro Cloud Lab

# Assignment should not be

**FILL IT.**    **SHUT IT.**    **FORGET IT.**





# General Instructions

1. Always use note book for writing important points and for solving problems
2. Use chat box for writing subject related questions
3. Do not repeat the questions on chat box. Questions will be answered during last 10 minutes of the session
4. Unanswered questions will be put up on the eLearn discussion forum
5. In case of query:
  - Subject related queries (Evaluation, Assignments, Quiz, discrepancy in Marks etc.) → Prof. Vaibhav Jain ([vaibhav.jain@wilp.bits-pilani.ac.in](mailto:vaibhav.jain@wilp.bits-pilani.ac.in))
  - Operation related queries ( Admission, registration, Exam related matters, grade sheets etc.) → [support@wilp.bits-pilani.ac.in](mailto:support@wilp.bits-pilani.ac.in)



# Today's Session

Contact Hour	List of Topic Title	Text/Ref Book/external resource
1-2	<b>Introduction to Computer Systems</b> <ul style="list-style-type: none"><li>• Hardware Organization of a computer</li><li>• Basic uniprocessor architecture</li><li>• Instruction Cycle State Diagram</li><li>• Operating System role in Managing Hardware</li><li>• Running a Hello Program</li></ul>	T1



# Definition of a Computer

- Is a complex system
- Is a programmable device
- Must be able to **process** data
- Must be able to **store** data
- Must be able to **move** data
- Must be able to **control** above three functions

# Computer System



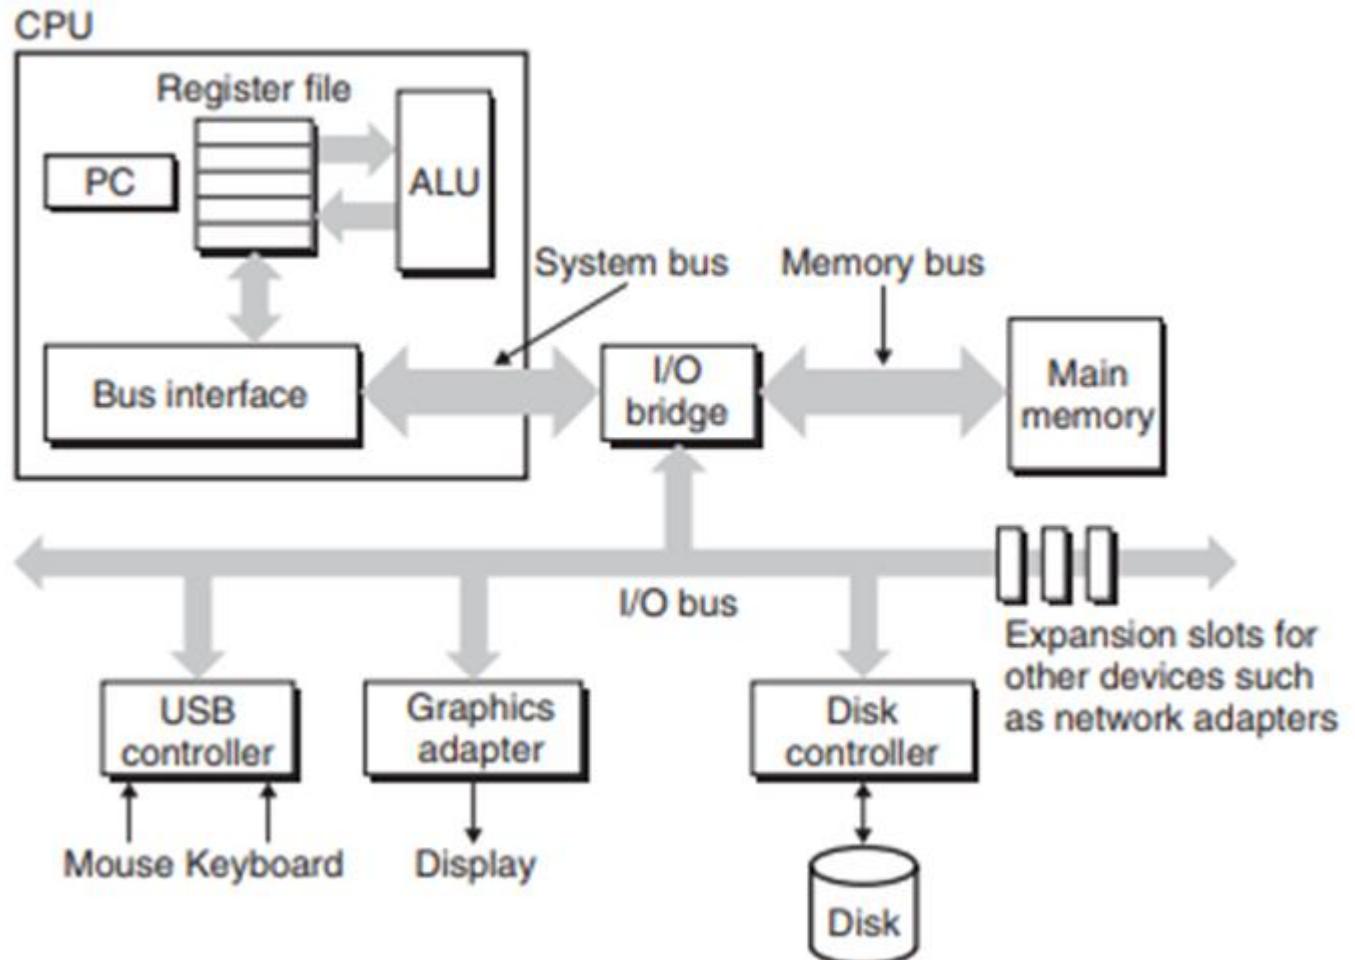
- **Hardware**

- Central Processing Unit (CPU)
- Memory
- I/O devices

- **Software**

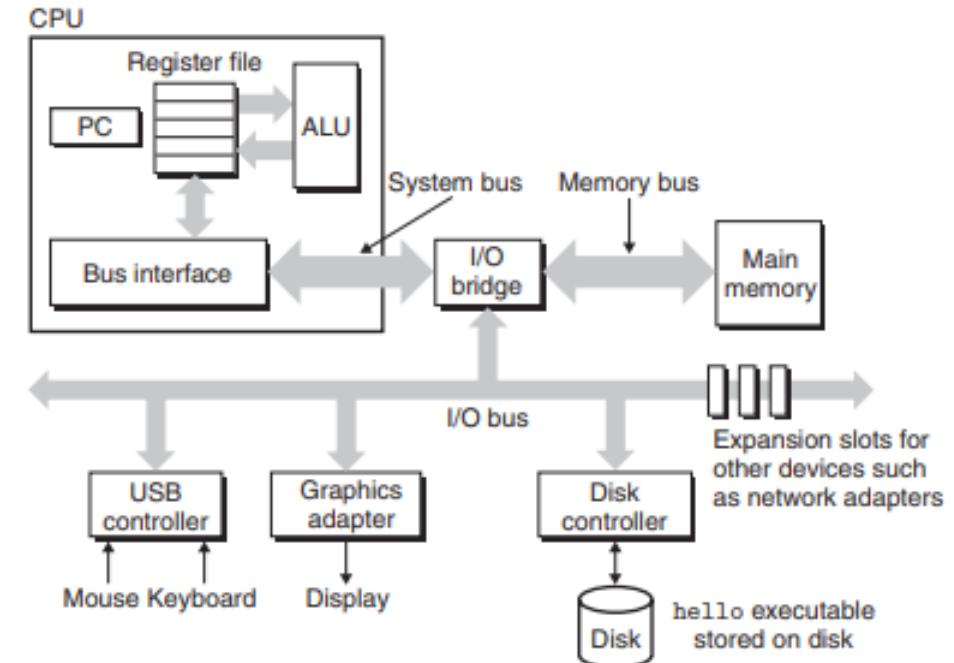
- System Software
  - System Management Software
  - Tools and Utilities for Developing the software
- Application Software
  - General Purpose Software
  - Specific Purposed Software

# Hardware Organization of a computer



# Von Neumann Architecture

- Three key concepts:
  - Data and instructions are stored in a single read - write memory
  - The contents of this memory are addressable by location, without regard to the type of data contained there
  - Execution occurs in a sequential fashion (unless explicitly modified) from one instruction to the next



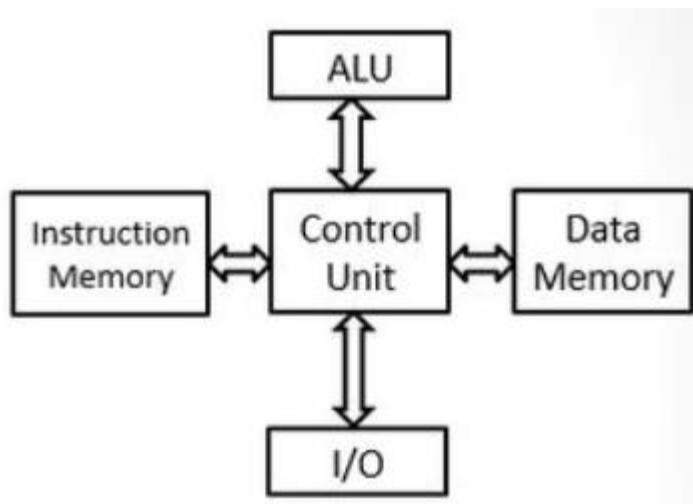


# Von Neumann Architecture...

- Stored-program computers have the following characteristics:
  - Three hardware systems:
    - A central processing unit (CPU)
    - A main memory system
    - An I/O system
  - The capacity to carry out sequential instruction processing.
  - A single path between the CPU and main memory.
    - This single path is known as the *von Neumann bottleneck*.
    - Side effect : reduced throughput (Data Rate)

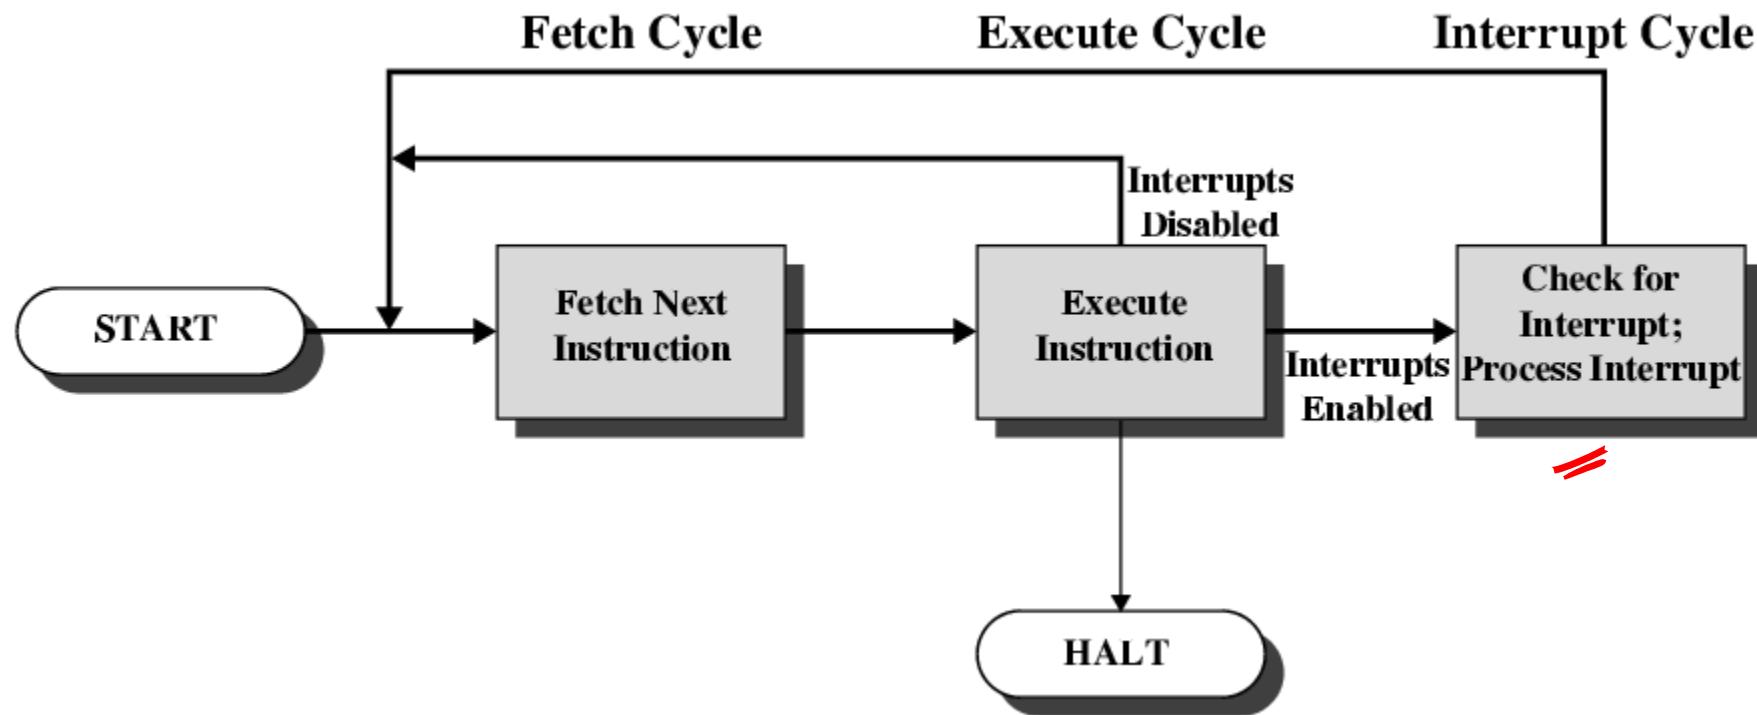
# Harvard Architecture

- Uses two memory systems and two separate busses
  - Instruction Memory
  - Data Memory



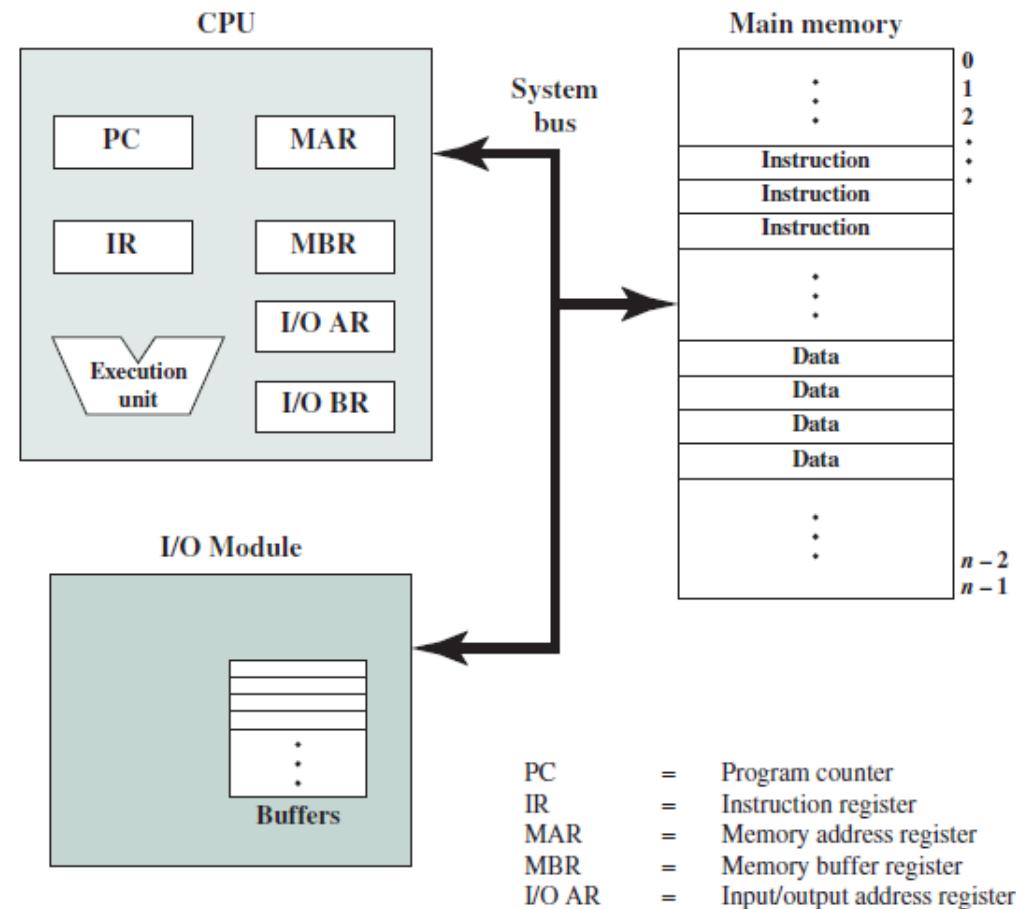
# Instruction Cycle Diagram

- Instruction execution : Two steps:
  - Fetch
  - Execute
- Interrupt: Interrupt is checked at the end of Instruction cycle



# Fetch Cycle

- Program Counter (PC) holds address of next instruction to be fetched
- Processor fetches instruction from memory location pointed to by PC
- Instruction loaded into Instruction Register (IR)
- Processor interprets instruction and performs required actions during execution cycle
- Increment PC
  - Unless told otherwise





# Execute Cycle

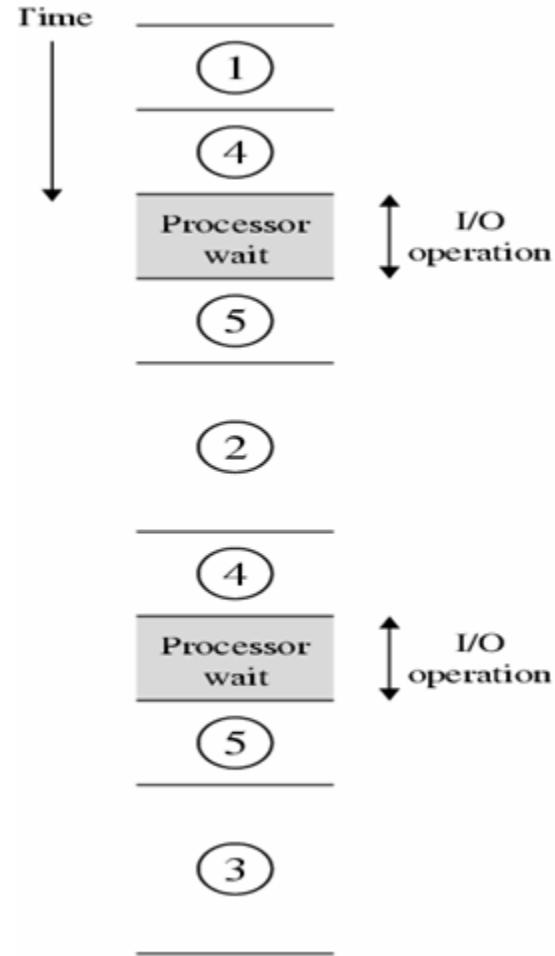
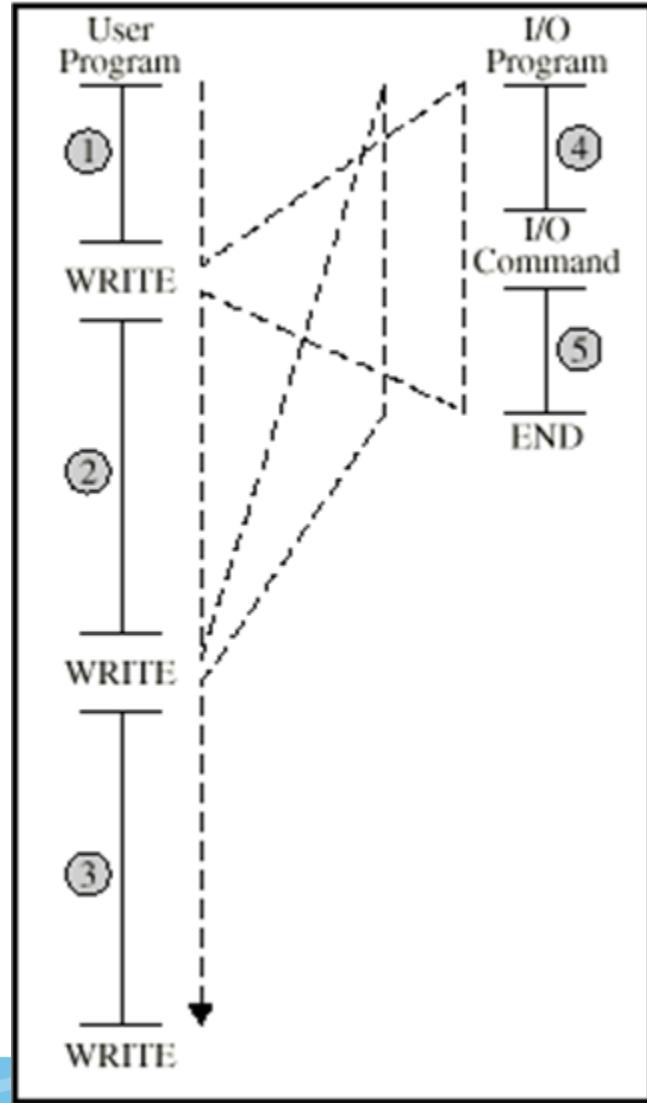
- Processor - memory
  - Data transfer between CPU and main memory
- Processor - I/O
  - Data transfer between CPU and I/O module
- Data processing
  - Some arithmetic or logical operation on data
- Control
  - Alteration of sequence of operations
    - e.g. jump
- Combination of above



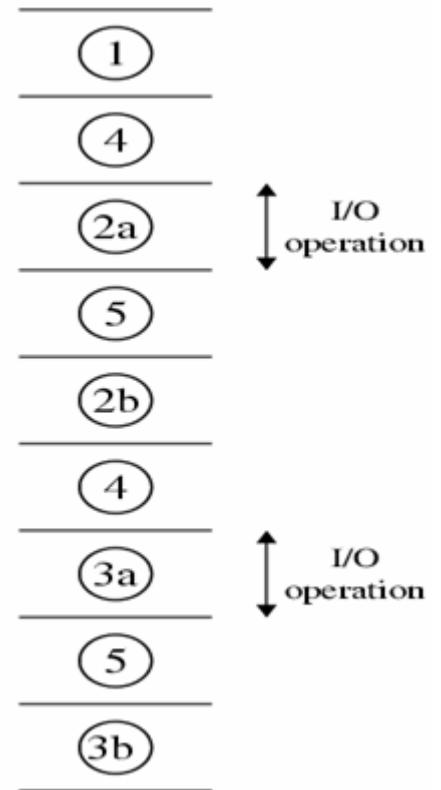
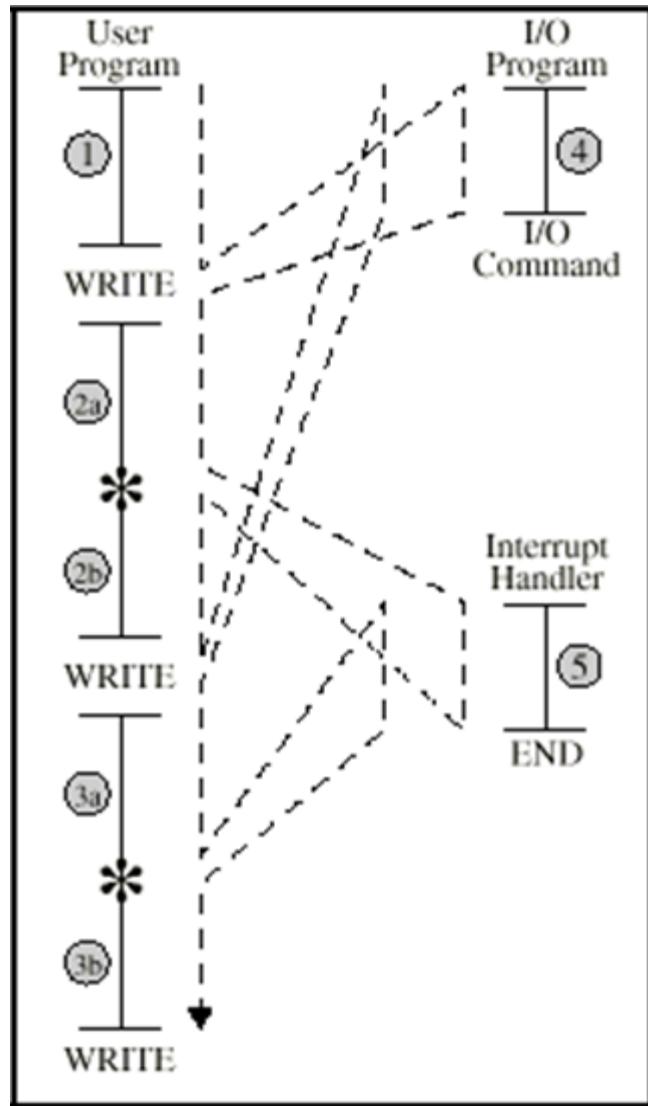
# Interrupt Cycle

- Interrupts: Mechanism by which other modules (e.g. I/O) may interrupt normal sequence of processing
- Interrupts enhances processing efficiency

# Program Flow Control (No Interrupts)



# Program Flow Control (With Interrupts)





# Types of Interrupts

- Types of interrupts:
  - Program
    - e.g. overflow, division by zero
  - Timer
    - Generated by internal processor timer
    - Used in pre-emptive multi-tasking
  - I/O
    - from I/O controller
  - Hardware failure
    - e.g. memory parity error

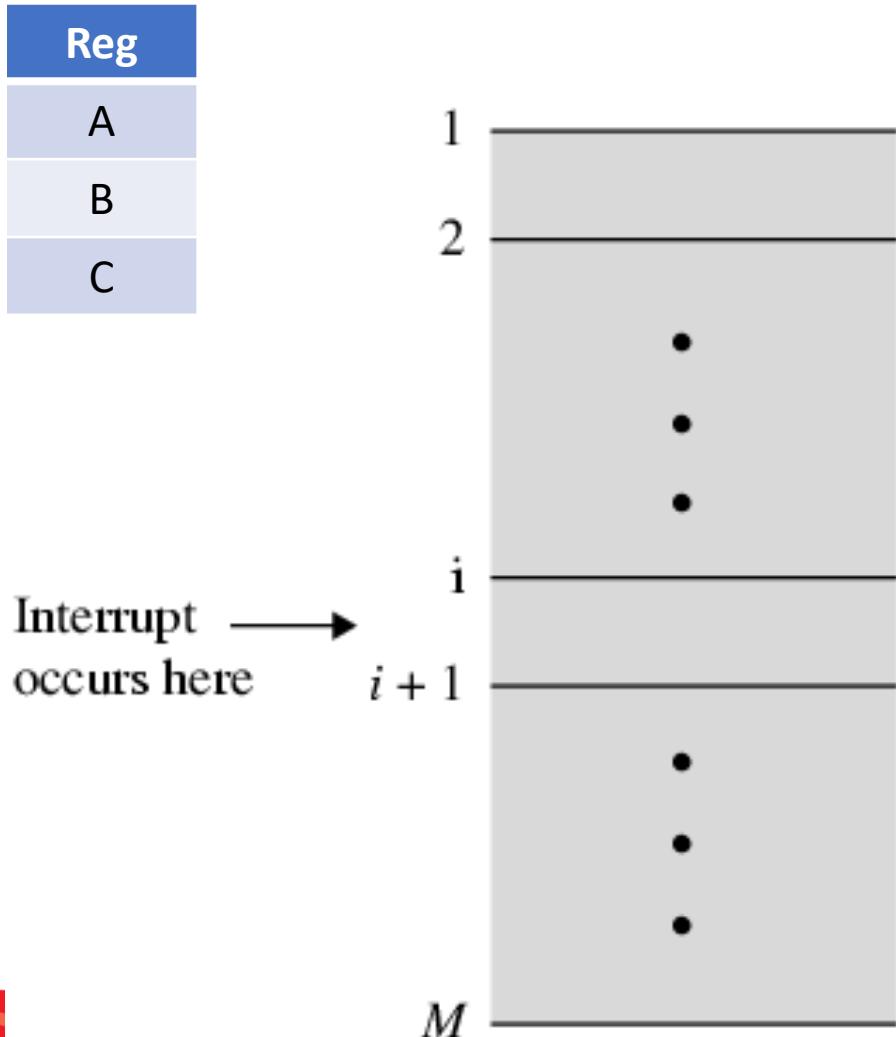


# Interrupt Cycle

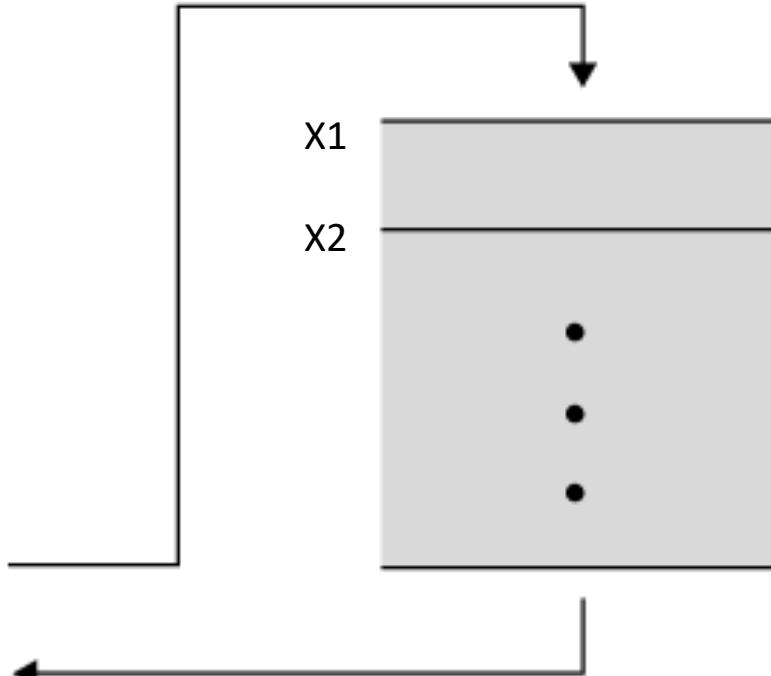
- Processor checks for interrupt
  - Indicated by an interrupt signal
- If no interrupt, fetch next instruction
- If interrupt pending:
  - Suspend execution of current program
  - Save context
  - Set PC to start address of interrupt handler routine
  - Process interrupt
  - Restore context and continue interrupted program

# Transfer of Control via Interrupts

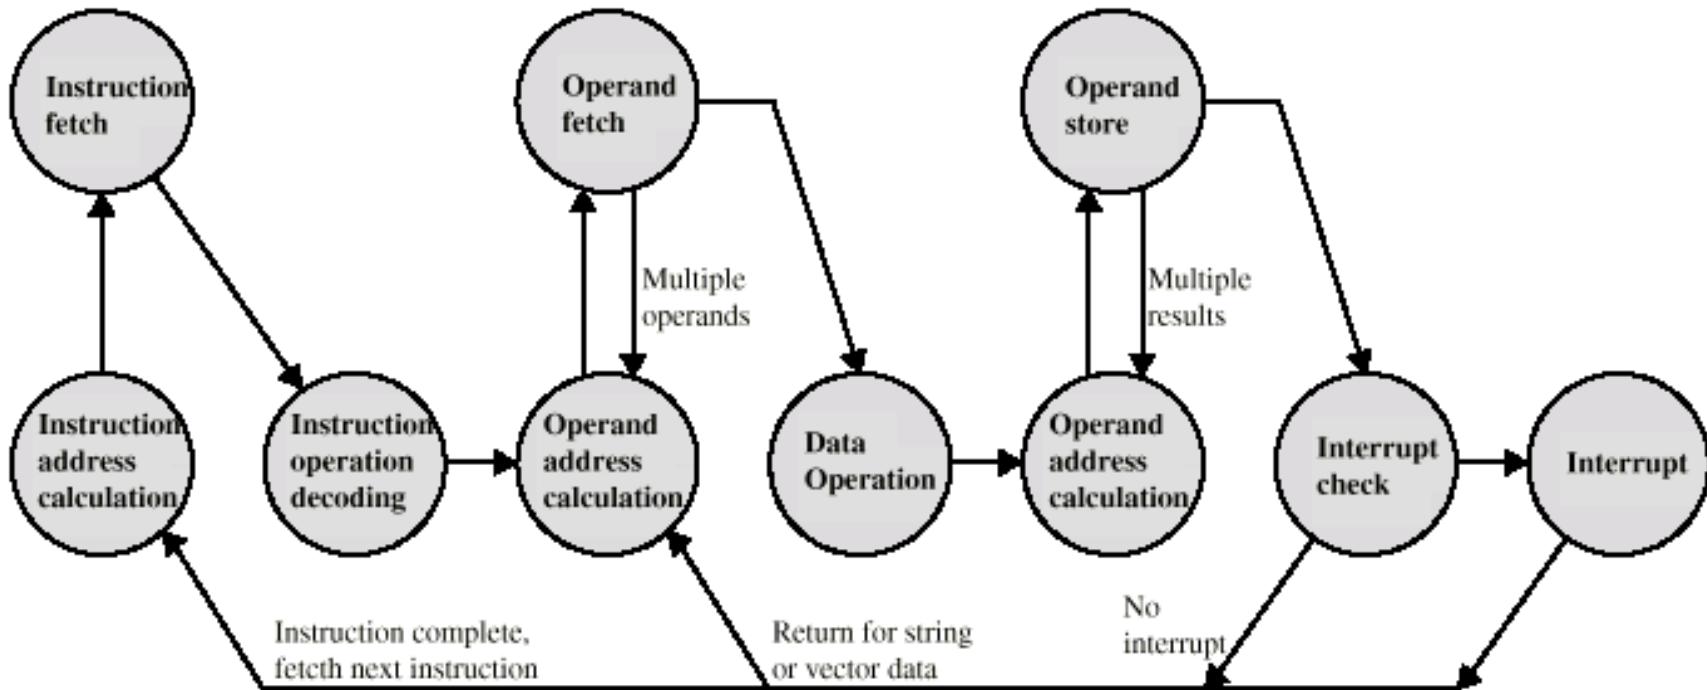
User Program



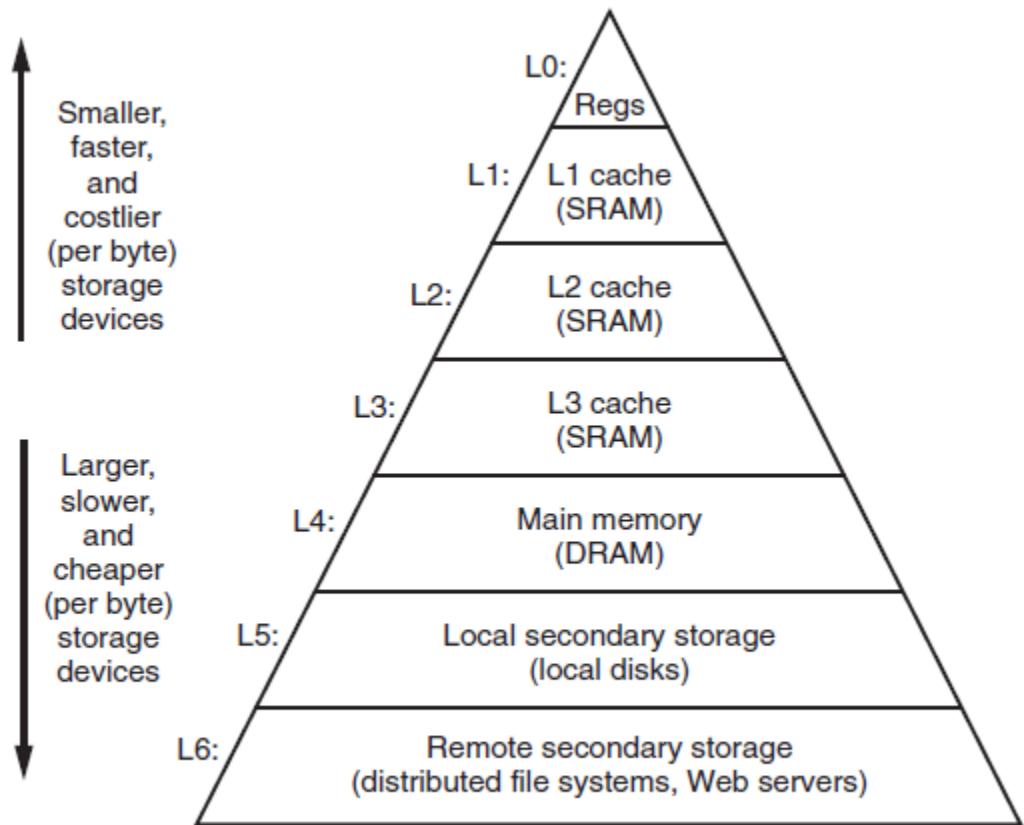
Interrupt Handler



# Instruction Cycle - State Diagram

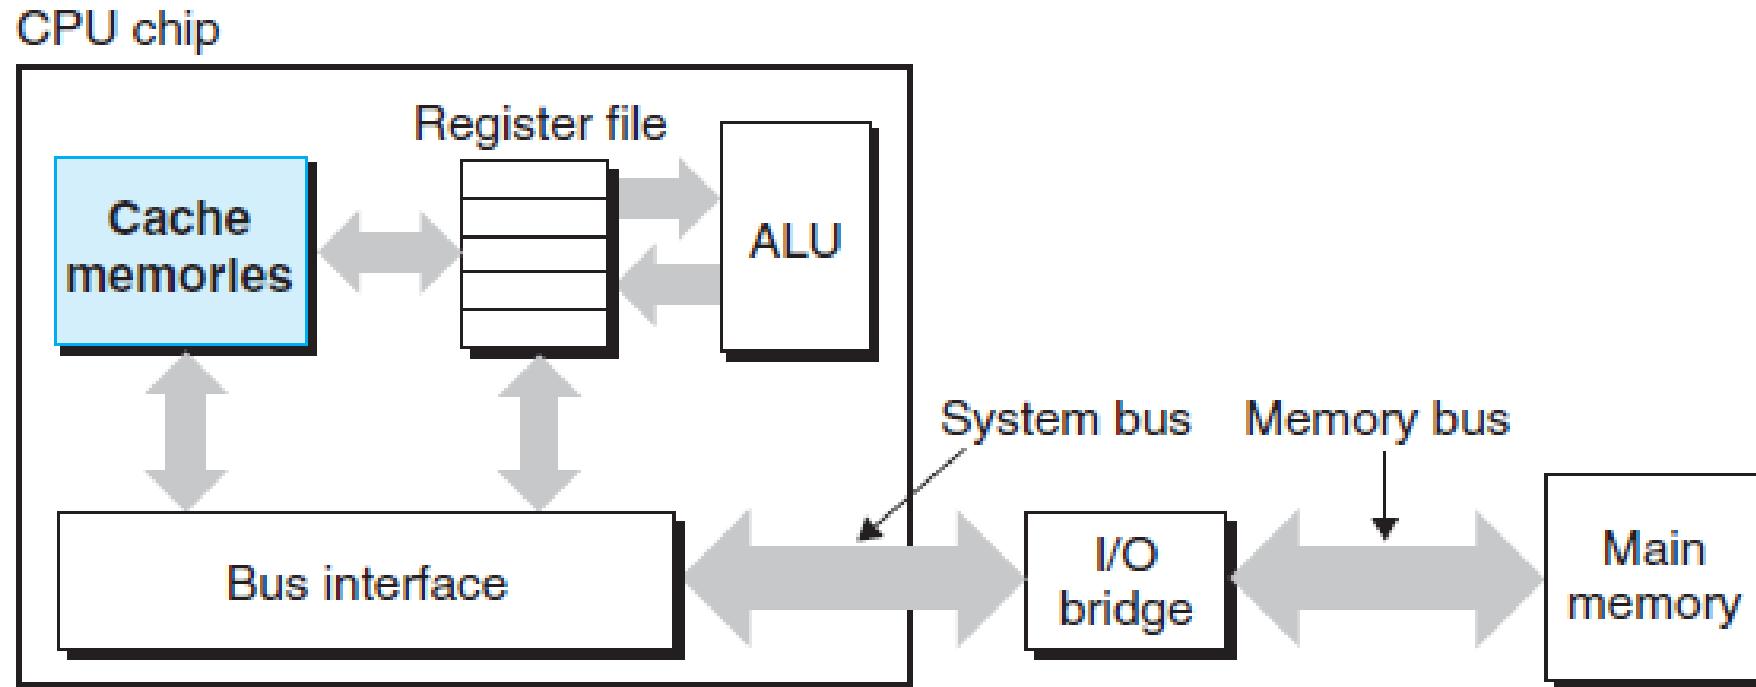


# Memory Hierarchy



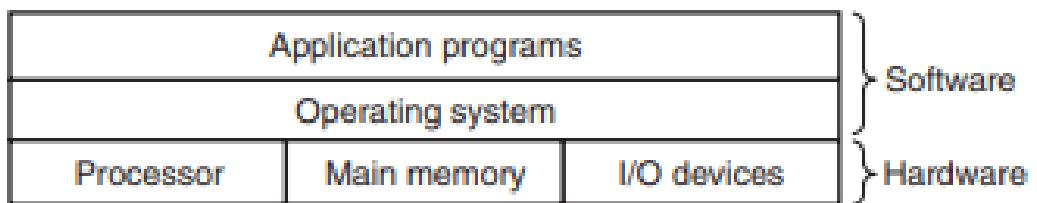
An example of a memory hierarchy.

# Role of Cache Memory



# Operating System

- collection of software/ Program that acts as an intermediary between an user of a computer and the computer hardware.
- is a program that helps to run all the other programs
- Three main functions:
  - Resource management
  - Establish an user interface
  - Execute and provide services for application software



Layered view of a  
computer system.



# Main objectives

- Convenience
- Efficiency
- Ability to evolve and offer new services
- Maximize System performance
- Protection and access control
- Footprint of OS should be small



# Important Note

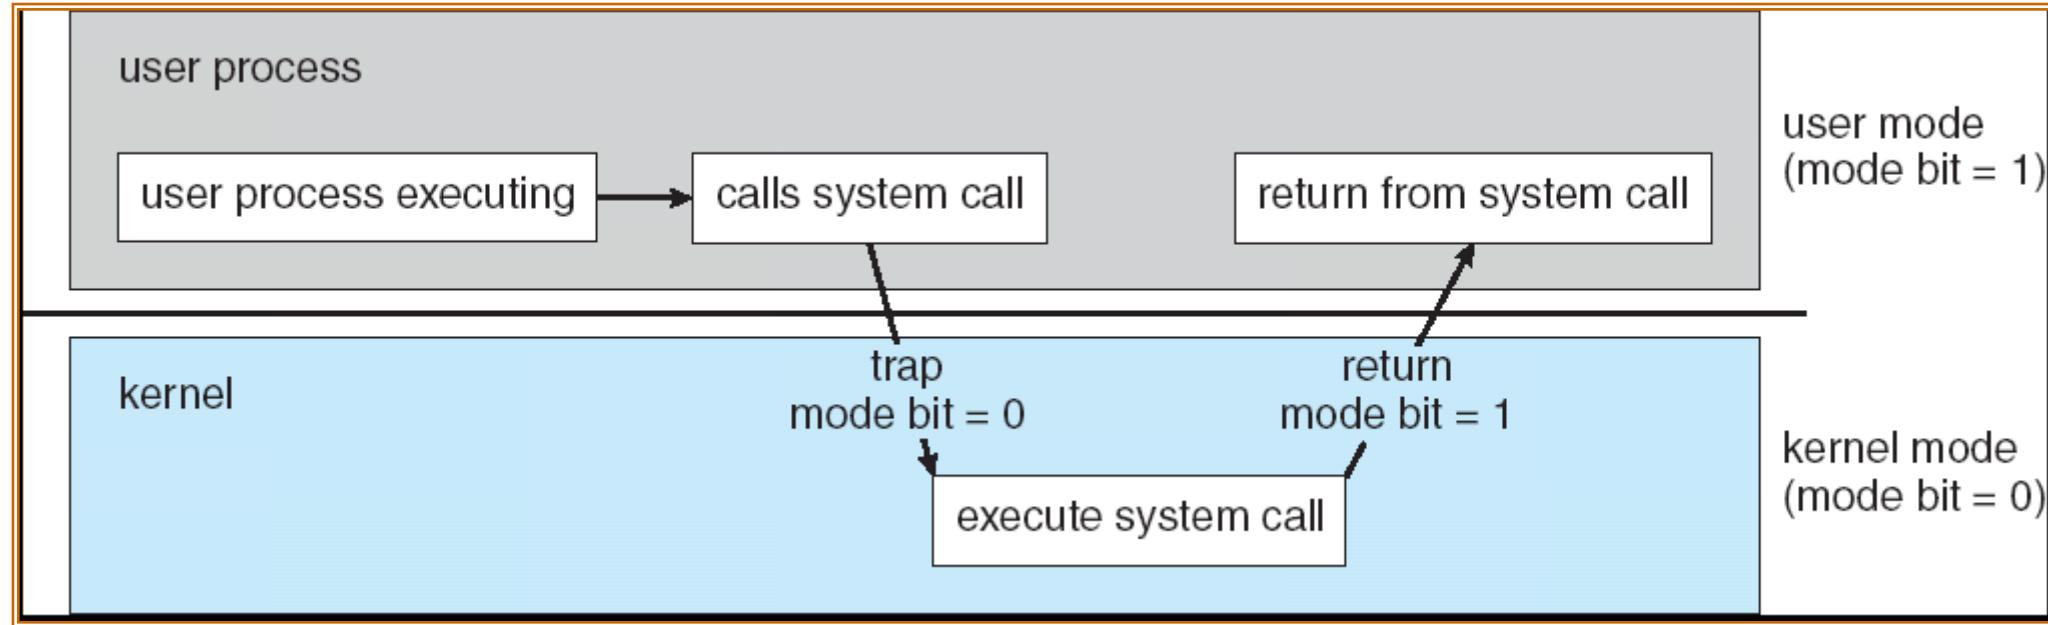
- **bootstrap program** is loaded at power-up or reboot
  - Typically stored in ROM or EPROM, generally known as **firmware**
  - Initializes all aspects of system
  - Loads operating system kernel and starts execution
- “The one program running at all times on the computer” is the **kernel**. Everything else is either a system program (ships with the operating system) or an application program



# Operating System Operations

- Dual-mode operation
  - User mode
  - Kernel mode ( also known as System Mode / Supervisor mode/ privileged mode )
- User mode(1):
  - user program executes in user mode
  - certain areas of memory are protected from user access
  - certain privileged instructions may not be executed
- Kernel Mode (0)
  - privileged instructions may be executed
  - protected areas of memory may be accessed

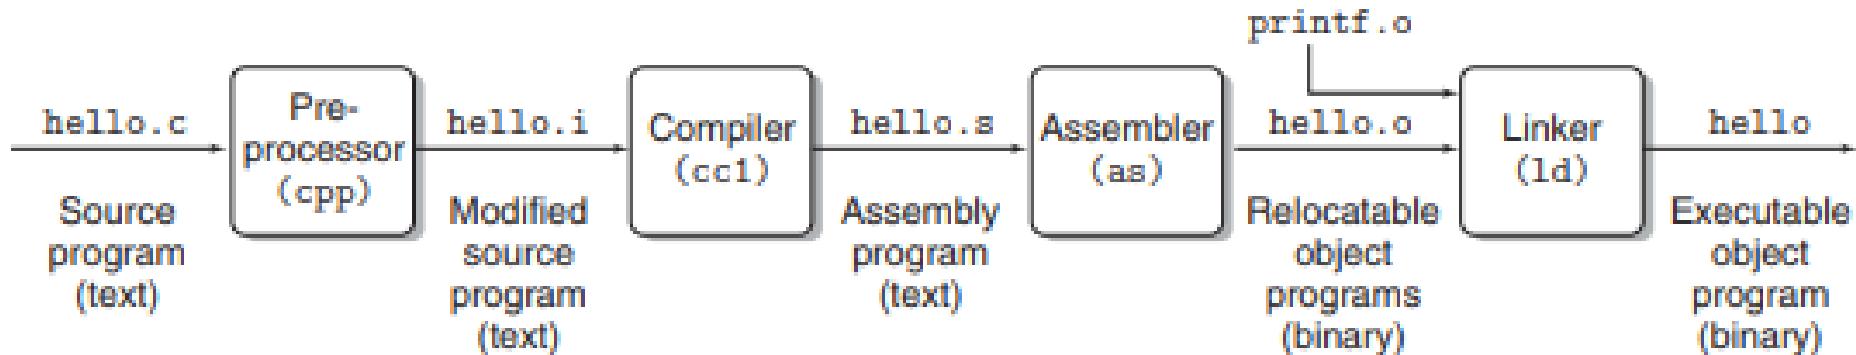
# Transition from user to kernel mode



# Running a Hello.c Program

```
#include <stdio.h>

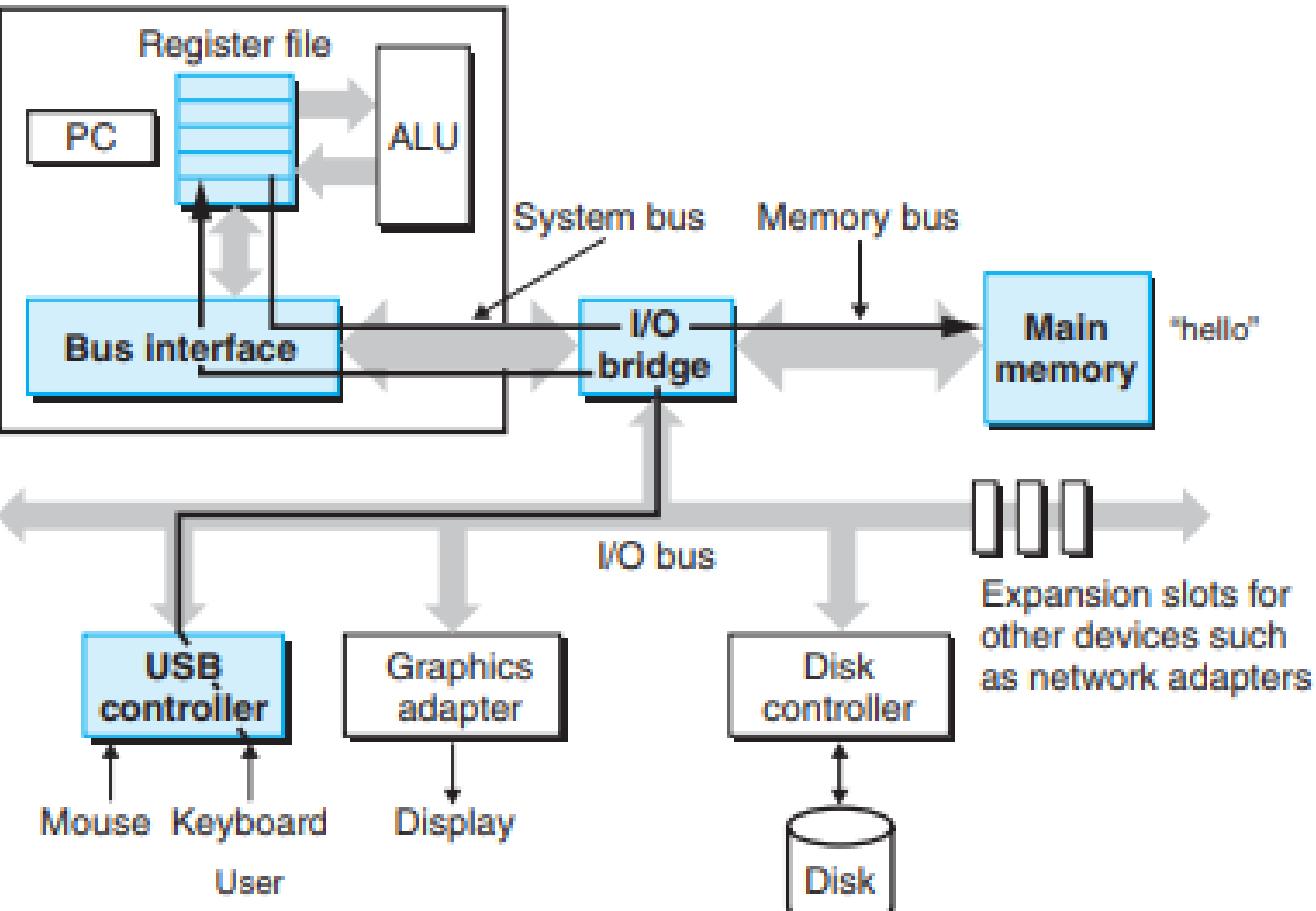
int main()
{
    printf("hello, world\n");
}
```



The compilation system.

# Reading ./hello command from Keyboard

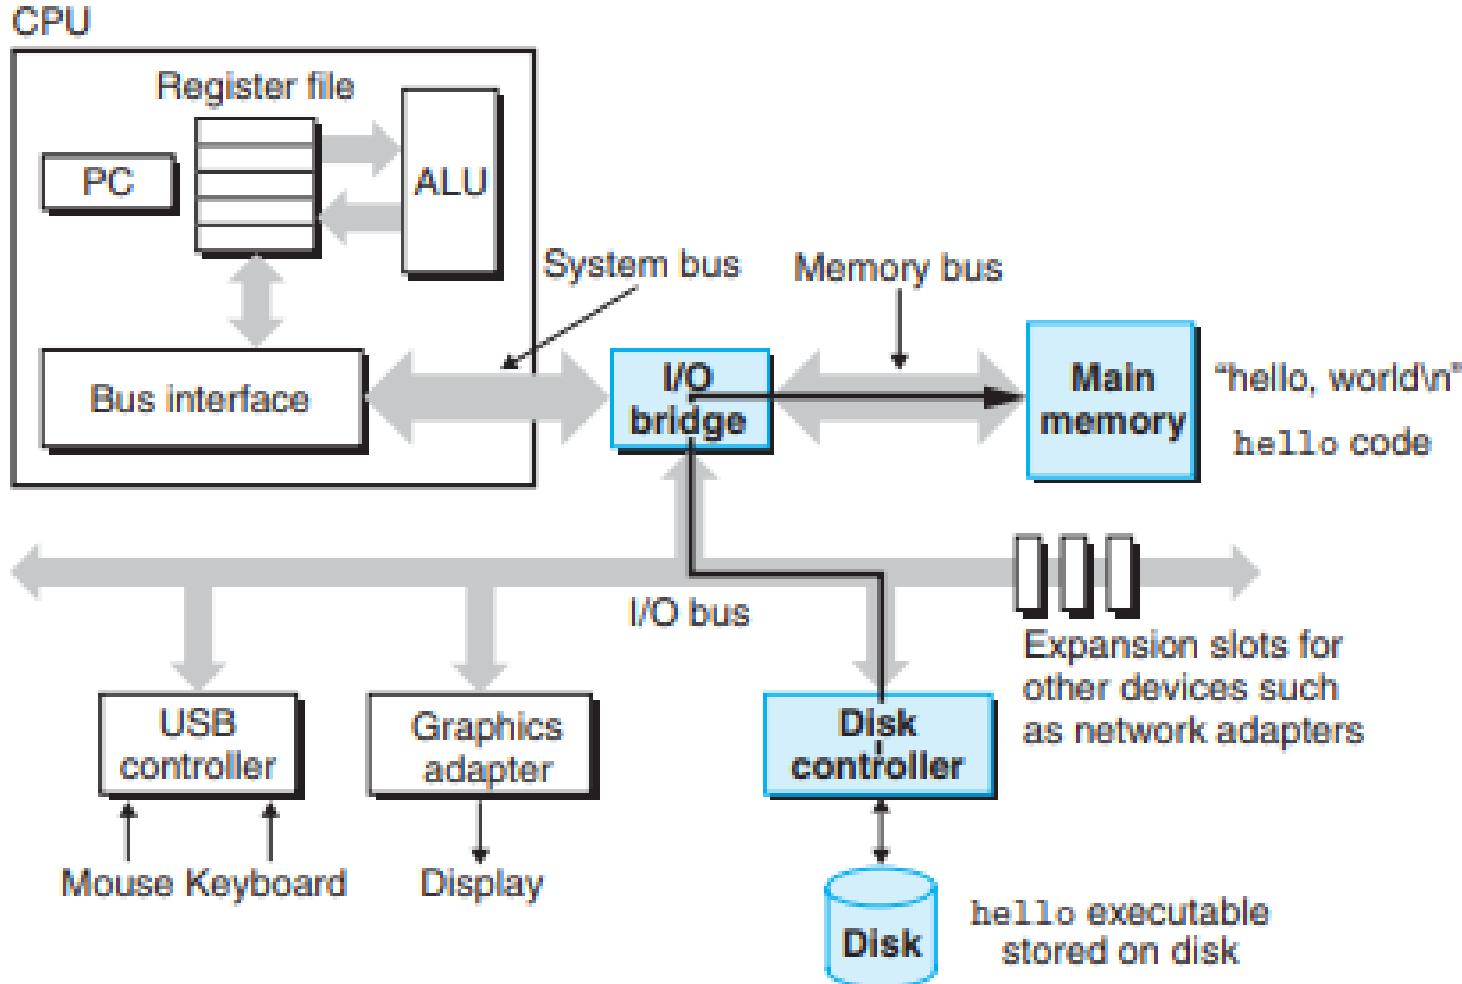
CPU



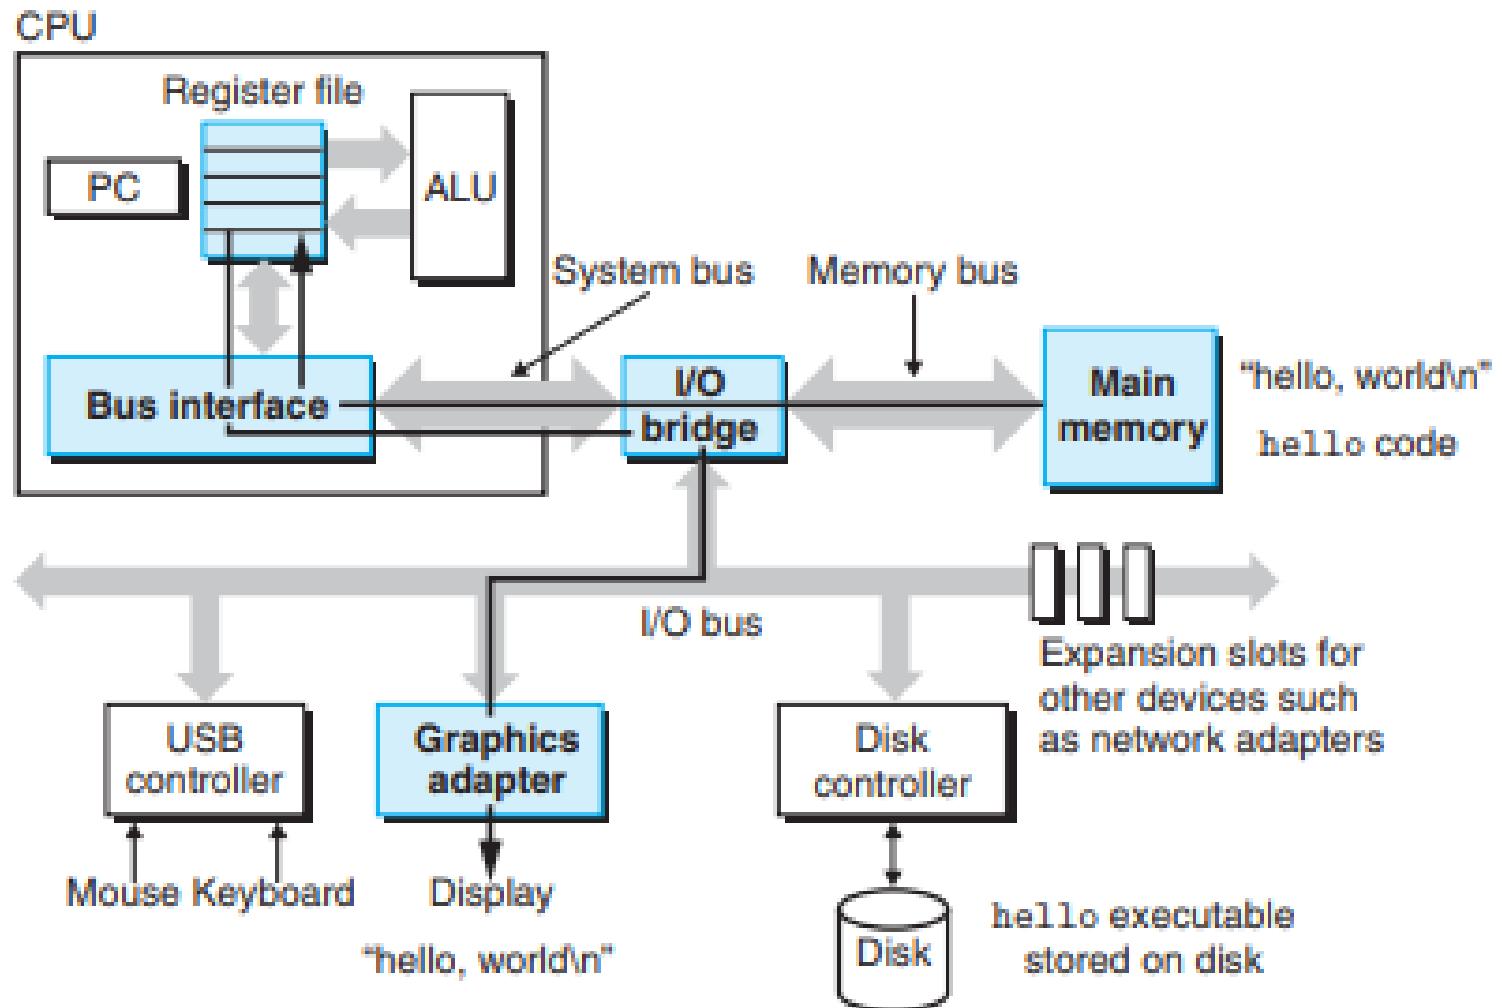
```
#include <stdio.h>
```

```
int main()
{
    printf("Hello, world\n");
}
```

# Loading the executable from disk into main memory



# Writing the output string from memory to the display





# Why do we need to know how compilation works?

- Optimizing program performance.
- Understanding link-time errors
- Avoiding security holes.



# Computer Organization and Software Systems

Contact Session 2

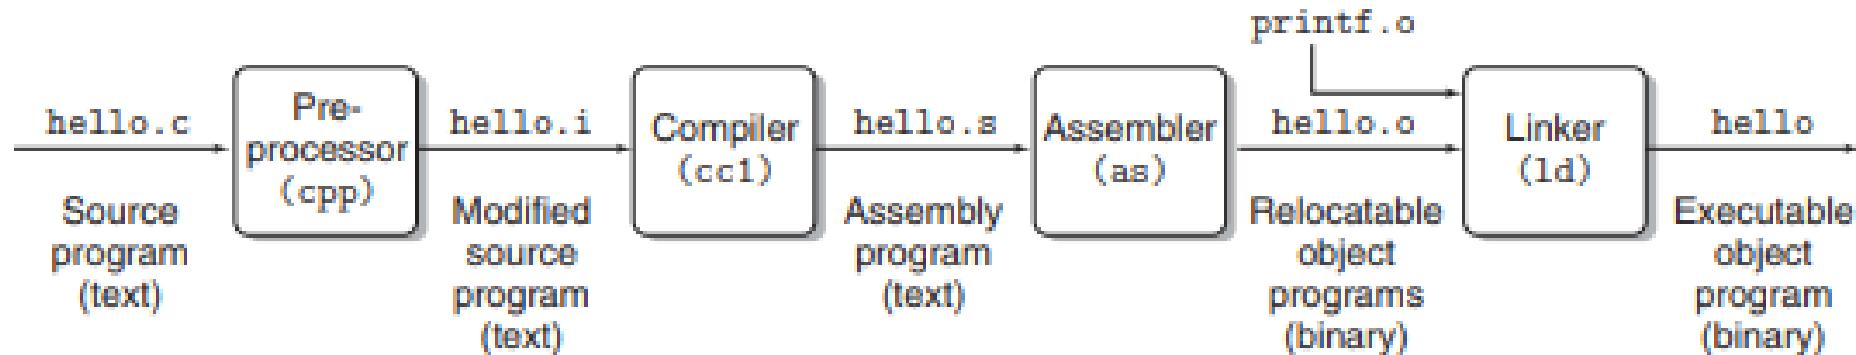
Dr. Lucy J. Gudino



# Running a Hello.c Program

```
#include <stdio.h>

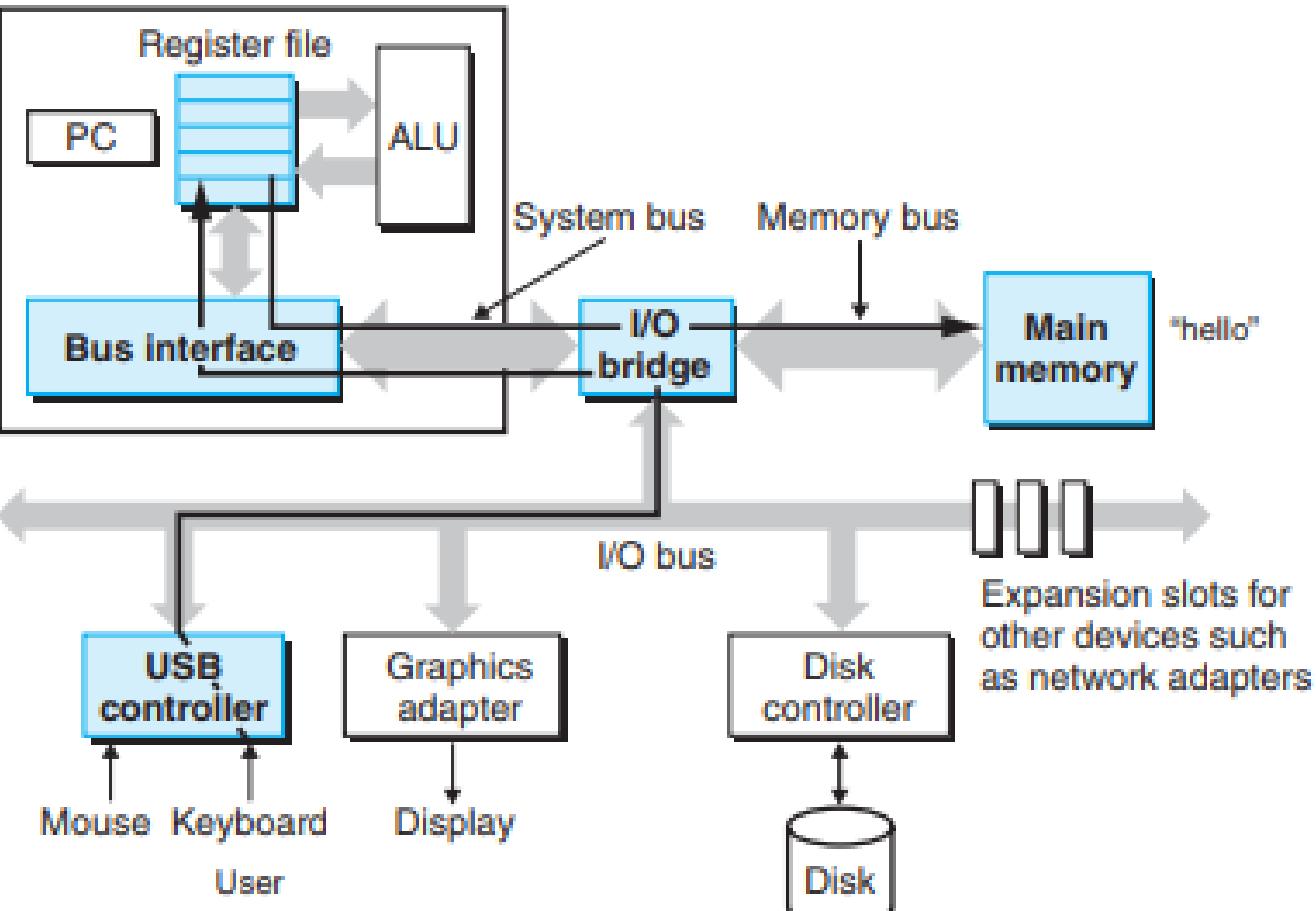
int main()
{
    printf("hello, world\n");
}
```



The compilation system.

# Reading ./hello command from Keyboard

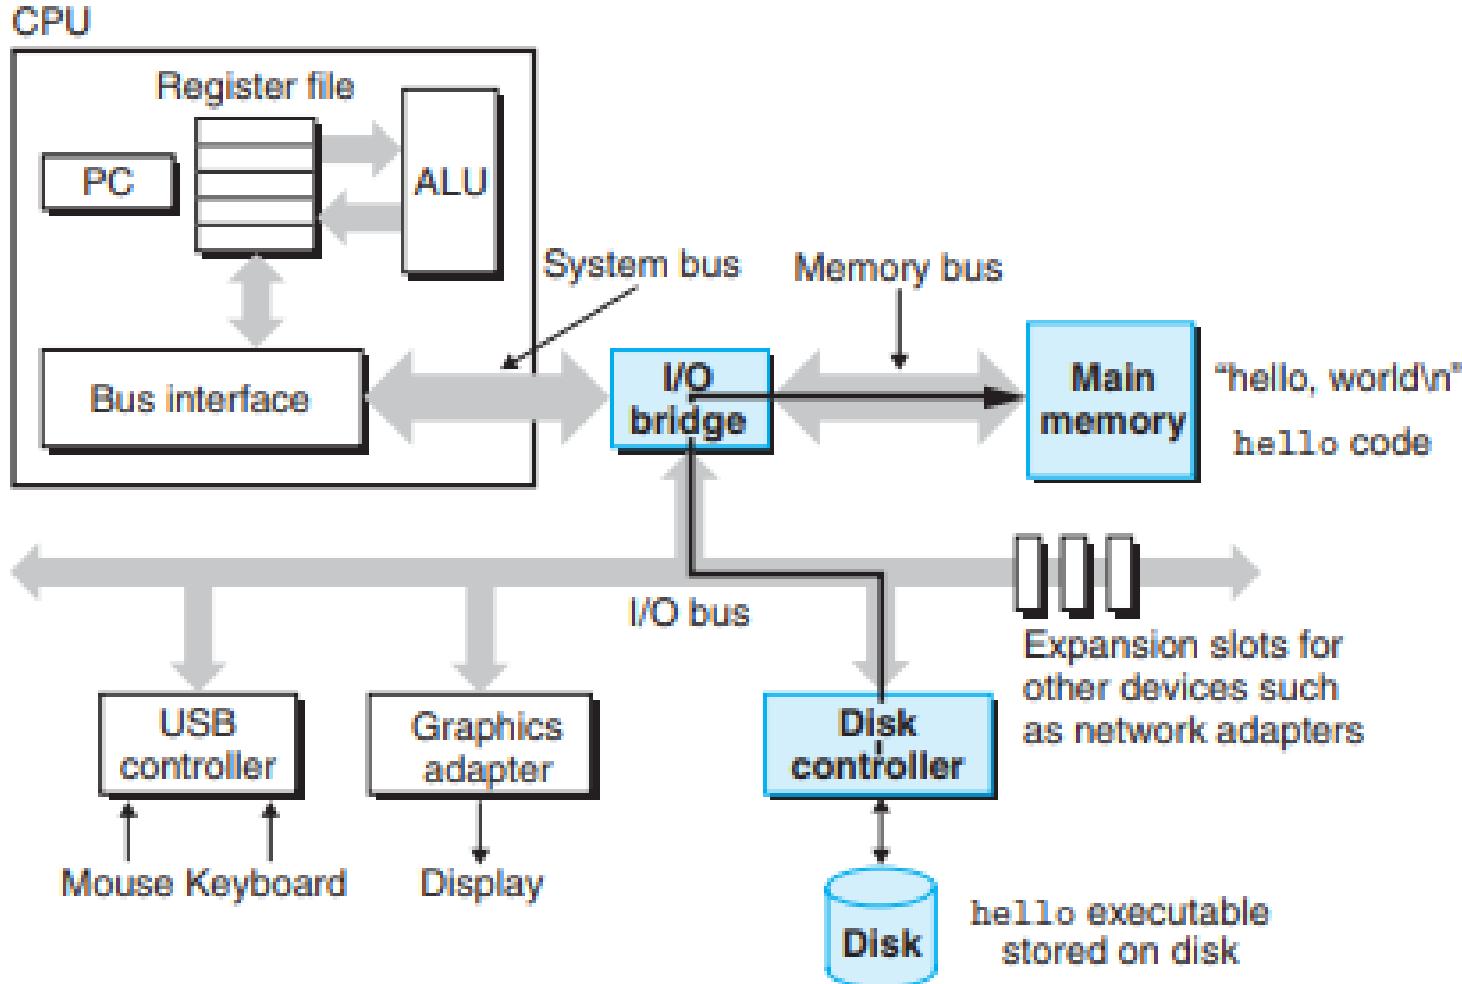
CPU



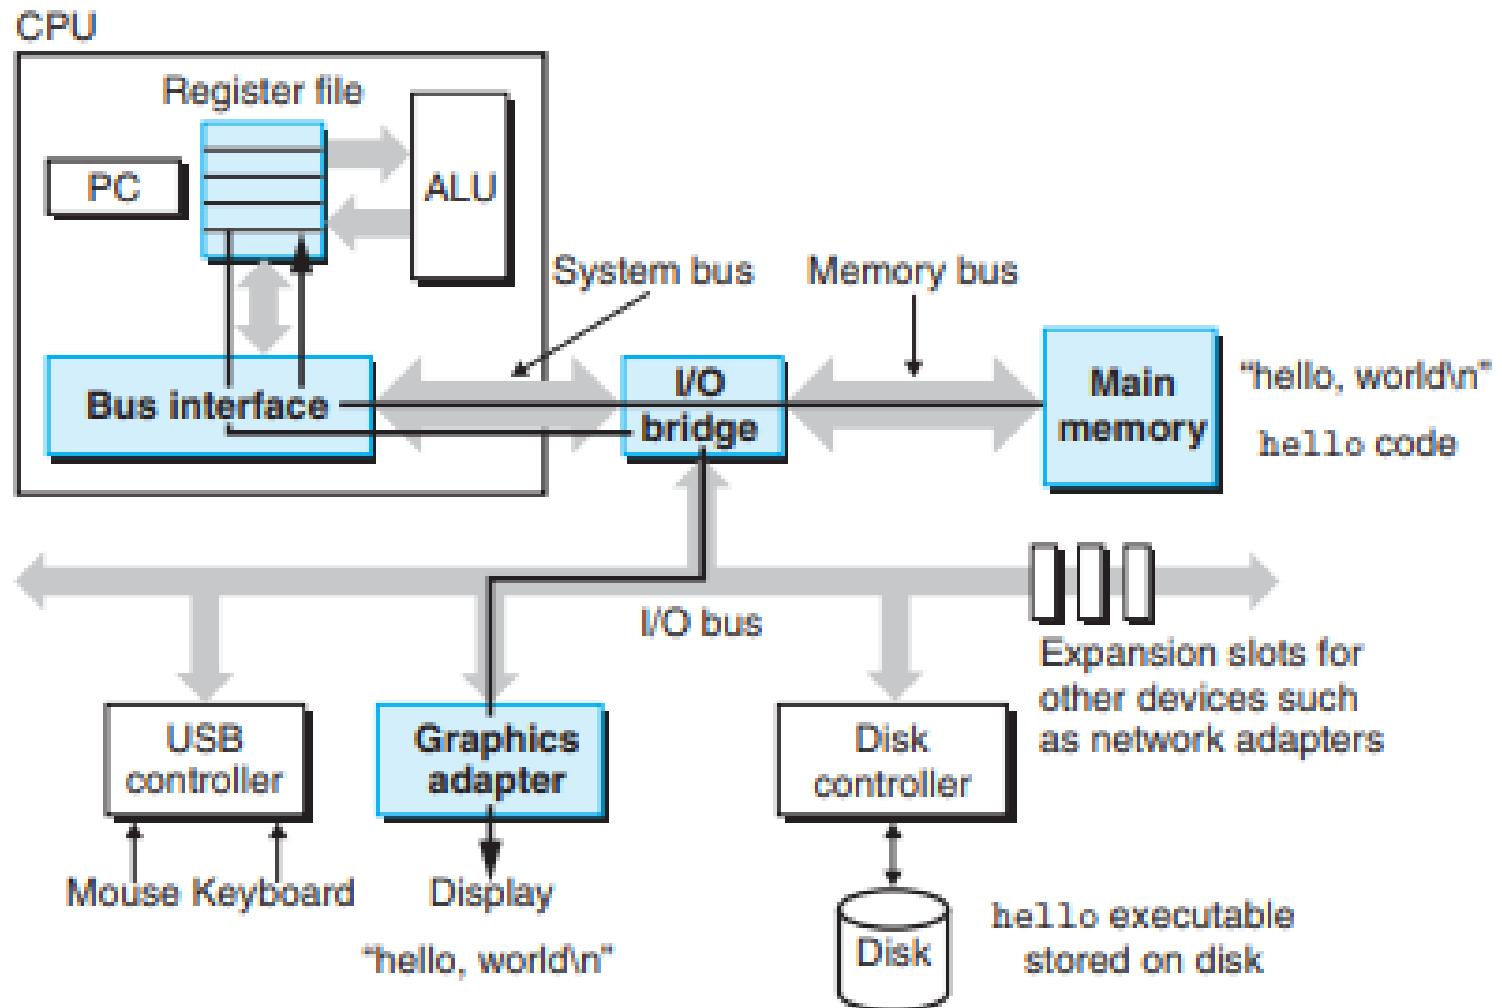
```
#include <stdio.h>
```

```
int main()
{
    printf("Hello, world\n");
}
```

# Loading the executable from disk into main memory



# Writing the output string from memory to the display





# Why do we need to know how compilation works?

- Optimizing program performance.
- Understanding link-time errors
- Avoiding security holes.

# Today's Class



Contact Hour	List of Topic Title	Text/Ref Book/external resource
3-4	<b>Memory Organization</b> - Internal Memory - External Memory	T1, R2

# Memory Organization

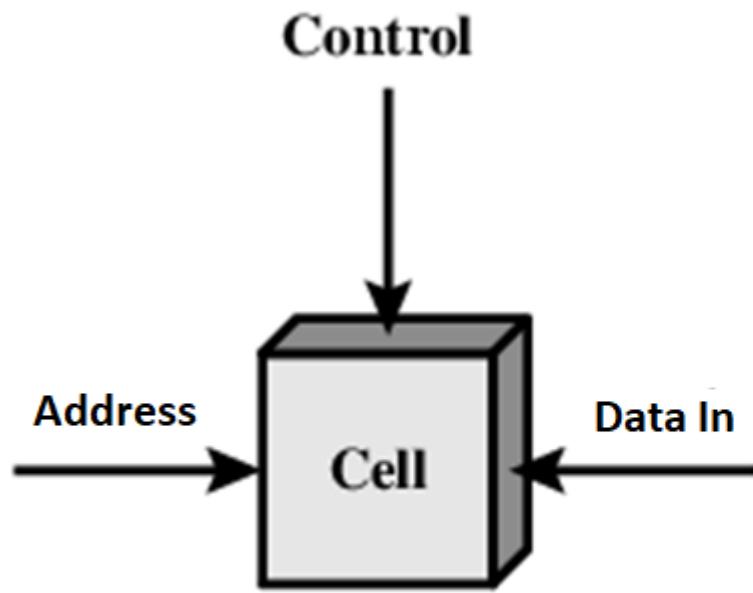




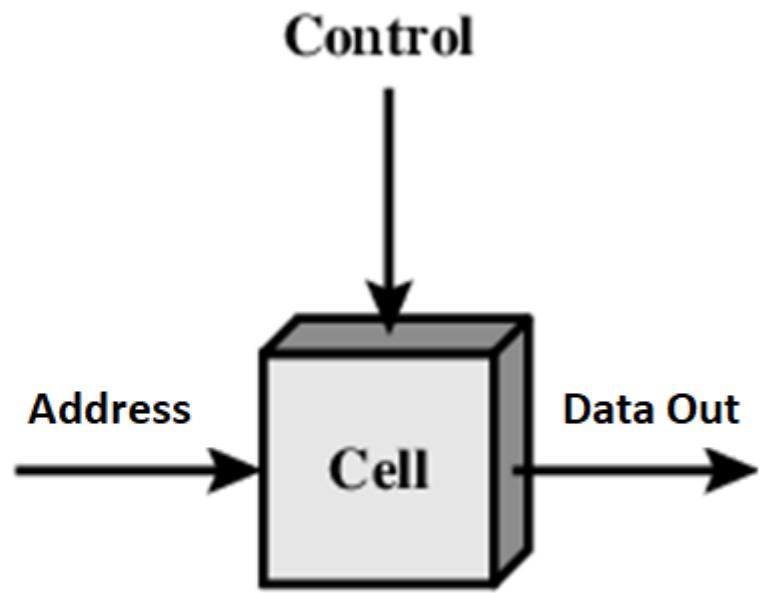
# Internal Memory Organization

- Internal Memory
  - Also known as main memory or Primary Memory
  - Small data storage but quick access.
  - Examples : RAM, ROM
- External Memory
  - Also Known as secondary Memory
  - Huge data stored persistently
  - Examples: hard disk, solid state drives, USB flash drives etc.

# Semiconductor Memory

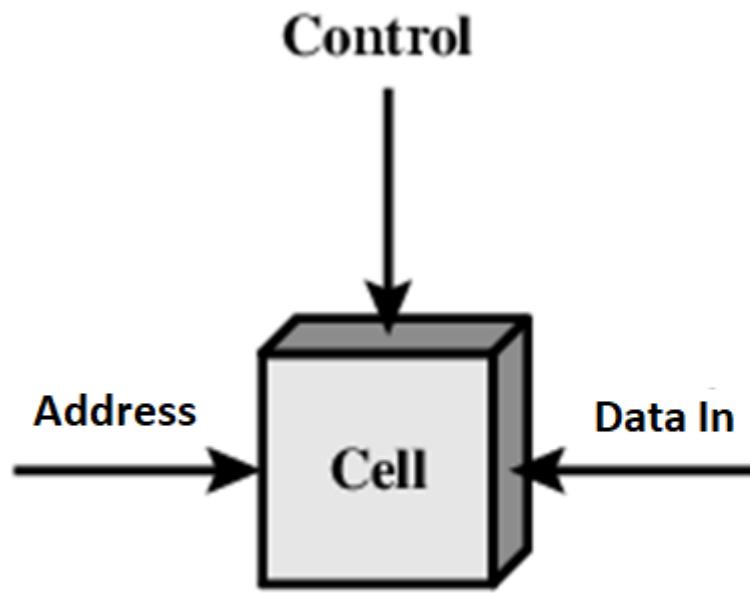


(a) Write

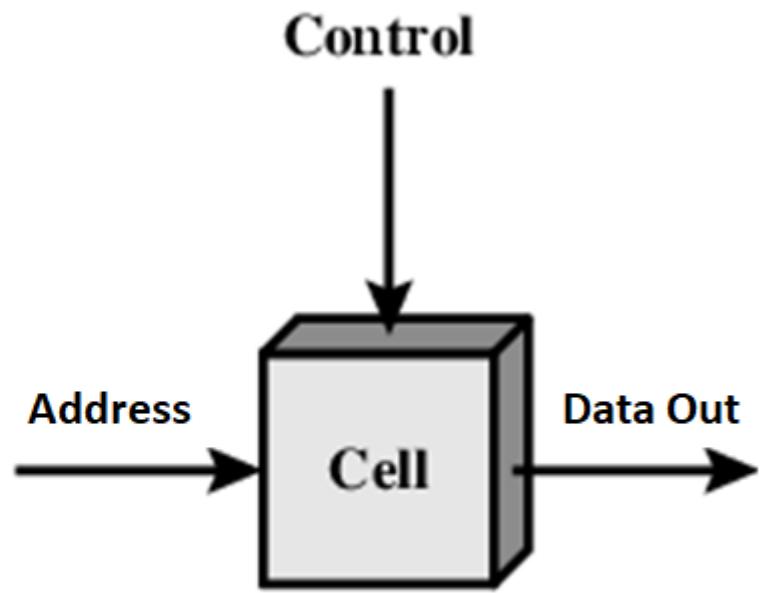


(b) Read

# Semiconductor Memory

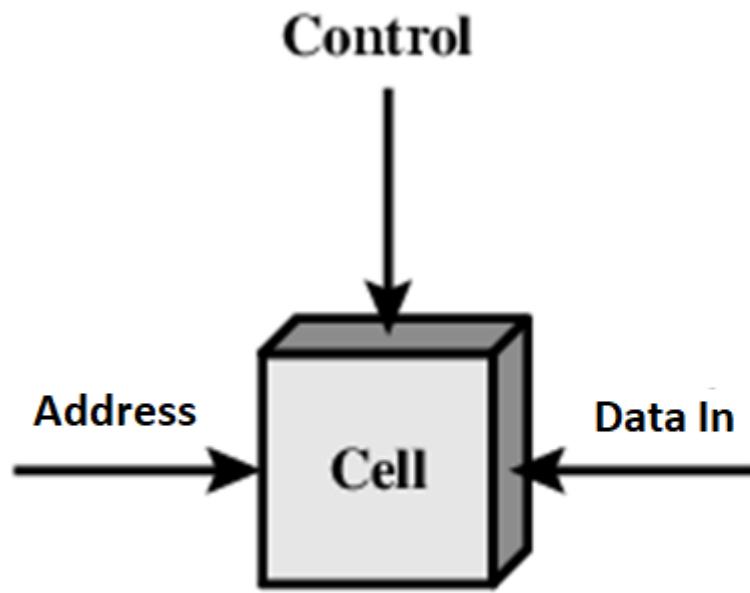


(a) Write

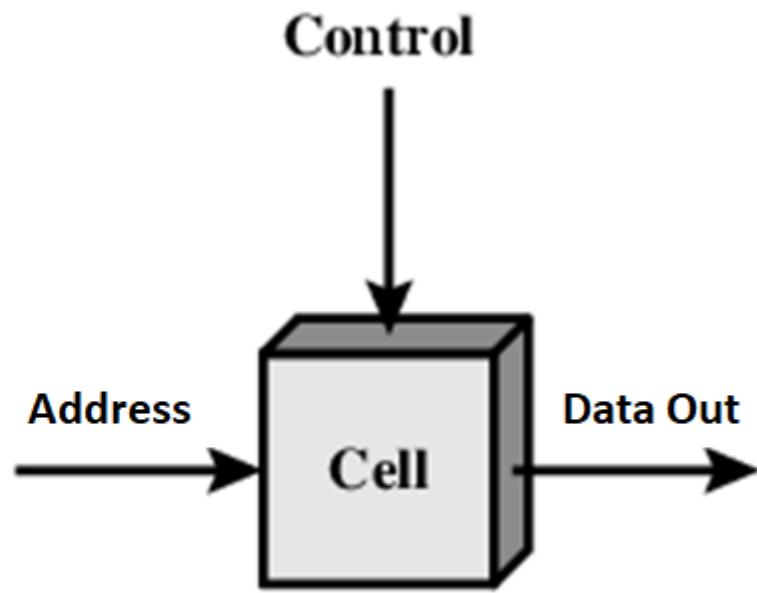


(b) Read

# Semiconductor Memory



(a) Write



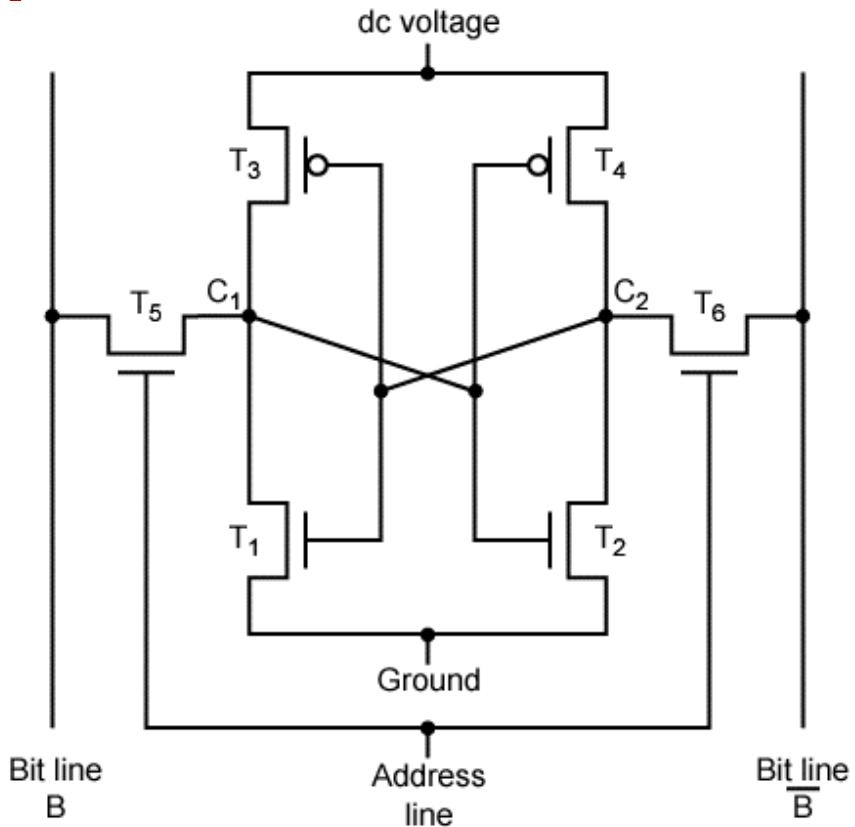
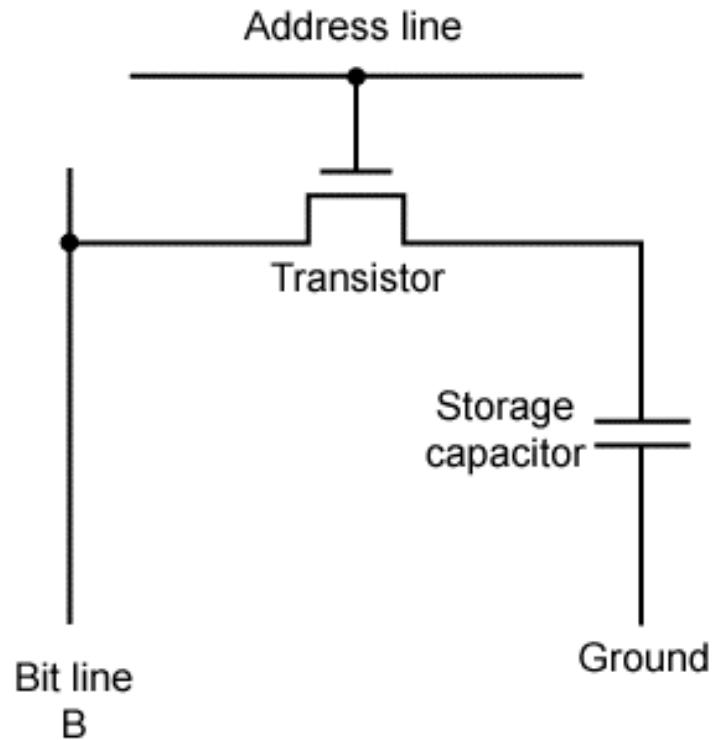
(b) Read



# Random-Access Memory (RAM)

- Key features
  - RAM is traditionally packaged as a chip.
  - Basic storage unit is normally a cell (one bit per cell).
  - Multiple RAM chips form a memory.
- RAM comes in two varieties:
  - SRAM (Static RAM)
  - DRAM (Dynamic RAM)
- SRAM and DRAM are volatile memories
  - Lose information if powered off.

# DRAM vs SRAM



	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 to 6	1X	No	Maybe	100x	Cache
DRAM	1	10X	Yes	Yes	1X	Main memories, frame buffers



# Read Only Memory

- Permanent Storage and Nonvolatile Memories
- Read Only Memory Variants:
  - Read-only memory (**ROM**): programmed during production
  - Programmable ROM (**PROM**): can be programmed once
  - Erasable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
  - Electrically erasable PROM (**EEPROM**): electronic erase capability
  - Flash memory: EEPROMs. with partial (block-level) erase capability
    - Wears out after about 100,000 erasing
- Firmware

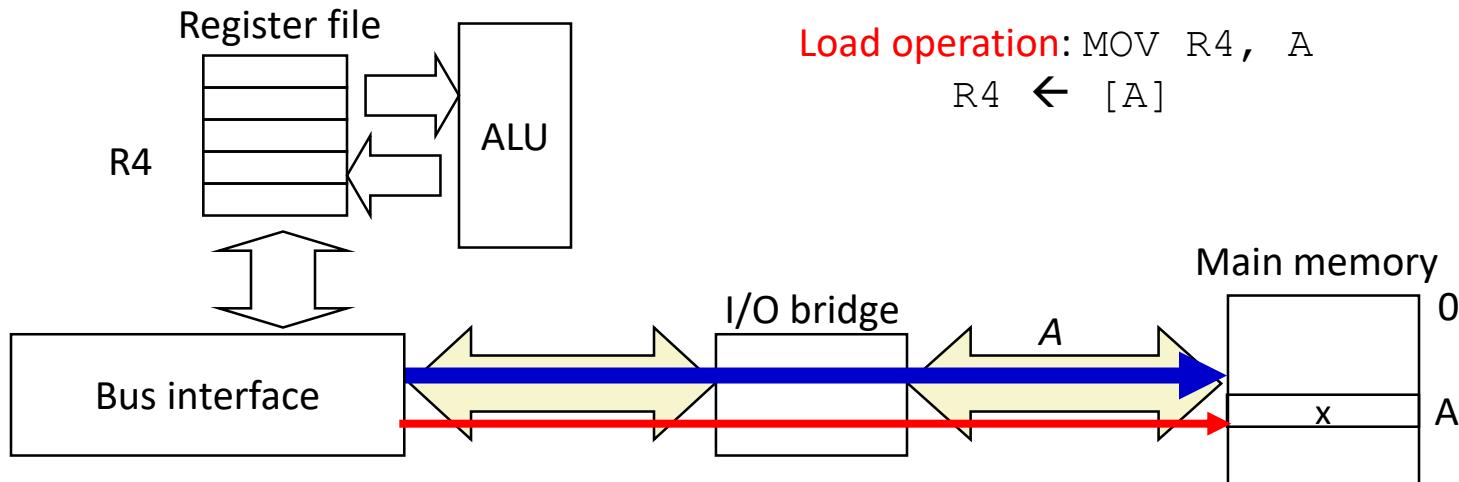
# Applications



- Storing fonts for printers
- Storing sound data in musical instruments
- Video game consoles
- Implantable Medical devices.
- High definition Multimedia Interfaces(HDMI)
- BIOS chip in computer
- Program storage chip in modem, video card and many electronic gadgets, controllers for disks, network cards, ....

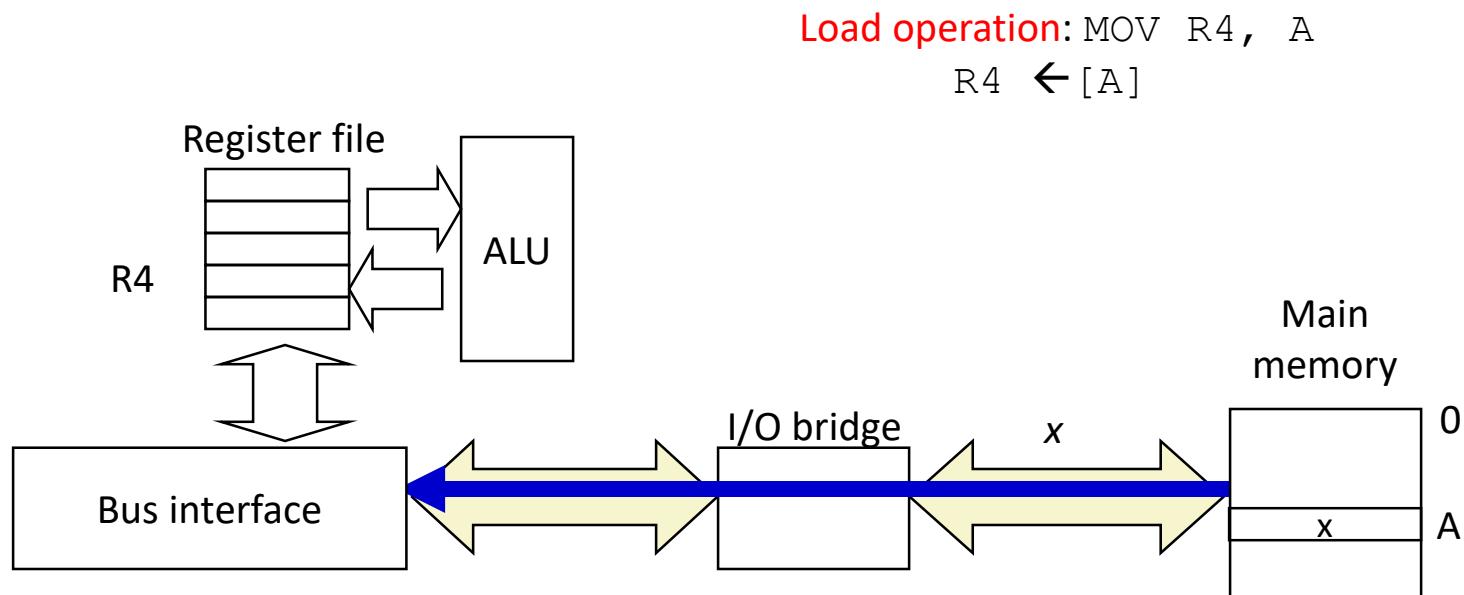
# Memory Read Operation (1)

- CPU places address  $A$  and then read control signal on the memory bus



# Memory Read Operation (2)

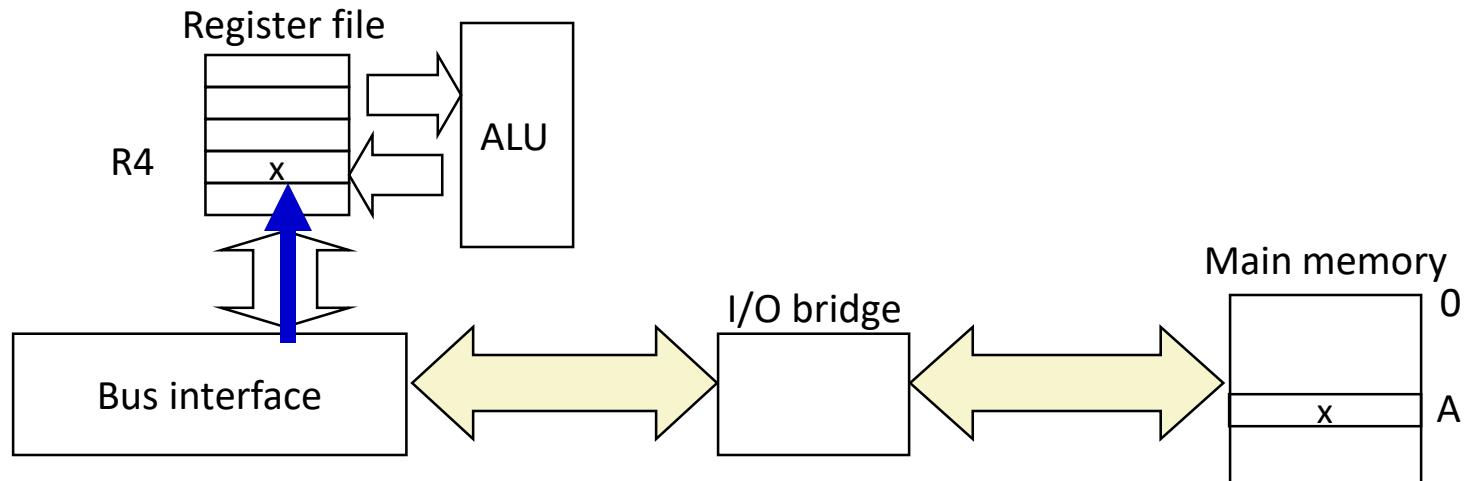
- Main memory reads  $A$  from the memory bus, retrieves word  $x$ , and places it on the bus



# Memory Read Operation (3)

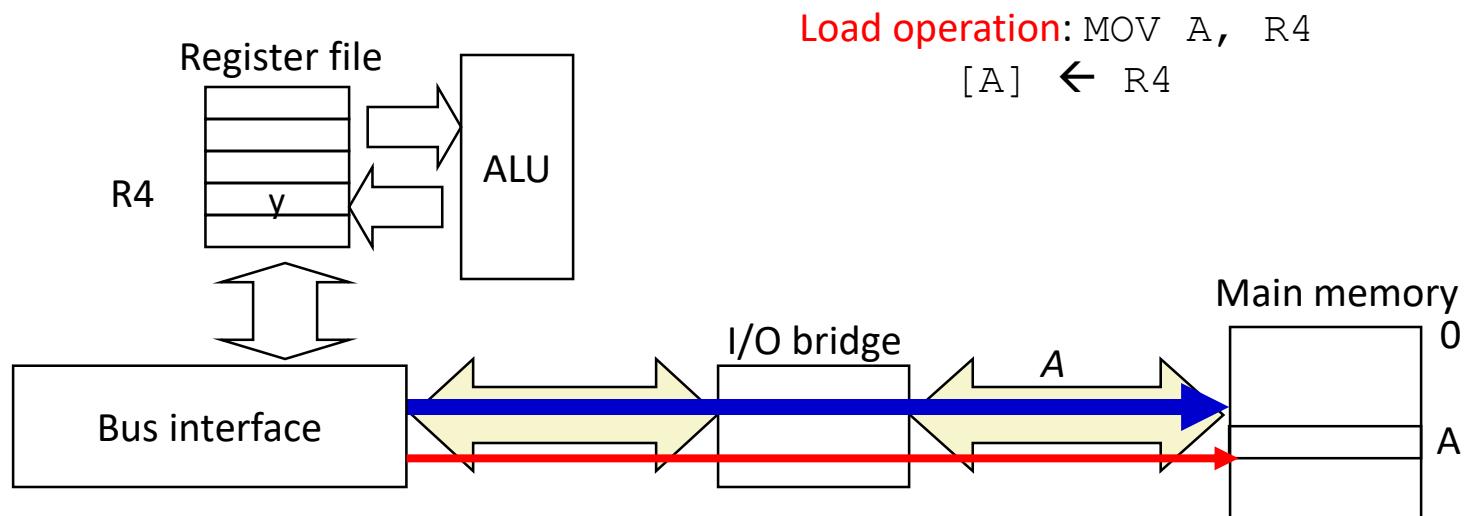
- CPU read word  $x$  from the bus and copies it into register R4.

Load operation:  $\text{MOV } R4, A$   
 $R4 \leftarrow [A]$



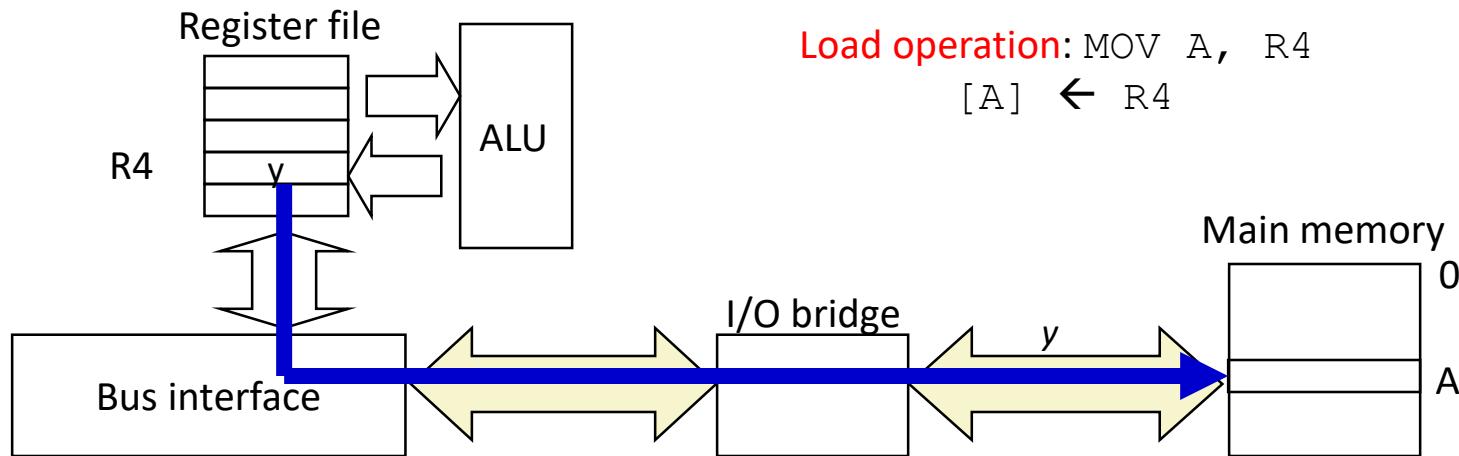
# Memory Write Operation (1)

- CPU places address  $A$  and WRITE control signal on bus. Main memory reads them and waits for the corresponding data word to arrive.



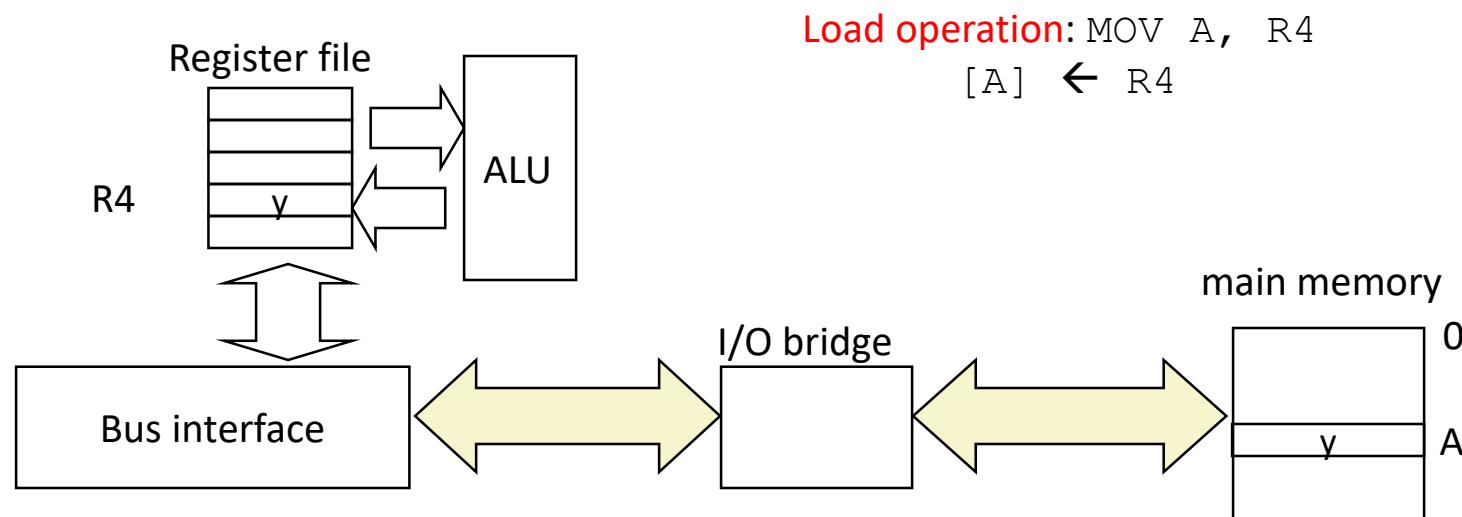
# Memory Write Operation (2)

- CPU places data word  $y$  on the bus



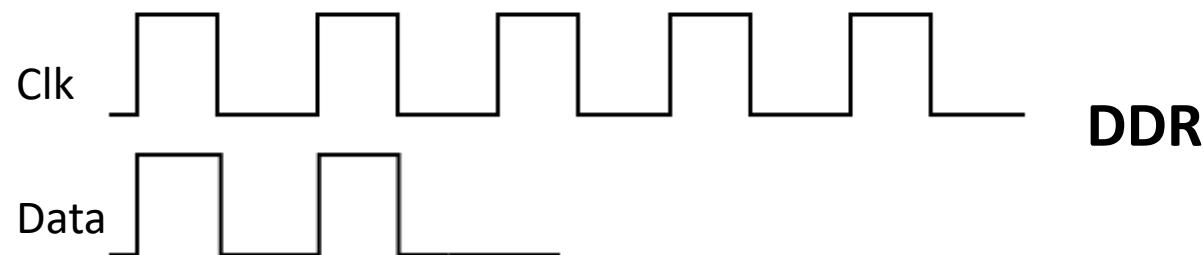
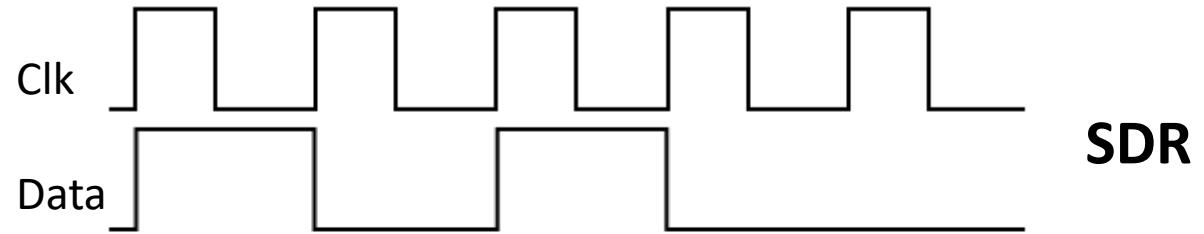
# Memory Write Operation (3)

- Main memory reads data word  $y$  from the bus and stores it at address  $A$ .



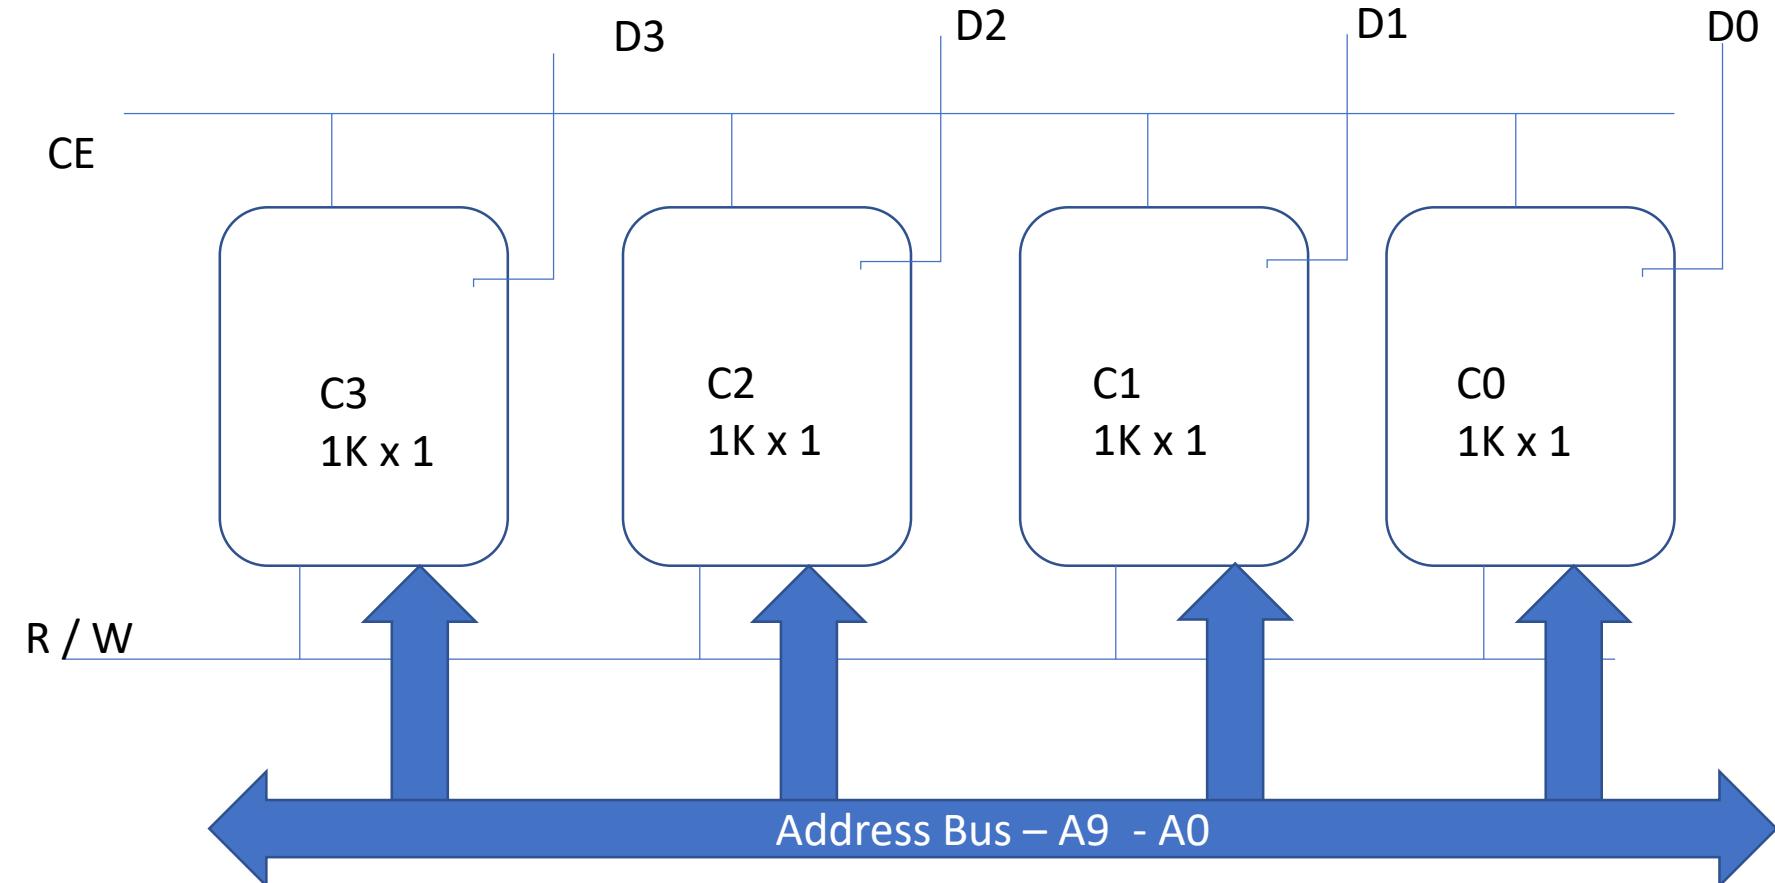
# SDR Vs DDR

- SDR → Single Data Rate
- DDR → Double Data Rate



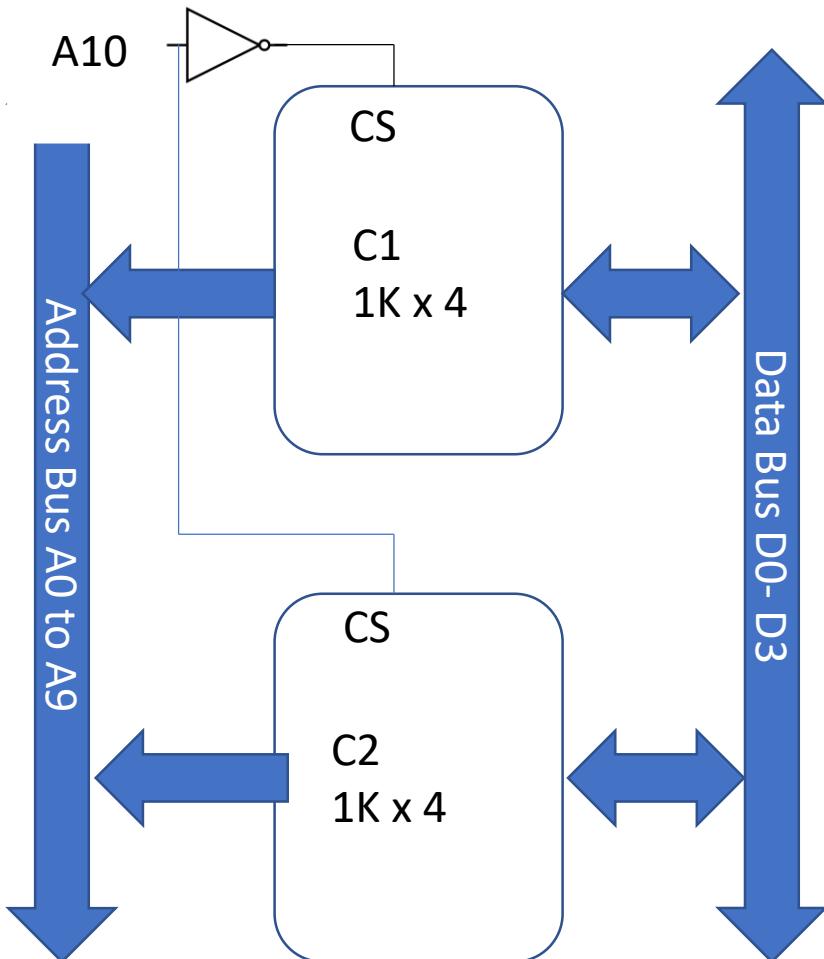
# Typical Memory Connection Examples

- Construct 1K X4 bit memory using 1Kx1 bit chip



# Typical Memory Connection Examples

- Construct 2K X4 bit memory using 1Kx4 bit chip



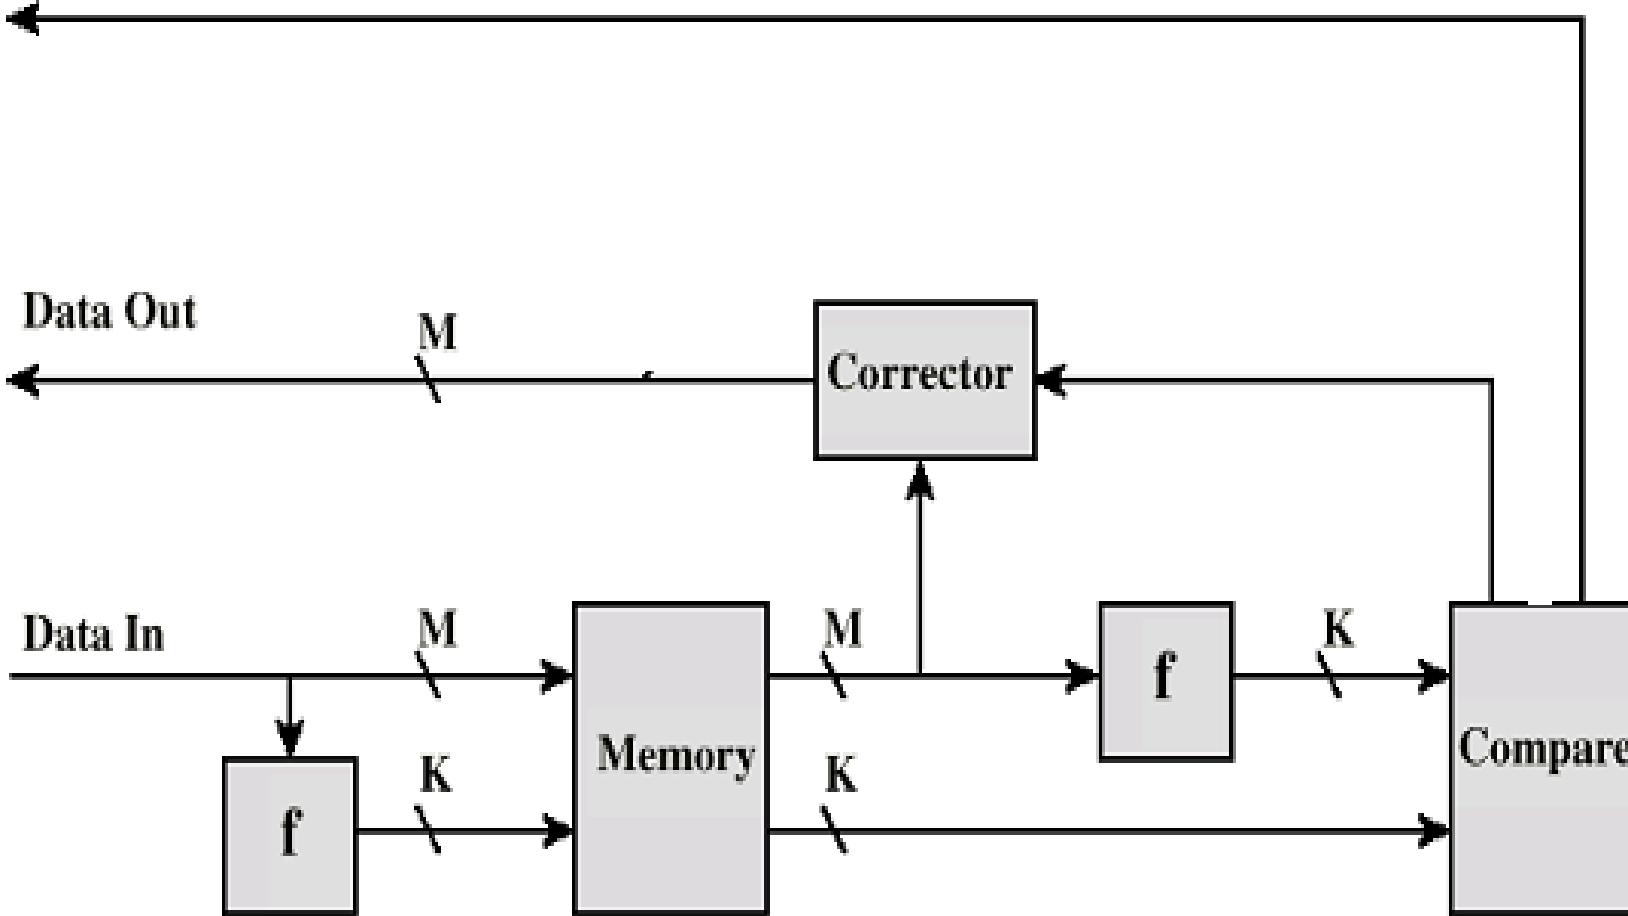


# Error Correction

- Hard Failure
  - Caused by harsh environmental abuse or manufacturing defects or wear
  - Memory cell is permanently stuck at 0 or 1
  - Permanent defect
- Soft Error
  - Random, non-destructive
  - alters the contents of one or more memory cells without damaging the memory.
  - No permanent damage to memory
  - Caused by power supply problems
- Detected using Hamming error correcting code

# Error Correcting Code Function

Error Signal





# Error Correcting Code Function

- The comparison logic receives as input two K-bit values. A bit-by-bit comparison is done by taking the exclusive-OR of the two inputs. The result is called the *syndrome word*.
- The comparison yields one of three results
  - No errors are detected.
  - An error is detected and it is possible to correct the error
  - An error is detected but it is not possible to correct it.



# Hamming Code....

- What should be the length of the code K ?
- Result of comparison is known as syndrome word
- length of the syndrome word is K bits, Length of Data is “M” bits
- Length of K should satisfy
- $2^k - 1 \geq M + K$



# Hamming code....

- Generate 4-bit syndrome for an 8 bit data word with following characteristics
  - If the syndrome contains all 0's, no error has been detected.
  - If the syndrome contains one and only one bit set to one, then an error has occurred in one of the 4 check bits. No correction is needed
  - If the syndrome contains more than one bit set to 1, then the numerical value of the syndrome indicates the position of the data bit in error. This data bit is inverted to correction



# Layout of Data and Check bits

Bit position	12	11	10	9	8	7	6	5	4	3	2	1
Position Number	1100	1011	1010	1001	1000	0111	0110	0101	0100	0011	0010	0001
Data bit	D8	D7	D6	D5		D4	D3	D2		D1		
Check Bit					C4				C3		C2	C1

# Layout of Data and Check bits

Bit position	12	11	10	9	8	7	6	5	4	3	2	1
Position Number	1100	1011	1010	1001	1000	0111	0110	0101	0100	0011	0010	0001
Data bit	D8	D7	D6	D5		D4	D3	D2		D1		
Check Bit					C4				C3		C2	C1

- $C_1 = D_1 \oplus D_2 \oplus D_4 \oplus D_5 \oplus D_7$
- $C_2 = D_1 \oplus D_3 \oplus D_4 \oplus D_6 \oplus D_7$
- $C_3 = D_2 \oplus D_3 \oplus D_4 \oplus D_8$
- $C_4 = D_5 \oplus D_6 \oplus D_7 \oplus D_8$

:

# Layout of Data and Check bits

Bit position	12	11	10	9	8	7	6	5	4	3	2	1
Position Number	1100	1011	1010	1001	1000	0111	0110	0101	0100	0011	0010	0001
Data bit	D8	D7	D6	D5		D4	D3	D2		D1		
Check Bit					C4				C3		C2	C1

- $C_1 = D_1 \oplus D_2 \oplus D_4 \oplus D_5 \oplus D_7$
- $C_2 = D_1 \oplus D_3 \oplus D_4 \oplus D_6 \oplus D_7$
- $C_3 = D_2 \oplus D_3 \oplus D_4 \oplus D_8$
- $C_4 = D_5 \oplus D_6 \oplus D_7 \oplus D_8$

:



# Problem 1

Consider the data + code k is as follows:

110101011101

Find out if there is an error. If so which bit is having error?

# Problem 1 - Solution



- $C_1 = D_1 \oplus D_2 \oplus D_4 \oplus D_5 \oplus D_7$
- $C_2 = D_1 \oplus D_3 \oplus D_4 \oplus D_6 \oplus D_7$
- $C_3 = D_2 \oplus D_3 \oplus D_4 \oplus D_8$
- $C_4 = D_5 \oplus D_6 \oplus D_7 \oplus D_8$

Consider the data + check bit is as follows:

110101011101

Find out if there is an error. If so which bit is having

Bit position	12	11	10	9	8	7	6	5	4	3	2	1
Position Number	1100	1011	1010	1001	1000	0111	0110	0101	0100	0011	0010	0001
Data bit	D8	D7	D6	D5		D4	D3	D2		D1		
Check Bit					C4				C3		C2	C1

- $C_1 = D_1 \oplus D_2 \oplus D_4 \oplus D_5 \oplus D_7$
- $C_2 = D_1 \oplus D_3 \oplus D_4 \oplus D_6 \oplus D_7$
- $C_3 = D_2 \oplus D_3 \oplus D_4 \oplus D_8$
- $C_4 = D_5 \oplus D_6 \oplus D_7 \oplus D_8$



## Problem 2

Data : 10101100 ( $M = 8$ )

Compute Check Bits

Bit position	12	11	10	9	8	7	6	5	4	3	2	1
Position Number	1100	1011	1010	1001	1000	0111	0110	0101	0100	0011	0010	0001
Data bit	D8	D7	D6	D5		D4	D3	D2		D1		
Check Bit					C4				C3		C2	C1



# Computer Organization and Software Systems

Contact Session 3

Dr. Lucy J. Gudino



# Last Class



Contact Hour	List of Topic Title	Text/Ref Book/external resource
3-4	<b>Memory Organization</b> - Internal Memory - External Memory (HDD)	T1, R2

# Today's Session



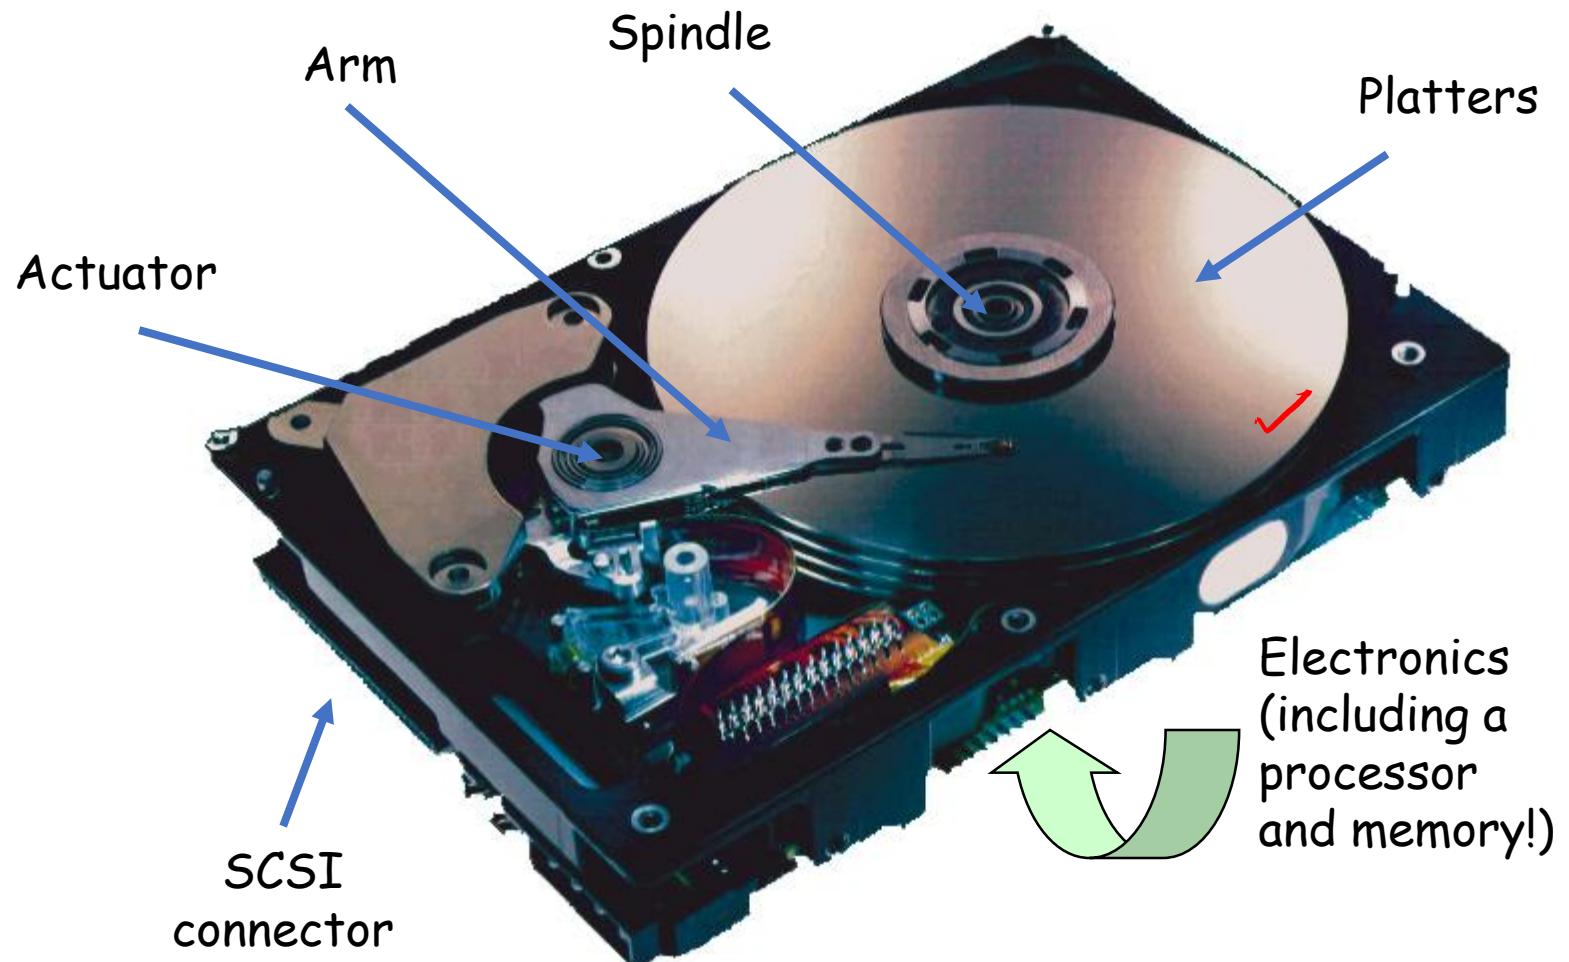
Contact Hour	List of Topic Title	Text/Ref Book/external resource
5-6	<b>Memory Organization</b> - External Memory (RAID, SSD) <b>Cache Memory Organization</b> - Locality - Locality of Reference to Program Data - Locality of instruction fetches	T1, R2



# Types of External memory

- Magnetic Disk
  - RAID Memories
  - Removable Disks
- Optical
  - CD-ROM
  - CD-Recordable (CD-R)
  - CD-R/W
  - DVD
- Magnetic Tape

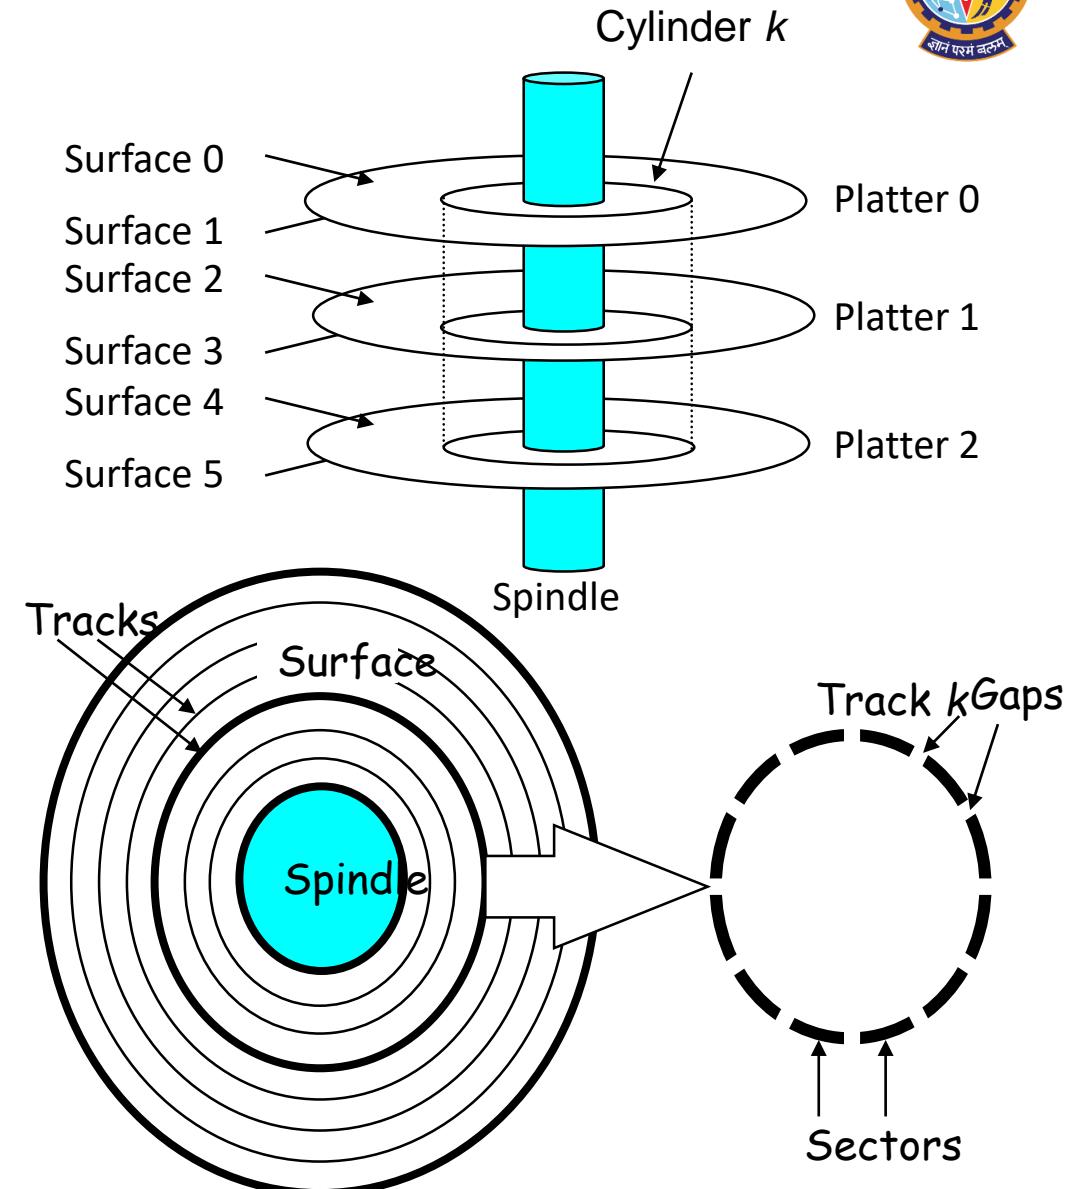
# Magnetic Disk Drive



*Image courtesy of Seagate Technology*

# Disk Geometry

- Disks consist of **platters**, each with two **surfaces**.
- Each surface consists of concentric rings called **tracks**
- Aligned tracks form a cylinder
- Each track consists of **sectors** separated by **gaps**



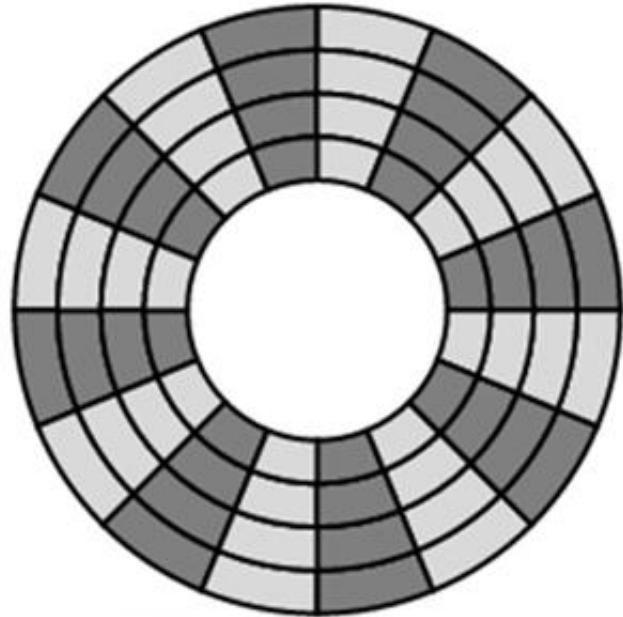


# Disk Capacity

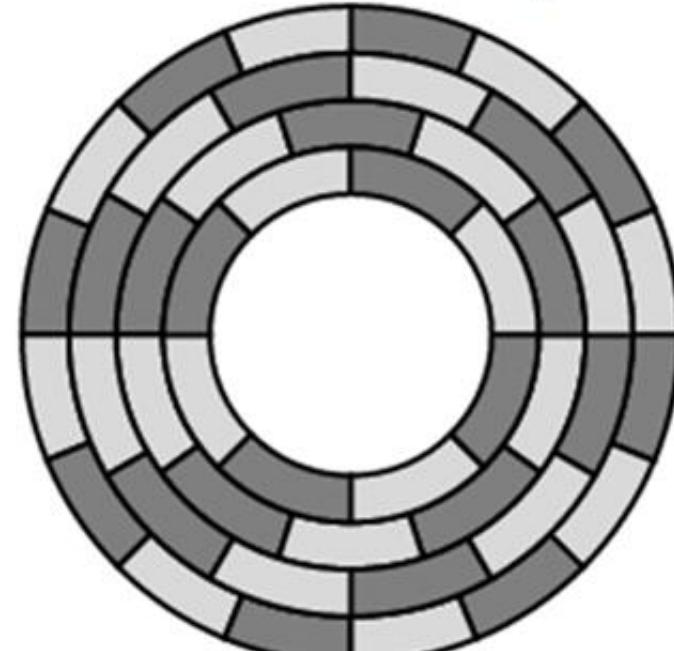
- **Capacity:** maximum number of bits that can be stored.
  - Vendors express capacity in units of gigabytes (GB /TB), where  $1\text{ GB} = 2^{30}\text{ Bytes}$ ,  $1\text{ TB} = 2^{40}\text{ Bytes}$ ,
- Capacity is determined by these technology factors:
  - **Recording density** (bits/in): number of bits that can be squeezed into a 1 inch segment of a track.
  - **Track density** (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment.
  - **Areal density** (bits/in<sup>2</sup>): product of recording and track density.

# Recording zones

- Modern disks partition tracks into disjoint subsets called **recording zones**
  - Each track in a zone has the same number of sectors, determined by the circumference of innermost track.
  - Each zone has a different number of sectors/track, outer zones have more sectors/track than inner zones.
  - So we use **average** number of sectors/track when computing capacity.



Without Recording Zones



With Recording Zones



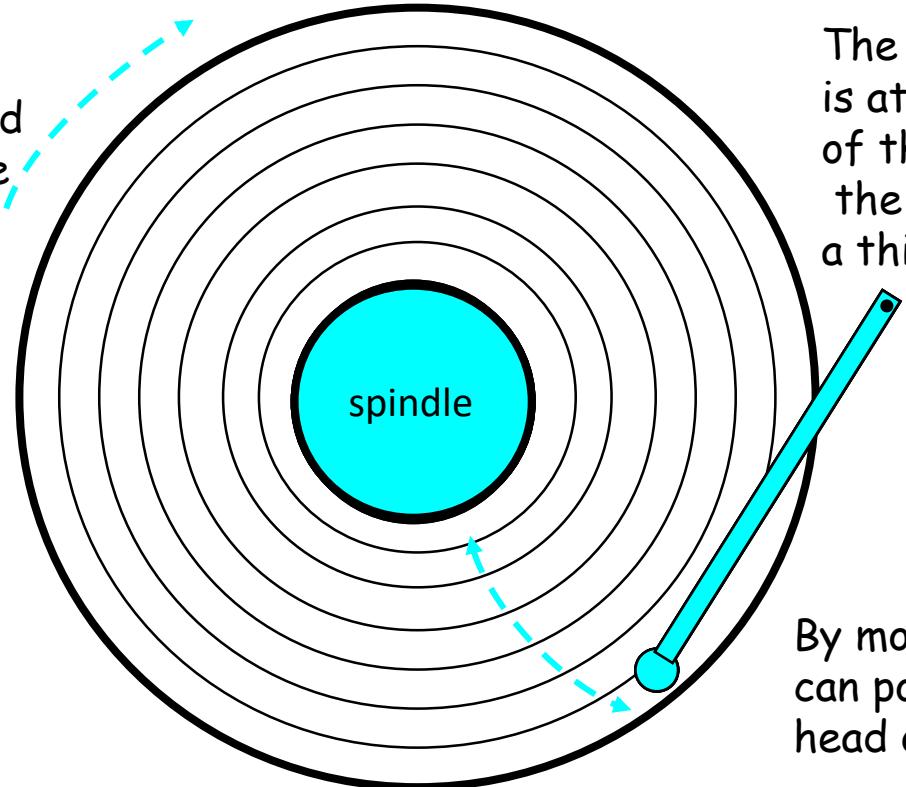
# Computing Disk Capacity

- Capacity = (# bytes/sector) × (avg. # sectors/track) ×  
(# tracks/surface) × (# surfaces/platter) ×  
(# platters/disk)
- Example:
  - 512 bytes/sector
  - 300 sectors/track (on average)
  - 20,000 tracks/surface
  - 2 surfaces/platter
  - 5 platters/disk
- Capacity =  $512 \times 300 \times 20000 \times 2 \times 5$   
= 30,720,000,000  
= 28.61 GB

# Disk Operation (Single-Platter View)



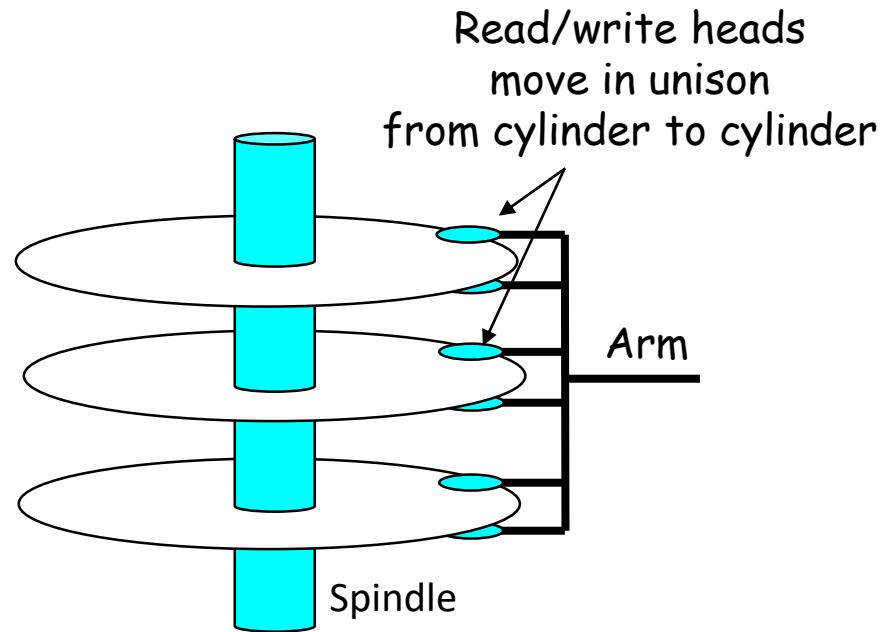
The disk surface spins at a fixed rotational rate



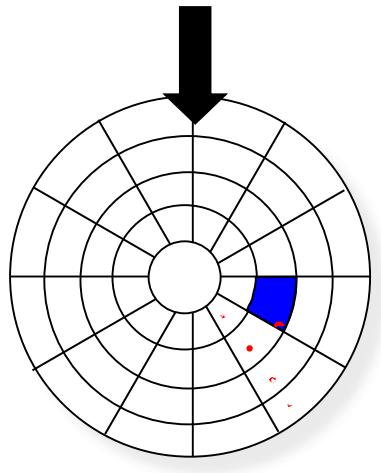
The read/write head is attached to the end of the arm and flies over the disk surface on a thin cushion of air.

By moving radially, the arm can position the read/write head over any track.

# Disk Operation (Multi-Platter View)

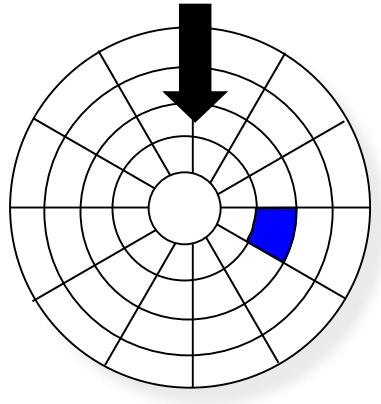


# Disk Access



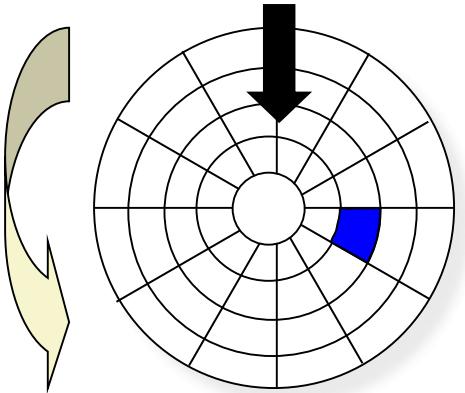
Need to access a sector  
colored in blue

# Disk Access



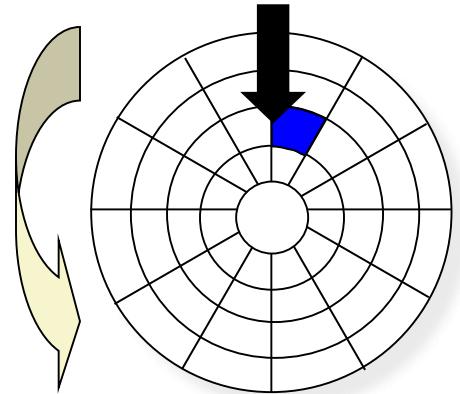
Head in position above a track

# Disk Access



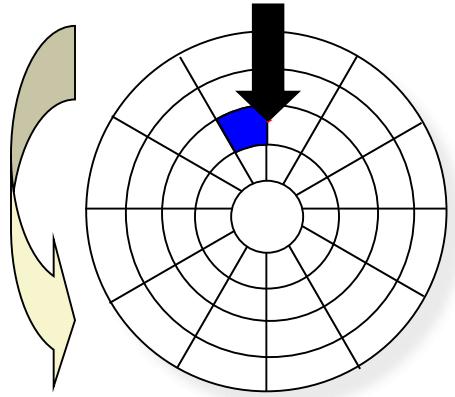
Rotate the platter in counter-clockwise direction

# Disk Access - Read



About to read blue sector

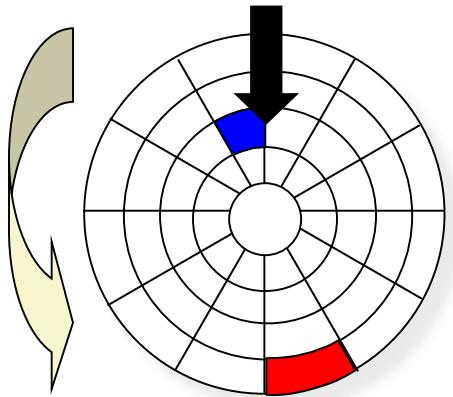
# Disk Access - Read



After BLUE  
read

After reading blue sector

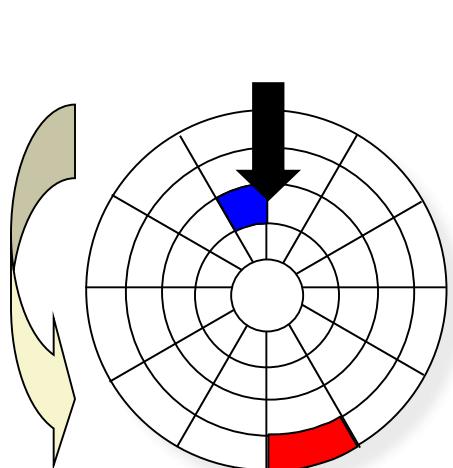
# Disk Access - Read



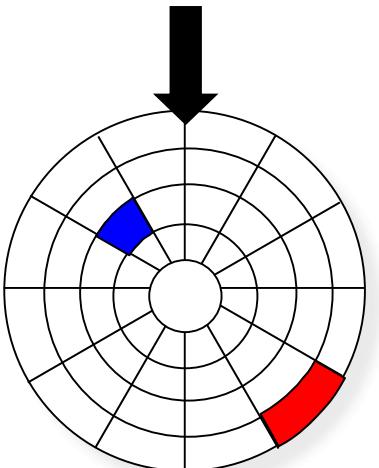
After BLUE  
read

Red request scheduled next

# Disk Access - Seek

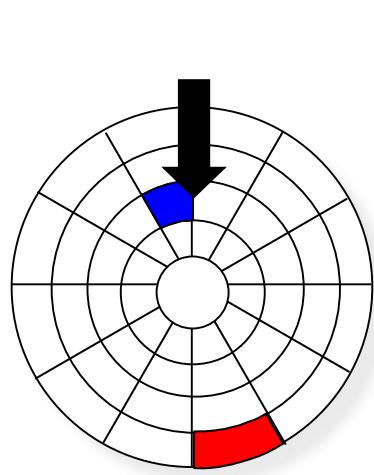


After BLUE  
read

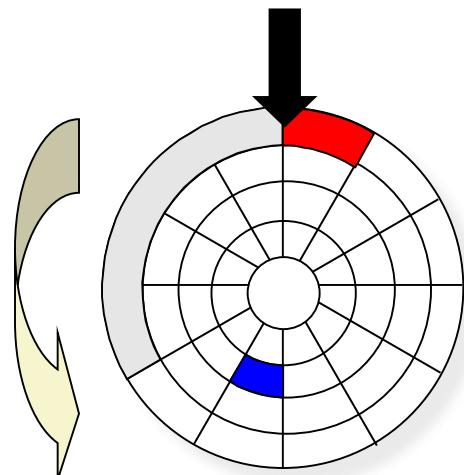
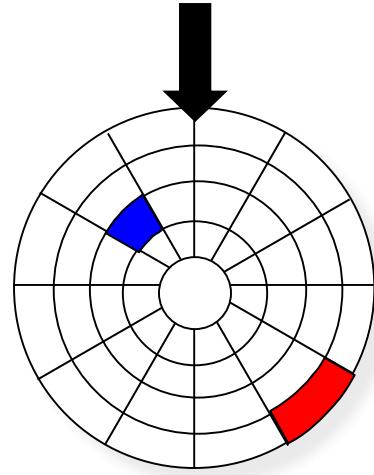


Seek for RED

# Disk Access – Rotational Latency



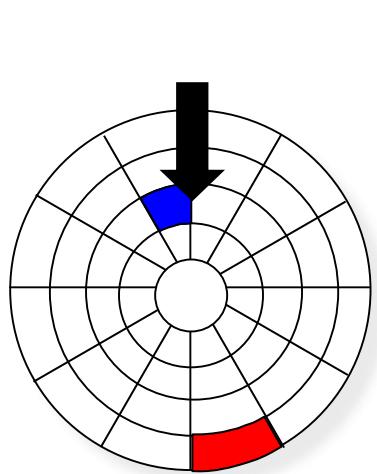
After BLUE Seek for RED  
read



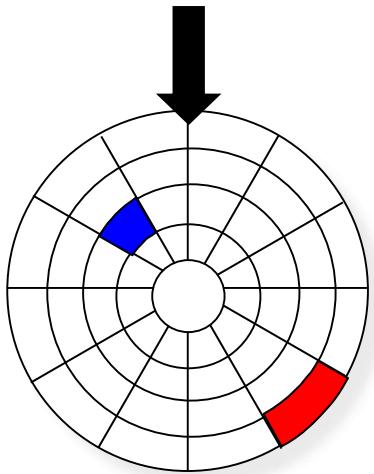
Rotational latency

Wait for red sector to rotate around

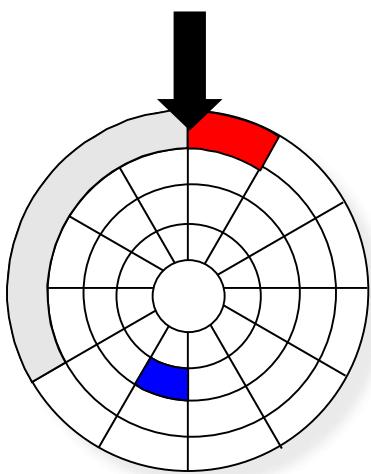
# Disk Access – Read



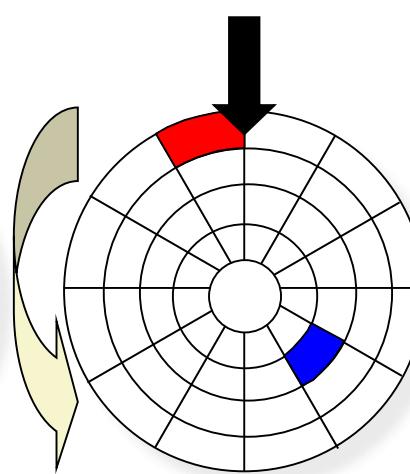
After BLUE  
read



Seek for RED

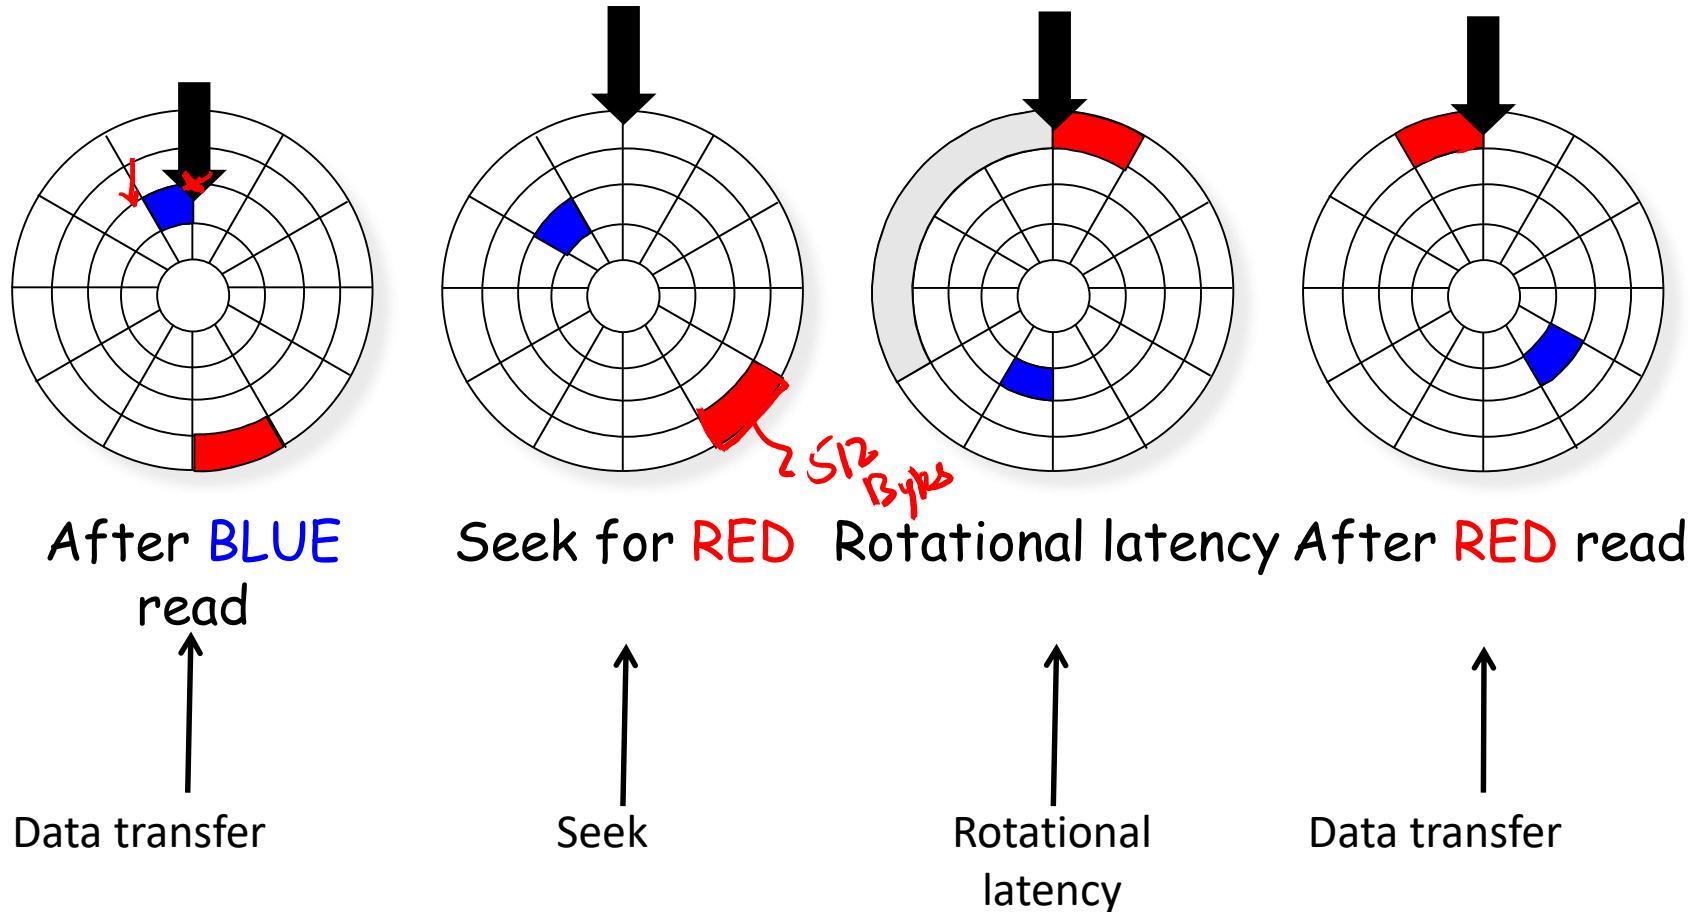


Rotational latency After RED read



Complete read of red

# Disk Access – Access Time Components





# Today's Class

Contact Hour	List of Topic Title	Text/Ref Book/external resource
5-6	<b>Memory Organization</b> - External Memory (RAID, SSD) <b>Cache Memory Organization</b> - Locality - Locality of Reference to Program Data - Locality of instruction fetches	T1, R2

# RAID



- RAID - Redundant Array of Independent Disks
- Variety of ways in which the data can be organized
- Need for RAID:
  - Parallel I/O
  - Reliability
  - Redundancy through multiple inexpensive disks



# Common Characteristics

1. RAID is a set of physical disk drives viewed by the operating system as a single logical drive.
2. Data are distributed across the physical drives of an array in a scheme known as striping.
3. Redundant disk capacity is used to store parity information, which guarantees data recoverability in case of a disk failure.

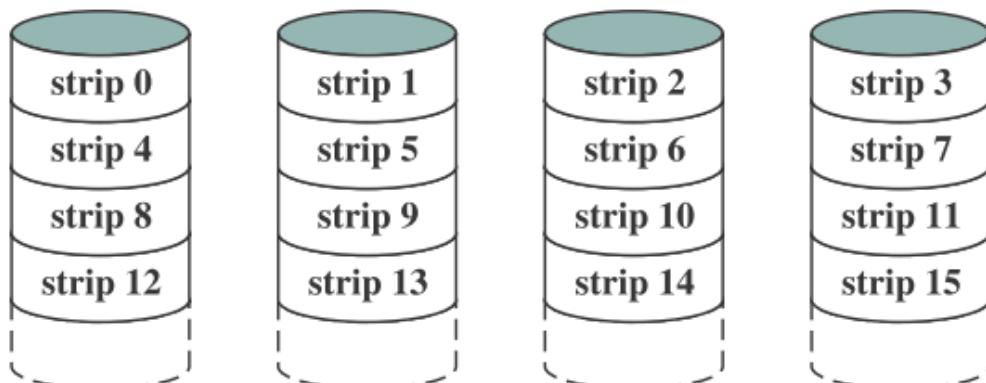


# Categories

- RAID category
  - Striping (Level 0)
  - Mirroring (Level 1)
  - Parallel access (Level 2,3)
  - Independent access (Level 4,5,6)

# RAID Level 0

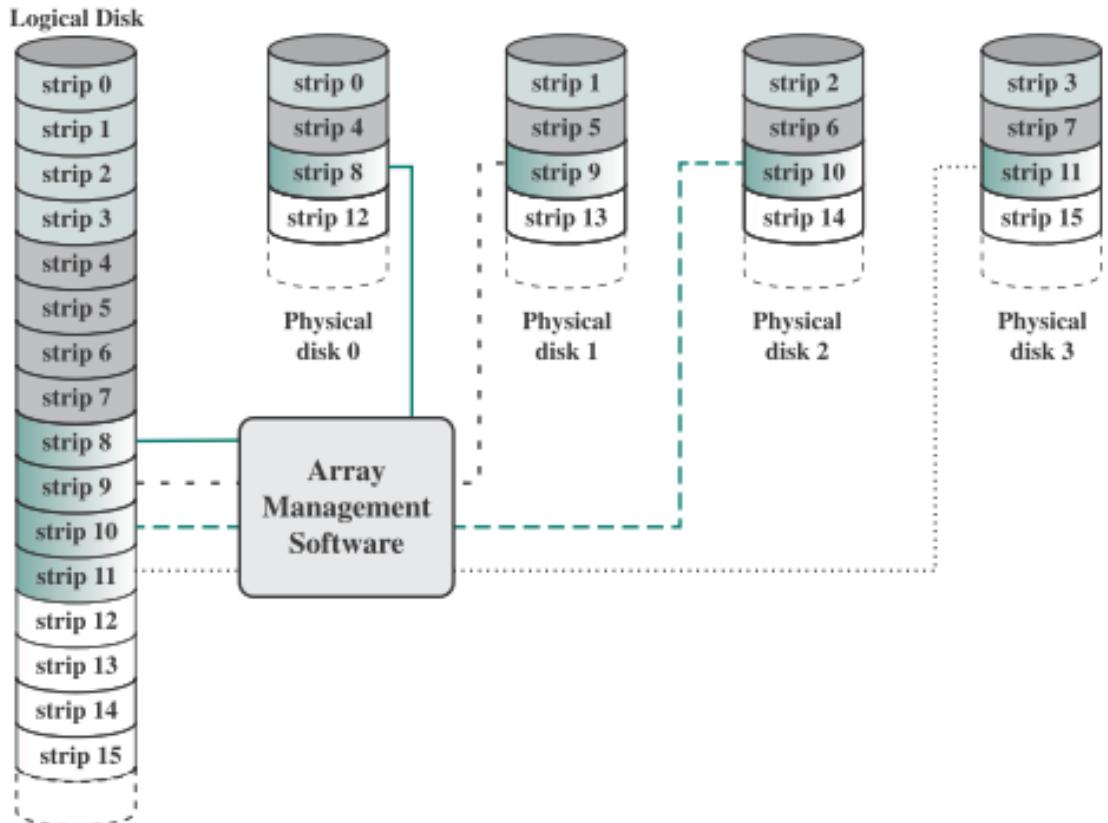
- Is not a true member of RAID family – No Redundancy
- The data are striped across the available disks
- Data are distributed across all of the disks in the array
- IO requests can be issued in Parallel
- Reduce IO queuing time



(a) RAID 0 (Nonredundant)

# RAID level 0 configuration

- RAID 0 for
  - high data transfer capacity
  - high I/O request rate



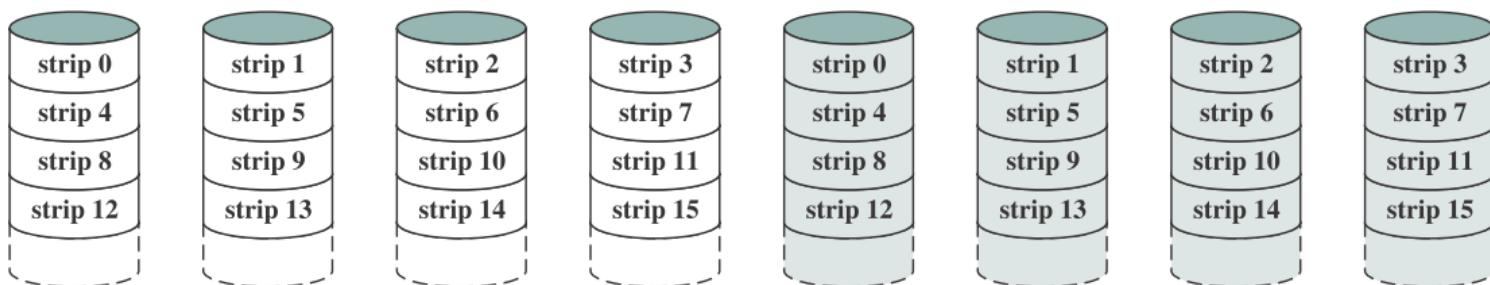
# RAID Level 0



Level	Advantages	Disadvantages	Applications
0	<ul style="list-style-type: none"><li>• Very simple design</li><li>• Easy to implement</li><li>• I/O performance is greatly improved by spreading the I/O load across many drives</li><li>• No parity calculation overhead is involved</li></ul>	The failure of just one drive will result in all data in an array being lost	<ul style="list-style-type: none"><li>• Video production and editing</li><li>• Image Editing</li><li>• Any application requiring high bandwidth</li></ul>

# RAID Level 1

- Mirroring: redundancy is achieved by duplicating all the data
- Each logical strip is mapped to two separate physical disks
- Every disk in the array has a mirror disk that contains the same data



(b) RAID 1 (Mirrored)

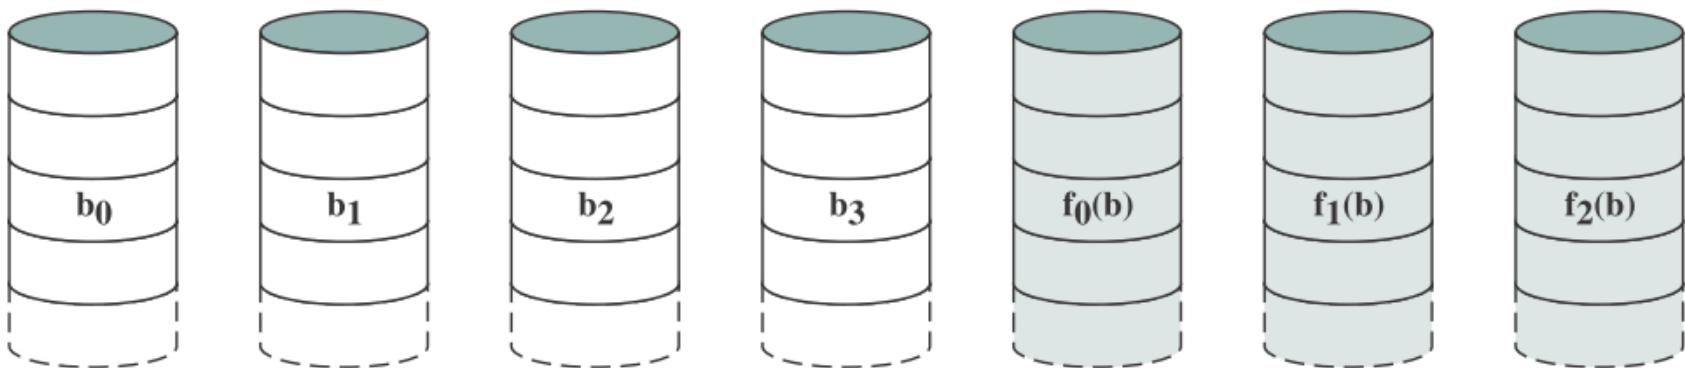
# RAID Level 1



Level	Advantages	Disadvantages	Applications
1	<ul style="list-style-type: none"><li>• Simplest RAID storage subsystem design</li><li>• 100% redundancy of data</li><li>• RAID 1 can sustain multiple simultaneous drive failures</li></ul>	<ul style="list-style-type: none"><li>• Highest disk overhead of all RAID types</li><li>• inefficient</li></ul>	<ul style="list-style-type: none"><li>• Accounting</li><li>• Payroll</li><li>• Financial</li><li>• Any application requiring very high availability</li></ul>

# RAID Level 2

- Use parallel access technique
- In a parallel access array, all member disks participate in the execution of every I/O request
- Spindles of the individual drives are synchronized



(c) RAID 2 (Redundancy through Hamming code)

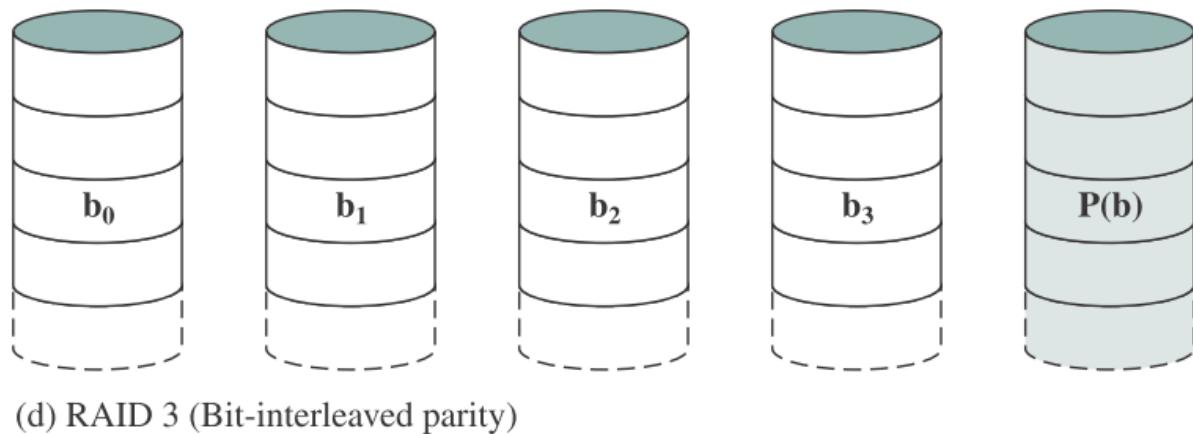
# RAID Level 2



Level	Advantages	Disadvantages	Applications
2	<ul style="list-style-type: none"><li>• Extremely high data transfer rates possible</li><li>• The higher the data transfer rate required, the better the ratio of data disks to ECC disks</li><li>• Relatively simple controller design compared to RAID levels 3, 4, &amp; 5</li></ul>	<ul style="list-style-type: none"><li>• Very high ratio of ECC disks to data disks with smaller word sizes— inefficient</li><li>• Entry level cost very high— requires very high transfer rate requirement to justify</li></ul>	<ul style="list-style-type: none"><li>• No commercial implementations exist/not commercially viable</li></ul>

# RAID Level 3

- RAID 3 is organized in a similar fashion to RAID 2
- RAID 3 requires only a single redundant disk
- RAID 3 employs parallel access, with data distributed in small strips
- Parity bit is computed for the set of individual bits in the same position on all of the data disks



# RAID Level 3

- Redundancy: Event of a drive failure - parity drive is accessed and data is reconstructed from the remaining devices
  - Missing data can be restored on the new drive
- Data reconstruction is simple

$$X4(i) = X3(i) \oplus X2(i) \oplus X1(i) \oplus X0(i)$$

- Suppose drive X1 fails

$$X1(i) = X4(i) \oplus X3(i) \oplus X2(i) \oplus X0(i)$$

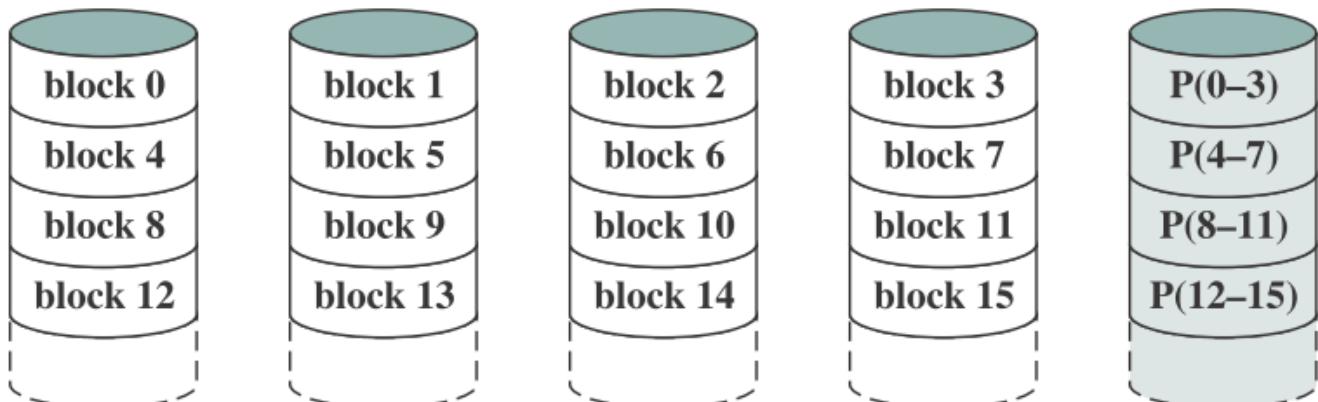
# RAID Level 3



Level	Advantages	Disadvantages	Applications
3	<ul style="list-style-type: none"><li>• Very high read data transfer rate</li><li>• Very high write data transfer rate</li><li>• Disk failure has an insignificant impact on throughput</li><li>• Low ratio of ECC (parity) disks to data disks means high efficiency</li></ul>	<ul style="list-style-type: none"><li>• In case of small size files it performs slowly.</li><li>• Controller design is fairly Complex</li></ul>	<ul style="list-style-type: none"><li>• Video production and live streaming</li><li>• Image editing</li><li>• Video editing</li><li>• Prepress applications</li><li>• Any application requiring high throughput</li></ul>

# RAID Level 4

- RAID levels 4 through 6 make use of an independent access technique
- Each member disk operates independently, so that separate I/O requests run in parallel



(e) RAID 4 (Block-level parity)

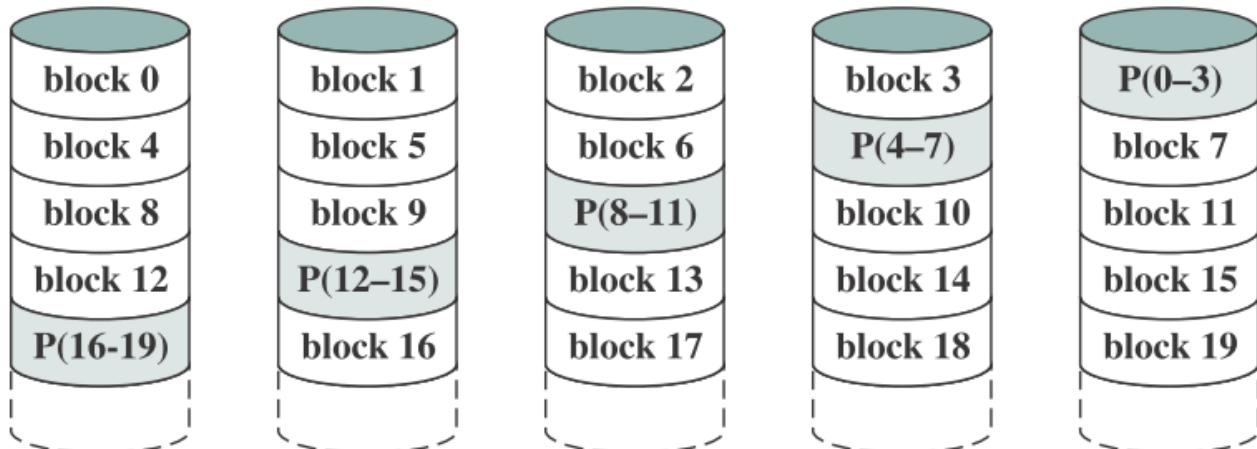
# RAID Level 4



Level	Advantages	Disadvantages	Applications
4	<ul style="list-style-type: none"><li>• Low ratio of ECC (parity) disks to data disks means high efficiency</li><li>• Very high Read data transaction rate</li></ul>	<ul style="list-style-type: none"><li>• Worst write transaction rate and Write aggregate transfer rate</li><li>• Quite complex controller design</li><li>• Difficult and inefficient data rebuild in the event of disk failure</li></ul>	<ul style="list-style-type: none"><li>• No commercial implementations exist/not commercially viable</li></ul>

# RAID Level 5

- RAID 5 is organized in a similar fashion to RAID 4.
- The difference is that RAID 5 distributes the parity strips across all disks



(f) RAID 5 (Block-level distributed parity)

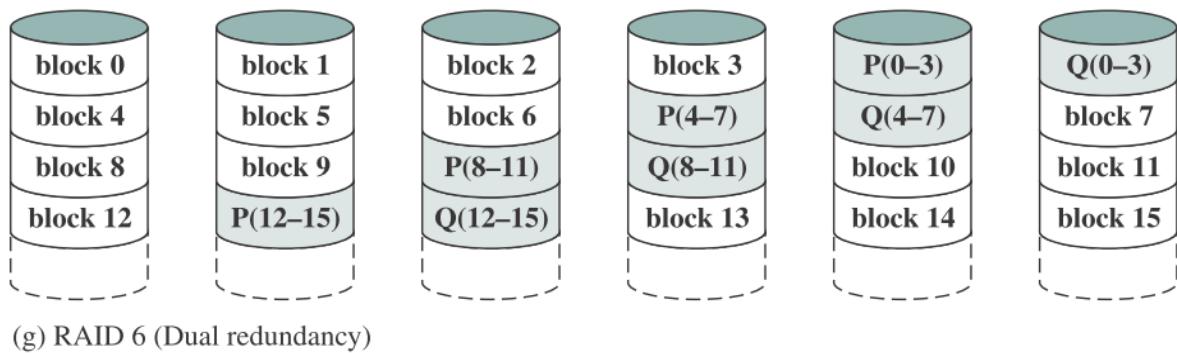
# RAID Level 5



Level	Advantages	Disadvantages	Applications
5	<ul style="list-style-type: none"><li>Highest Read data transaction rate</li><li>Low ratio of ECC (parity) disks to data disks means high efficiency</li><li>Better reliability</li></ul>	<ul style="list-style-type: none"><li>Most complex controller design</li><li>Difficult to rebuild in the event of a disk failure (as compared to RAID level 1)</li></ul>	<ul style="list-style-type: none"><li>File and application servers</li><li>Database servers</li><li>Web, e-mail, and news servers</li><li>Intranet servers</li></ul>

# RAID Level 6

- Two different parity calculations are carried out and stored in separate blocks on different disks
- RAID 6 array whose user data require N disks consists of N + 2 disks
- One of the two is the exclusive- OR calculation used in RAID 4 and 5
- Other is an independent data check algorithm
- Possible to regenerate data even if two disks containing user data fail



# RAID Level 6



Level	Advantages	Disadvantages	Applications
6	<ul style="list-style-type: none"><li>Provides for an extremely high data fault tolerance and can sustain multiple simultaneous drive failures</li></ul>	<ul style="list-style-type: none"><li>More complex controller design</li><li>Controller overhead to compute parity is extremely high</li></ul>	<ul style="list-style-type: none"><li>Perfect solution for mission critical applications</li></ul>

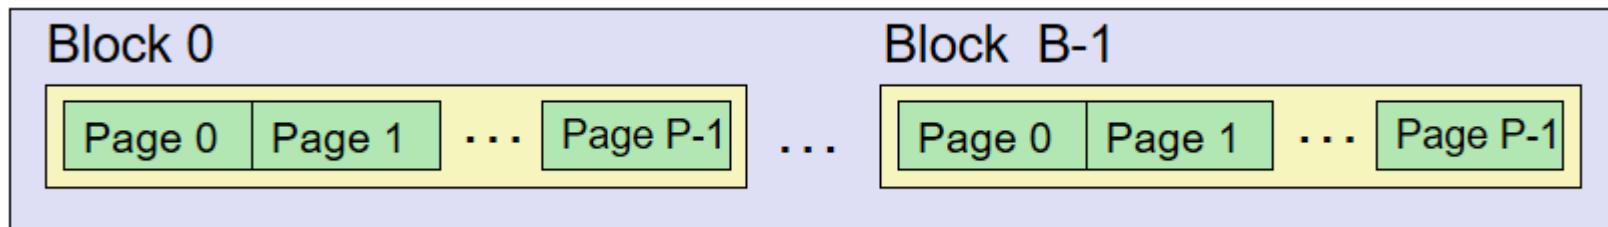


# Solid state drives

- SSDs have the following advantages over HDDs:
  - High- performance input/output operations per second (IOPS): Significantly increases performance I/O subsystems.
  - Lower access times and latency rates: Over 10 times faster than the spinning disks in an HDD.
  - Lower power consumption: SSDs use considerably less power than comparable- size HDDs.
  - Durability: Less susceptible to physical shock and vibration.
  - Longer lifespan: SSDs are not susceptible to mechanical wear.
  - Quieter and cooler running capabilities: Less space required, lower energy costs, and a greener enterprise.

# Solid State Disks (SSDs)

- SSDs contain Blocks and Pages
- Pages: 512KB to 4KB, Blocks: 32 to 128 pages
- Data read/written in units of pages.
- Page can be written only after its block has been erased
- A block wears out after about 100,000 repeated writes



# SSD Organization

- SSD contains the following components:
  - Controller: Provides SSD device level interfacing and firmware execution.
  - Addressing: Logic that performs the selection function across the flash memory components.
  - Data buffer/cache: High speed RAM memory components used for speed matching and to increased data throughput.
  - Error correction: Logic for error detection and correction.
  - Flash memory components: Individual NAND flash chips.

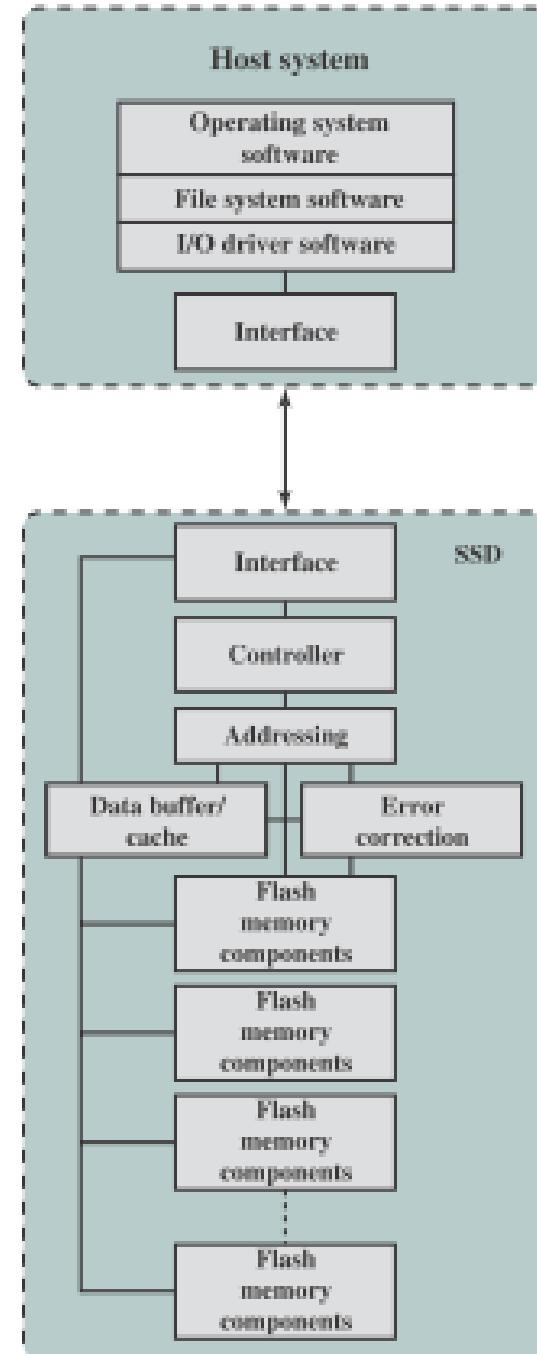


Figure 6.8 Solid State Drive Architecture



# Issues with SSD

- Two practical issues:
  - SSD performance has a tendency to slow down as the device is used
  - Memory becomes unusable after a certain number of writes

# First Issue



- Files are stored on disk as a set of pages, typically 4 KB in length
- Pages are not necessarily stored as a contiguous set of pages on the disk
- Flash memory is accessed in blocks, block size of 512 KB, (128 pages per block)
- Write a page onto a flash memory:
  1. The entire block must be read from the flash memory and placed in a RAM buffer. Then the appropriate page in the RAM buffer is updated.
  2. Before the block can be written back to flash memory, the entire block of flash memory must be erased— it is not possible to erase just one page of the flash memory.
  3. The entire block from the buffer is now written back to the flash memory.



# Second Issue

- Flash memory becomes unusable after a certain number of writes.
- Typical limit is 100,000 writes



# Comparison of Solid State Drives and Disk Drives

	NAND Flash Drives	Seagate Laptop Internal HDD
File copy/write speed	200–550 Mbps	50–120 Mbps
Power draw/battery life	Less power draw, averages 2–3 watts, resulting in 30+ minute battery boost	More power draw, averages 6–7 watts and therefore uses more battery
Storage capacity	Typically not larger than 512 GB for notebook size drives; 1 TB max for Desktops	Typically around 500 GB and 2 TB max for notebook size drives; 4 TB max for desktops
Cost	Approx. \$0.50 per GB for a 1-TB drive	Approx. \$0.15 per GB for a 4-TB drive



# Computer Organization and Software Systems

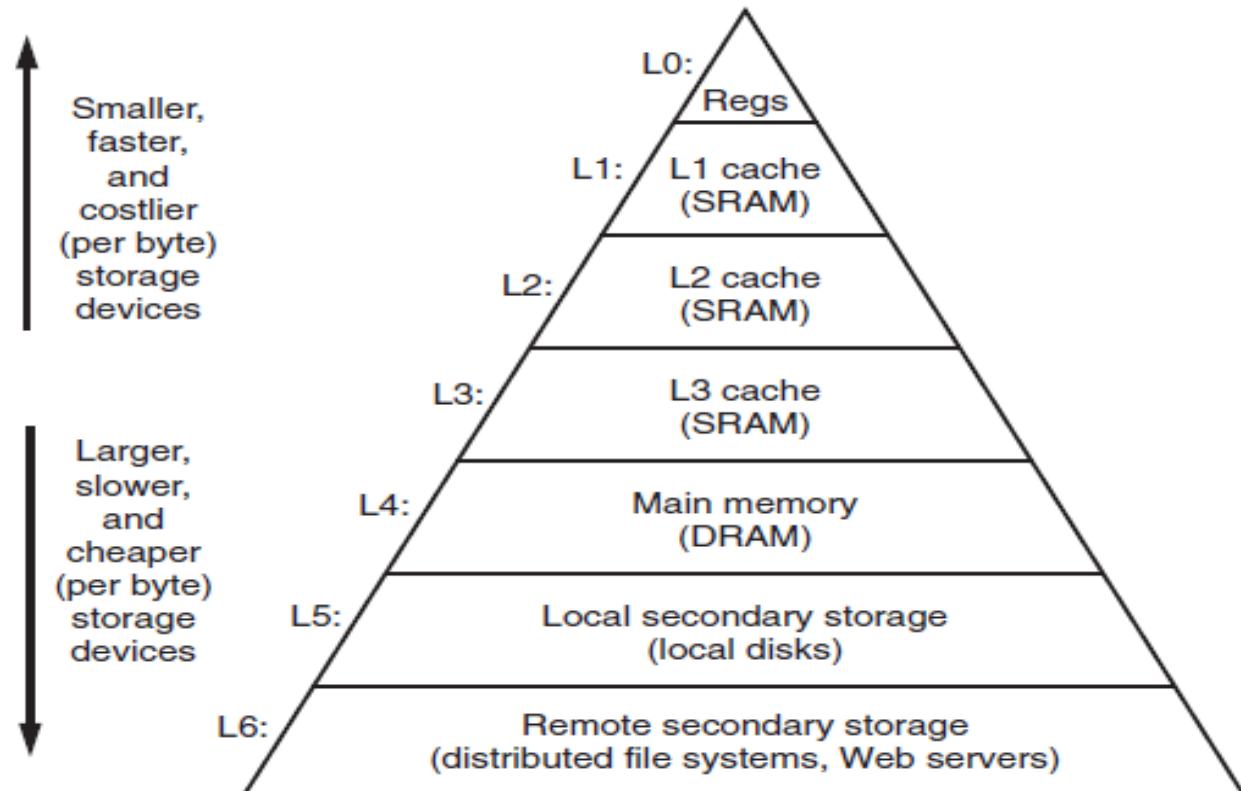
Contact Session 4

Dr. Lucy J. Gudino



# The Bottom Line

- How much?
  - Capacity
- How fast?
  - Time is money
- How expensive?



An example of a memory hierarchy.

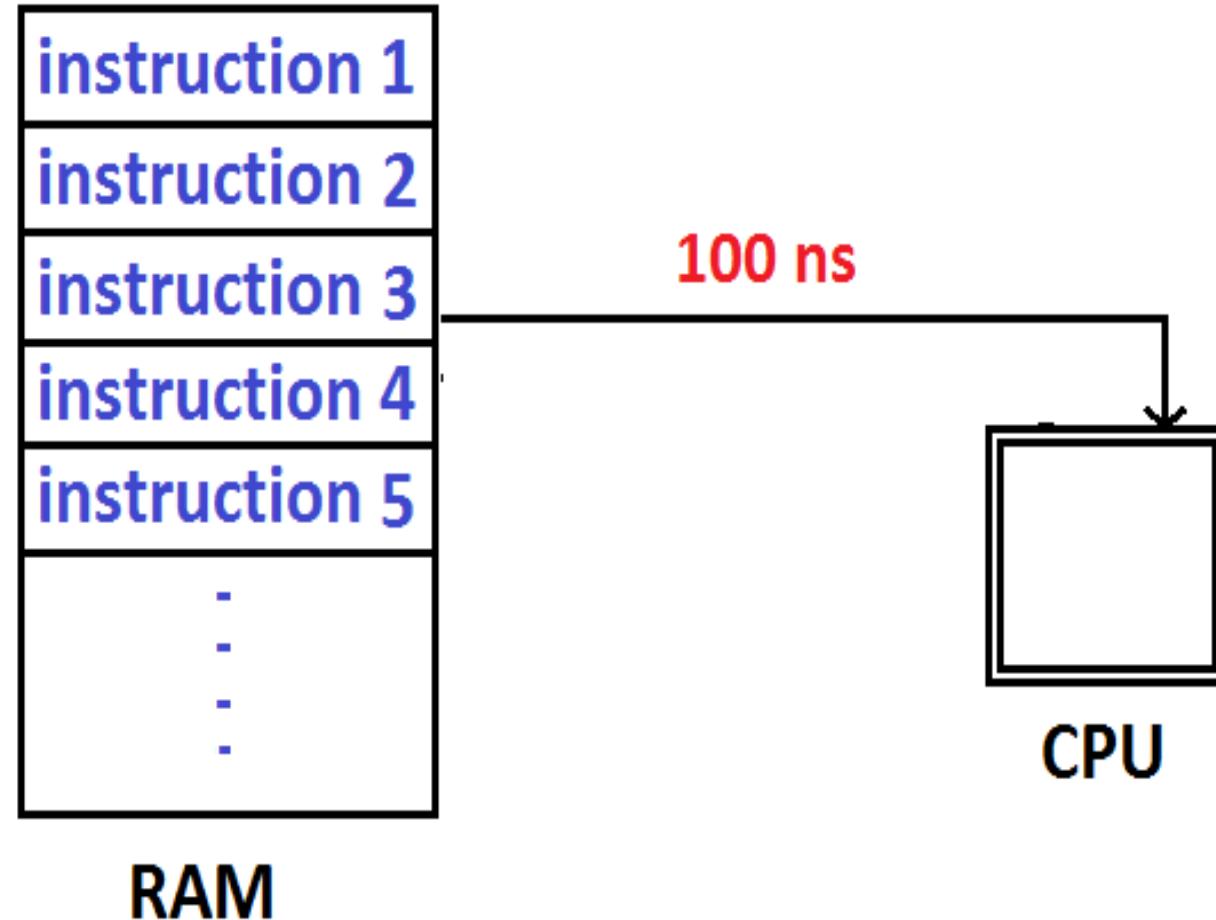
- Faster access time , ----- cost per bit
- Greater capacity, ----- cost per bit
- Greater capacity, -----access time



# Memory Hierarchy

- Registers
  - In CPU
- Internal or Main memory
  - May include one or more levels of cache
  - "RAM"
- External memory
  - Backing store

# Performance enhancement - Motivation

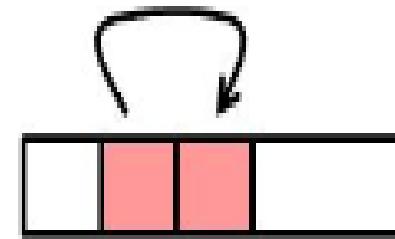
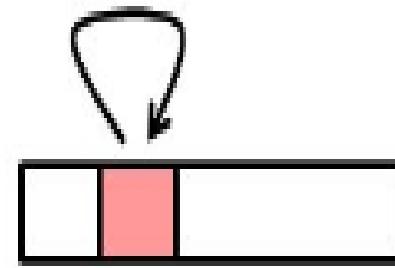


# Performance enhancement - Motivation

## Locality of Reference

During the course of the execution of a program, memory references tend to cluster

- **Temporal locality:** Locality in time
  - If an item is referenced, it will tend to be referenced again soon
- **Spatial locality:** Locality in space
  - If an item is referenced, items whose addresses are close by will tend to be referenced soon.



# Example



```
product = 1;  
for ( i = 0; i < n-1; i++)  
    product = product * a[i] ;
```

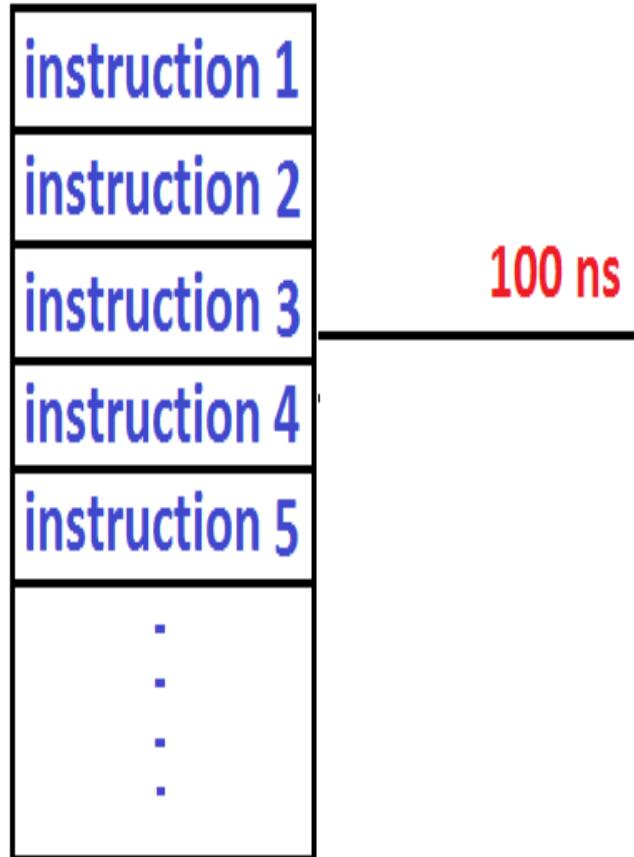
- Data :

- Access array elements in succession - spatial locality
- Reference to "product" "i" and "n" in each iteration - Temporal locality

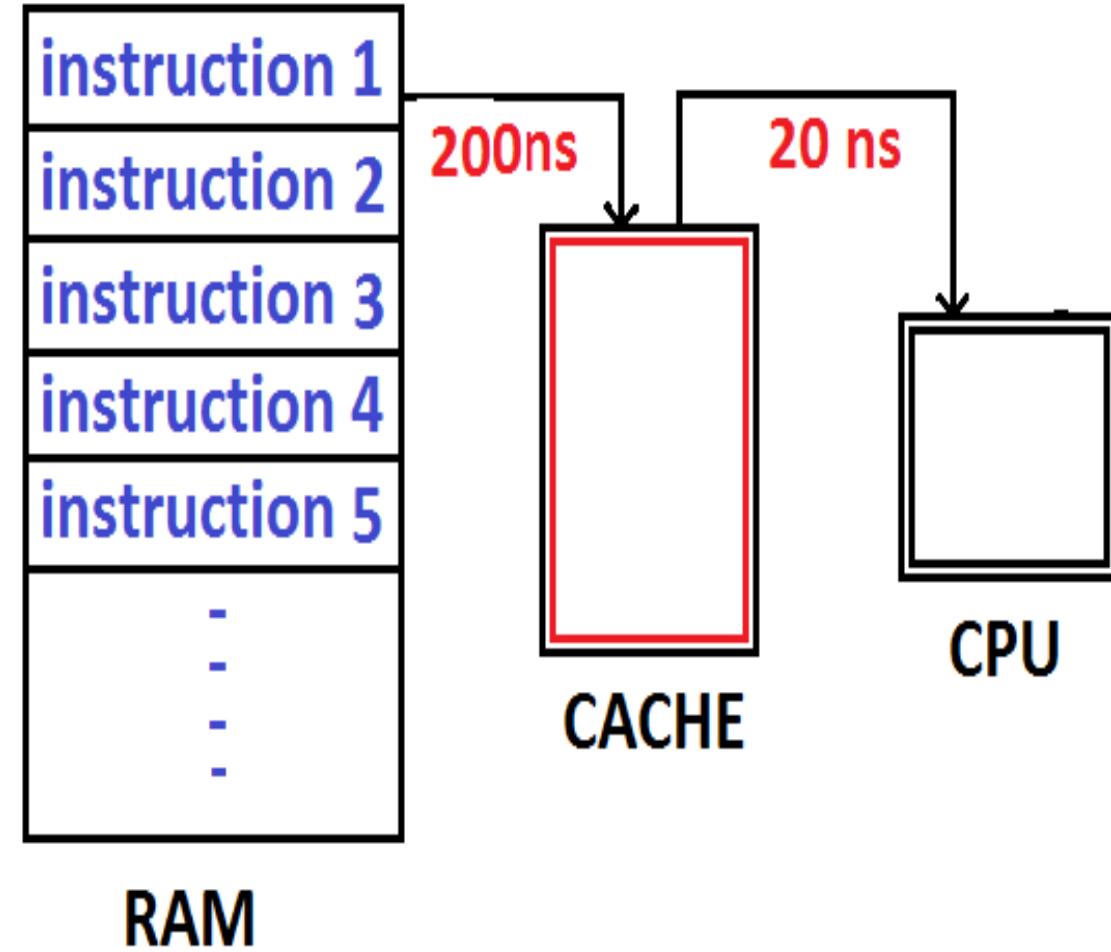
- Instructions :

- Reference instructions in sequence : Spatial locality
- Looping through : Temporal locality

# Performance enhancement -Motivation



RAM



CACHE

CPU

RAM

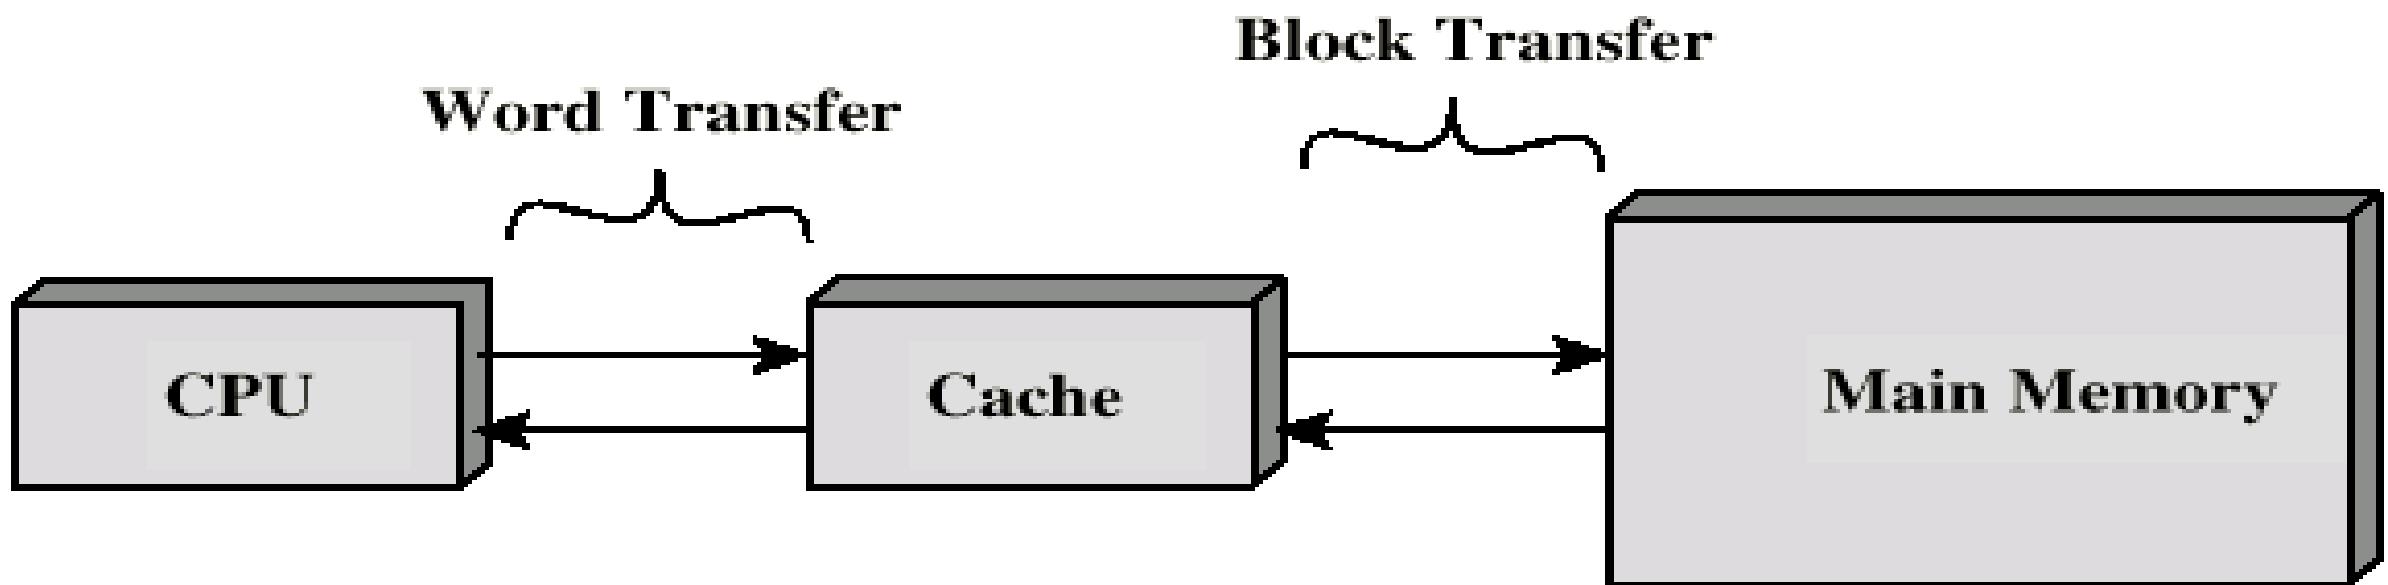
# Cache



# Cache



- Small, fast memory
- Sits between normal main memory and CPU
- May be located on CPU chip or separate module



# Cache and Main Memory Structure



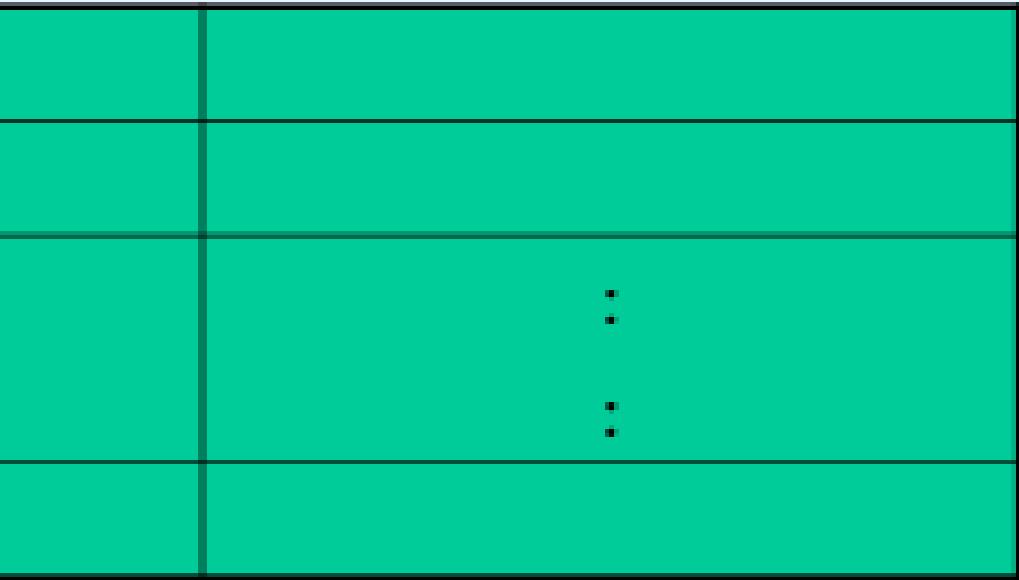
Line  
No:

Tag

Block

0  
1

C-1



Cache

Main Memory  $2^{n-1}$

←Word length →

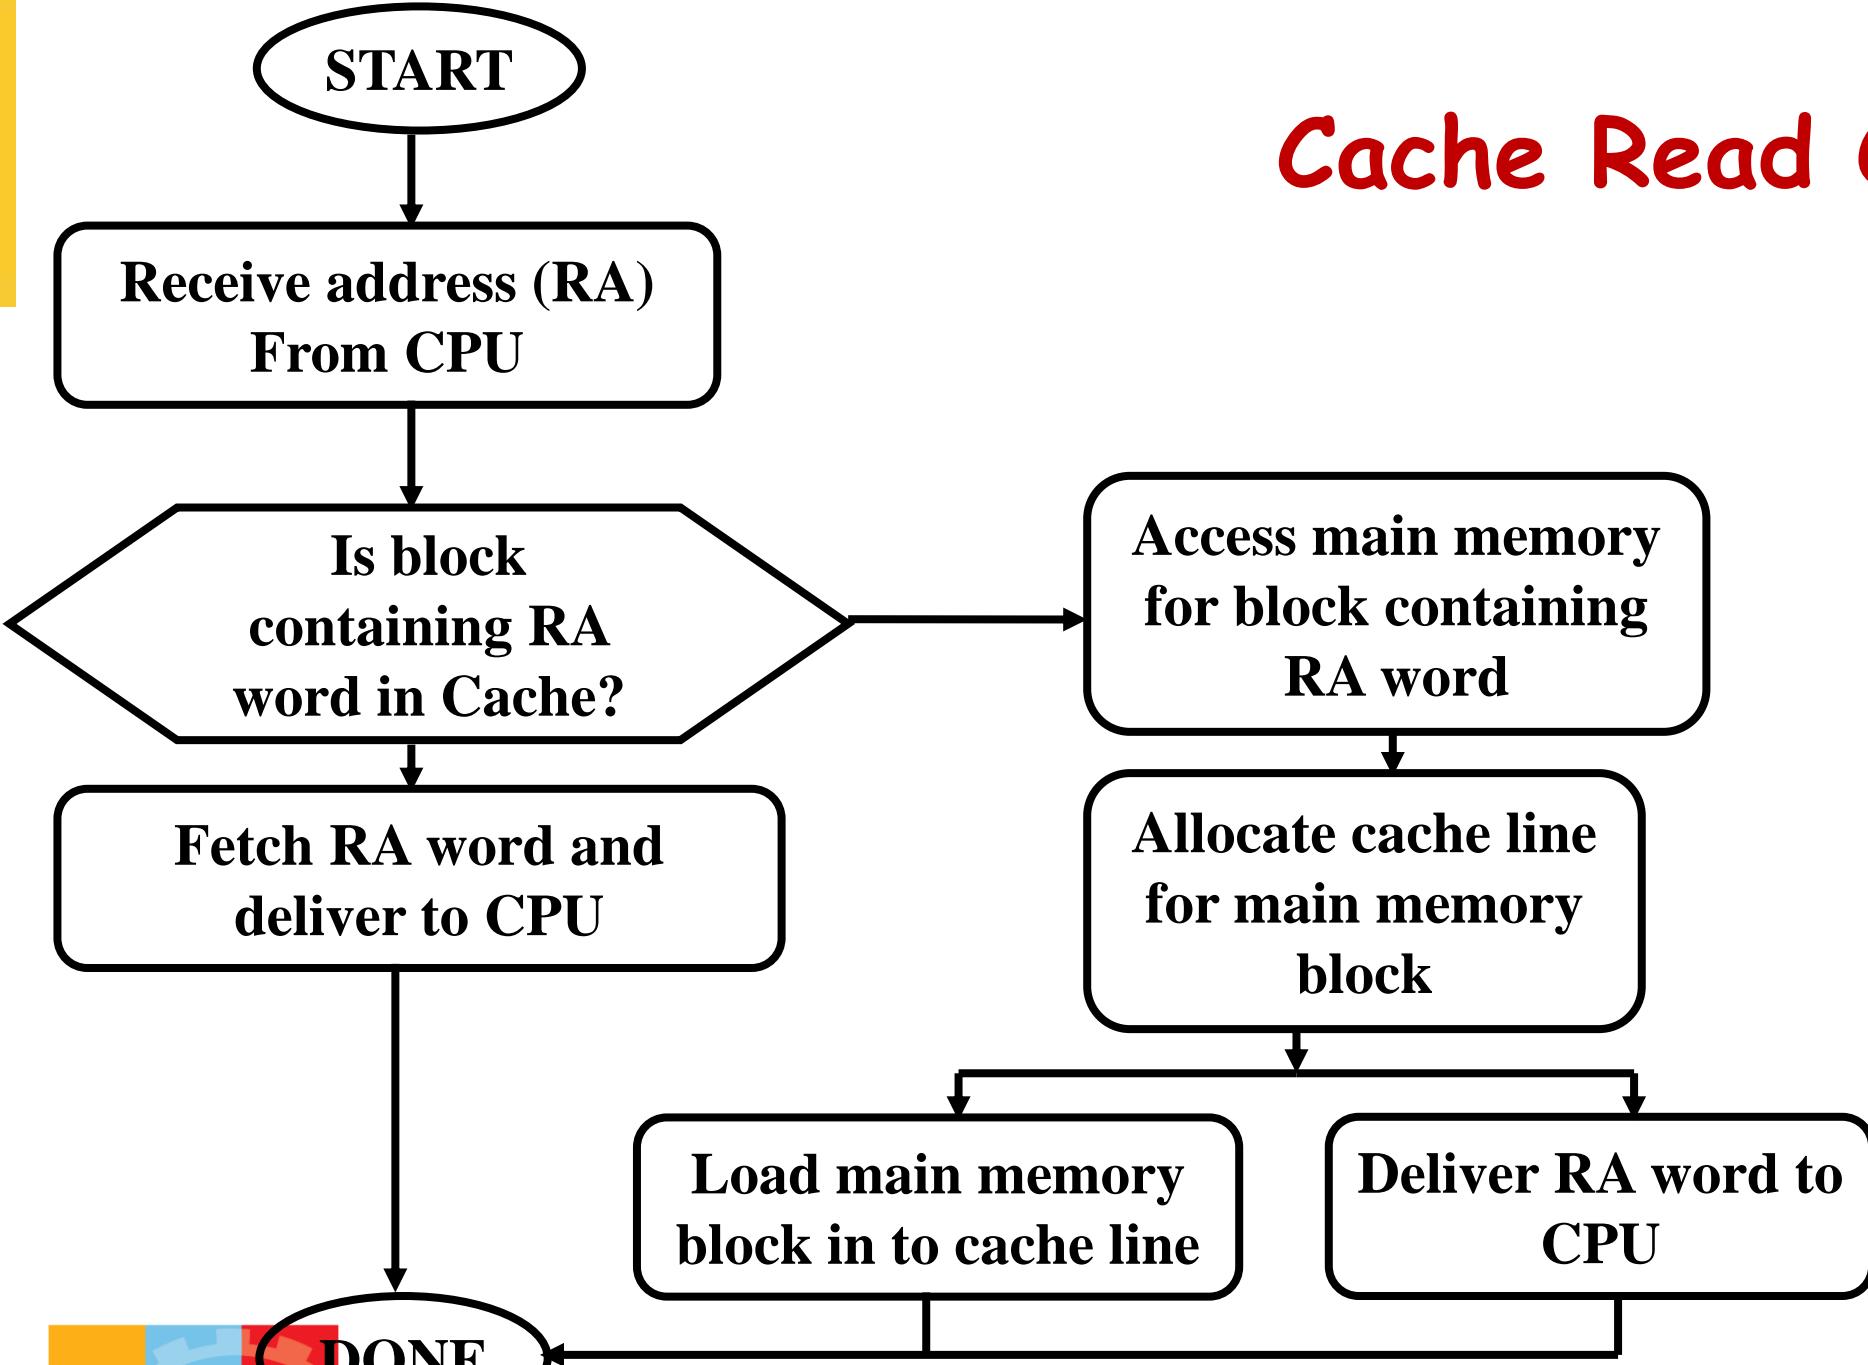
0  
1  
2

Block

⋮  
⋮  
⋮  
⋮

Block

# Cache Read Operation

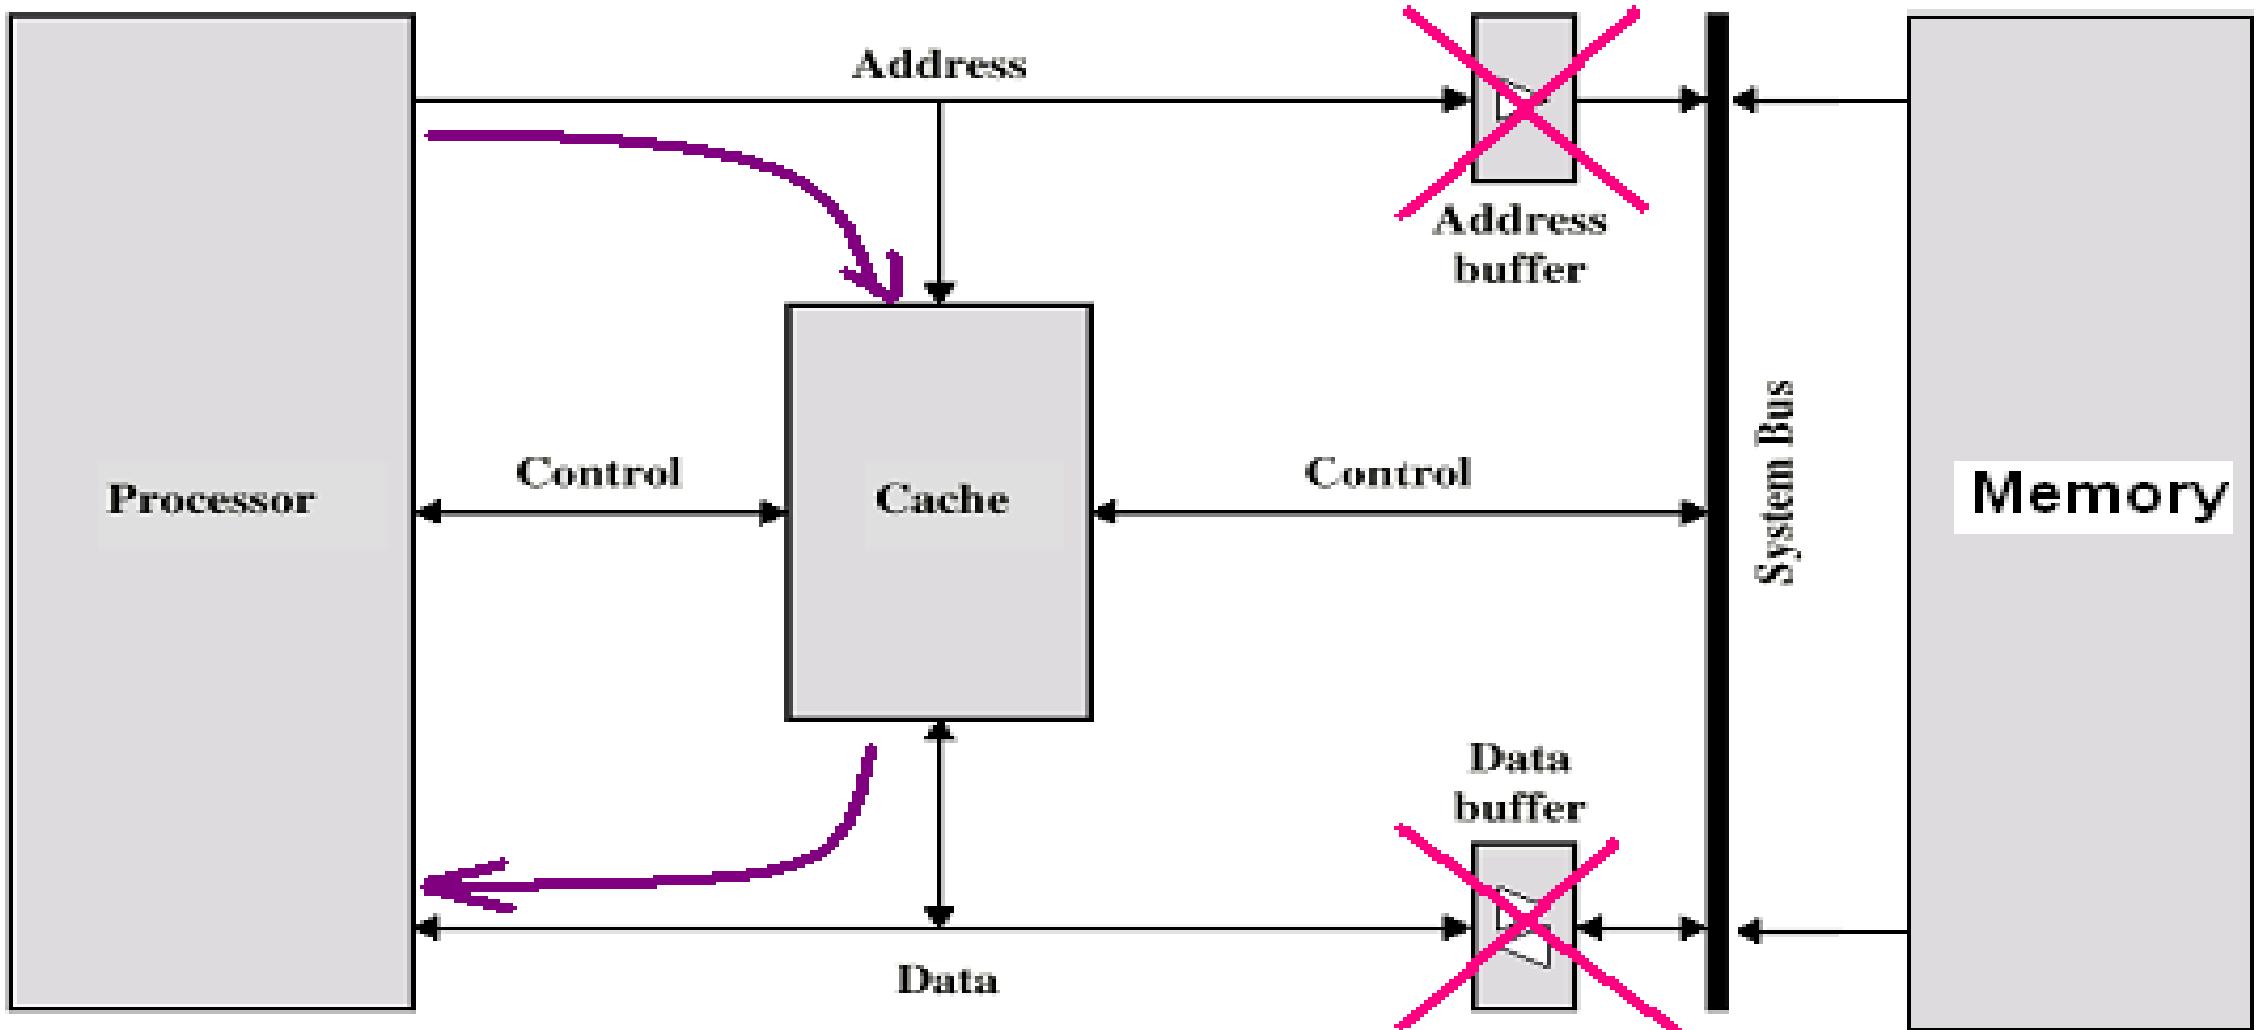


# Performance of cache

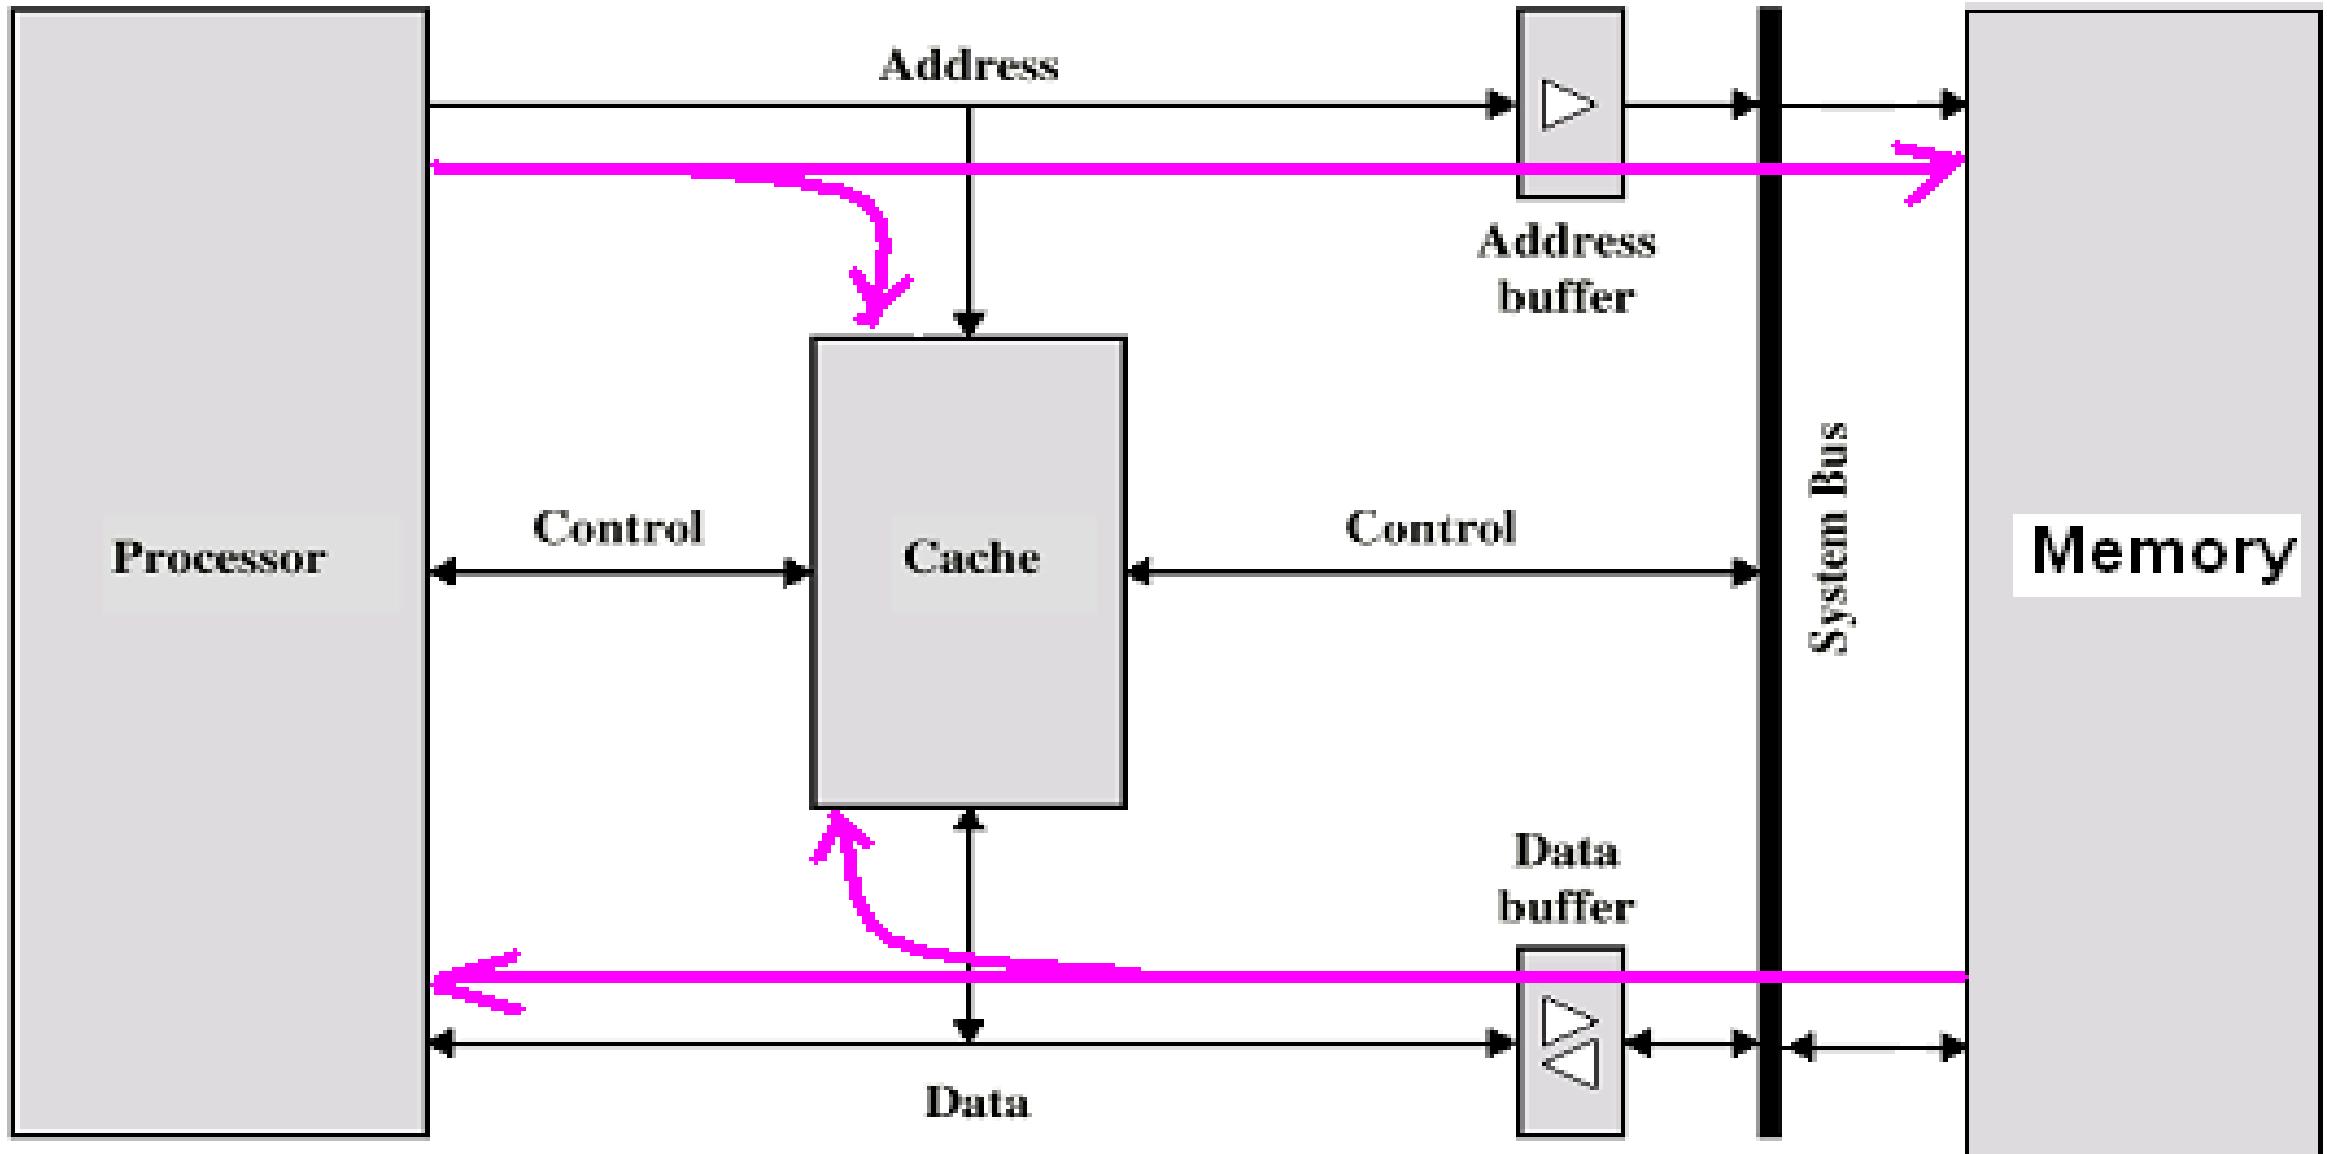


- Hit ratio : Number of Hits / total references to memory
- Hit
- Miss

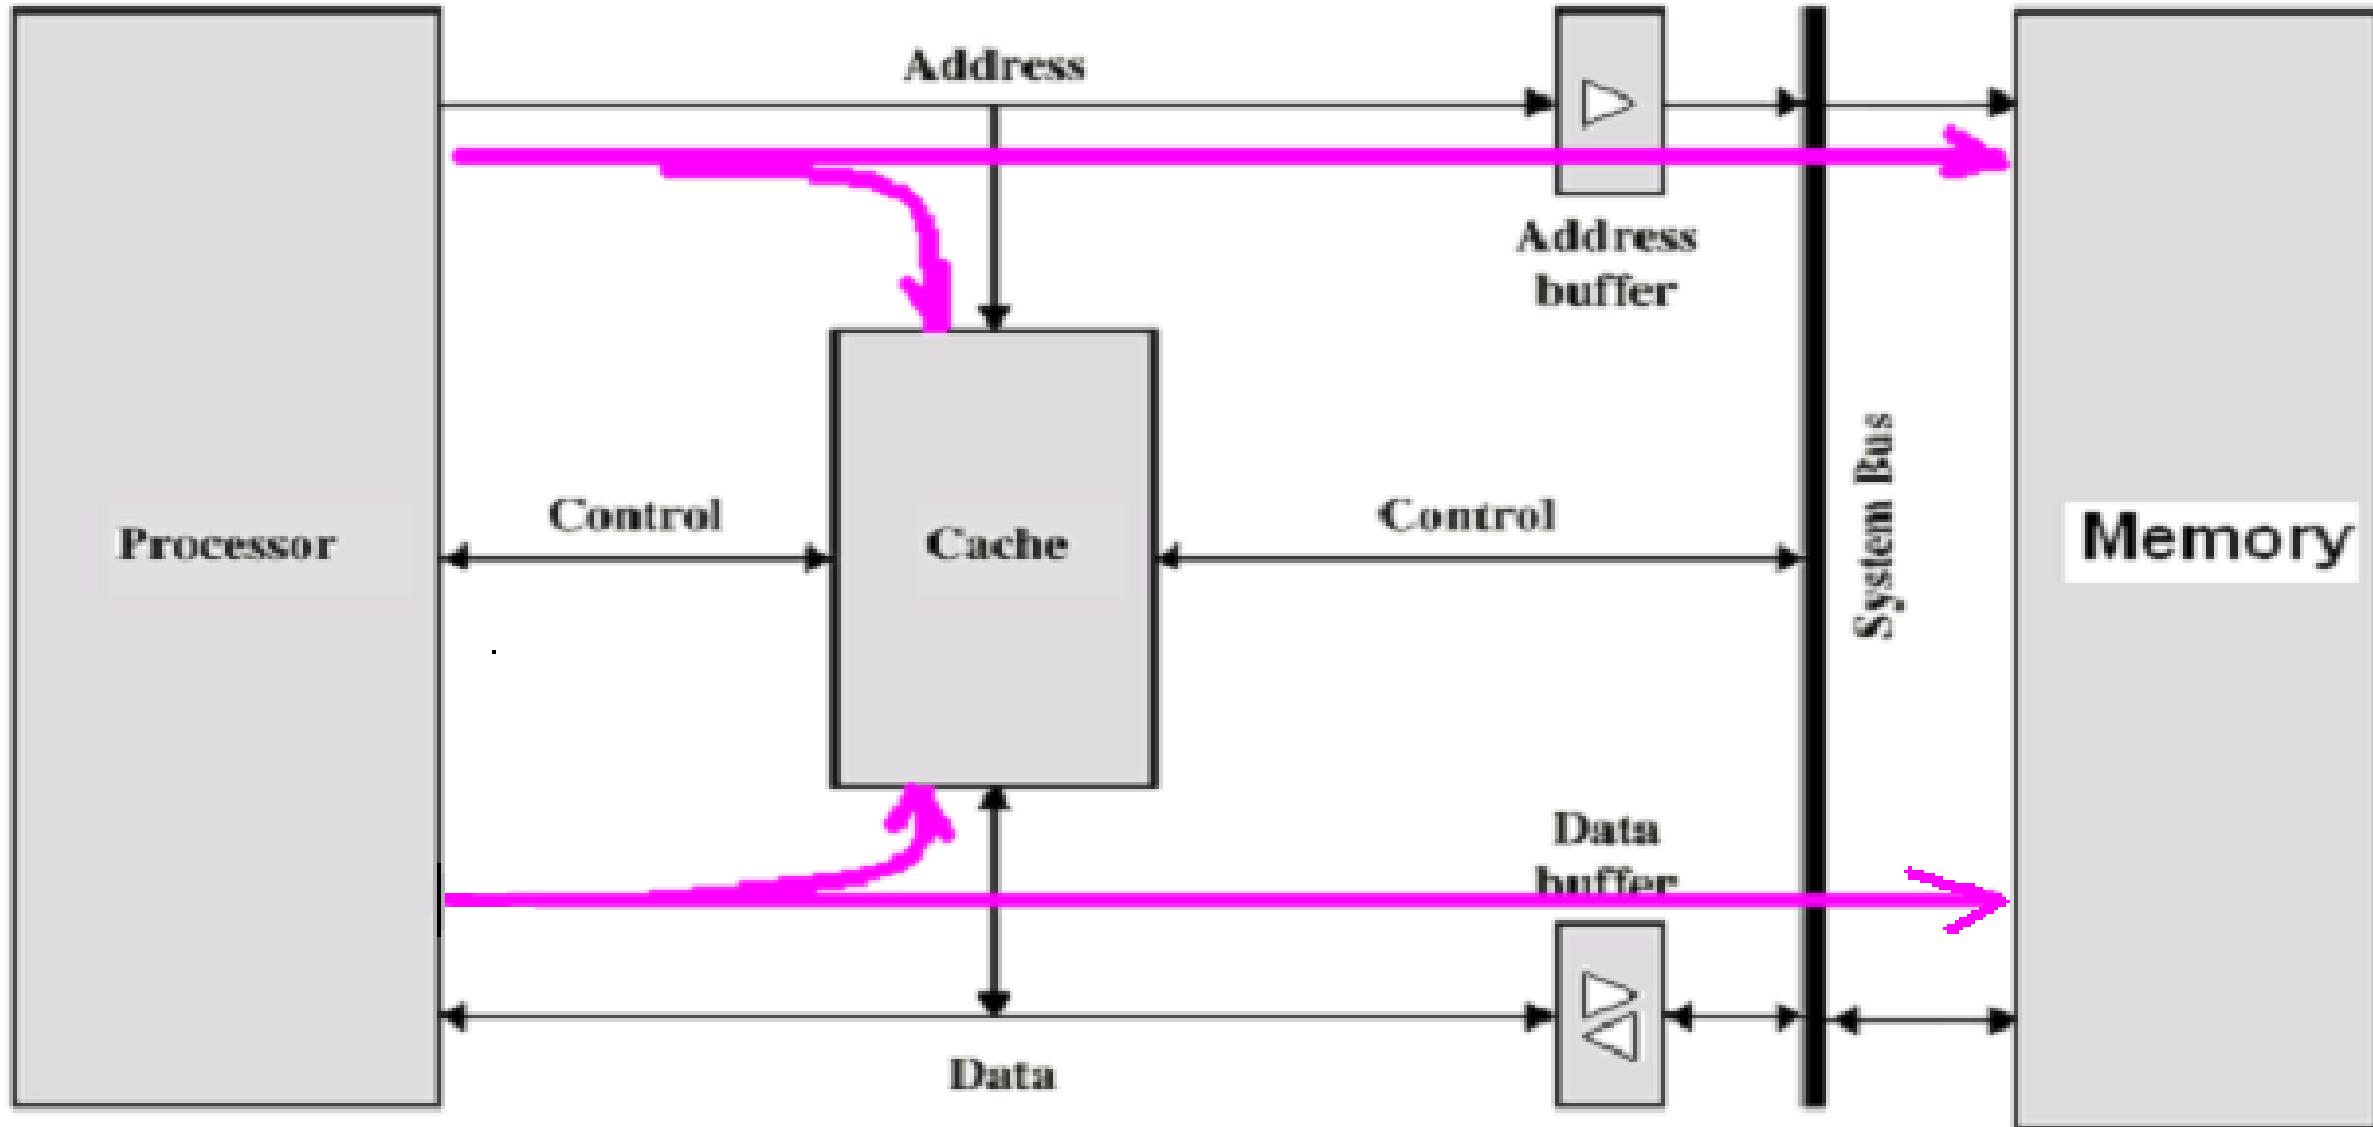
# Read Hit



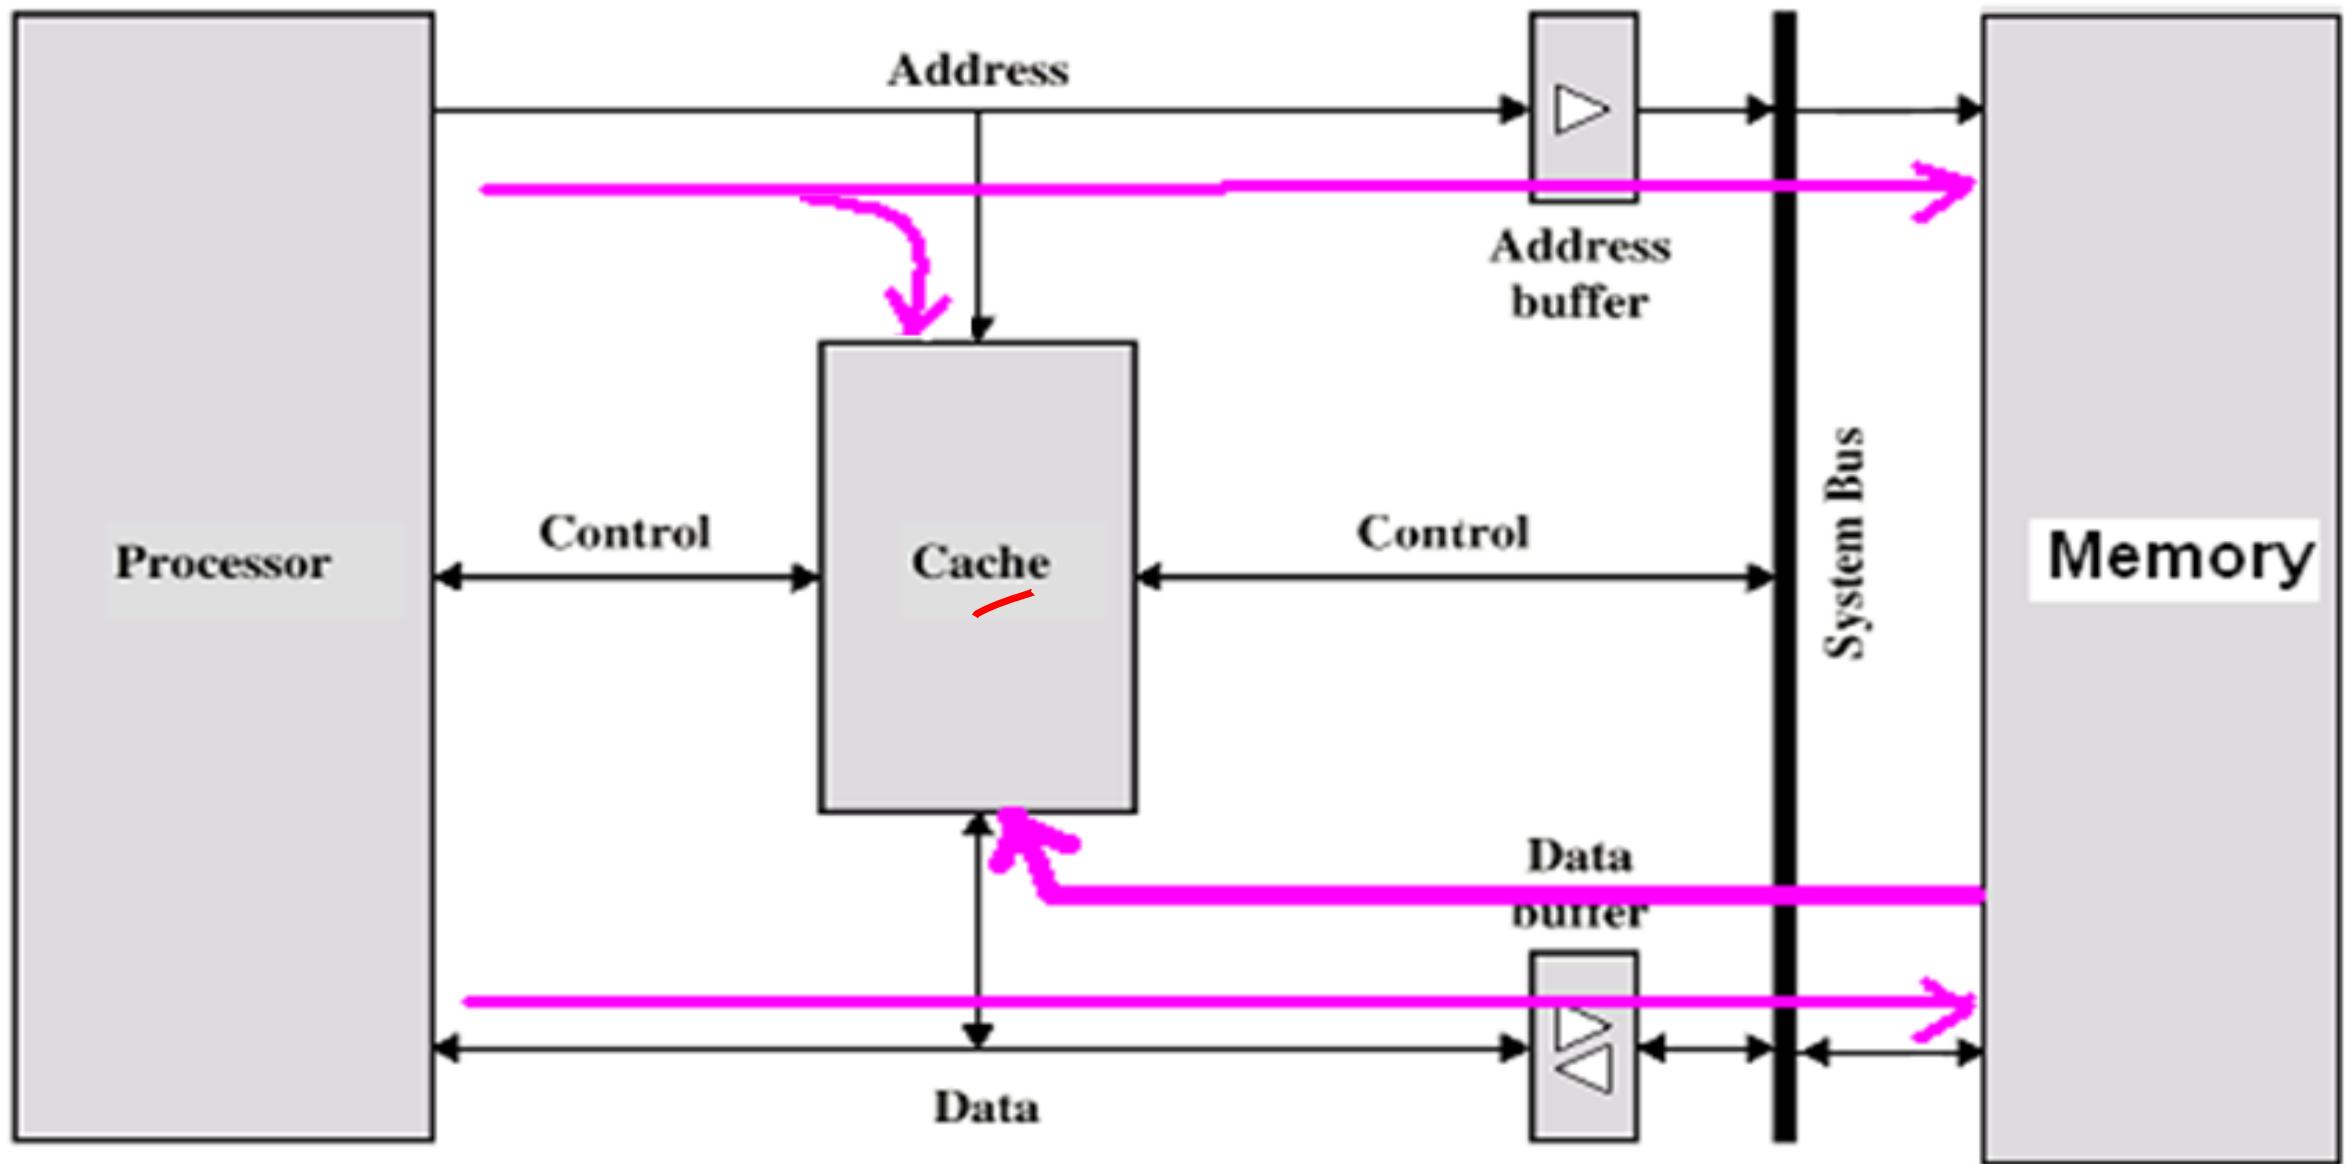
# Read Miss



# Write hit



# Write miss





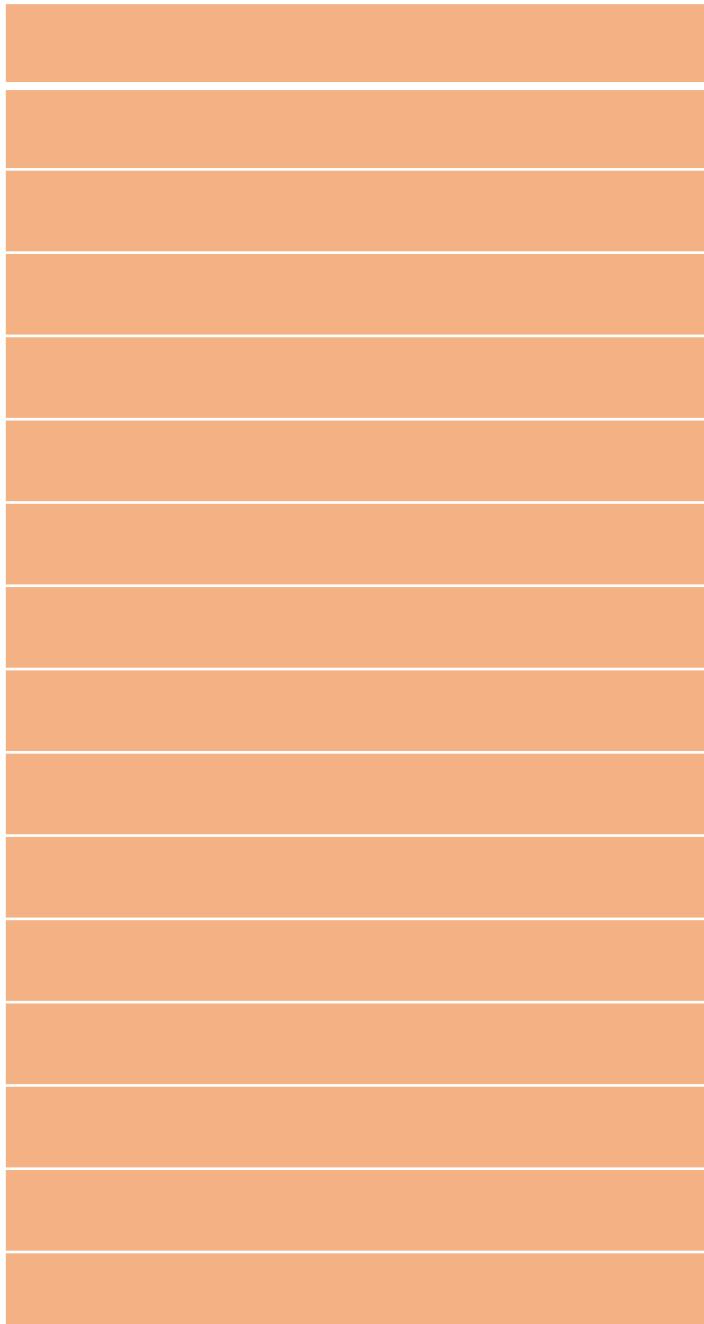
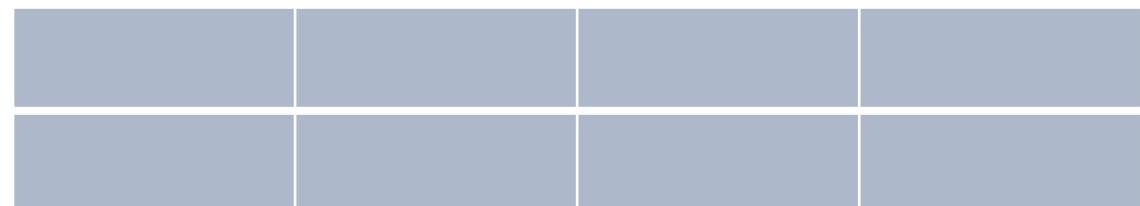
# Mapping Function

- How memory blocks are mapped to cache lines
- Three types
  - Direct mapping
  - Associative mapping
  - Set Associative mapping

# Direct Mapped Cache



- 16 Bytes main memory
  - How many address bits are required?
- Memory block size is 4 bytes
- Cache of 8 Byte
  - How many cache lines?
  - cache contains 2 lines ( 4 bytes per Line)





# Direct Mapped Cache

- Each block of main memory maps to only one cache line
  - i.e. if a block is in cache, it must be in one specific place
  - $i = j \text{ modulo } m$

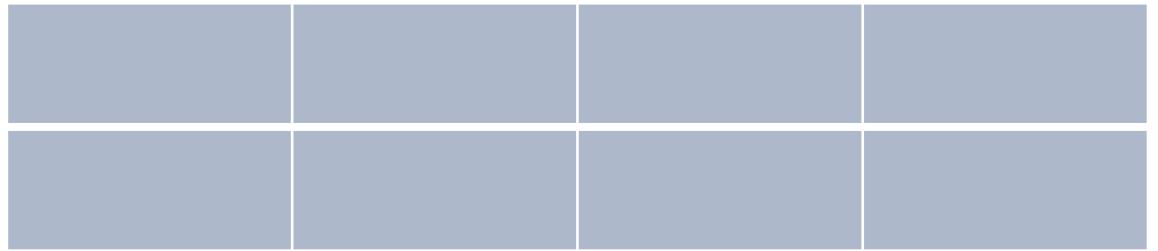
where  $i$  = cache line number

$j$  = main memory block no.

$m$  = no.of lines in the  
cache

# Direct Mapped Cache

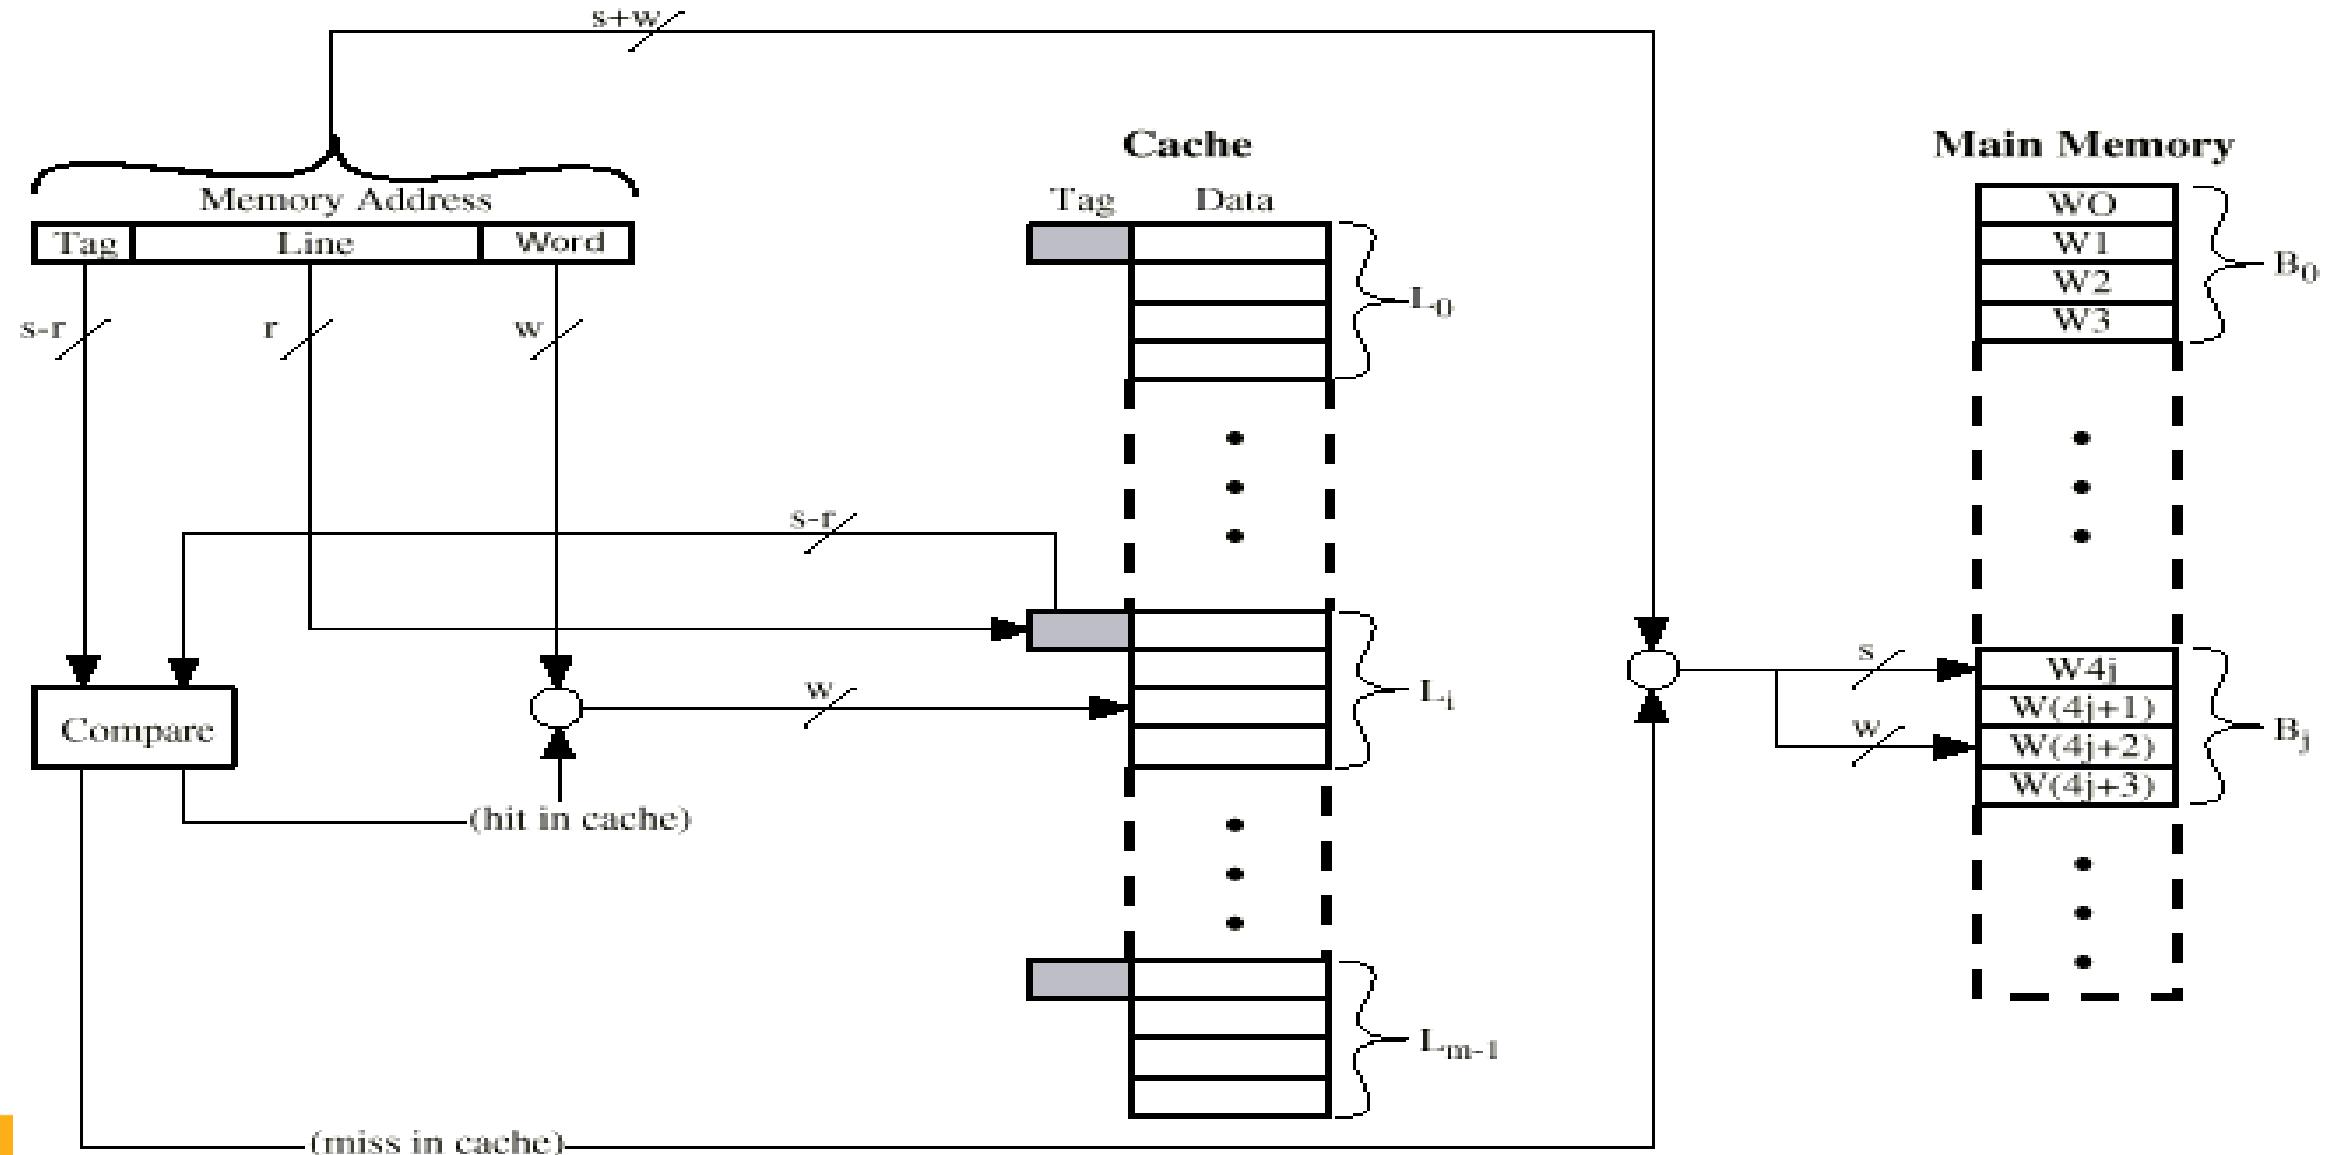
$$i = j \bmod m$$



- Address is split in three parts:
  - Tag
  - Line
  - Word



# Direct Mapped Cache Organization





# Direct mapped cache- Summary

- Address length =  $(s+w)$  bits
- Number of addressable units =  $2^{s+w}$  words or bytes
- Block size = line size =  $2^w$  words or bytes
- Number of blocks in main memory =  $2^{s+w} / 2^w = 2^s$
- Number of lines in cache =  $m = 2^r$
- Size of tag =  $(s-r)$  bits



# Direct mapped cache- Summary



Tag s-r

Line or Slot r

Word w



# Direct mapped cache-pros & cons

- Simple
- Inexpensive
- Fixed location for given block
  - If a program accesses 2 blocks that map to the same line repeatedly, cache misses are very high



# Problem 1 : Direct Mapped Cache

• Given :

- Cache of 64KByte, Cache block of 4 bytes
- 16MBytes main memory

• Find out

- a) Number of bits required to address the main memory
- b) Number of blocks in main memory
- c) Number of cache lines
- d) Number of bits required to identify a word (byte) in a block
- e) Number of bits to identify a block
- f) Tag, Line, Word

# Solution 1 $2^n = \text{capacity}$

Given :

- Cache of 64kByte, Cache block of 4 bytes
- 16MBytes main memory

$\Downarrow$   
MM Block Size

$$\begin{aligned} 2^{10} &\Rightarrow K \\ 2^{20} &\Rightarrow M \\ 2^{30} &- 9 \\ 2^{40} &- T \end{aligned}$$

$$\begin{aligned} 2^0 &= 1 \\ 2^1 &= 2 \\ 2^2 &= 4 \\ 2^3 &= 8 \\ 2^4 &= 16 \\ 2^5 &= 32 \\ 2^6 &= 64 \\ 2^7 &= 128 \\ 2^8 &= 256 \\ 2^9 &= 512 \\ 2^{10} &= 1024 = 1K \end{aligned}$$

Find out

a) Number of bits required to address the main memory

$$\begin{aligned} 2^n &= 16M \\ &= 2^4 \cdot 2^{20} \Rightarrow 2^{24} \quad \therefore [n = 24 \text{ bits}] \end{aligned}$$

b) Number of blocks in main memory

$$\frac{\text{Capacity}}{\text{MM Block Size}} = \frac{16MB}{4B} = \frac{2^{24}}{2^2} = 2^{24-2} = 2^{22} = 2^2 \cdot 2^{20} = [4M \text{ Blocks}]$$

c) Number of cache lines

$$\frac{\text{Cache Size}}{\text{Line Size}} = \frac{64KB}{4B} = \frac{2^6 \cdot 2^{10} B}{2^2 B} = 2^{16-2} \Rightarrow 2^{14} = 16K \cancel{X} \text{ Lines}$$

# Solution 1

Given :

- Cache of 64kByte, Cache block of 4 bytes
- 16MBytes main memory

Find out

- d) Number of bits required to identify a word (byte) in a block?

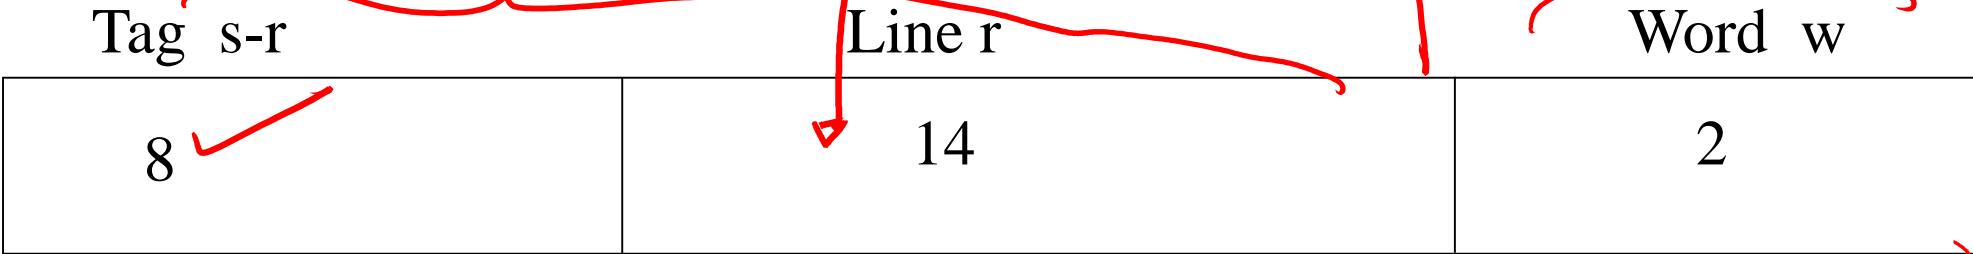
$$2^n = 4 = 2^2$$

$$n = 2 \text{ bits}$$

- e) Number of bits required to identify a block

$$2^n = 4M \Rightarrow 2^{22} \Rightarrow n = 22 \text{ bits}$$

- f) Tag, Line, Word



# Cache lines = 16K lines

$$2^n = 16 \text{ K}$$

$$- = 2^4$$

$$n = 14$$

Tag = # address bits  
 - line-word  
 $= 24 - 14 - 2 \Rightarrow 8$

2 bits

Word w

## Problem 2



Consider a machine with a byte addressable main memory of  $2^{16}$  bytes and block size of 8 bytes. Assume that a direct mapped cache consisting of 32 lines is used with this machine.

- How is a 16-bit memory address divided into tag, line number, and byte number?
- Into what line would bytes with each of the following addresses be stored?

0001 0001 0001 1011

1100 0011 0011 0100

1101 0000 0001 1101

1010 1010 1010 1010

- Suppose the byte with address 0001 1010 0001 1010 is stored in the cache. What are the addresses of the other bytes stored along with it?
- How many total bytes of memory can be stored in the cache?
- Why is the tag also stored in the cache?

## Solution 2

6432168421 ←  
 0 0 1 1 0  
 0 0 0 1 1  
 1 0 1 0 1



Consider a machine with a byte addressable main memory of  $2^{16}$  bytes and block size of 8 bytes. Assume that a direct mapped cache consisting of 32 lines is used with this machine.

a. How is a 16-bit memory address divided into tag, line number, and byte number?

$$\text{Block size} = 8 \text{ bytes}$$

$$2^w = 8 \Rightarrow 2^3$$

$$2^r = 32 \Rightarrow 2^5$$

$$\text{word} = 3 \text{ bits}$$

$$\text{line} = 5 \text{ bits}$$

$$\begin{aligned} \text{Tag} &= \# \text{ address bits} - r - w \\ &= 16 - 5 - 3 = 8 \text{ bits} \end{aligned}$$

b. Into what line would bytes with each of the following addresses be stored?

	<u>Tag</u>	<u>line</u>	<u>word</u>	
0001	0001	0001	1011	$\Rightarrow 00011 \Rightarrow l_3 \Rightarrow \text{line } 3$
1100	0011	0011	0100	$\Rightarrow 00110 \Rightarrow l_6 \Rightarrow \text{line } 6$
1101	0000	0001	1101	$\Rightarrow 00011 \Rightarrow l_3 \Rightarrow \text{line } 3$
1010	1010	1010	1010	$\Rightarrow 10101 \Rightarrow l_{21} \Rightarrow \text{line } 21$

## Solution 2

Consider a machine with a byte addressable main memory of  $2^{16}$  bytes and block size of 8 bytes. Assume that a direct mapped cache consisting of 32 lines is used with this machine.

- c. Suppose the byte with address 0001 1010 0001 1010 is stored in the cache. What are the addresses of the other bytes stored along with it?

Tag	line		000
0001	1010	0001	1000
0001	1010	0001	1001
0001	1010	0001	1010 : given in the problem
0001	1010	0001	1011
0001	1010	0001	1100
0001	1010	0001	1101
0001	1010	0001	1110
0001	1010	0001	1111



## Solution 2

Consider a machine with a byte addressable main memory of  $2^{16}$  bytes and block size of 8 bytes. Assume that a direct mapped cache consisting of 32 lines is used with this machine.

- d. How many total bytes of memory can be stored in the cache?

$$\text{# lines} \times \text{block size} = 32 \times 8 \text{ bytes} \Rightarrow 256 \text{ bytes}$$

- e. Why is the tag also stored in the cache?

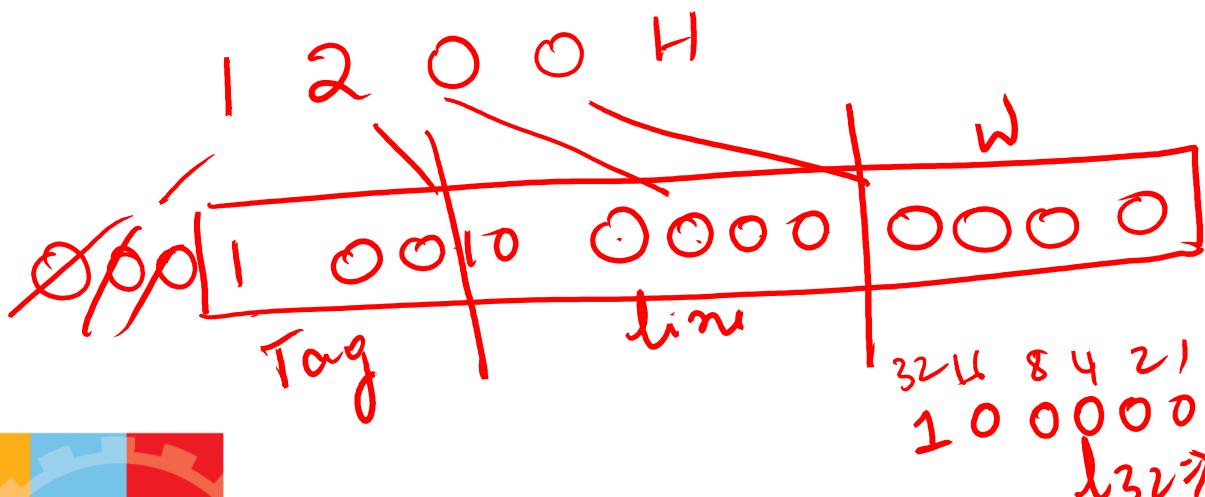
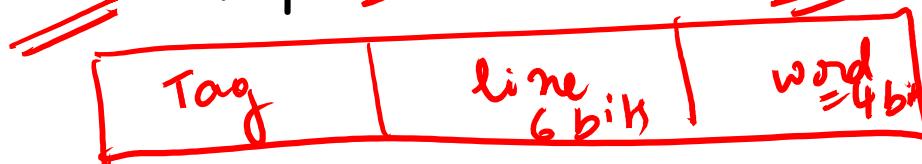
## Problem 3

$$\text{MM Size} = 8K$$

$$2^n = 8K \Rightarrow 2^{\frac{n}{3}} \cdot 2^{10} = 2^{13}$$

$$n = 13 \text{ bits}$$

- Consider a direct-mapped cache with 64 cache lines and a block size of 16 bytes and main memory of 8K (Byte addressable memory). To what line number does byte address 1200H map?  $\rightarrow$  line # 32



Hexadecimal	Binary	Decimal
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
A	1010	10
B	1011	11
C	1100	12
D	1101	13
E	1110	14
F	1111	15

$$\text{Block size} = 16 \text{ bytes}$$

$$2^W = 16 \Rightarrow 2^4 \Rightarrow W = 4 \text{ bits}$$

$$\text{Cache lines} = 64$$

$$2^R = 64 \Rightarrow R = 6$$

# Problem 4

- The system uses a L1 cache with direct mapping and 32-bit address format is as follows:

	Tag s-r	Line r	Word w
	17	11	4

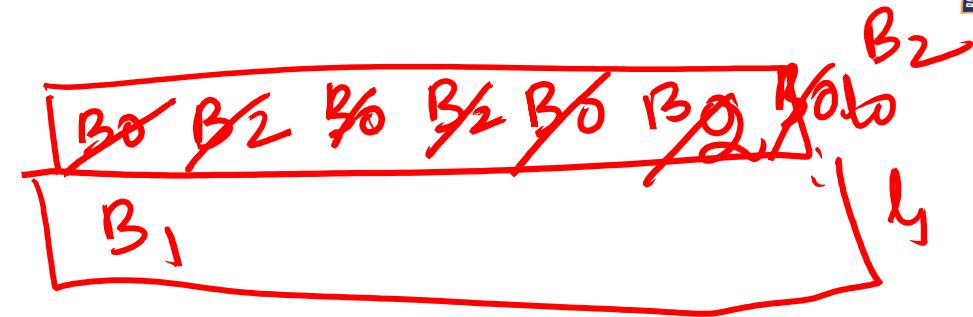
bits 0 - 3 = offset (word) = 4 bits  $\Rightarrow 3-0+1=4$   
 bits 4 - 14 = index bits (Line)  $14-4+1=11$   
 bits 15 - 31 = tag  $31-15+1=17$   
 a) What is the size of cache line?  $\Rightarrow$  word field = 16 bytes  
 b) How many Cache lines are there?  $\Rightarrow$  Line field  $\Rightarrow 2^{11} \Rightarrow 2^2 = 2K \Rightarrow 2048$  lines  
 c) How much space is required to store the tags in the L1 cache?  
 $\Rightarrow \# \text{Cache lines} \times \text{tag size}$   
 $2048 \times 17 = \underline{\hspace{2cm}}$  bits  
 d) What is the total Capacity of cache including tag storage?  
 $= 2048 \times 4 B + 2048 \times 17 \text{ bits}$   
 $= 2048 \times 4 \times 8 + 2048 \times 17 \text{ bits}$   
 $= \underline{\hspace{2cm}} \times 8 = \underline{\hspace{2cm}} \text{ Bytes.}$

# Problem 5



- 16 Bytes main memory, Memory block size is 4 bytes, Cache of 8 Byte (cache is 2 lines of 4 bytes each )
- Block access sequence :

$B_0 \ B_1 \ . \ . \ . \ - \ - \ - \ B_1$   
0 2 0 2 2 0 0 2 0 0 0 2 1



$$\begin{aligned} B_0 &\Rightarrow 0 \ 1 \ . \ 2 = l_0 \\ B_1 &\Rightarrow 1 \ 1 \ . \ 2 = l_1 \\ B_2 &\Rightarrow 2 \ 1 \ . \ 2 = l_0 \end{aligned}$$

Find out hit ratio.

0 2 0 2 2 0 0 2 0 0 0 2 1  
M M M M H H M H M M H H M M

4 / 13



# Problem 5 - Direct Mapped Cache

- 16 Bytes main memory, Memory block size is 4 bytes, Cache of 8 Byte (cache is 2 lines of 4 bytes each )

Block access sequence :

0 2 0 2 2 0 0 2 0 0 0 2 1

Find out hit ratio.

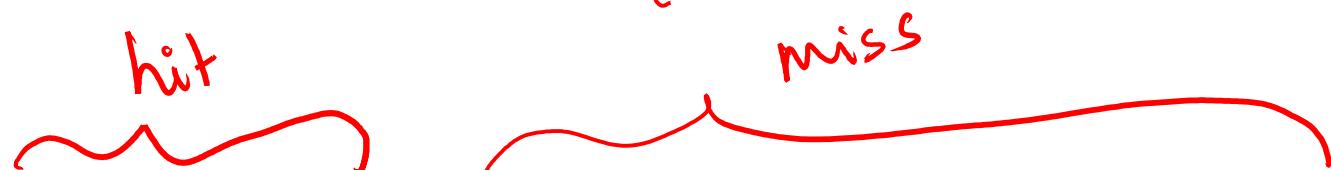
0 2 0 2 2 0 0 2 0 0 0 2 1

# Problem 6

- Suppose a 1024-byte cache has an access time of 0.1 microseconds and the main memory stores 1 Mbytes with an access time of 1 microsecond. A referenced memory block that is not in cache must be loaded into cache .
- Answer the following questions:
- a) What is the number of bits needed to address the main memory?

$$2^n = 1MB \Rightarrow 2^{20} \therefore n = 20 \text{ bits}$$

- a) If the cache hit ratio is 95%, what is the average access time for a memory reference?



$$\text{Avg access time} = \text{hit ratio} * \text{cache access} + (1 - \text{hit ratio}) * (\text{cache access} + \text{memory access})$$

$$\begin{aligned}
 &= 0.95 * 0.1\mu s + (0.05) * (0.1\mu s + 1\mu sec) \\
 &=
 \end{aligned}$$



# Computer Organization and Software Systems

Contact Session 5

Dr. Lucy J. Gudino

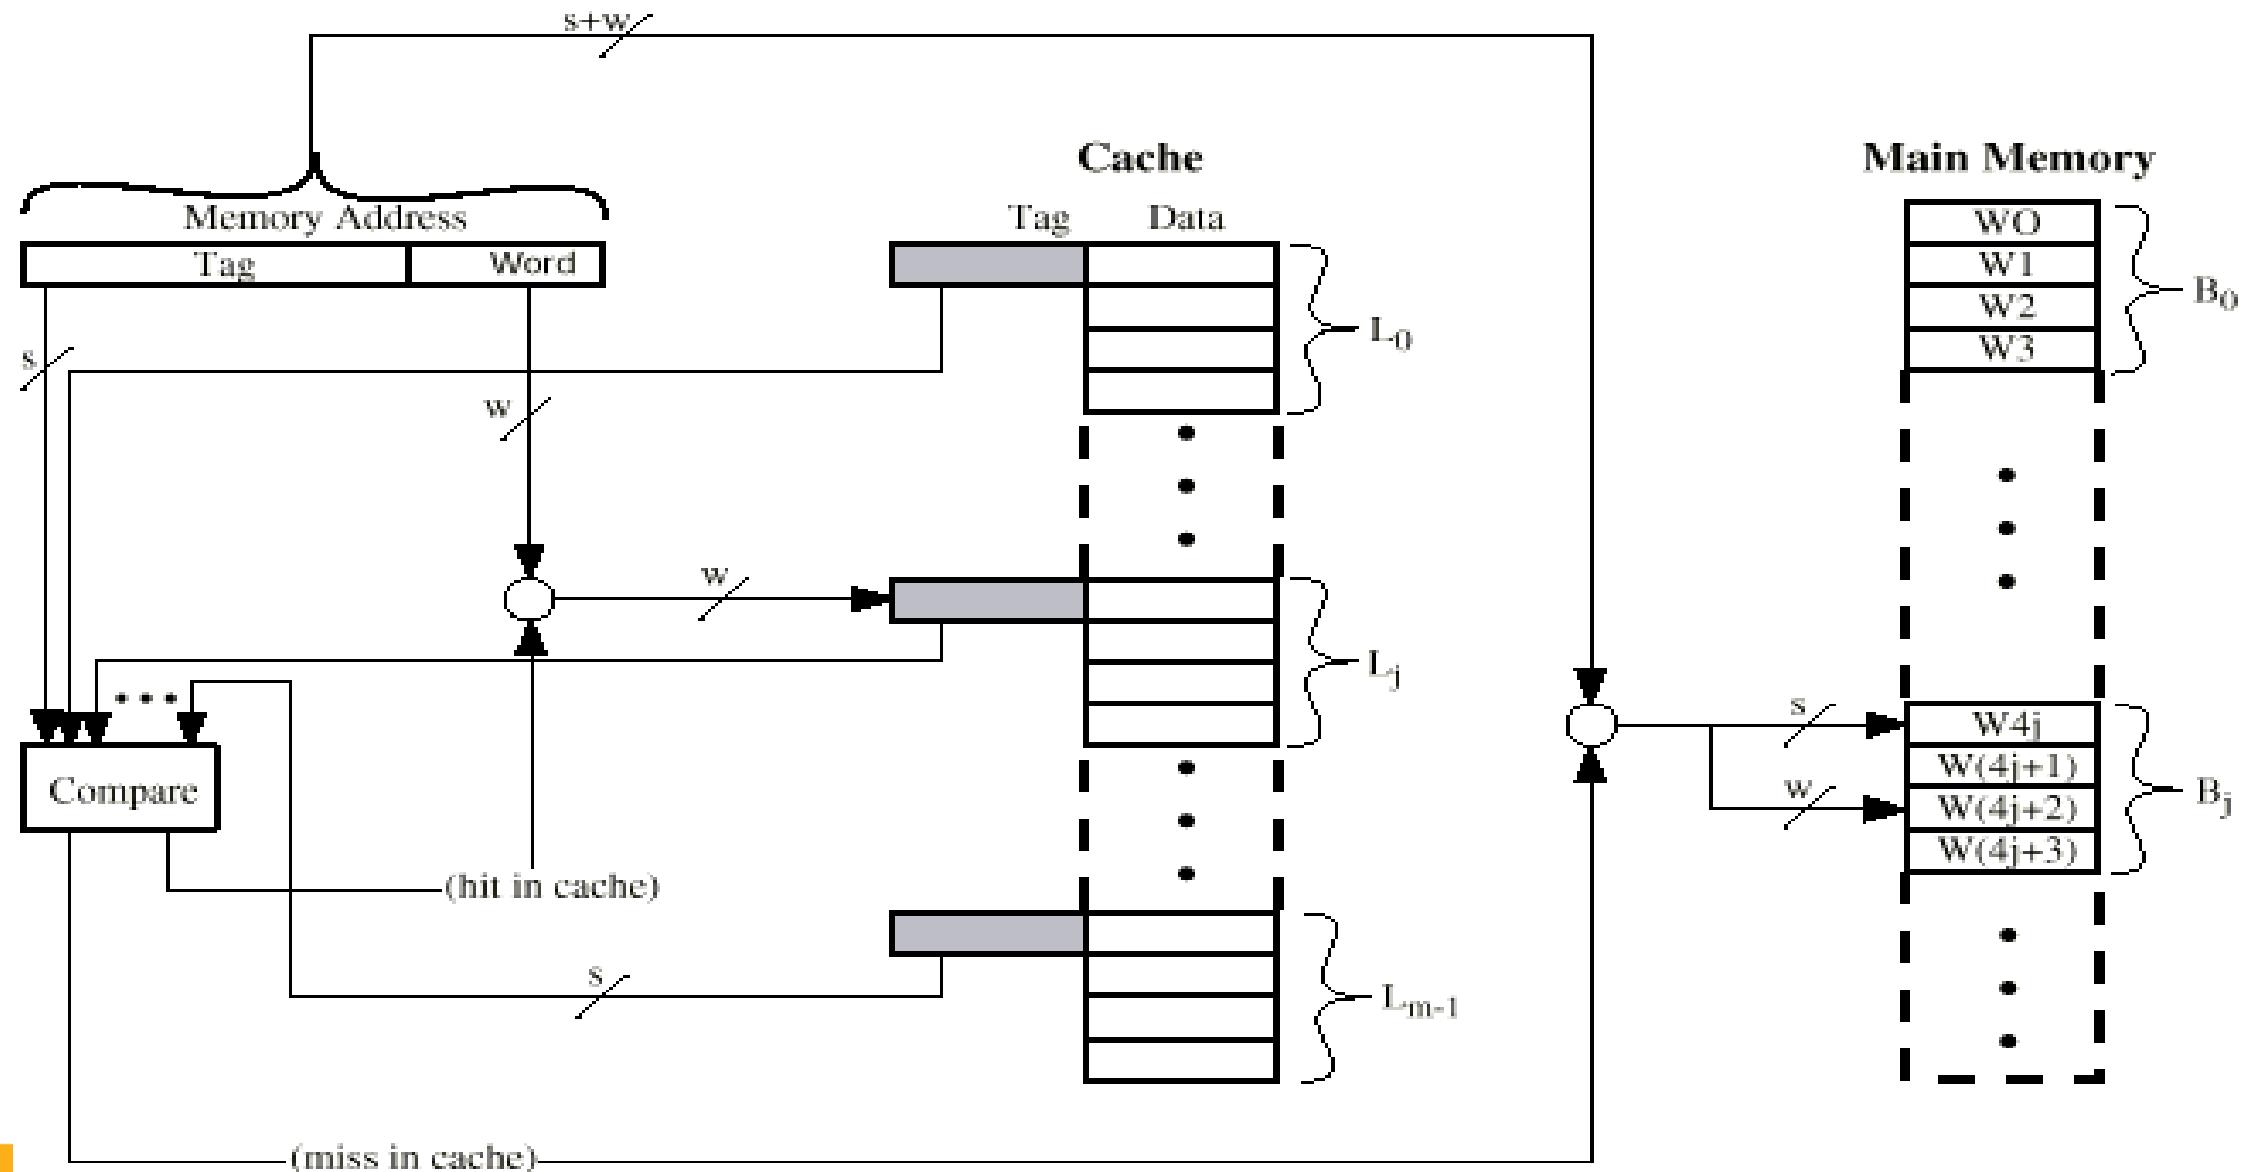


# Associative Mapping



- A main memory block can load into any line of cache
- Memory address is interpreted as tag and word
- Tag uniquely identifies block of memory
- Every line's tag is examined for a match
- Cache searching gets expensive

# Associative Cache Organization





# Associative Mapping Summary

- Address length =  $(s + w)$  bits
- Number of addressable units =  $2^{s+w}$  words or bytes
- Block size = line size =  $2^w$  words or bytes
- Number of blocks in main memory =  $2^{s+w}/2^w = 2^s$
- Number of lines in cache = undetermined
- Size of tag =  $s$  bits

## Problem 7

$$\begin{aligned}
 (a) \quad 2^n &= 32 \text{ MB} \\
 &= 2^5 \cdot 2^{20} \text{ B} \\
 &= 2^{25} \text{ B} \\
 \boxed{n = 25 \text{ bits}}
 \end{aligned}$$

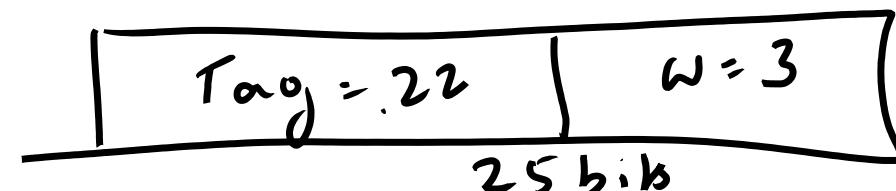
(b) # blocks = Main capacity

$$\begin{aligned}
 &\text{Block size} \\
 &= \frac{2^{25} \text{ B}}{2^3 \text{ B}} = 2^{22} = 2^{22} \text{ blocks} \\
 &= 4 \text{ MB} \text{ blocks}
 \end{aligned}$$

(c) # cache lines

$$\begin{aligned}
 &= \frac{\text{Cache capacity}}{\text{Block size}} \\
 &= \frac{128 \text{ KB}}{8 \text{ B}} = \frac{2^7 \cdot 2^{10} \text{ B}}{2^3 \text{ B}} \\
 &= 2^4 = 2^4 \cdot 2^{10} \text{ bytes} \\
 &= 16 \text{ K bytes}
 \end{aligned}$$

$$(d) \quad 2^w = 2^3 \text{ B} \Rightarrow \boxed{w = 3}$$



Given :

- Cache of 128KByte, Cache block of 8 bytes
- 32 MBytes main memory

Find out

a) Number of bits required to address the memory

b) Number of blocks in main memory

c) Number of cache lines

d) Number of bits required to identify a word (byte) in a block?

e) Tag, Word ✓





## Problem 8

- Cache of 64KByte, Cache block of 4 bytes , 16 M Bytes main memory and associative mapping.

Fill in the blanks:

Number of bits in main memory address = \_\_\_\_\_

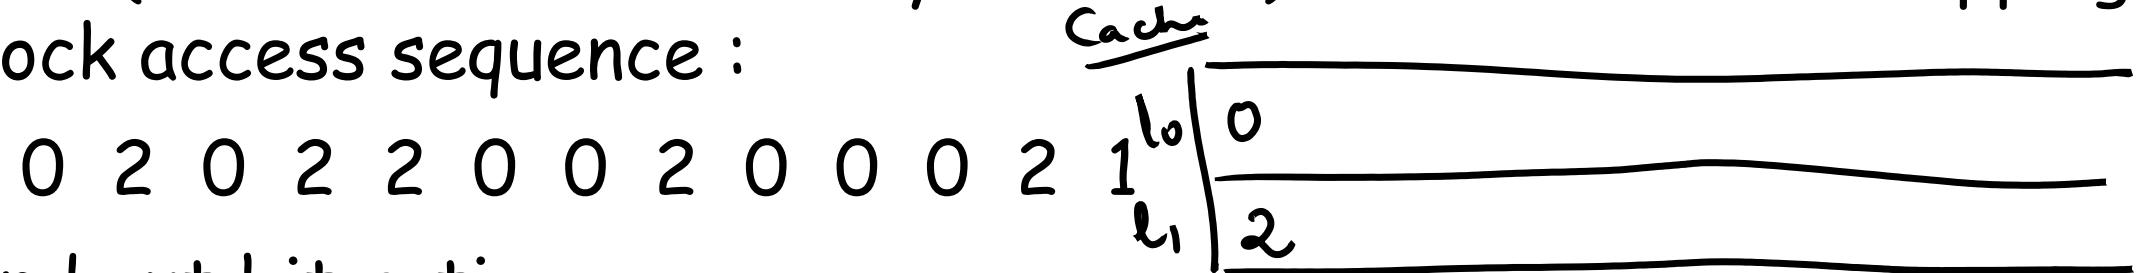
Number of lines in the cache memory = \_\_\_\_\_

Word bits = \_\_\_\_\_

Tag bits = \_\_\_\_\_

## Problem 9 ~~# blocks~~ = 4 Blocks $\Rightarrow B_0, B_1, B_2, B_3$

- 16 Bytes main memory, Memory block size is 4 bytes, Cache of 8 Byte (cache is 2 lines of 4 bytes each) and associative mapping.
- Block access sequence :



Find out hit ratio.

0 2 0 2 2 0 0 2 0 0 0 2 1  
 M M H H H H H H H H H M

$$\text{Hit ratio} = \frac{10}{13}$$

$$DM \Rightarrow \frac{4}{13}$$

# Set Associative Mapping

- Cache is divided into a number of sets ( $v$  sets each with  $k$  lines)  $v = 10 \text{ sets}$
- $m = v * k$   $k = 2$
- $i = j \text{ modulo } v$   $m = v * k$   
 $= 10 * 2$   
 $= 20$
- where  $i$  = cache set number
- $j$  = main memory block number
- $v$  = number of sets in the cache
- Each set contains ' $k$ ' number of lines
- A given block maps to any line in a given set
  - e.g. Block B can be in any line of set  $i$
- ~~$k$ -way set associative cache~~
  - 2 way set associative mapping  $\rightarrow$  2 lines per set
  - A given block can be in one of 2 lines in only one set

# Example

- 16 Bytes main memory, Block Size is 2 Bytes,
  - Cache of 8 Bytes, 2 way set associative cache
    - # address bits  $\Rightarrow$  4 bits
    - Cache line size  $\Rightarrow$  2 B
    - # main memory blocks  $\frac{\text{main mem}}{\text{BS}} = 16$
    - # Number of cache lines
    - # lines per set  $\frac{\text{Cache size}}{\text{Line size}}$
    - # of sets

memory blocks :  $\frac{mm\text{ cap}}{8S} = \frac{15}{2} = 8$

f cache lines

- # Number of cache lines

- # lines per set

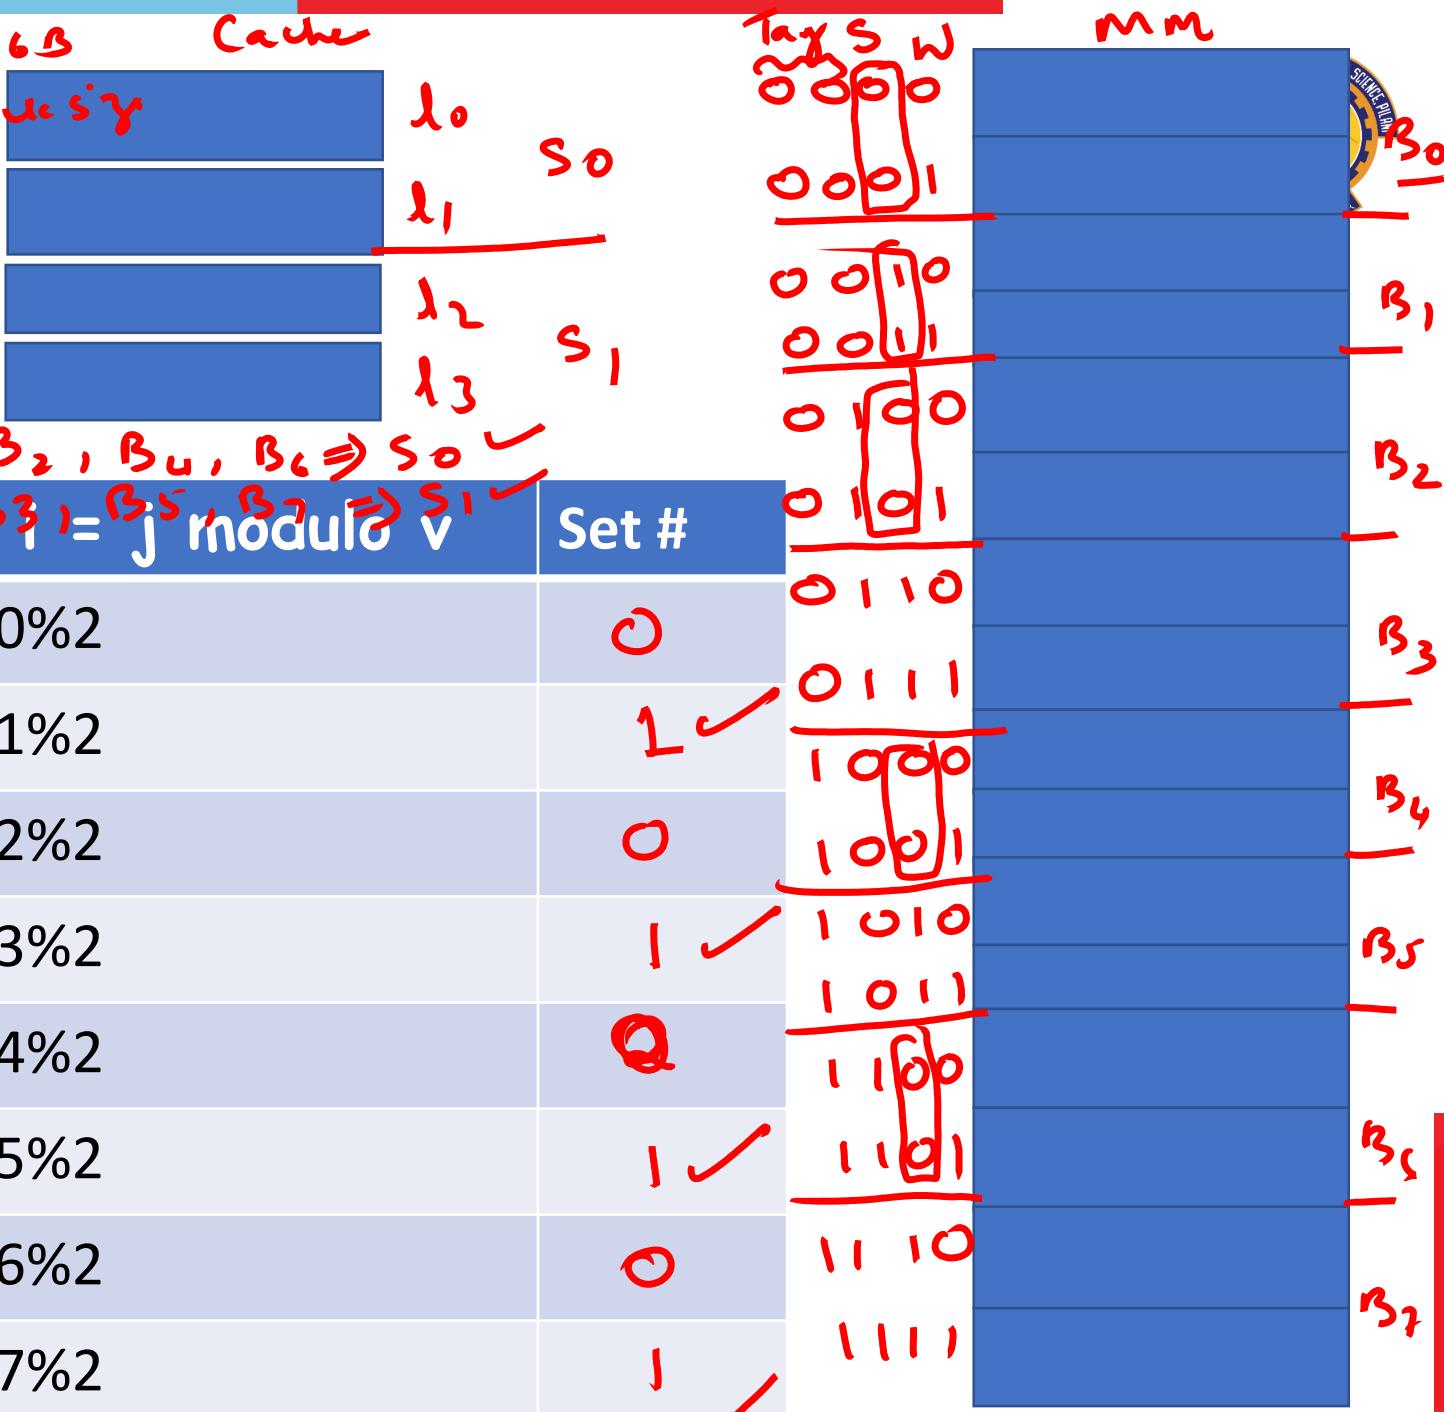
- # of sets

$$\Rightarrow 2 \text{ hours / set}$$

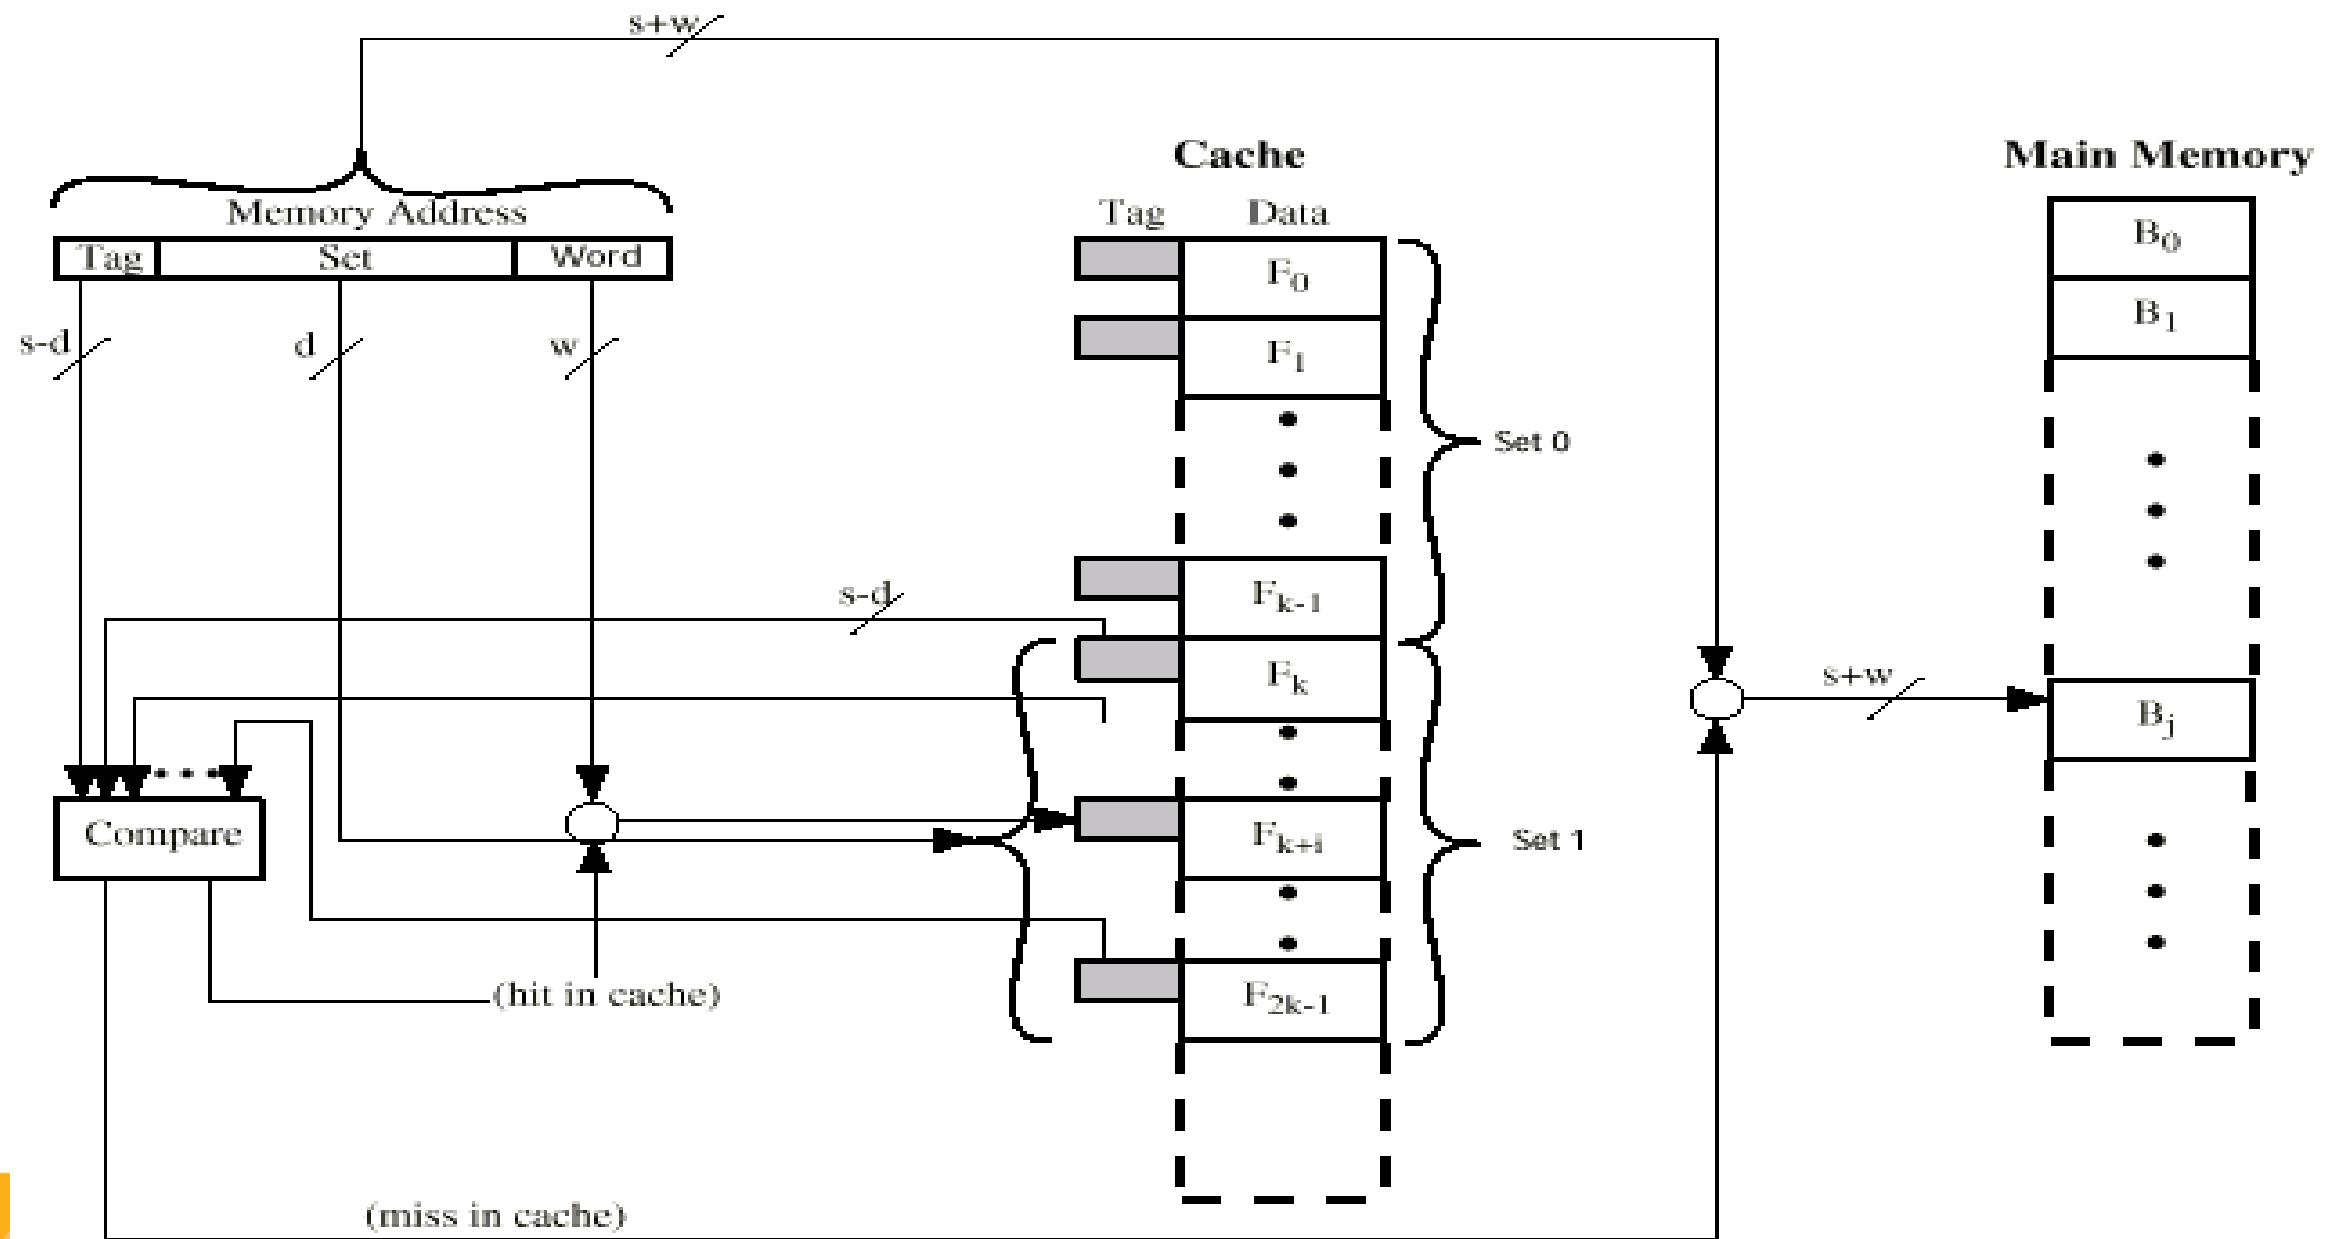
$\dots + \sqrt{+} \kappa$

$$t = \sqrt{2} + 2 \approx 4.12 \text{ s}$$

6362843505 MM=16B Cache  
Cache line size: Block size



# Two-Way Set Associative Cache Organization





# Set Associative Mapping Summary

- Address length =  $(s + w)$  bits
- Number of addressable units =  $2^{s+w}$  words or bytes
- Block size = line size =  $2^w$  words or bytes
- Number of blocks in main memory =  $2^d$
- Number of lines in set =  $k$
- Number of sets =  $v = 2^d$
- Number of lines in cache =  $kv = k * 2^d$
- Size of tag =  $(s - d)$  bits

# Problem 1



Given data

$$\# \text{ of lines} = 64$$

$$m \Rightarrow$$

$$K = 4$$

↓  
4 way set associative cache

$$MM \text{ capacity} = 4K \text{ Blocks}$$

$$\text{Block size} = 128B$$

$$2^W = 128 \Rightarrow 2^7 B$$

$$2^S = 16 \Rightarrow 2^4$$

$$= \# \text{ of address bits} = W - S$$

$$= 19 - 7 - 4$$

$$= 8 \text{ bits}$$

$$\text{a) Capacity of MM} = 4K \times 128B = 2^7 \cdot 2^10 \cdot 2^7 B$$

$$= 2^{19} B$$

$$= 512 KB$$

$$\text{b) Cache Memory capacity} = \# \text{ lines} \times \text{Block size}$$

$$= 64 \times 128B$$

$$= 2^6 \cdot 2^7 B = 2^{13} B$$

$$= 8 KB$$

$$\text{c) } \# \text{ Sets} = \frac{64}{4} = \frac{m}{K} = 16 \text{ sets}$$

(d)





## Problem 2: Home work

- A computer has an 8 GByte memory with 64 bit word sizes. Each block of memory stores 16 words. The computer has a direct-mapped cache of 128 blocks. The computer uses **word level addressing**. What is the address format? If we change the cache to a 4- way set associative cache, what is the new address format?



# Replacement Algorithms (1/3)

## Direct mapped cache

- No choice
- Each block maps to one line and replace that line



# Replacement Algorithms (2/3)

- Needed in Associative & Set Associative mapped cache
- Hardware implemented algorithm (speed)
- Methods:
  - Least Recently Used (LRU)
  - Least Frequently Used (LFU)
  - First In First Out (FIFO)
  - Random



# Replacement Algorithms (3/3)

- **Least Recently used (LRU):** Replace the block that has been in the cache longest with no reference to it
  - e.g. 2 way set associative
  - Uses "USE" bits
  - Most effective method
- **Least frequently used:** Replace block which has had fewest hits
  - Uses counter with each line
- **First in first out (FIFO):** Replace block that has been in cache longest
  - Round robin or circular buffer technique
- **Random**

# Problem 3



- Consider a reference pattern that accesses the sequence of blocks 0, 4, 0, 2, 1, 8, 0, 1, 2, 3, 0, 4. Assuming that the cache uses associative mapping, find the hit ratio for a cache with four lines
  - a) LRU
  - b) LFU
  - c) FIFO

## Problem 2 - LRU

→

Ref	0	4	0	2	1	8	0	1	2	3	0	4
time	0	1	2	3	4	5	6	7	8	9	10	11
L0	0 <sub>0</sub>	0 <sub>0</sub>	0 <sub>2</sub>	0 <sub>2</sub>	0 <sub>2</sub>	0 <sub>2</sub>	0 <sub>6</sub>	0 <sub>6</sub>	0 <sub>6</sub>	0 <sub>10</sub>	0 <sub>10</sub>	
L1		4 <sub>1</sub>	4 <sub>1</sub>	4 <sub>1</sub>	(4 <sub>1</sub> )	8 <sub>5</sub>	8 <sub>5</sub>	8 <sub>5</sub>	(8 <sub>5</sub> )	3 <sub>9</sub>	3 <sub>9</sub>	3 <sub>9</sub>
L2				2 <sub>3</sub>	2 <sub>3</sub>	2 <sub>3</sub>	2 <sub>3</sub>	2 <sub>3</sub>	2 <sub>8</sub>	2 <sub>8</sub>	2 <sub>8</sub>	2 <sub>8</sub>
L3					1 <sub>4</sub>	1 <sub>4</sub>	1 <sub>4</sub>	1 <sub>7</sub>	1 <sub>7</sub>	1 <sub>7</sub>	1 <sub>7</sub>	4 <sub>11</sub>
H/M	M	M	H	M	M	M	H	H	H	M	H	M

$$\text{Hit ratio} \geq 5/12$$

## Problem 2 - LFU



Ref	0	4	0	2	1	8	0	1	2	3	0	4
L0	O <sub>1</sub>	O <sub>1</sub>	O <sub>2</sub>	O <sub>2</sub>	O <sub>2</sub>	O <sub>2</sub>	O <sub>3</sub>	O <sub>3</sub>	O <sub>3</sub>	O <sub>3</sub>	O <sub>4</sub>	O <sub>4</sub>
L1	4 <sub>1</sub>	4 <sub>1</sub>	4 <sub>1</sub>	4 <sub>1</sub>	(4 <sub>1</sub> )	8 <sub>1</sub>	8 <sub>1</sub>	8 <sub>1</sub>	(8 <sub>1</sub> )	3 <sub>1</sub>	(3 <sub>1</sub> )	4 <sub>1</sub>
L2			2 <sub>1</sub>	2 <sub>1</sub>	2 <sub>1</sub>	2 <sub>1</sub>	2 <sub>1</sub>	2 <sub>1</sub>	2 <sub>2</sub>	2 <sub>2</sub>	2 <sub>2</sub>	2 <sub>2</sub>
L3				1 <sub>1</sub>	1 <sub>1</sub>	1 <sub>1</sub>	1 <sub>2</sub>	1 <sub>2</sub>	1 <sub>2</sub>	1 <sub>2</sub>	1 <sub>2</sub>	1 <sub>2</sub>
H/M	M	M	H	M	M	M	H	H	H	M	H	M

Hit ratio = 5/12

L<sub>0</sub> ⇒ B<sub>0</sub>

L<sub>1</sub> ⇒ B<sub>4</sub>

L<sub>2</sub> ⇒ B<sub>2</sub>

L<sub>3</sub> ⇒ B<sub>1</sub>

## Problem 2 - FIFO

⇒

Ref	0	<del>4</del>	0	2	1	8	0	1	2	3	0	4
time	0	1	2	3	4	5	6	7	8	9	10	11
L0	0 <sub>0</sub>	8 <sub>5</sub>										
L1	4 <sub>1</sub>	0 <sub>6</sub>										
L2			2 <sub>3</sub>	3 <sub>9</sub>	3 <sub>9</sub>	3 <sub>9</sub>						
L3				1 <sub>4</sub>	4 <sub>11</sub>							
H/M	M	M	H	M	M	M	M	H	H	H	H	M

Hit ratio : 4/12

L<sub>0</sub> ⇒ B<sub>8</sub>

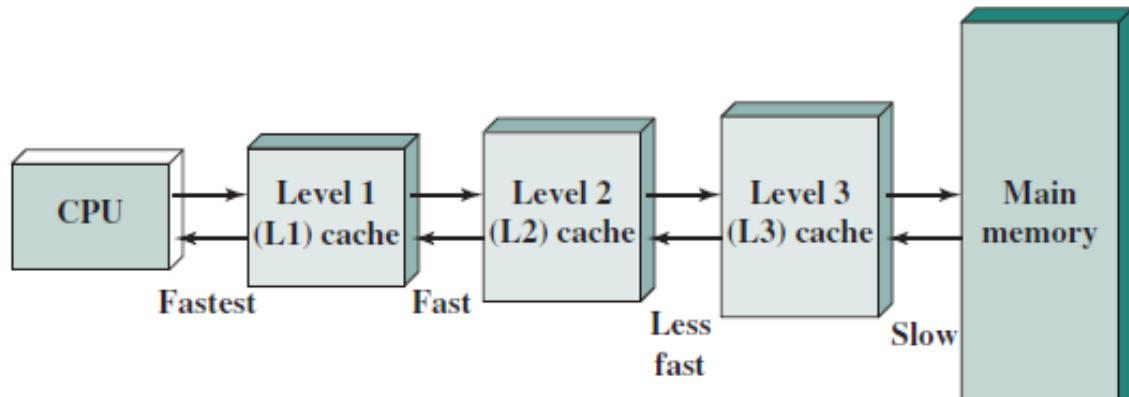
L<sub>1</sub> ⇒ B<sub>0</sub>

L<sub>2</sub> ⇒ B<sub>3</sub>

L<sub>3</sub> ⇒ B<sub>4</sub>

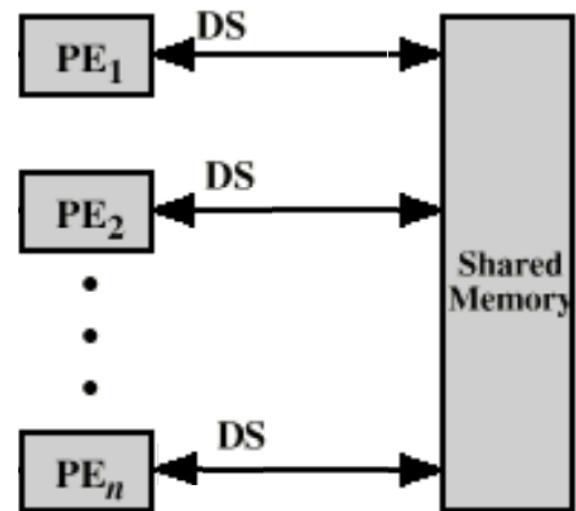
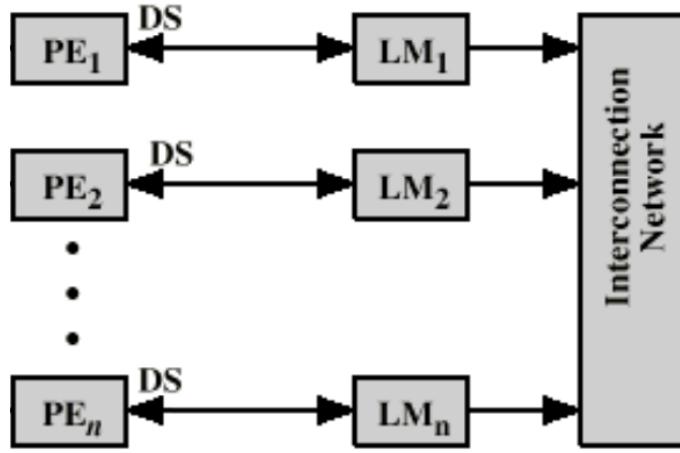
# Issues with Writes

- Multiple copies of data exist:
  - L1, L2, L3, Main Memory, Disk
- What to do on a write-hit?
  - Write-through (write immediately to memory)
  - Write-back (defer write to memory until replacement of line)
    - Need a dirty bit (line different from memory or not)

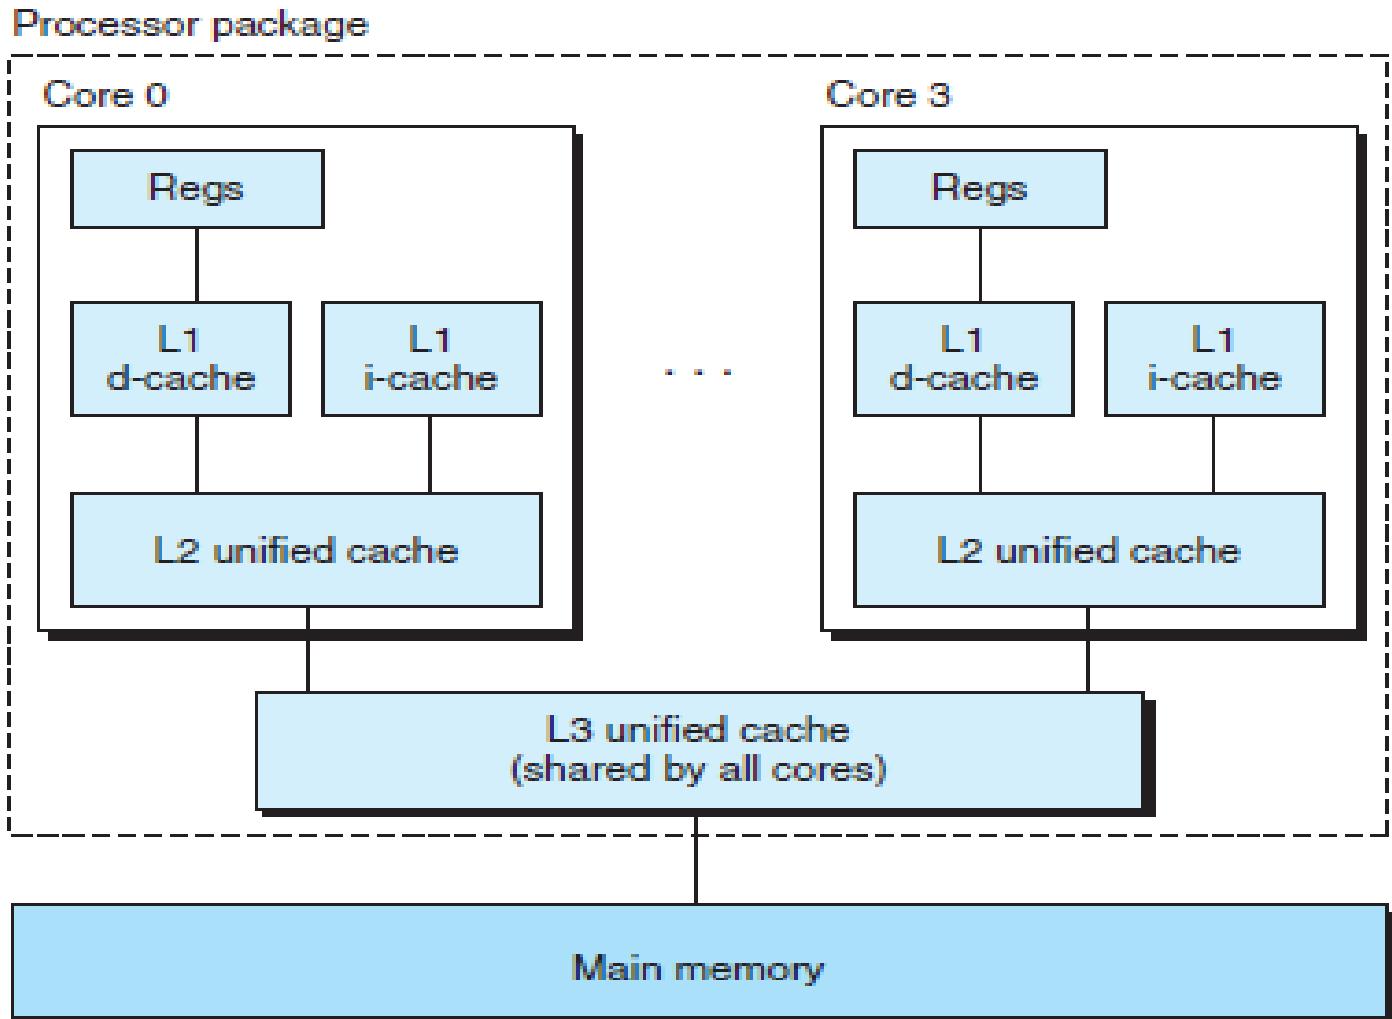


(b) Three-level cache organization

Cache and Main Memory



# Intel Core i7 Cache Hierarchy





# Intel Core i7 Cache Hierarchy

Cache type	Access time (cycles)	Cache size ( $C$ )	Assoc. ( $E$ )	Block size ( $B$ )	Sets ( $S$ )
L1 i-cache	4	32 KB	8	64 B	64
L1 d-cache	4	32 KB	8	64 B	64
L2 unified cache	11	256 KB	8	64 B	512
L3 unified cache	30–40	8 MB	16	64 B	8192

Characteristics of the Intel Core i7 cache hierarchy.



# Performance Impact of Cache Parameters

- Associativity :
  - higher associativity → more complex hardware
  - Higher Associativity → Lower miss rate
  - Higher Associativity → reduces average memory access time (AMAT)
- Cache Size
  - Larger the cache size → Lower miss rate
  - Larger the cache size → reduces average memory access time (AMAT)
- Block Size:
  - Smaller blocks do not take maximum advantage of spatial locality.



# Computer Organization and Software Systems

Contact Session 6

Dr. Lucy J. Gudino



# Revisiting Locality of reference

```
1 int sumvec(int v[N])
2 {
3     int i, sum = 0;
4
5     for (i = 0; i < N; i++)
6         sum += v[i];
7     return sum;
8 }
```

Does this function have good locality?



N=8							
Address	0	1	2	3	4	5	6
Contents	v[0]	v[1]	v[2]	v[3]	v[4]	v[5]	v[6]
Access Order	1	2	3	4	5	6	7

# Stride k – reference pattern



Byte Addressable  
memory and word  
length is 1 byte

i	Address
0	0000
1	0001
2	0002
3	0003
4	0004
5	0005
6	0006
7	0007
8	0008
9	0009
10	000A
11	000B
12	000C

Stride 1

Address difference  
**Stride = -----**  
Word Length

i	Address
0	0000
1	0001
2	0002
3	0003
4	0004
5	0005
6	0006
7	0007
8	0008
9	0009
10	000A
11	000B
12	000C

Stride 2

# Stride k – reference pattern



Byte Addressable  
memory and word  
length is 2 bytes

i	Address
0	0000
1	0002
2	0004
3	0006
4	0008
5	000A
6	000C
7	000E
8	0010
9	0012
10	0014
11	0016
12	0018

Stride 1

Address difference  
 $\text{Stride} = \frac{\text{Address difference}}{\text{Word Length}}$

i	Address
0	0000
1	0002
2	0004
3	0006
4	0008
5	000A
6	000C
7	000E
8	0010
9	0012
10	0014
11	0016
12	0018

Stride 2

# Revisiting Locality of reference

```

1 int sumarrayrows(int a[M][N])
2 {
3     int i, j, sum = 0;
4
5     for (i = 0; i < M; i++)
6         for (j = 0; j < N; j++)
7             sum += a[i][j];
8
9     return sum;
}
  
```

Does this function have good locality?

MxN	[0]	[1]	[2]
[0]	1	2	3
[1]	4	5	6

M = 2, N=3

Address(Index)	0	1	2	3	4	5
Contents	A[0][0]	A[0][1]	A[0][2]	A[1][0]	A[1][1]	A[1][2]
Access Order	1	2	3	4	5	6

# Revisiting Locality of reference

```

1 int sumarraycols(int a[M] [N])
2 {
3     int i, j, sum = 0;
4
5     for (j = 0; j < N; j++)
6         for (i = 0; i < M; i++)
7             sum += a[i][j];
8
9     return sum;
}
  
```

Does this function have good locality?

M = 2, N=3						NxM	[0]	[1]	
Address(Index)	0	1	2	3	4	5	[0]	1	2
Contents	A[0][0]	A[1][0]	A[2][0]	A[1][0]	A[1][1]	A[2][1]	[1]	3	4
Access Order	1	3	5	2	4	6	[2]	5	6



# Writing Cache Friendly Code

- Make the common case go fast
  - Focus on the inner loops of the core functions
- Minimize the misses in the inner loops
  - Repeated references to variables are good (**temporal locality**)
  - Stride-1 reference patterns are good (**spatial locality**)



# Example 1

```
int sumarrayrows(int a[4][4])
{
    int i, j, sum = 0;
    for (i = 0; i < 4; i++)
        for (j = 0; j < 4; j++)
            sum += a[i][j];
    return sum;
}
```

## Assumption:

- The cache has a block size of 4 words each, 2 cache lines
- Word size 4 bytes.
- C stores arrays in row-major order

# Example 1(Contd..)

```
int sumarrayrows(int a[4][4])
```

```
{  
    int i, j, sum = 0;  
    for (i = 0; i < 4; i++)  
        for (j = 0; j < 4; j++)  
            sum += a[i][j];  
    return sum;  
}
```



A[i][j]	J = 0	J = 1	J = 2	J = 3
i = 0	W0 M	W1 H	W2 H	W3 H
i = 1	W4 M	W5 H	W6 H	W7 H
i = 2	W8 M	W9 H	W10 H	W11 H
i = 3	W12 M	W13 H	W14 H	W15 H

a[0][0] →	W0
a[0][1] →	W1
→ a[0][2] →	W2
→ a[0][3] →	W3
<u>a[0][3]</u>	<u>W3</u>
a[1][0] →	W4
a[1][1] →	W5
a[1][2] →	W6
a[1][3] →	W7
<u>a[1][3]</u>	<u>W7</u>
<u>a[2][0]</u>	<u>W8</u>
a[2][1] →	W9
a[2][2] →	W10
a[2][3] →	W11
<u>a[2][3]</u>	<u>W11</u>
<u>a[3][0]</u>	<u>W12</u>
a[3][1] →	W13
a[3][2] →	W14
a[3][3] →	W15
<u>a[3][3]</u>	<u>W15</u>



# Example 2

```
int sumarraycols(int a[4][4])
{
    int i, j, sum = 0;
    for (j = 0; j < 4; j++)
        for (i = 0; i < 4; i++)
            sum += a[i][j];
    return sum;
}
```

## Assumption:

- The cache has a block size of 4 words each, 2 cache lines
- Word size 4 bytes.
- C stores arrays in row-major order

## Example 2(Contd..)

```
int sum_array(int a[4][4])
```

```
{
```

```
    int i, j, sum = 0;
```

```
    for (j = 0; j < 4; j++)
```

```
        for (i = 0; i < 4; i++)
```

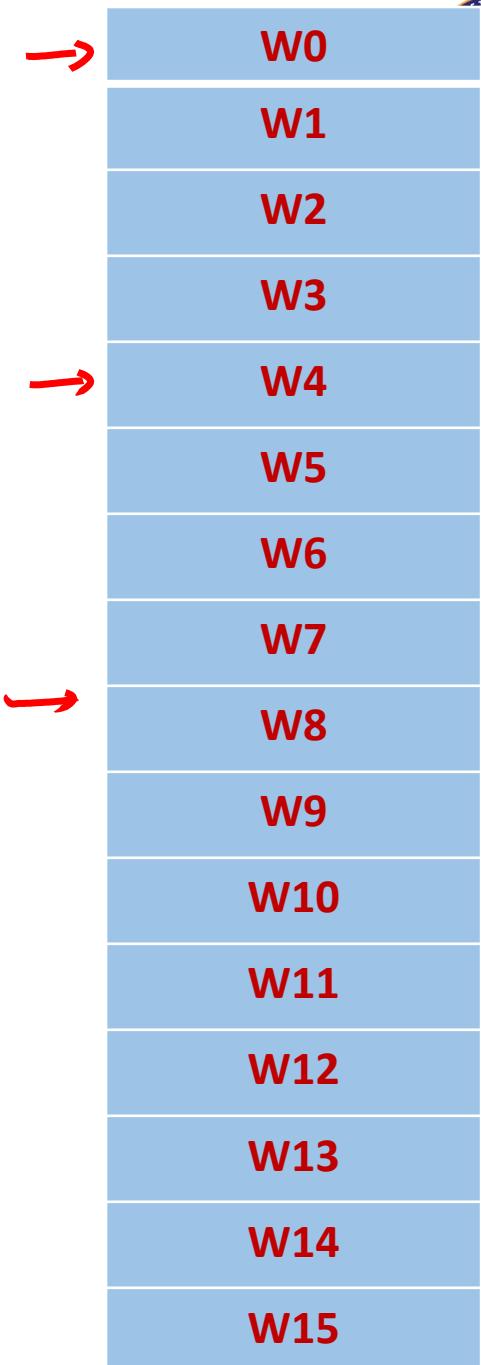
```
            sum += a[i][j];
```

```
    return sum;
```

```
}
```



A[i][j]	J = 0	J = 1	J = 2	J = 3
i = 0	w <sub>0</sub> m	w <sub>1</sub> m	w <sub>2</sub> m	w <sub>3</sub> m
i = 1	w <sub>4</sub> m	w <sub>5</sub> m	w <sub>6</sub> m	w <sub>7</sub> m
i = 2	w <sub>8</sub> m	w <sub>9</sub> m	w <sub>10</sub> m	w <sub>11</sub> m
i = 3	w <sub>12</sub> m	w <sub>13</sub> m	w <sub>14</sub> m	w <sub>15</sub> m





# Home Work – Which one is better ?

Program 1:

```
for (int i = 0; i < n; i++) {  
    z[i] = x[i] - y[i];  
    z[i] = z[i] * z[i];  
}
```

Program 2:

```
for (int i = 0; i < n; i++) {  
    z[i] = x[i] - y[i];  
}  
for (int i = 0; i < n; i++) {  
    z[i] = z[i] * z[i];  
}
```

# Today's Session



Contact Hour	List of Topic Title	Text/Ref Book/external resource
11-12	<ul style="list-style-type: none"><li>• <b>Instruction Set Architecture - CISC Vs RISC</b></li><li>• CISC Instruction Set (Intel x86 as an example)<ul style="list-style-type: none"><li>• Machine Instruction Characteristics</li><li>• Types of Operands</li><li>• Types of Operations</li><li>• Addressing Modes</li></ul></li><li>• Instruction Formats</li></ul>	T1

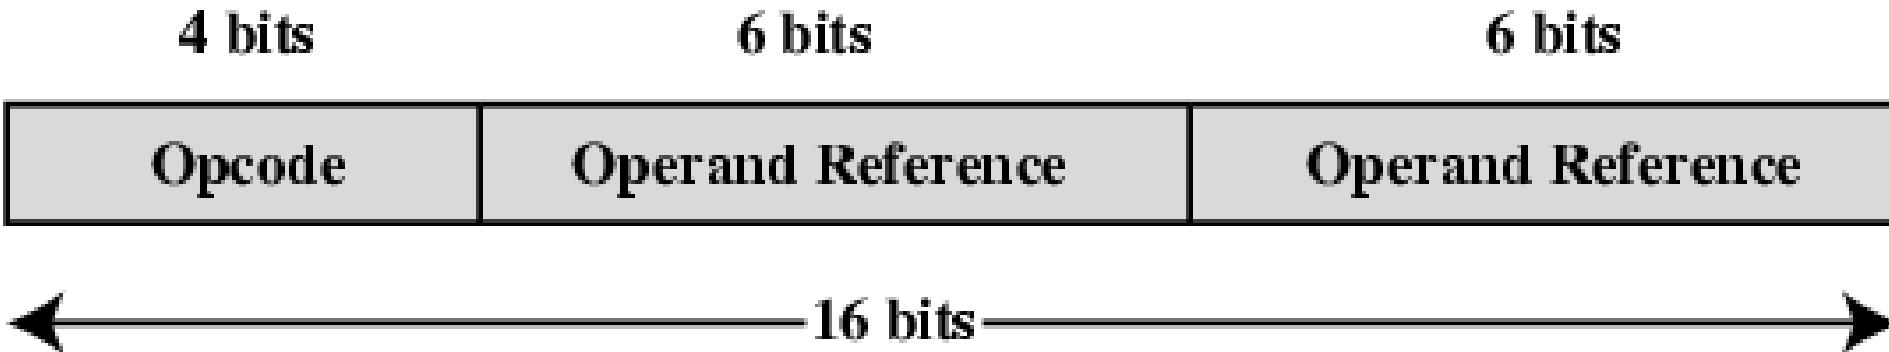


# Introduction

- What is an Instruction Set?
  - The complete collection of instructions that are understood by a CPU
- Elements of an Instruction
  - Operation code (Op code)
  - Source Operand reference
  - Result Operand reference
  - Next Instruction Reference
- Source and Destination Operands can be found in four areas
  - Main memory (or virtual memory or cache)
  - CPU register
  - I/O device
  - Immediate



# Simple Instruction Format



- During instruction execution, an instruction is read into an instruction register (IR) in the processor.
- The processor must be able to extract the data from the various instruction fields to perform the required operation.
- Opcodes are represented by abbreviations, called **mnemonics**  
Example: ADD AX, BX → Add instruction



# Instruction Types

- Data processing : Arithmetic and logic instructions
- Data storage (main memory) : Movement of data into or out of register and or memory locations
- Data movement (I/O) : I/O instructions
- Program flow control : Test and branch instruction



# Number of Addresses (1/2)

- 3 addresses
  - Result, Operand 1, Operand 2
  - $c = a + b$ ; add  $c$ ,  $a$ ,  $b$
  - May be a forth - next instruction (usually implicit)
  - Needs very long words to hold everything
- 2 addresses
  - One address doubles as operand and result
  - $a = a + b$  : add  $a$ ,  $b$
  - Reduces length of instruction
  - The original value of  $a$  is lost.



# Number of Addresses (2/2)

- 1 address
  - Implicit second address
  - Usually a register (accumulator)
  - Common on early machines
- 0 (zero) addresses
  - All addresses implicit
  - Uses a stack
  - e.g.  $c = a + b$

push a

push b

add

pop c

# Example

$$\text{Execute } Y = \frac{A - B}{C + (D \times E)}$$



<u>Instruction</u>	<u>Comment</u>
SUB Y, A, B	$Y \leftarrow A - B$
MPY T, D, E	$T \leftarrow D \times E$
ADD T, T, C	$T \leftarrow T + C$
DIV Y, Y, T	$Y \leftarrow Y \div T$

(a) Three-address instructions

<u>Instruction</u>	<u>Comment</u>
MOVE Y, A	$Y \leftarrow A$
SUB Y, B	$Y \leftarrow Y - B$
MOVE T, D	$T \leftarrow D$
MPY T, E	$T \leftarrow T \times E$
ADD T, C	$T \leftarrow T + C$
DIV Y, T	$Y \leftarrow Y \div T$

(b) Two-address instructions

<u>Instruction</u>	<u>Comment</u>
LOAD D	$AC \leftarrow D$
MPY E	$AC \leftarrow AC \times E$
ADD C	$AC \leftarrow AC + C$
STOR Y	$Y \leftarrow AC$
LOAD A	$AC \leftarrow A$
SUB B	$AC \leftarrow AC - B$
DIV Y	$AC \leftarrow AC \div Y$
STOR Y	$Y \leftarrow AC$

(c) One-address instructions



# How Many Addresses

- Fewer addresses
  - More Primitive instructions, shorter length instructions
  - Less complex instructions, hence requires less complex hardware
  - More instructions per program
    - Longer programs
    - More complex programs
    - Longer execution time
- Multiple address instructions
  - Lengthy instructions
  - More registers
    - Inter-register operations are quicker
  - Fewer instructions per program

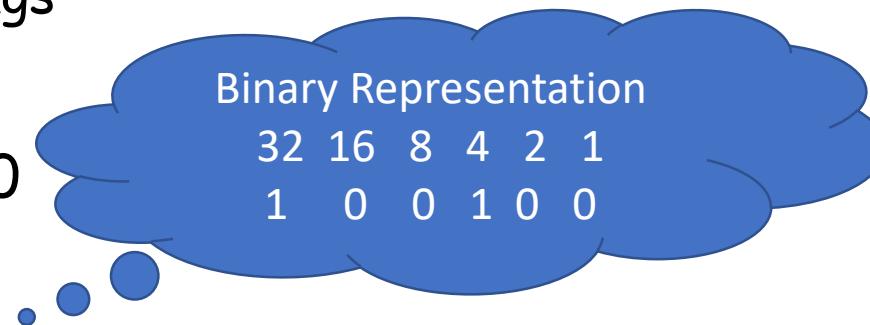


# Instruction set Design Decisions

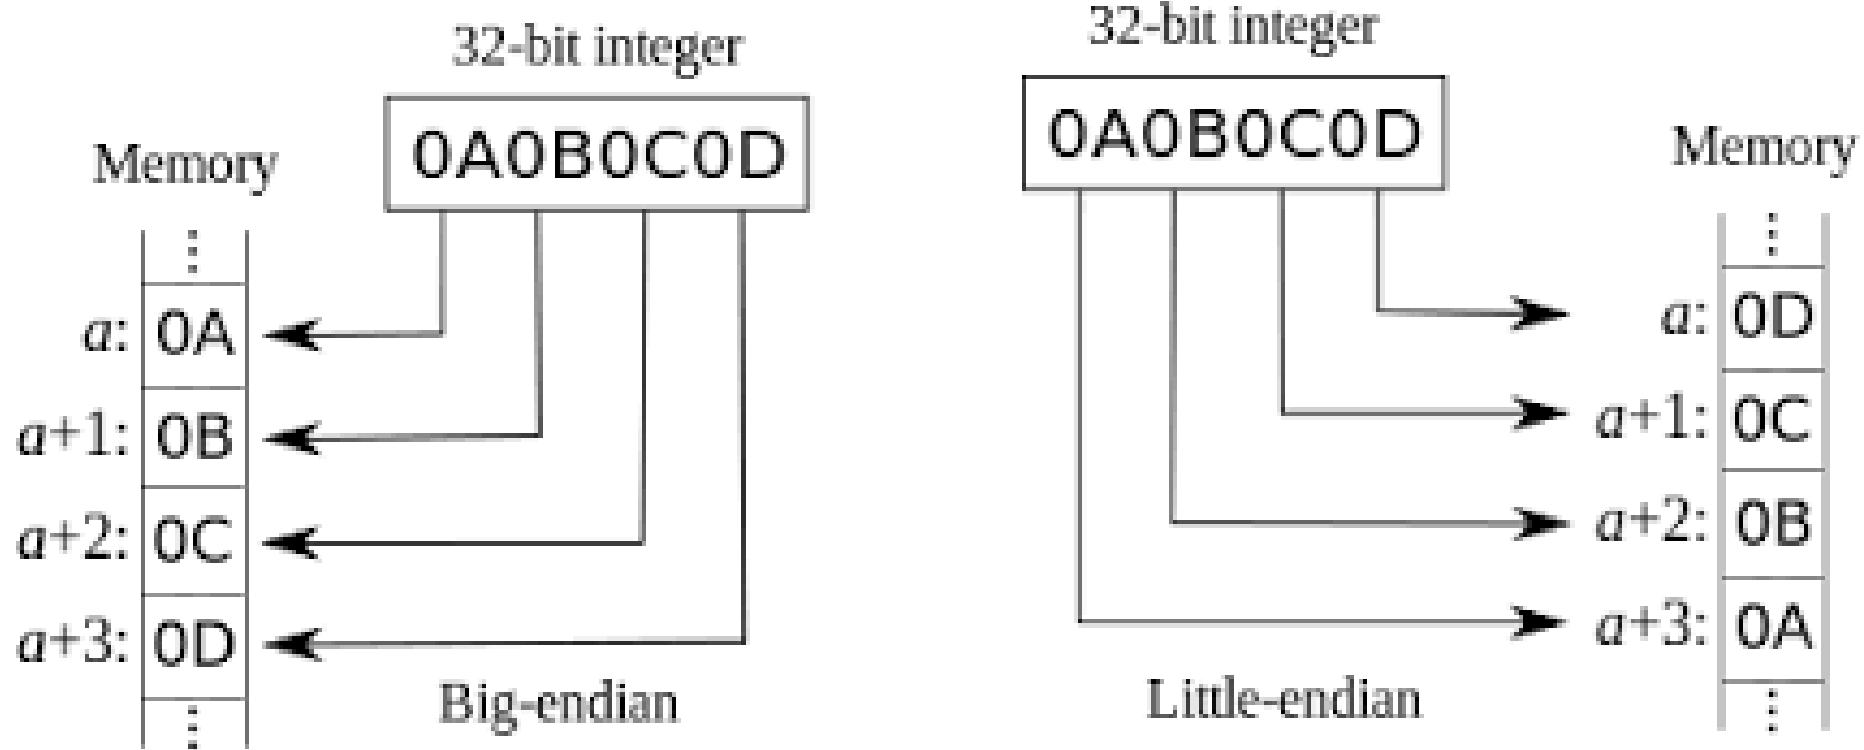
- Operation repertoire
  - How many ops?
  - What can they do?
  - How complex are they?
- Data types
- Instruction formats
  - Length of op code field
  - Number of addresses
- Registers
  - Number of CPU registers available
  - Which operations can be performed on which registers?
- Addressing modes

# Types of Operand

- Machine instructions operate on data
- General categories of data
  - Addresses
  - Numbers
    - Binary integer or binary fixed point, floating point, decimal
  - Characters
    - ASCII etc.
  - Logical Data
    - Bits or flags
- Packed Decimal
  - 36 : 0011 0110



# Byte Ordering



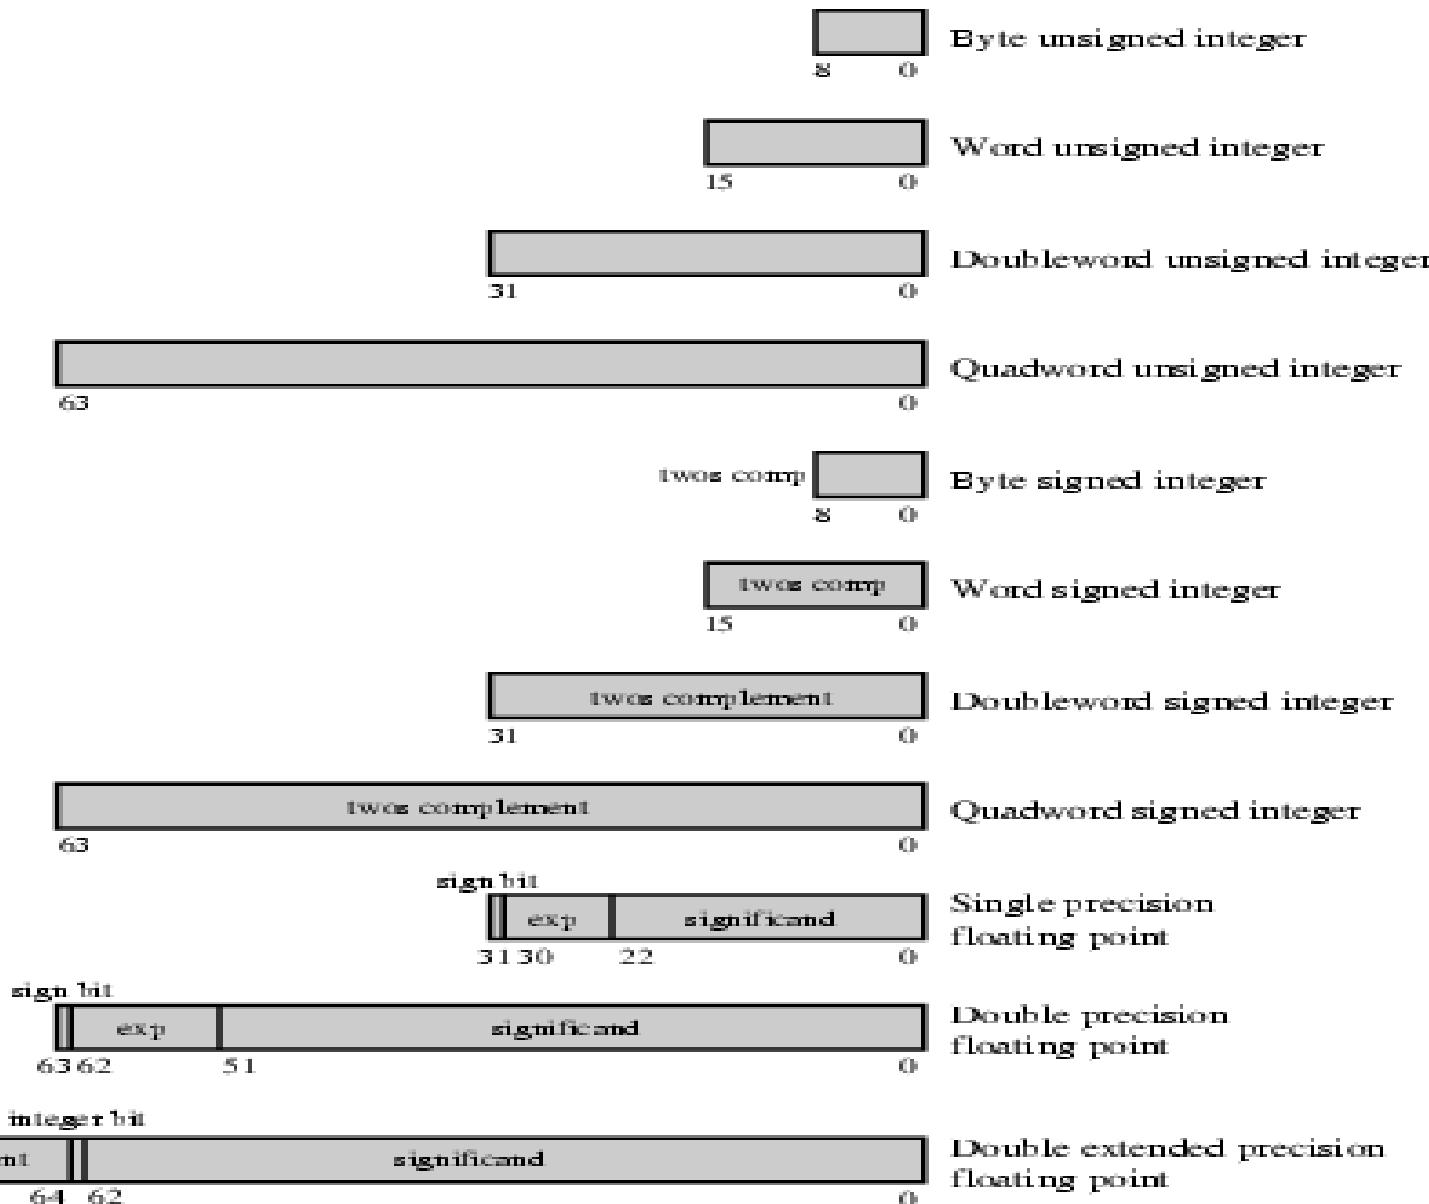


# x86 Data Types

- General - Byte, Word, double word, quadword, double quad word - arbitrary binary contents
- Integer - signed binary using two's complement representation
- Ordinal - unsigned integer
- Unpacked BCD - One digit per byte
- Packed BCD - 2 digits per byte
- Near Pointer - 16/32 bit offset within segment
- Far pointer - 16/32 bit offset outside segment
- Bit field : A contiguous sequence of bits in which the position of each bit is considered as an independent unit.
- Bit and Byte String
- Floating Point

Data Type	Description
General	Byte, word (16 bits), doubleword (32 bits), quadword (64 bits), and double quadword (128 bits) locations with arbitrary binary contents.
Integer	A signed binary value contained in a byte, word, or doubleword, using twos complement representation.
Ordinal	An unsigned integer contained in a byte, word, or doubleword.
Unpacked binary coded decimal (BCD)	A representation of a BCD digit in the range 0 through 9, with one digit in each byte.
Packed BCD	Packed byte representation of two BCD digits; value in the range 0 to 99.
Near pointer	A 16-bit, 32-bit, or 64-bit effective address that represents the offset within a segment. Used for all pointers in a nonsegmented memory and for references within a segment in a segmented memory.
Far pointer	A logical address consisting of a 16-bit segment selector and an offset of 16, 32, or 64 bits. Far pointers are used for memory references in a segmented memory model where the identity of a segment being accessed must be specified explicitly.
Bit field	A contiguous sequence of bits in which the position of each bit is considered as an independent unit. A bit string can begin at any bit position of any byte and can contain up to 32 bits.
Bit string	A contiguous sequence of bits, containing from zero to $2^{32} - 1$ bits.
Byte string	A contiguous sequence of bytes, words, or doublewords, containing from zero to $2^{32} - 1$ bytes.
Floating point	See Figure 10.4.
Packed SIMD (single instruction, multiple data)	Packed 64-bit and 128-bit data types

# x86 Numeric Data Formats





# Types of Operation

- Data Transfer
- Arithmetic
- Logical
- Conversion
- I/O
- System Control
- Transfer of Control

# Data Transfer



- Specify
  - Source
  - Destination
  - Amount of data
- Action:
  1. Calculate the memory address, based on the address mode
  2. If the address refers to virtual memory, translate from virtual to real memory address.
  3. Determine whether the addressed item is in cache.
  4. If not, issue a command to the memory



# Arithmetic

- Add, Subtract, Multiply, Divide
- May include
  - Absolute value ( $|a|$ )
  - Increment ( $a++$ )
  - Decrement ( $a--$ )
  - Negate ( $-a$ )
- Signed Integer



# Logical

- Bitwise operations
- AND, OR, NOT

## Basic Logical Operations

P	Q	NOT P	P AND Q	P OR Q	P XOR Q	P = Q
0	0	1	0	0	0	1
0	1	1	0	1	1	0
1	0	0	0	1	1	0
1	1	0	1	1	0	1

# Shift and Rotate Operations



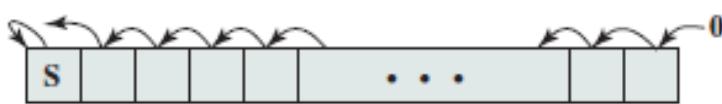
(a) Logical right shift



(b) Logical left shift



(c) Arithmetic right shift



(d) Arithmetic left shift



(e) Right rotate



(f) Left rotate

Input	Operation	Output
10101101	Logical right shift (3 bits)	10101101=> 00010101
10101101	Logical left shift (3 bits)	10101101=> <b>01101000</b>
10101101	Arithmetic right shift (3 bits)	10101101=> 11110101
10101101	Arithmetic left shift (3 bits)	10101101=> <b>11101000</b>
10101101	Right rotate (3 bits)	10101101=> <b>10110101</b>
10101101	Left rotate (3 bits)	10101101=> <b>01101101</b>

# Conversion



- E.g. Binary to Decimal



# Input/Output

- May be specific instructions (I/O-Mapped I/O)
- May be done using data movement instructions (memory mapped)
- May be done by a separate controller (DMA)



# Systems Control

- Privileged instructions
- CPU needs to be in specific state
  - User Mode
  - Kernel mode
- For operating systems use



# Transfer of Control

- Jump / Branch (Unconditional / Conditional)
  - e.g. jump to x if result is zero
- Skip (Unconditional / Conditional)
  - skip (unconditional) : Increment to skip next instruction
  - e.g. increment and skip if zero

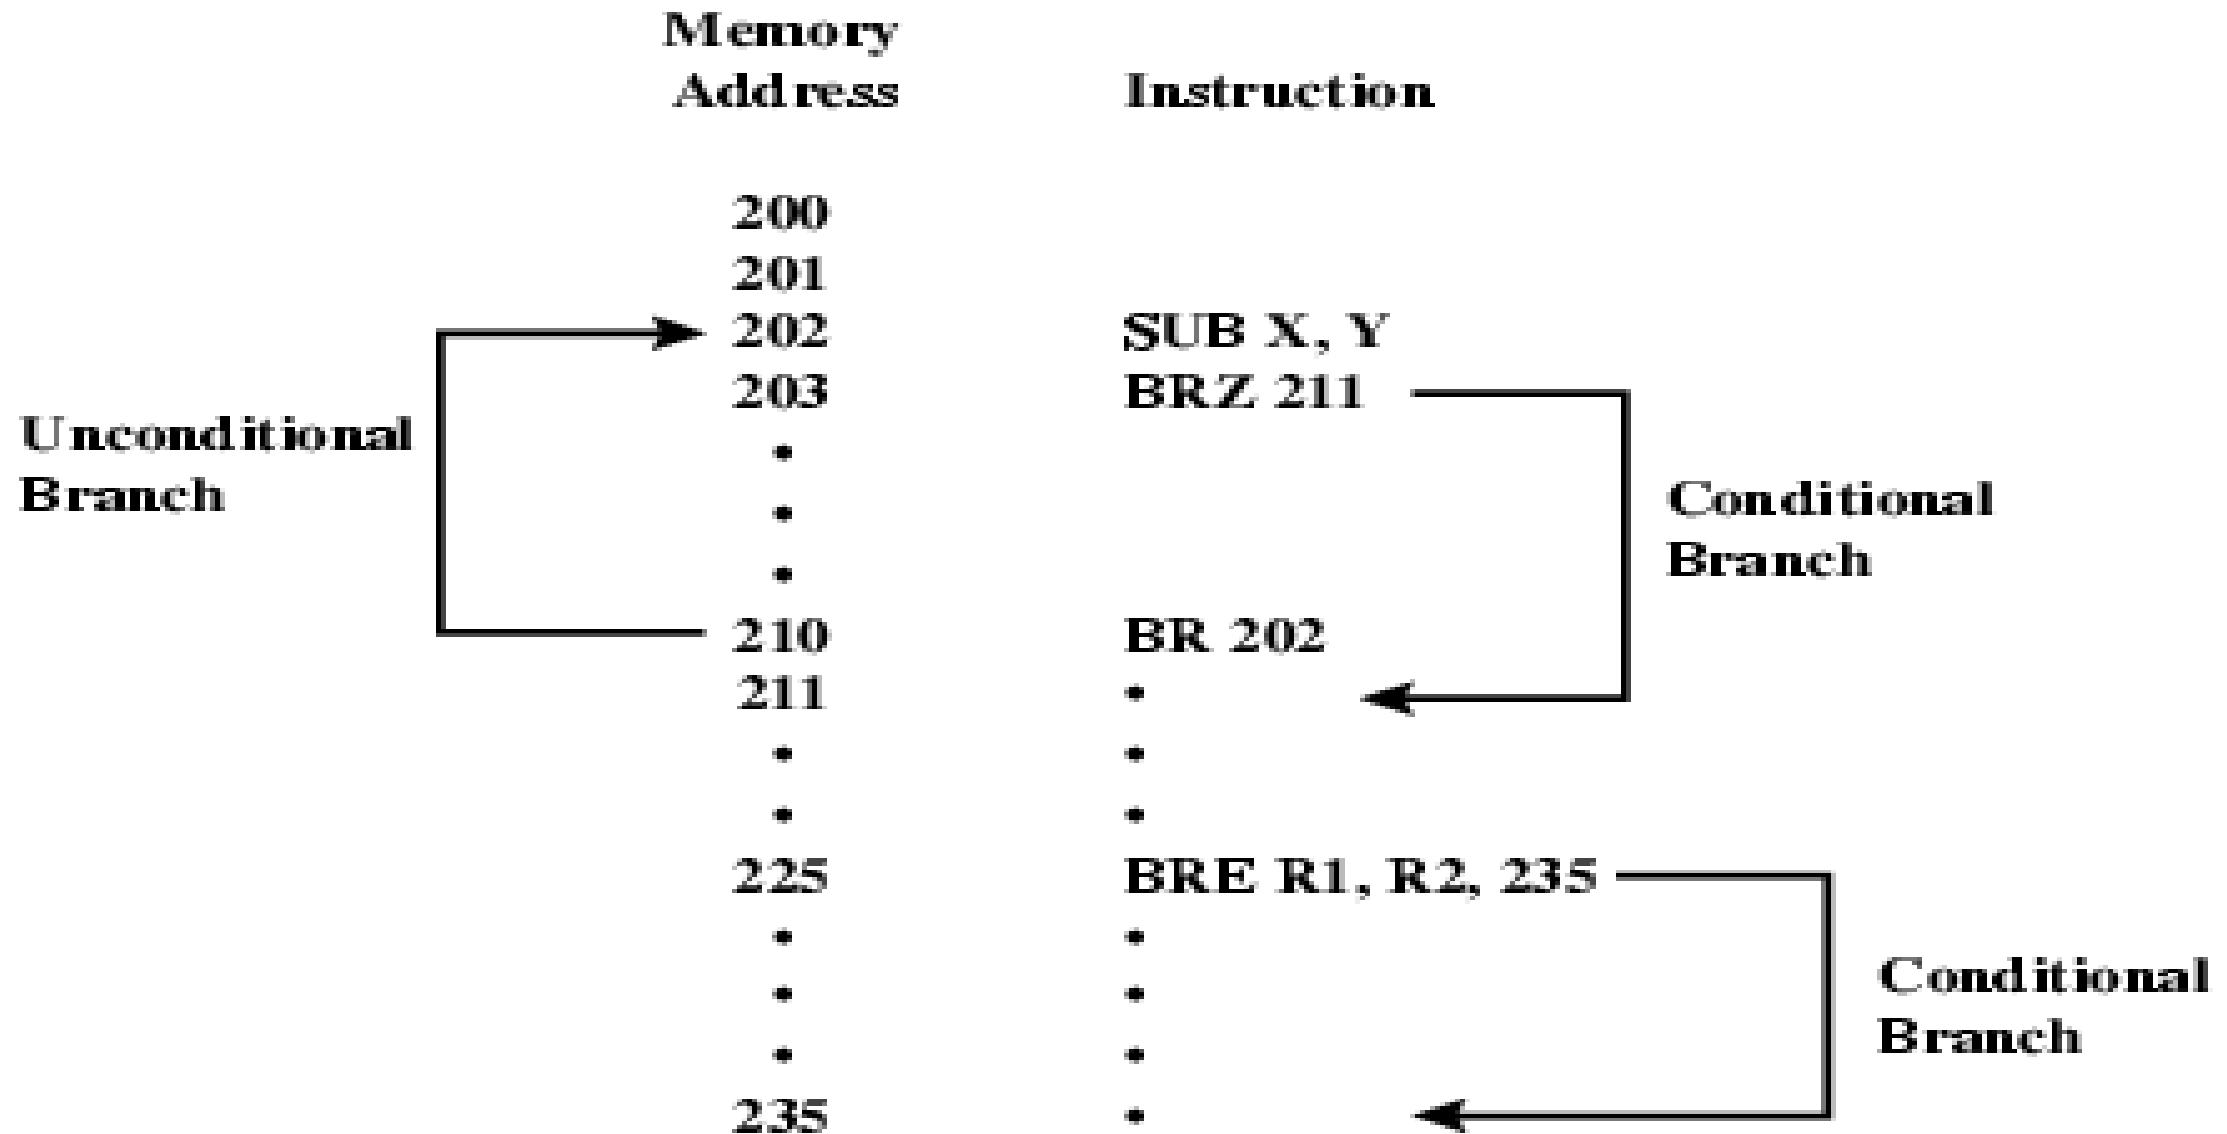
ISZ Register1

Branch xxxx

ADD A



# Branch / Jump Instruction

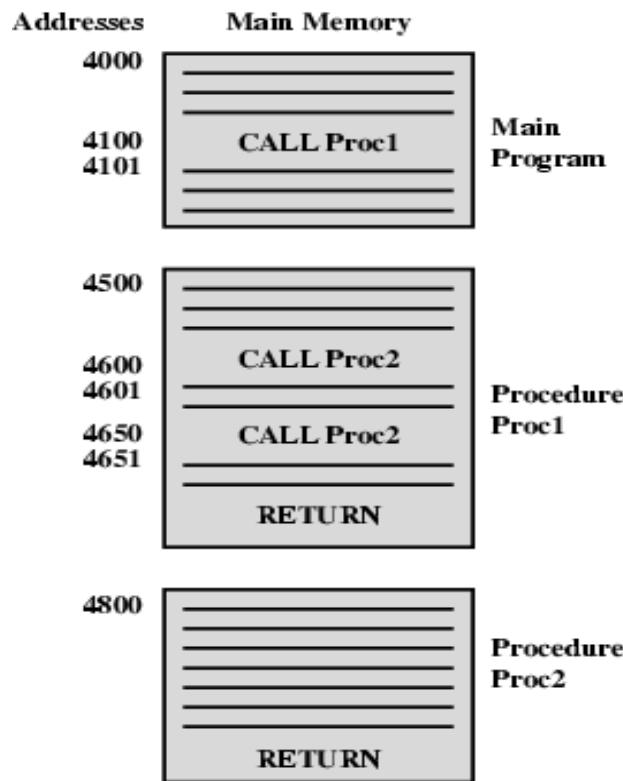




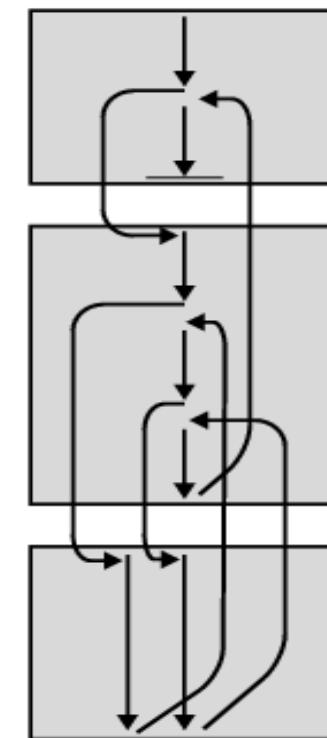
# Transfer of Control

- Jump / Branch (Unconditional / Conditional)
  - e.g. jump to x if result is zero
- Skip (Unconditional / Conditional)
  - skip (unconditional) : Increment to skip next instruction
  - e.g. increment and skip if zero
- ISZ Register1  
Branch xxxx  
ADD A
- Subroutine call
- interrupt call

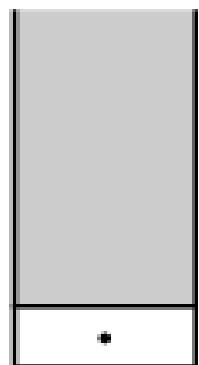
# Use of Stack



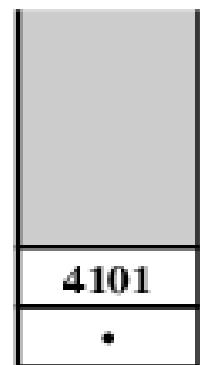
(a) Calls and returns



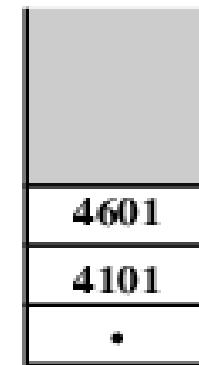
(b) Execution sequence



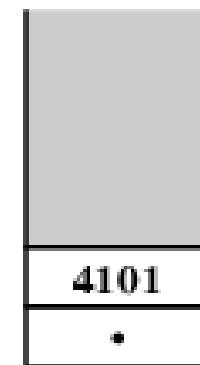
(a) Initial stack contents



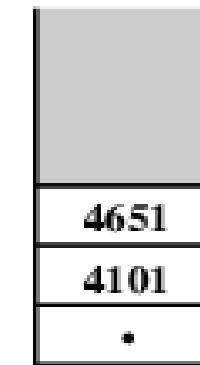
(b) After CALL Proc1



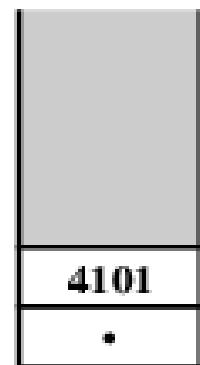
(c) Initial CALL Proc2



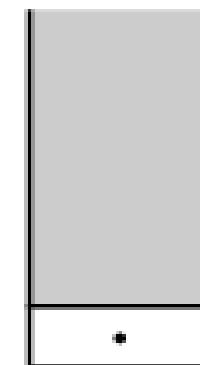
(d) After RETURN



(e) After CALL Proc2



(f) After RETURN



(g) After RETURN



# Computer Organization and Software Systems

Contact Session 7

Dr. Lucy J. Gudino





# Addressing Modes

- Addressing modes refers to the way in which the operand of an instruction is specified
- Types:
  - Immediate
  - Direct
  - Indirect
  - Register
  - Register Indirect
  - Displacement (Indexed)
  - Stack

# Immediate Addressing

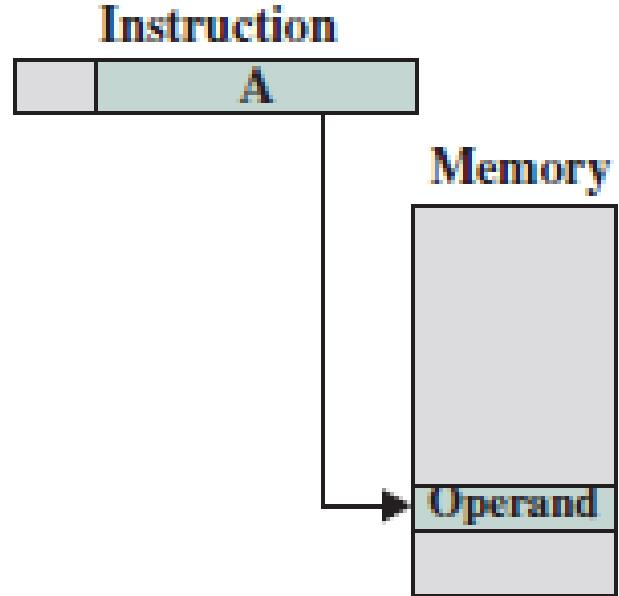
- Operand is specified in the instruction itself
- e.g. ADD #5
  - Add 5 to contents of accumulator
  - 5 is operand
- No memory reference to fetch data
- Fast
- Limited range

Instruction



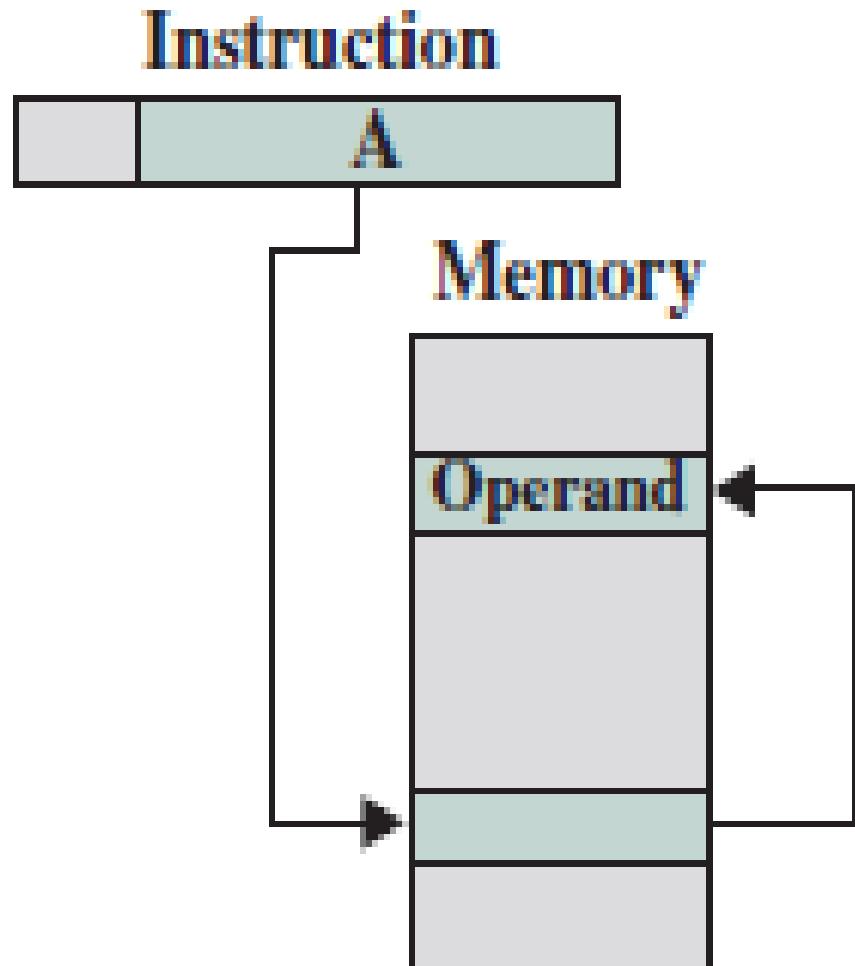
# Direct Addressing

- Address of the operand is specified in the instruction
- Effective address (EA) = address field (A)
- e.g. ADD A
  - Add contents of memory cell whose address is A to accumulator
  - Look in memory at address A for operand
- Single memory reference to access data
- No additional calculations to work out effective address
- Limited address space



# Indirect Addressing

- Memory cell pointed to by address field of the instruction contains the address of (pointer to) the operand
- $EA = (A)$ 
  - Look in A, find address and look there for operand
- e.g. ADD (A)
  - Add contents of cell pointed to by contents of A to accumulator



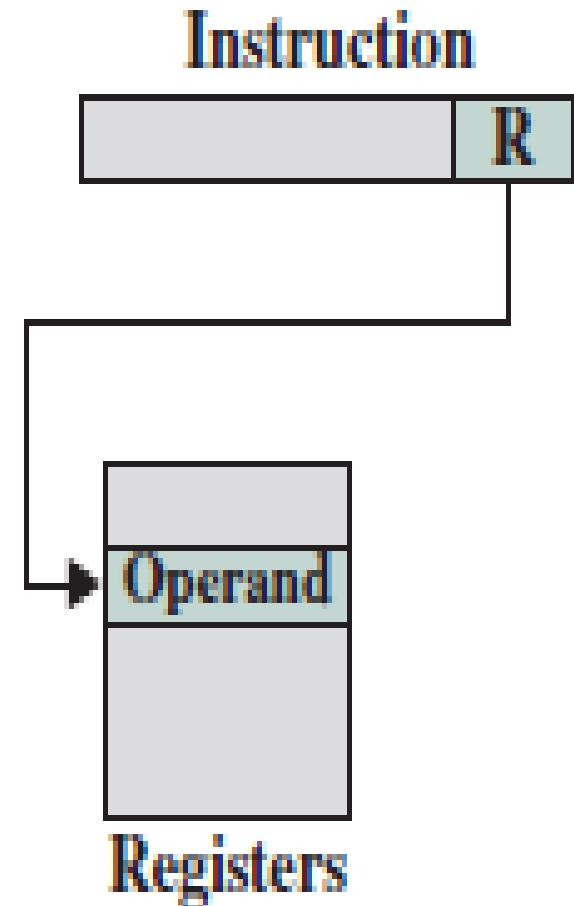


# Indirect Addressing...

- Large address space
- $2^n$  where n = word length
- May be nested, multilevel, cascaded
  - e.g. EA = (((A)))
- Multiple memory accesses to find operand
- Slower

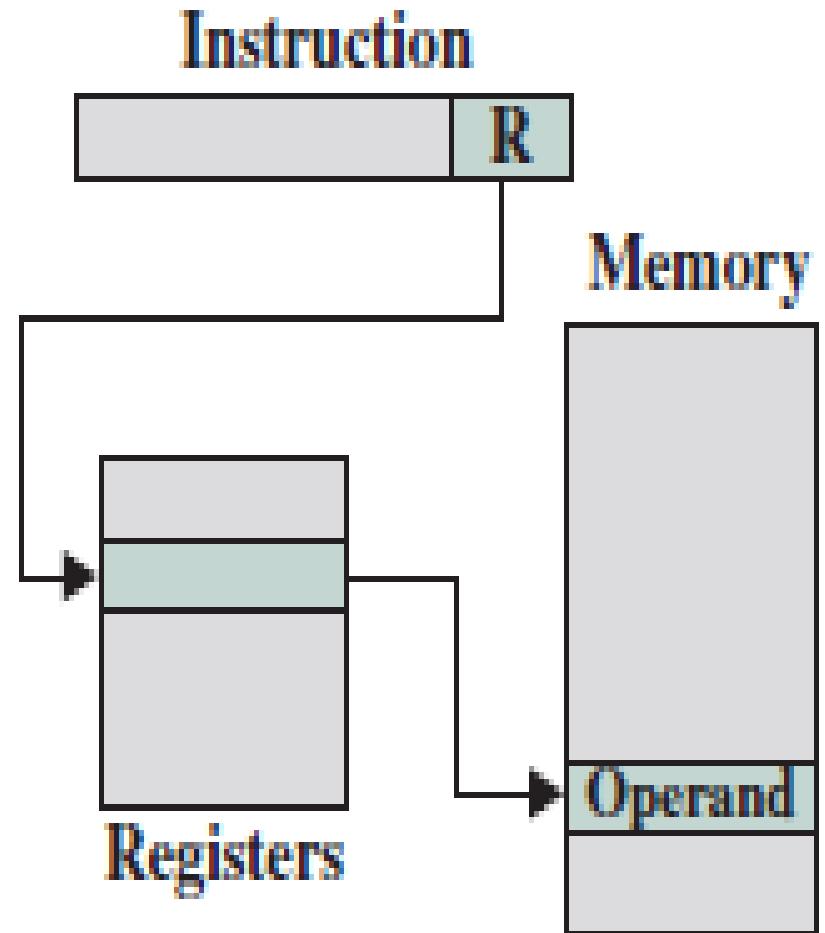
# Register Addressing

- Operand is held in register named in address filed
- $EA = R$
- Limited number of registers
- Very small address field needed
  - Shorter instructions
  - Faster instruction fetch
- No memory access hence Very fast execution but very limited address space
- Multiple registers helps in improving performance
  - Requires good assembly programming or compiler writing
  - C programming : `register int a;`



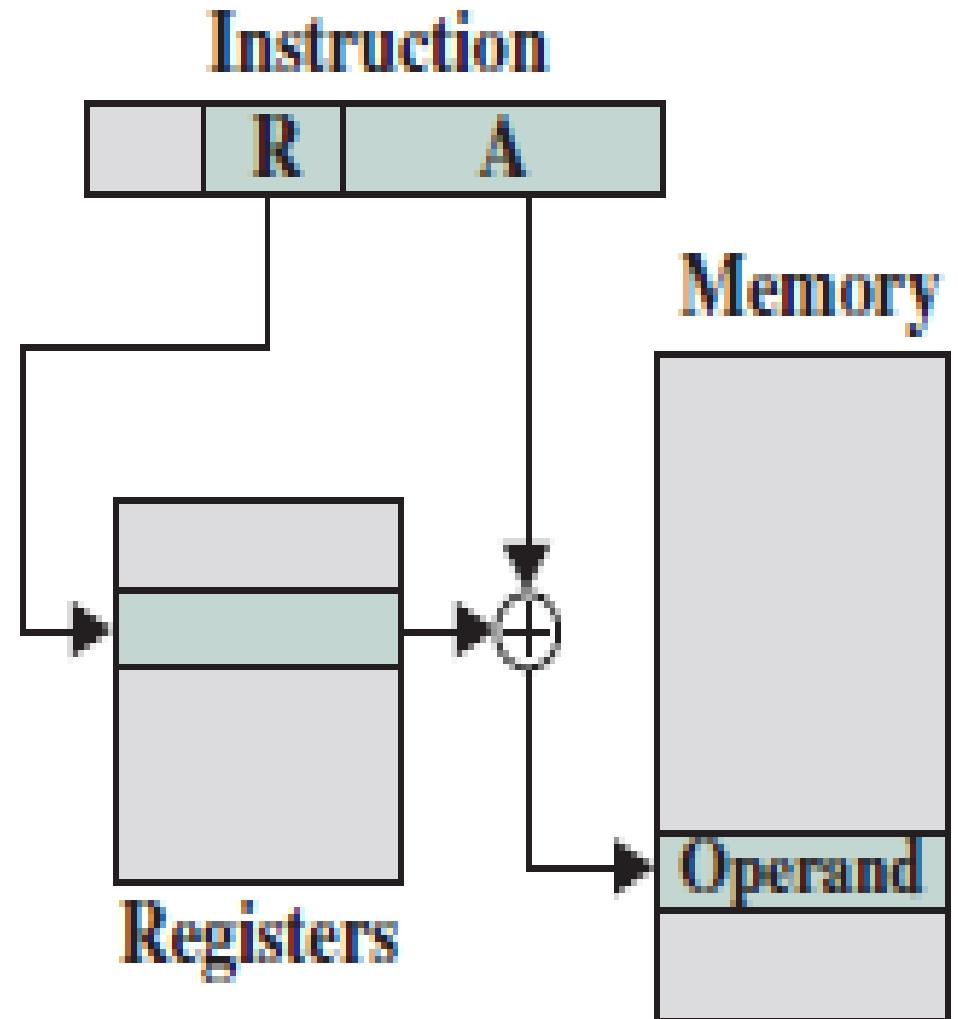
# Register Indirect Addressing

- Similar to indirect addressing
- $EA = (R)$
- Operand is in memory cell pointed to by contents of register R
- Large address space ( $2^n$ )
- One memory access compared indirect addressing



# Displacement Addressing

- $EA = A + (R)$
- Address field hold two values
  - $A$  = base value
  - $R$  = register that holds displacement
  - or vice versa
- Three variants:
  - Relative addressing
  - Base register addressing
  - Indexing





# Relative Addressing

- Also known as PC relative addressing
- A version of displacement addressing
- R = Program counter, PC
- $EA = A + (PC)$
- Relative addressing exploits the concept of locality
  - If most memory references are relatively near to the instruction being executed, then the use of relative addressing saves address bits in the instruction.



# Base-Register Addressing

- The referenced register "R" contains a main memory address
- address field contains a displacement A
- R may be explicit or implicit
- e.g. segment registers in 80x86

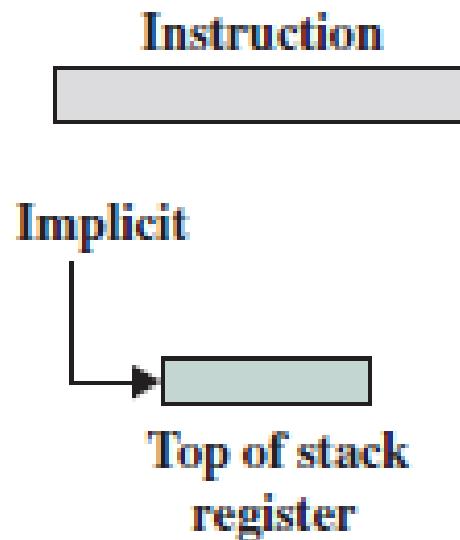


# Indexed Addressing

- The address field references a main memory address A
- The referenced register R contains a positive displacement from that address.
- $EA = A + R$
- Good for accessing arrays
  - $EA = A + R$
  - $R++$

# Stack Addressing

- Operand is (implicitly) on top of stack
- e.g.
  - ADD    Pop top two items from stack and add, push the result on stack top



# Instruction Formats



- Layout of bits in an instruction
  - Includes opcode
  - Includes (implicit or explicit) operand(s)
  - Usually more than one instruction format in an instruction set

# Instruction Length



- Affected by and affects:
  - - Memory size
  - - Memory organization
  - - Bus structure
  - - CPU complexity
  - - CPU speed
- Trade off between powerful instruction repertoire and saving space

# Allocation of Bits



- Number of addressing modes
- Number of operands
- Register versus memory
- Number of register sets
- Address range
- Address granularity

# CISC Vs RISC



#	RISC	CISC
1.	Reduced Instruction Set Computer.	Complex Instruction Set Computer.
2.	Fixed length instructions	Variable length instructions
3.	Instructions are executed in one clock cycle.	Takes one or more clock cycle
4.	Relatively simple to design.	Complex to design.
5.	Fewer, Simple addressing modes	Many addressing modes
6.	Hardwired Control Unit	Microprogrammed Control Unit
7.	Harvard Architecture	Von-Neumann Architecture
8.	Examples: SPARC, POWER PC.	Examples: Intel architecture(x86 and x64) AMD.



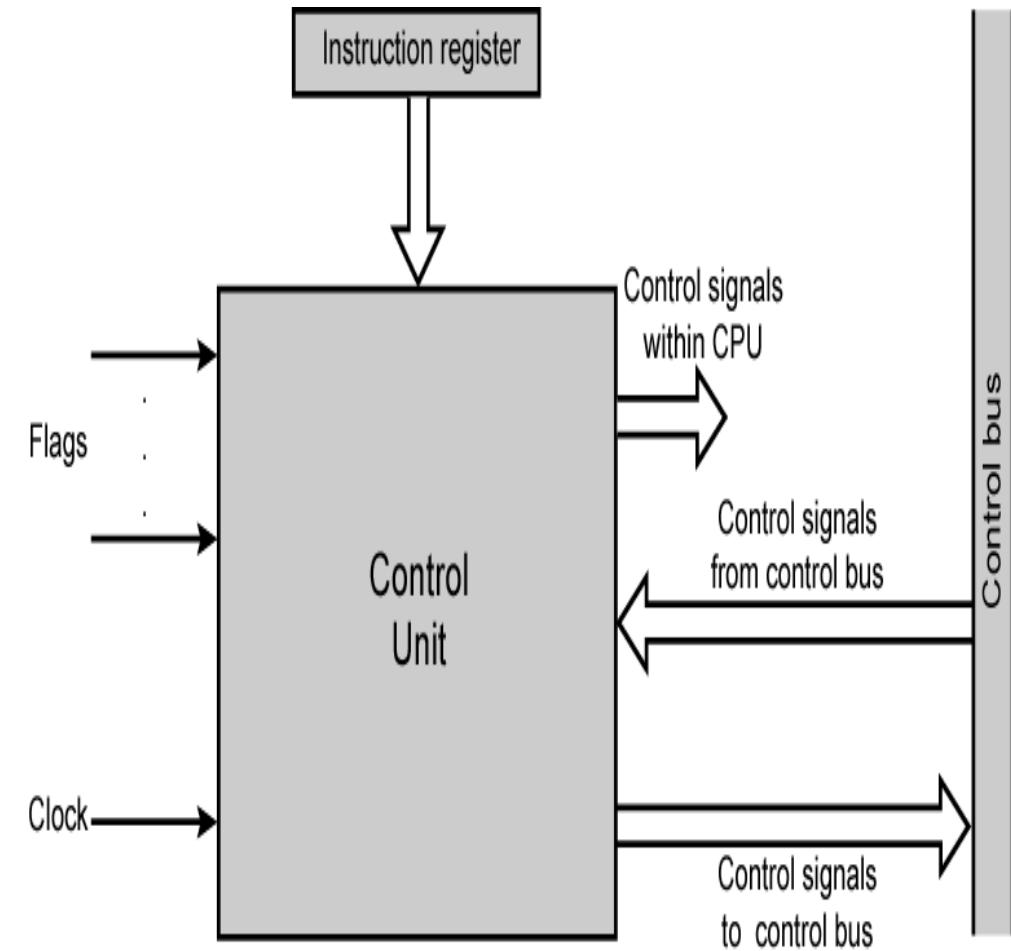
# Control Unit implementation

Hardwired control unit (RISC)

Microprogrammed control unit(CISC)

# Hardwired Implementation (1)

- Control unit inputs
  - Flags and control bus
    - Each bit means something
  - Instruction register
    - Op-code causes different control signals for each different instruction
    - Unique logic for each op-code
  - Clock
  - Time efficient





# Problems With Hardwired Designs

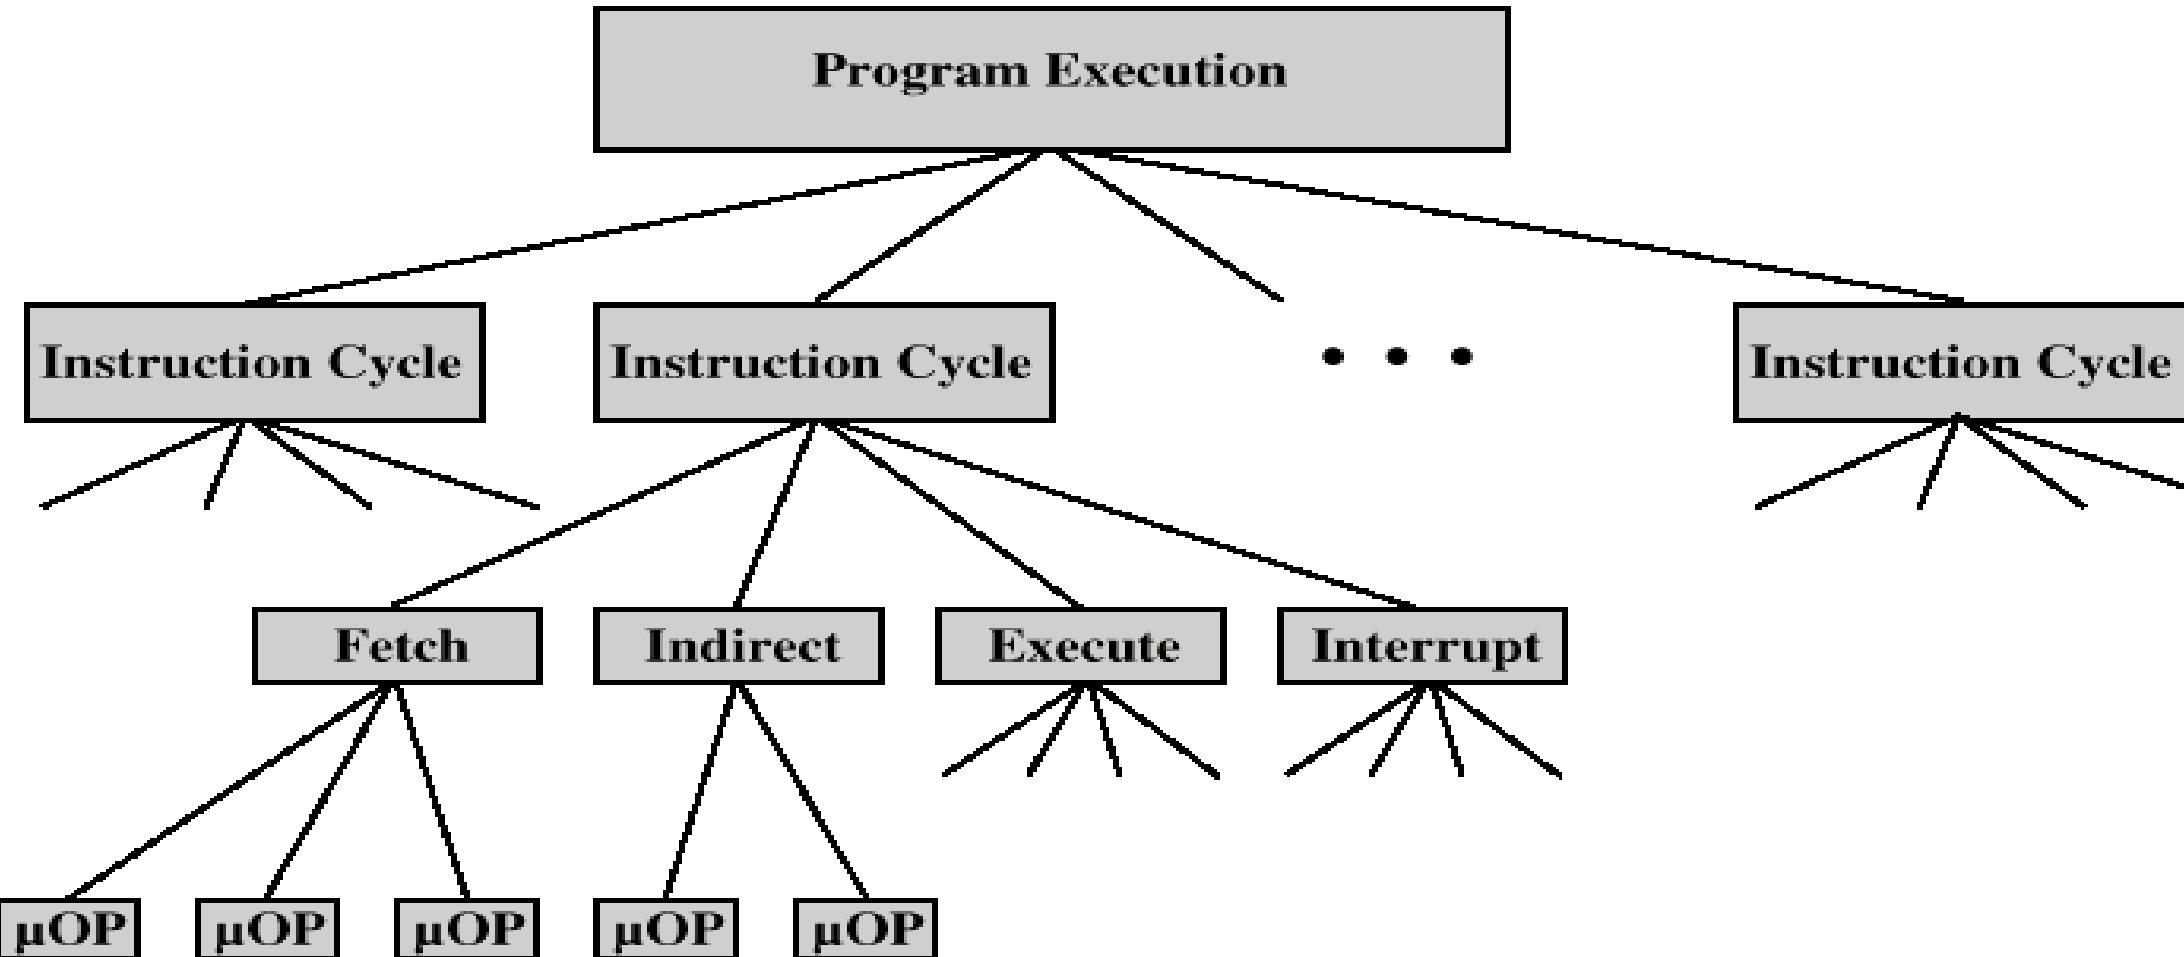
- Complex sequencing & micro-operation logic
- Difficult to design and test
- Inflexible design
- Difficult to add new instructions



# Microprogrammed Control Unit

- A computer executes a program
- Fetch/execute cycle
- Each cycle has a number of steps
  - Called micro-operations
  - Each step does very little
  - Atomic operation of CPU

# Constituent Elements of Program Execution



# Example: Fetch Sequence (ADD R1,X )

t1: MAR  $\leftarrow$  (PC)  
t2: MBR  $\leftarrow$  (memory)  
    PC  $\leftarrow$  (PC) +1  
t3: IR  $\leftarrow$  (MBR)

OR

t1: MAR  $\leftarrow$  (PC)  
t2: MBR  $\leftarrow$  (memory)  
t3: PC  $\leftarrow$  (PC) +1  
    IR  $\leftarrow$  (MBR)

# Example: Execute Cycle (ADD R1, X)

- Different for each instruction
- e.g. ADD R1,X - add the contents of location X to Register 1 , result in R1

t1: MAR  $\leftarrow$  (IR<sub>address</sub>)

t2: MBR  $\leftarrow$  (memory)

t3: R1  $\leftarrow$  R1 + (MBR)

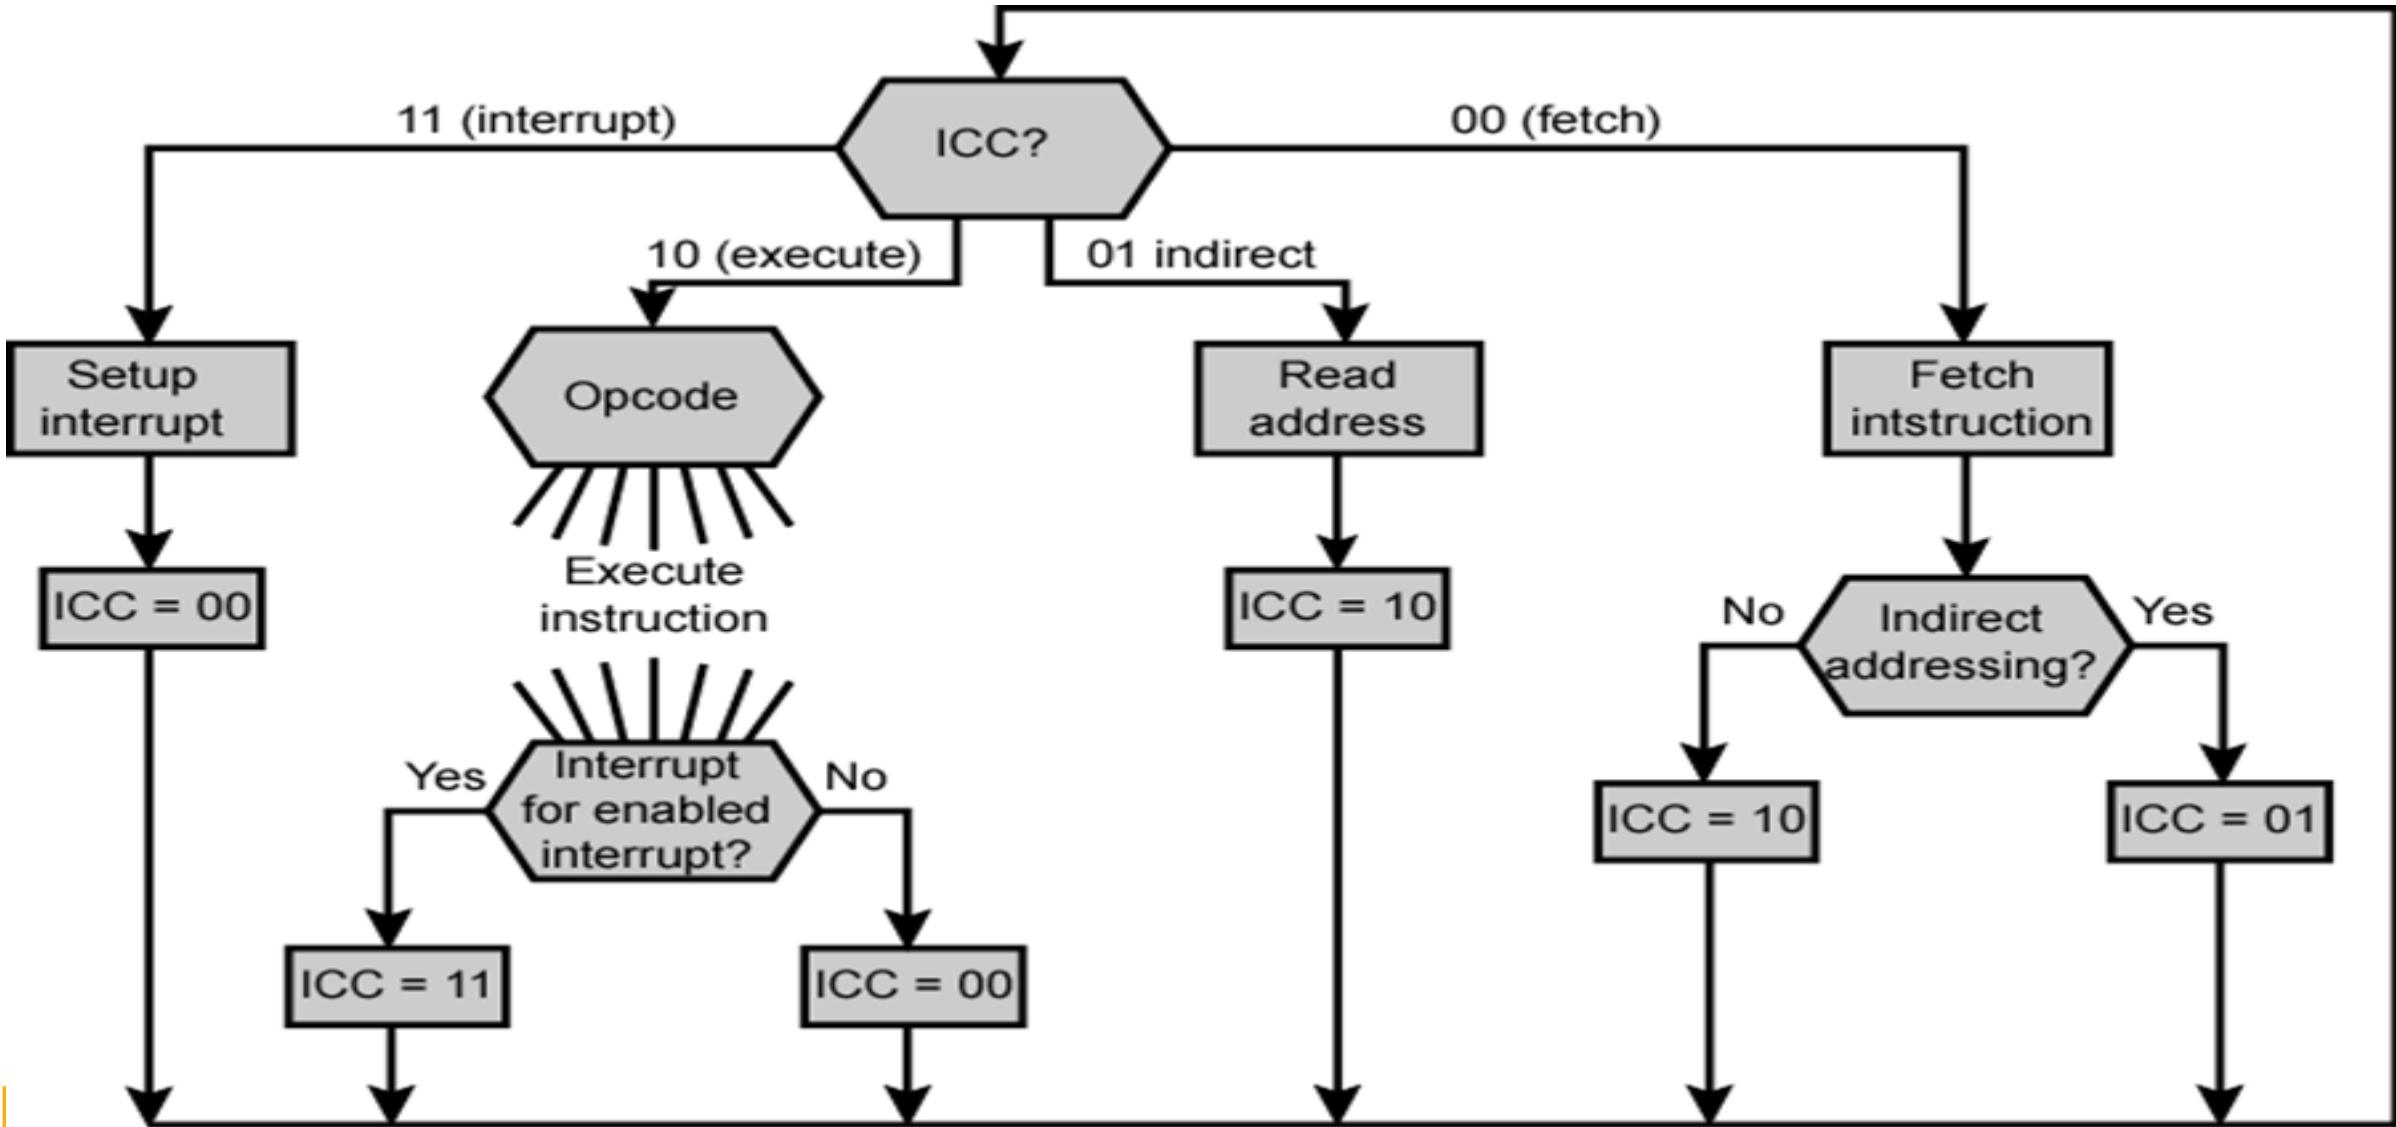
- Note no overlap of micro-operations



# Instruction Cycle

- Each phase decomposed into sequence of elementary micro-operations
- E.g. fetch, indirect, and interrupt cycles
- Execute cycle
  - One sequence of micro-operations for each opcode
- Need to tie sequences together
- Assume new 2-bit register
  - Instruction cycle code (ICC) designates which part of cycle processor is in
    - 00: Fetch
    - 01: Indirect
    - 10: Execute
    - 11: Interrupt

# Flowchart for Instruction Cycle



# Functions of Control Unit



- The control unit performs two basic tasks:
  - Sequencing
    - Causing the CPU to step through a series of micro-operations
  - Execution
    - Causing the performance of each micro-op
- This is done using Control Signals

Example: ADD R1, X

t1: MAR  $\leftarrow$  (PC)  
t2: MBR  $\leftarrow$  (memory)  
    PC  $\leftarrow$  (PC) +1  
t3: IR  $\leftarrow$  (MBR)

t4: MAR  $\leftarrow$  (IR<sub>address</sub>)  
t5: MBR  $\leftarrow$  (memory)  
t6: R1  $\leftarrow$  R1 + (MBR)



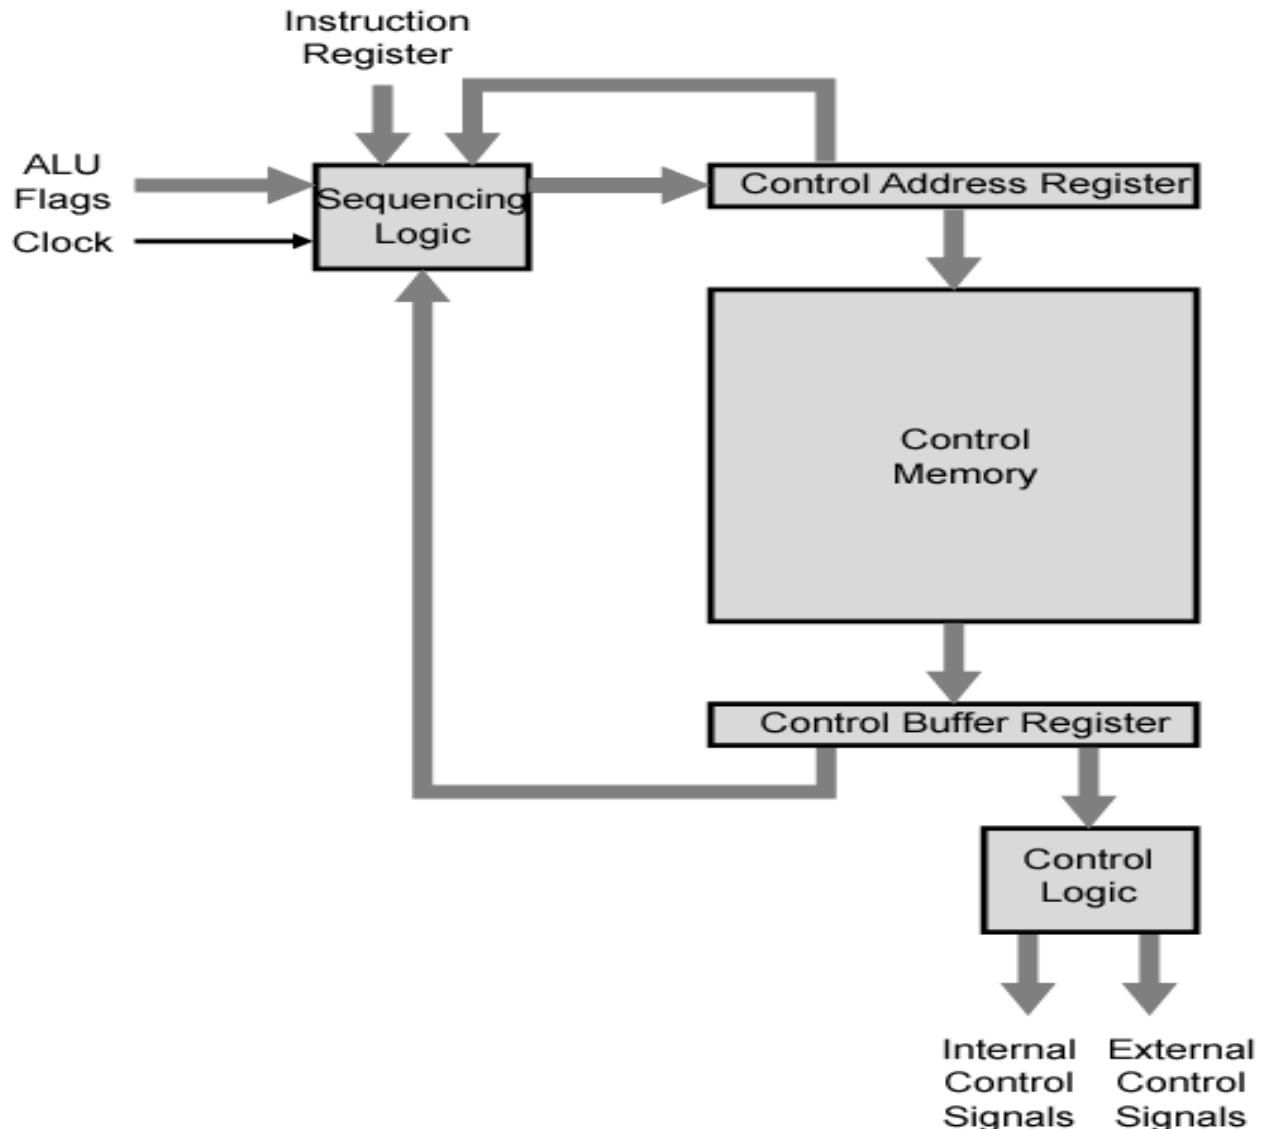
# Control Signals

- Inputs to the control unit
  - Clock
    - One micro-instruction (or set of parallel micro-instructions) per clock cycle
  - Instruction register
    - Op-code for current instruction
    - Determines which micro-instructions are performed
  - Flags
    - State of CPU
    - Results of previous operations
  - From control bus
    - Interrupts
    - Acknowledgements

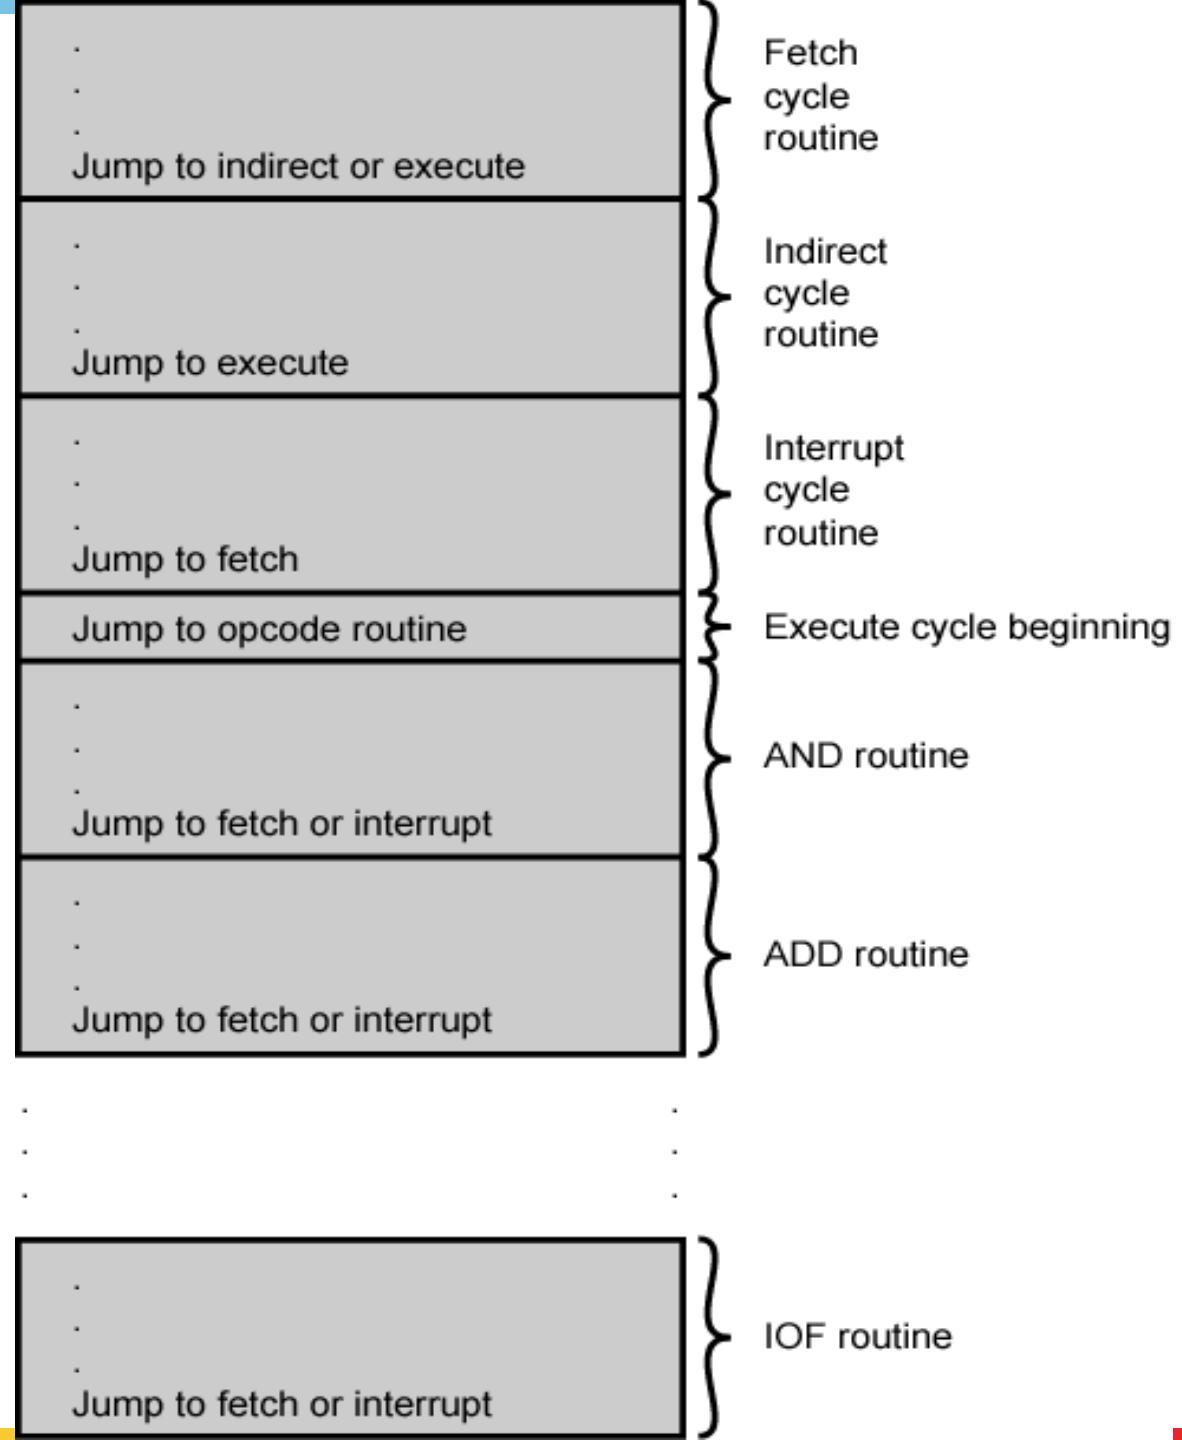
# Microprogrammed Control Unit



- Use sequences of micro-instructions to control complex operations called micro-programming or firmware
- Firmware is midway between hardware and software



# Organization of Control Memory



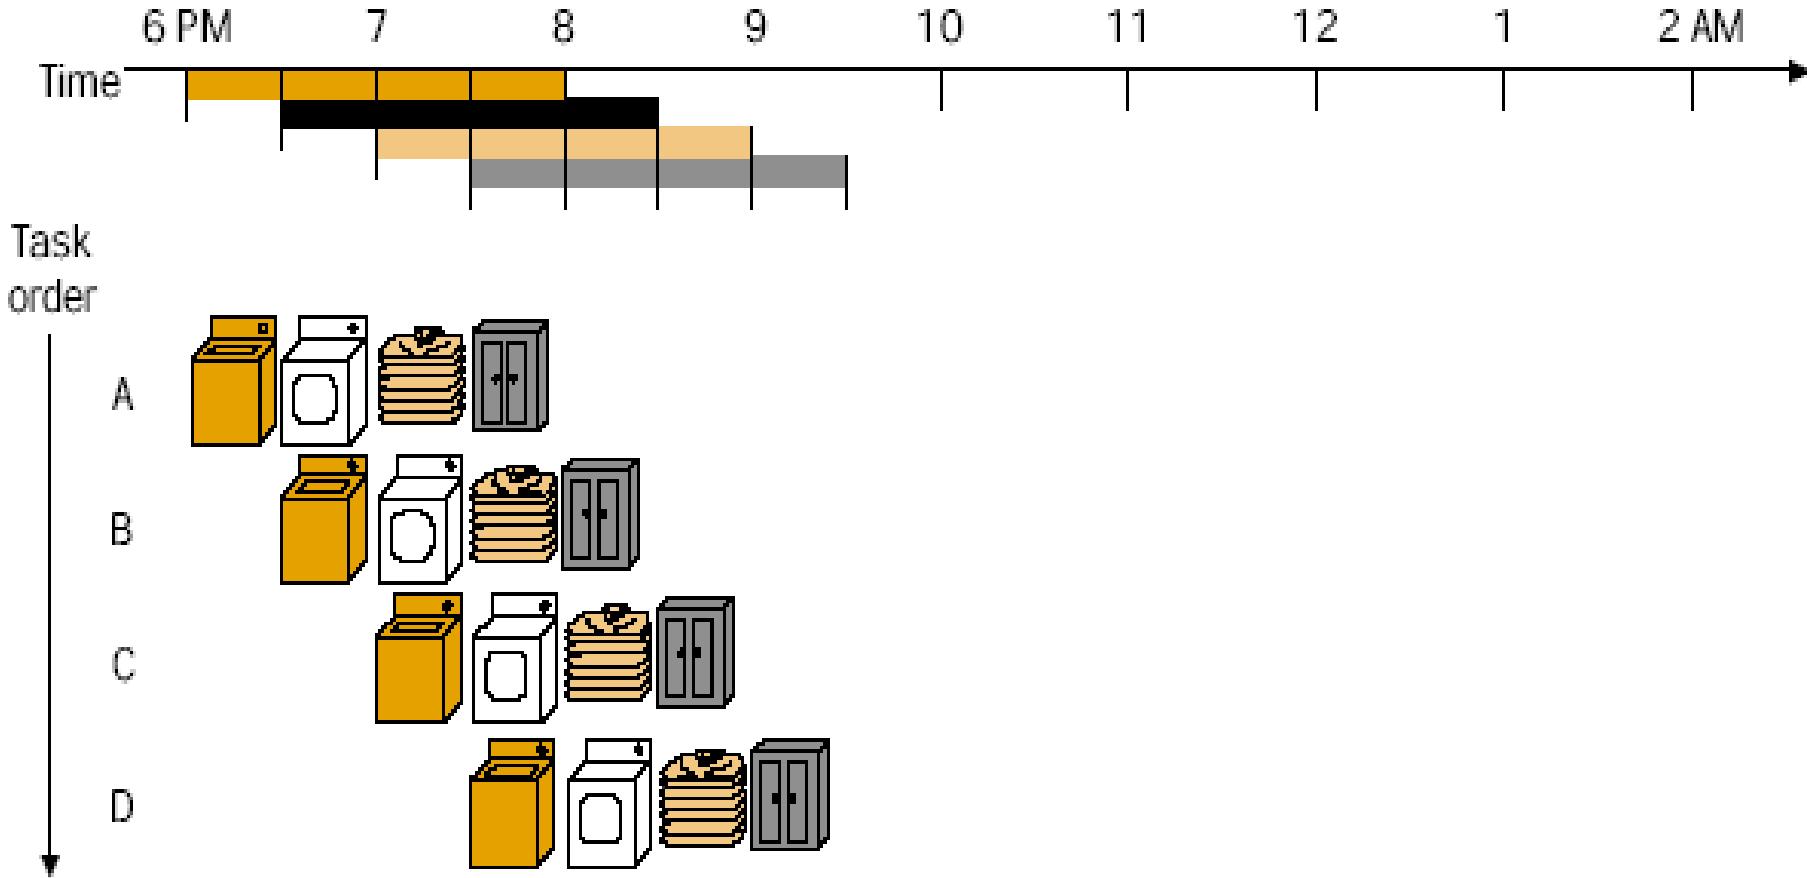


# Pipeline

# Laundry System



# Laundry System.....



# Pipelining



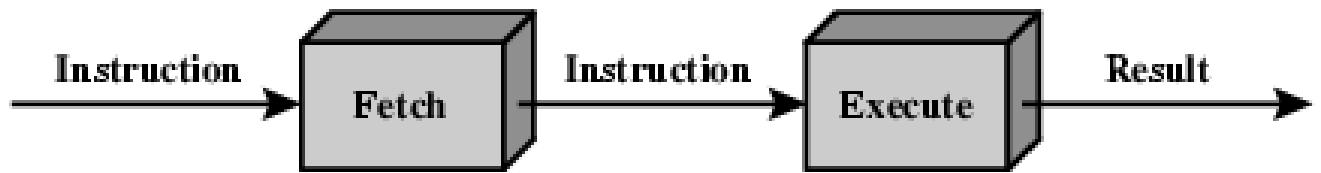
- An overlapped parallelism: overlapped execution of multiple operations
- Pipelining
  - Subdivide the input task into a sequence of subtasks
  - Specialized hardware stages
  - Concurrent operation of all the stages



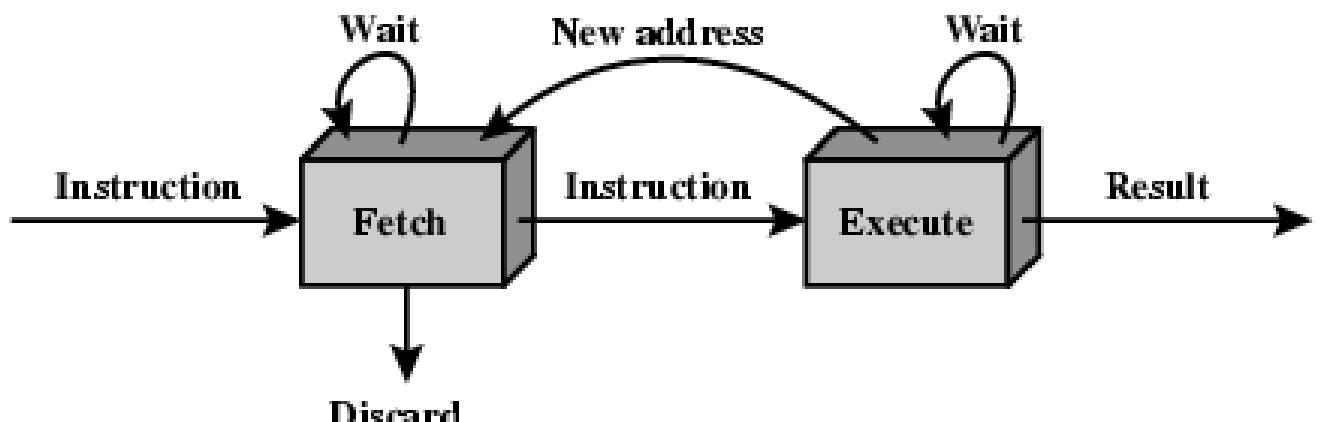
# Two segment instruction pipeline

- Contains
  - Instruction Fetch (IF)
  - Execute (EX)
- Example: 8086 microprocessor
- Instruction Fetch unit is implemented by means of first in first out buffer (Queue)

# Two Stage Pipeline...



(a) Simplified view



(b) Expanded view

# Issues



1. The execution time will generally be longer than the fetch time
2. A conditional branch instruction makes the address of the next instruction to be fetched unknown.



# Computer Organization and Software Systems

Contact Session 8

Dr. Lucy J. Gudino





# Pipeline



# Four Segment instruction pipeline

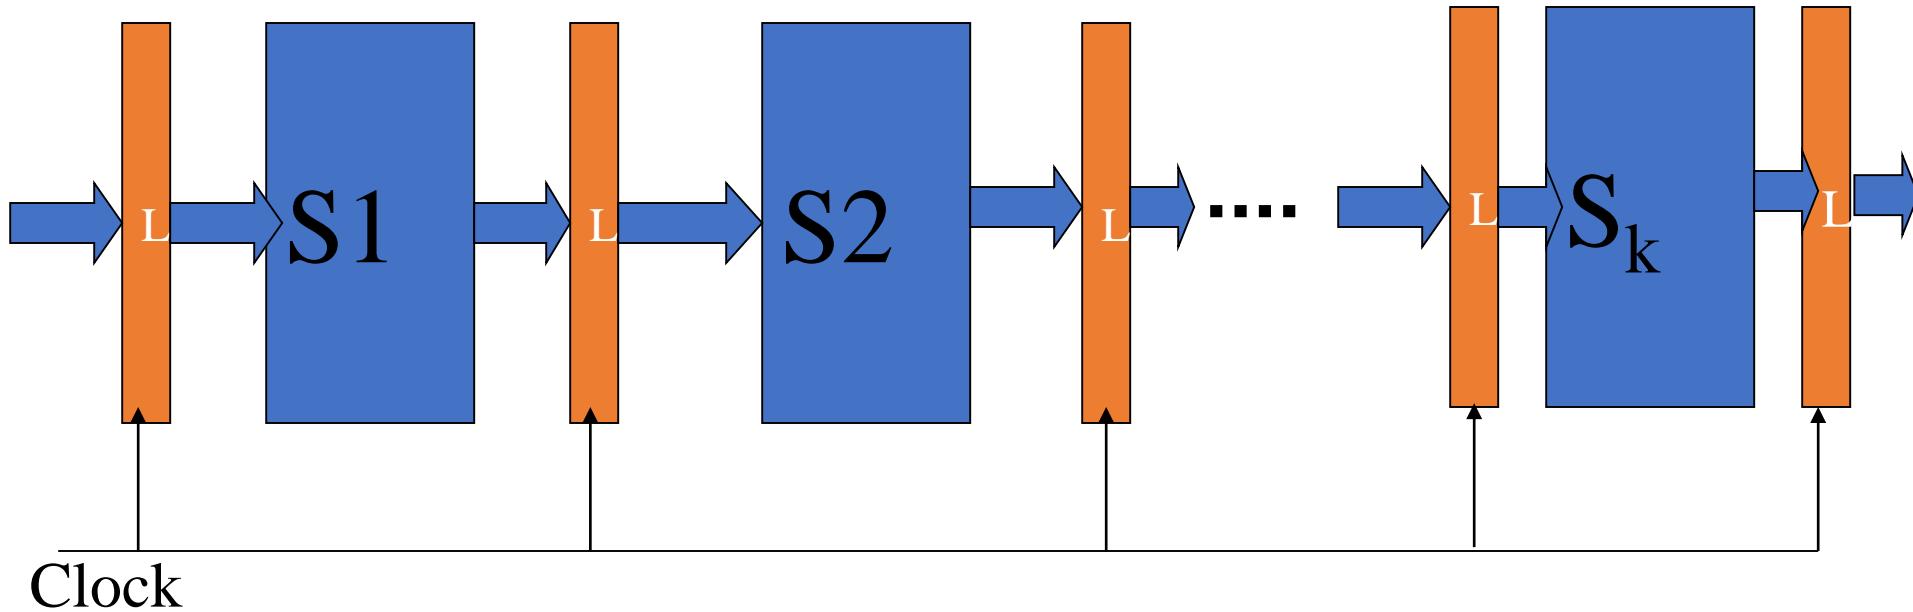
- Contains
  - FI : fetch instruction
  - DA : Decode instruction and calculate effective address
  - FO :Fetch operand
  - EX: Execute instruction



# Six stage pipeline

- Contains
  - FI : fetch instruction
  - DI : Decode instruction
  - CO : calculate effective address
  - FO :Fetch operand
  - EI : Execute instruction
  - WO : Write Operand

# Structure of a Pipeline



# Timing Diagram for Instruction Pipeline Operation



Time →

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction 1	FI	DI	CO	FO	EI	WO								
Instruction 2		FI	DI	CO	FO	EI	WO							
Instruction 3			FI	DI	CO	FO	EI	WO						
Instruction 4				FI	DI	CO	FO	EI	WO					
Instruction 5					FI	DI	CO	FO	EI	WO				
Instruction 6						FI	DI	CO	FO	EI	WO			
Instruction 7							FI	DI	CO	FO	EI	WO		
Instruction 8								FI	DI	CO	FO	EI	WO	
Instruction 9									FI	DI	CO	FO	EI	WO



# Classification

- Arithmetic pipelining
- Instruction pipelining
- Processor pipelining
- Unifunction and multifunction pipelining
- Static and Dynamic pipelining
- Scalar and Vector pipelining



# Arithmetic pipelining

- Arithmetic and logic units of a computer can be segmented for pipeline operations
- Usually found in high speed computers
- Example:
  - Star 100 → 4 stage
  - TI-ASC → 8 stage
  - Cray-1 → 14 stage
  - Cyber 205 → 26 stages
  - Intel Cooper Lake (3rd Gen Intel Xeon) = 14 stages
- Floating point adder pipeline

$$X = A * 2^a$$

$$Y = B * 2^b$$

# Floating point addition

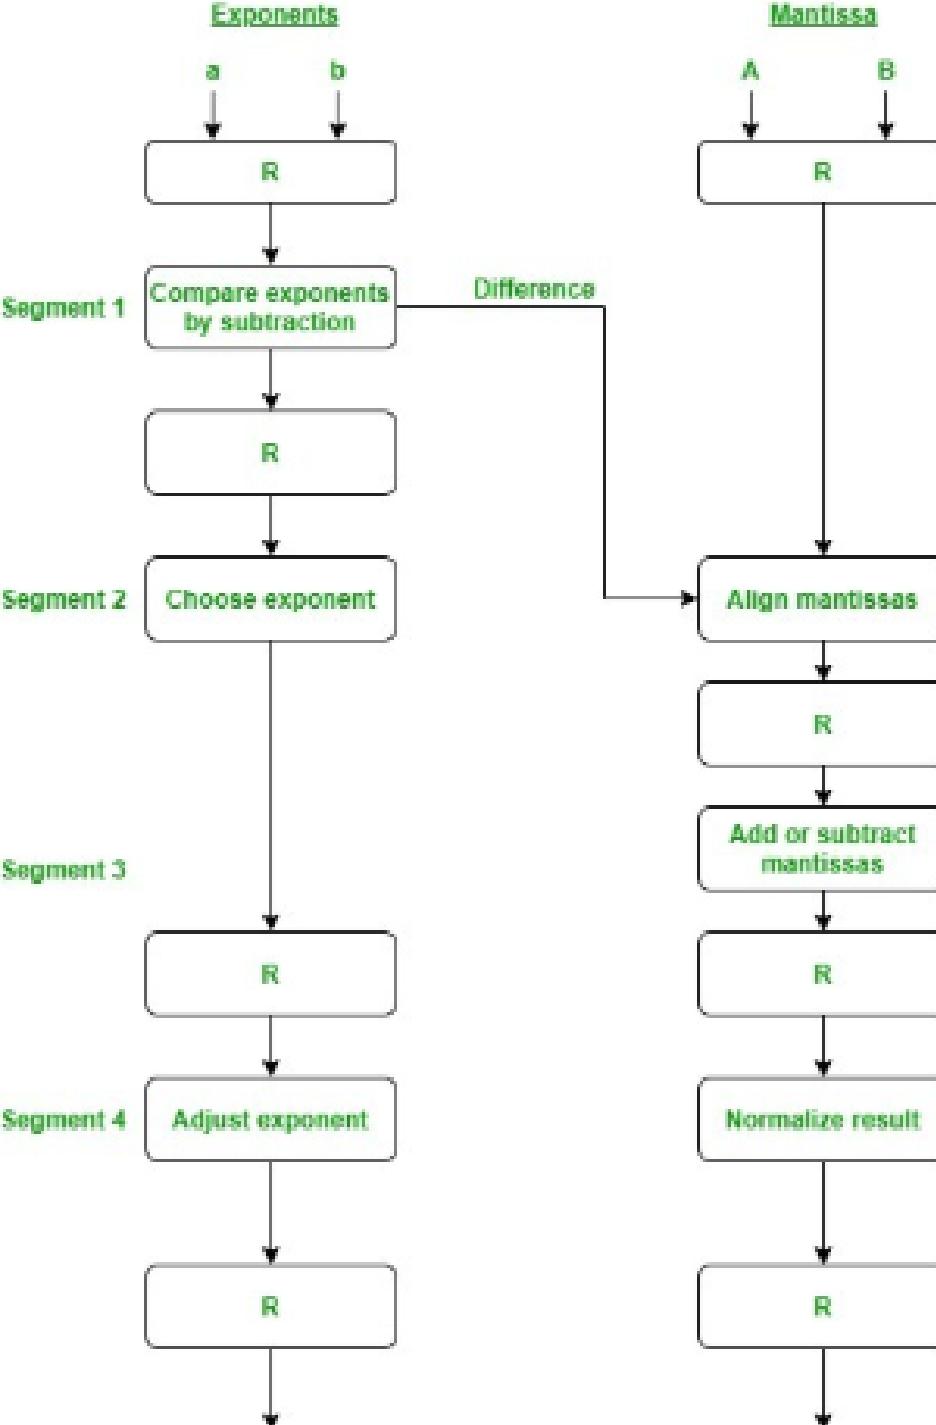


1. Compare the exponents
2. Align the mantissas
3. Add or subtract the mantissas
4. Normalize the result

Numerical Example:

$$X = 0.9504 \times 10^3$$

$$Y = 0.8200 \times 10^2$$



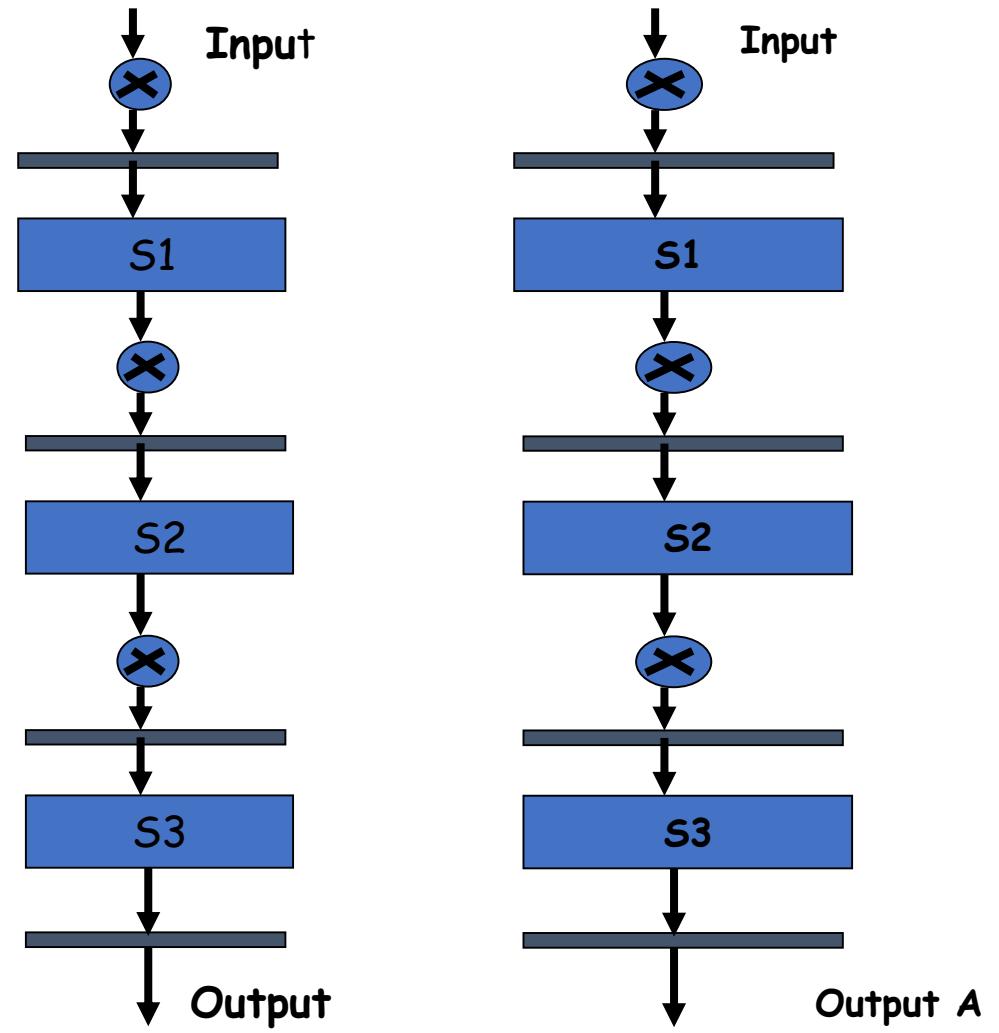


# Instruction pipelining

- The execution of a stream of instructions can be pipelined by overlapping the execution of the current instruction with the fetch, decode.....of subsequent instructions
- Sequence of steps followed in most general purpose computer to process instruction
  1. IF: Fetch the instruction from memory
  2. ID: Decode the instruction
  3. CO: Calculate the effective address
  4. FO: Fetch the operands from memory
  5. EI: Execute the instruction
  6. WO: Store the result in the proper place

# Unifunction and multifunction pipelining

- Unifunction
  - Pipeline with a fixed and dedicated function
  - Ex: Floating point adder
- Multifunction
  - Pipeline may perform different functions



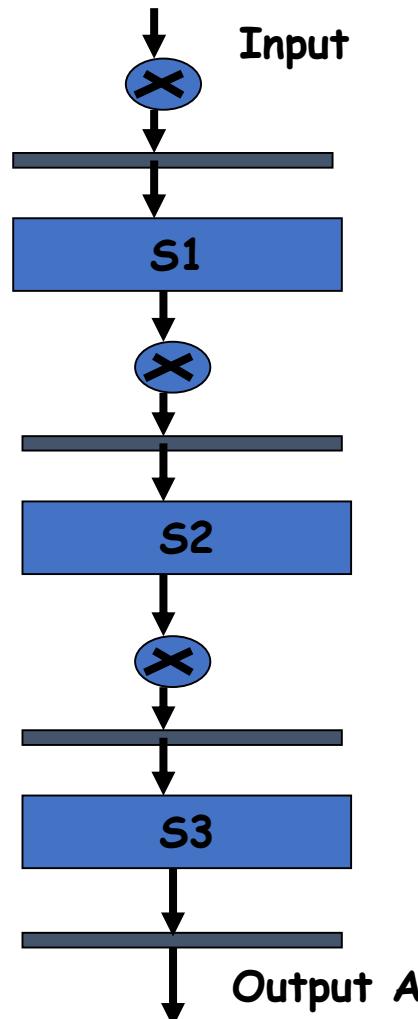
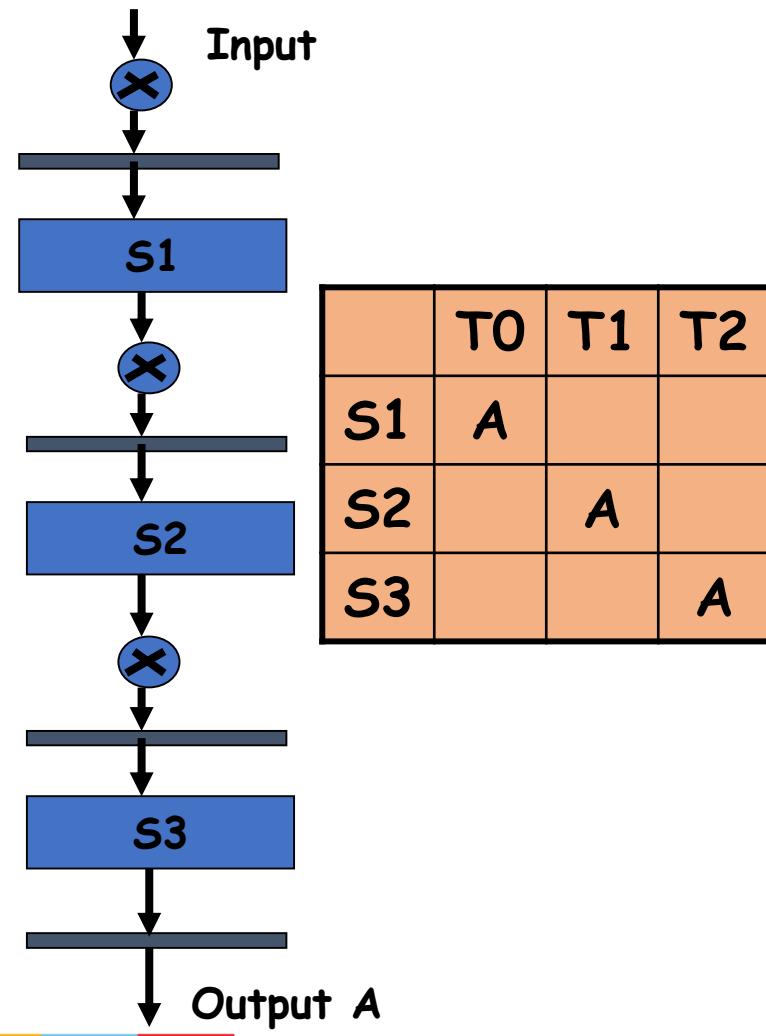
# Reservation Table



- Is a two dimensional chart
- Used to show how successive pipeline stages are utilized or reserved

# Uni-function

# Vs Multifunction



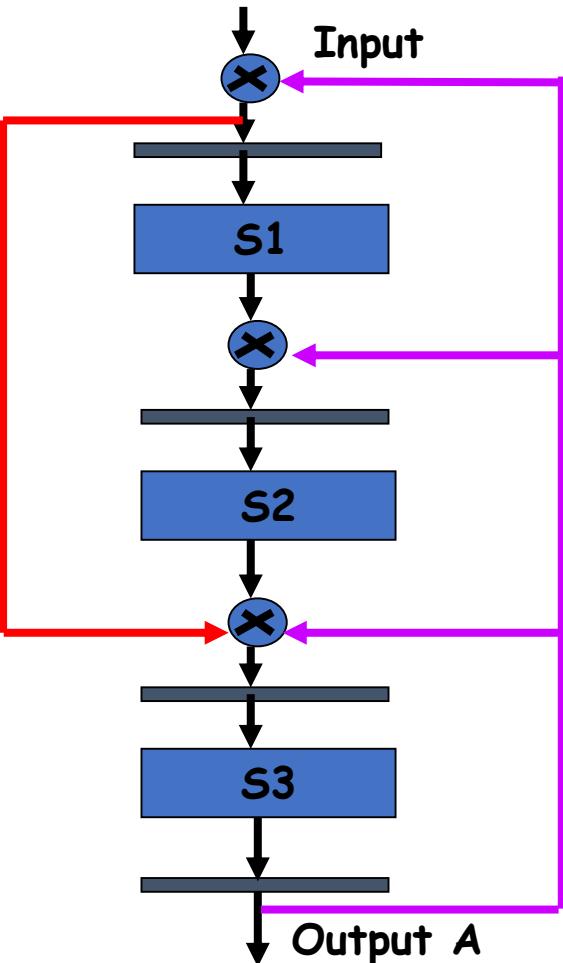
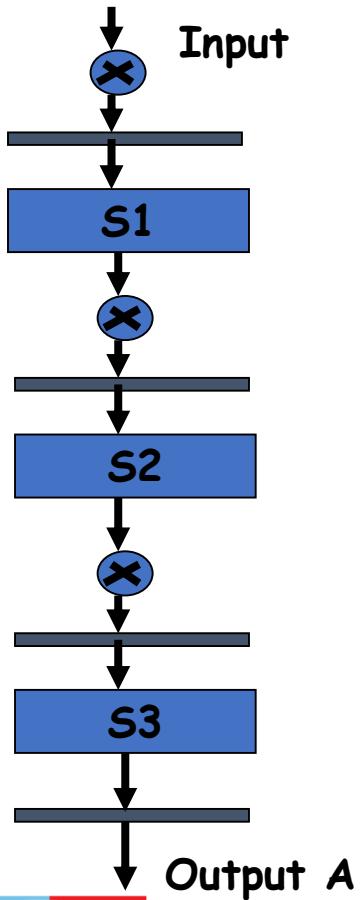
	T0	T1	T2
S1	A		
S2		A	
S3			A

	T0	T1
S1	B	
S2		B

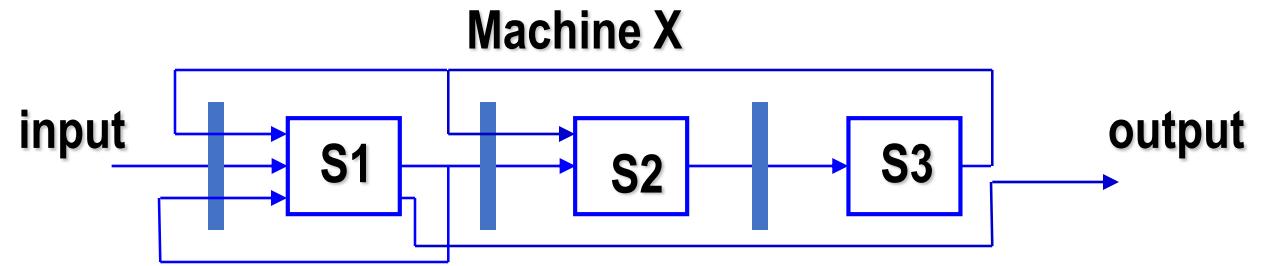


# Linear and Nonlinear Pipelines

- Linear Pipeline: Without feed forward and feed back connection
- Nonlinear Pipeline with feed forward and/or feed back connection



# Reservation Table

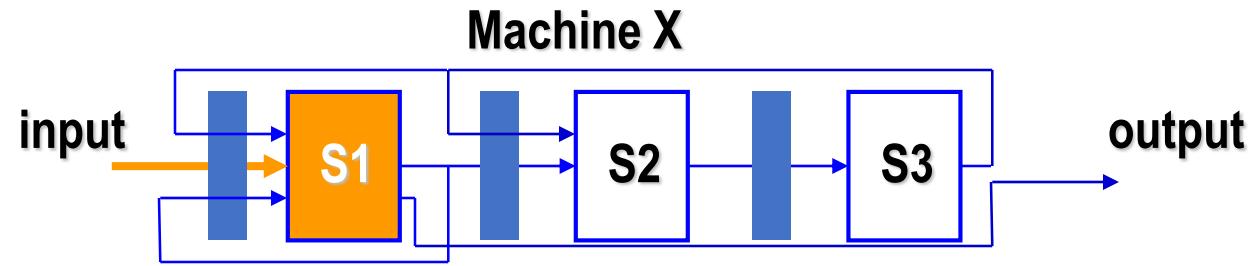


## Reservation Table

Time →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



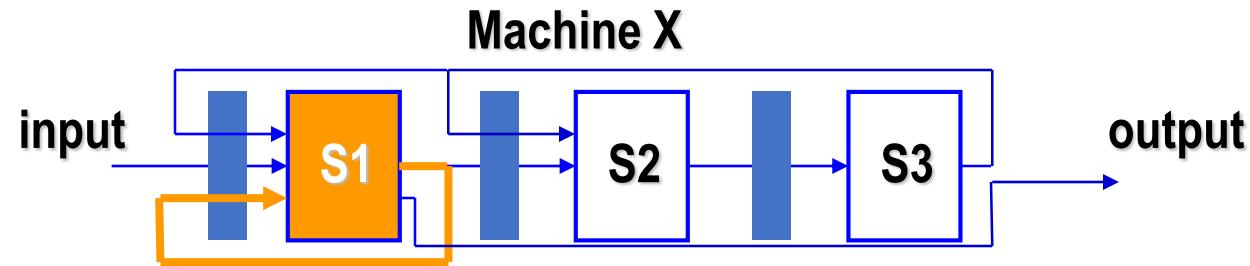
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



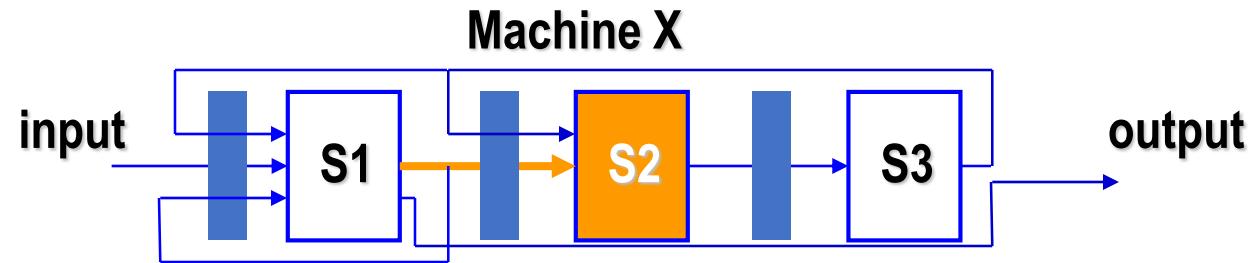
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



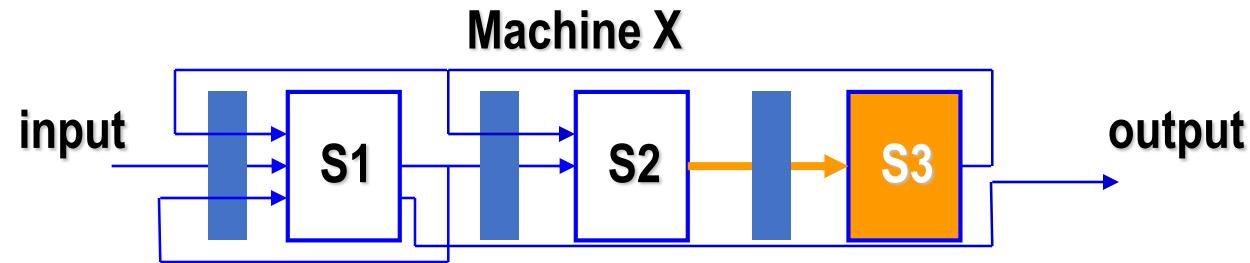
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



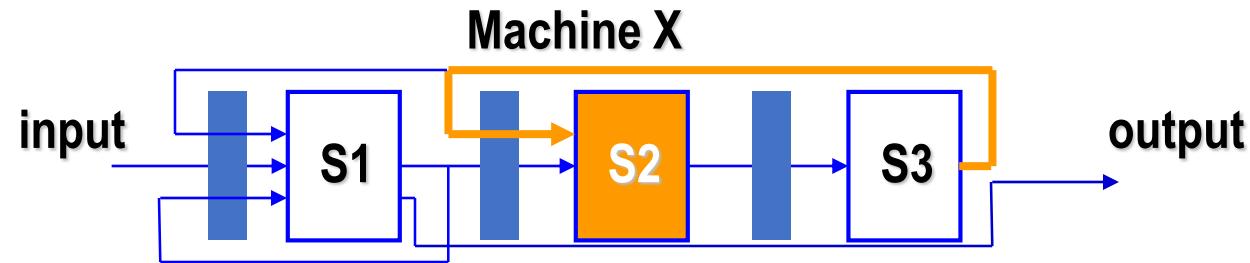
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table

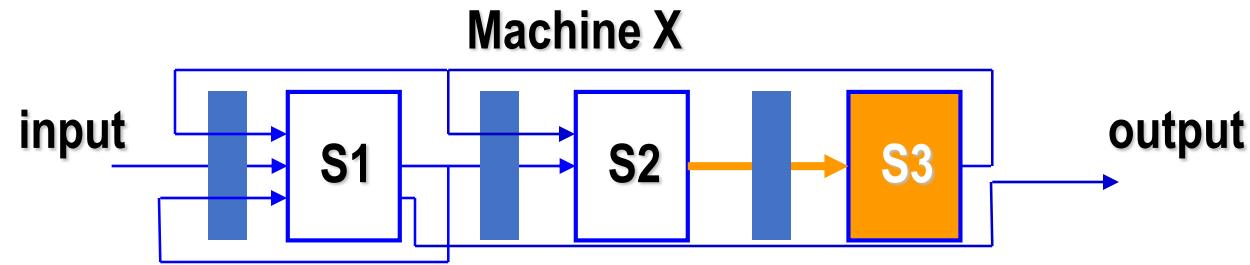


Reservation Table

Time →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



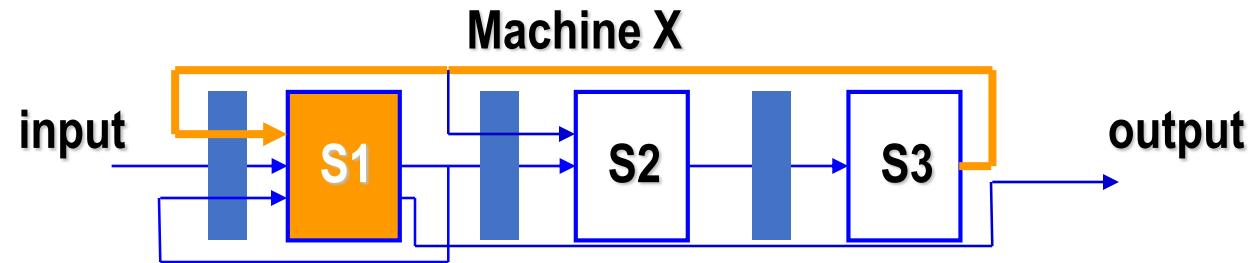
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



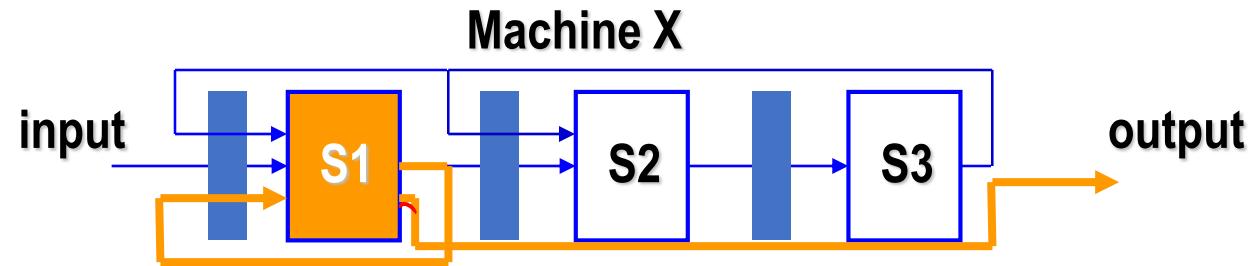
Reservation Table

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



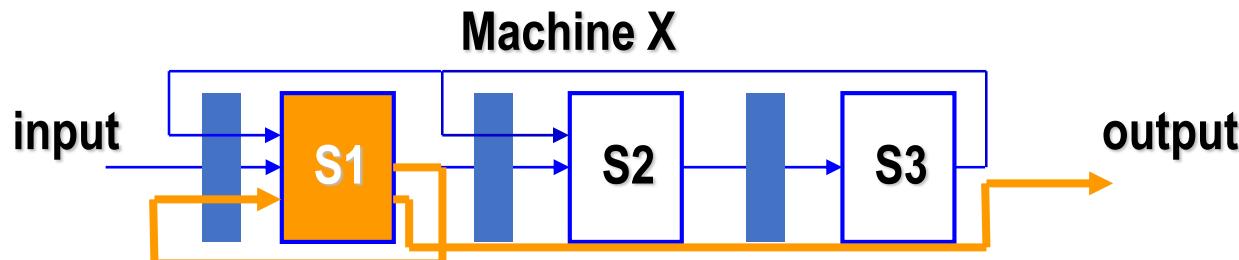
**Reservation Table**

Time →

Stage →

	0	1	2	3	4	5	6	7
S1	X	X					X	X
S2			X		X			
S3				X		X		

# Reservation Table



**Reservation Table**

Time →

	0	1	2	3	4	5	6	7
Stage →	S1	X	X				X	X
	S2			X		X		
	S3				X		X	

Time space diagram  
A, B, C

19 cycles

Time →

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
Stage →	S1	A	A	B	B	A	A	B	B	C	C			C	C					
	S2		A	A	B	B				C	C									
	S3		A	A	B	B				C	C	C								

✓ 24 cycles



# Static and Dynamic pipelining

- Dynamic pipeline allows more frequent changes in its configuration
- Require more elaborate sequencing and control mechanisms



# Scalar and Vector pipelining

- Based on the operand types or instruction type
- Scalar pipeline processes scalar operands
- Vector pipeline operate on vector data and instructions.

# Important Terms

Clock period:  $\tau$

$\tau_i$ : time delay of  $S_i$  stage

$\tau_L$ : time delay of latch

$$\tau = \max\{\tau_i\} + \tau_L$$

Pipeline processor frequency  $f = 1/\tau$

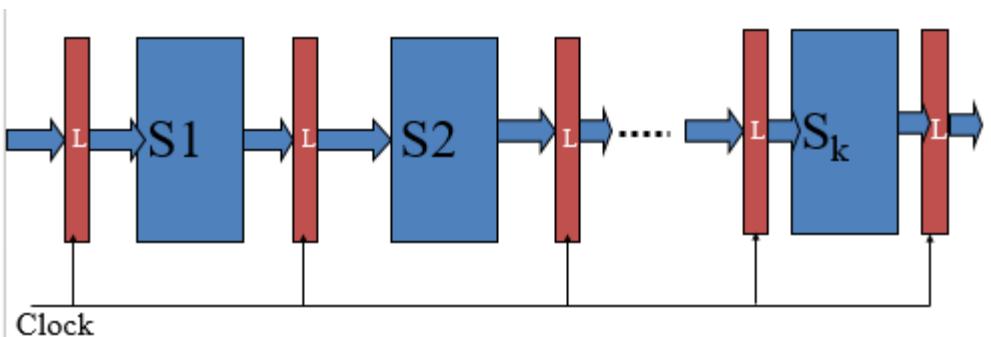


Fig: Structure of a pipeline

# Important Terms

Time taken to complete n tasks by k stage pipeline is

$$T_k = [k + (n-1)]\tau$$

Time taken by the nonpipelined processor ?

$$T_1 = k * n * \tau$$

Time →

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction 1	FI	DI	CO	FO	EI	WO								
Instruction 2		FI	DI	CO	FO	EI	WO							
Instruction 3			FI	DI	CO	FO	EI	WO						
Instruction 4				FI	DI	CO	FO	EI	WO					
Instruction 5					FI	DI	CO	FO	EI	WO				
Instruction 6						FI	DI	CO	FO	EI	WO			
Instruction 7							FI	DI	CO	FO	EI	WO		
Instruction 8								FI	DI	CO	FO	EI	WO	
Instruction 9									FI	DI	CO	FO	EI	WO



# Important Terms

- Speedup: speedup of a k-stage linear-pipeline over an equivalent nonpipelined processor

$$\begin{aligned} S_k &= \frac{T_1}{T_k} \\ &= \frac{n * k * \tau}{[k + (n - 1)]\tau} \\ &= \frac{n * k}{[k + (n - 1)]} \end{aligned}$$

- The maximum speedup is  $S_k \rightarrow k$  when  $n \rightarrow \text{INF}$
- Maximum speedup is very difficult to achieve because of data dependencies between successive tasks, program branches, interrupts etc.

# Important Terms

- Efficiency: the ratio of actual speedup to ideal speedup

$$\begin{aligned}\eta &= \frac{n.k}{k.[k+(n-1)]} \\ &= \frac{n}{[k+(n-1)]}\end{aligned}$$

- Maximum efficiency

$$\eta \rightarrow 1 \text{ as } n \rightarrow \infty$$

- Implies that the larger the number of tasks flowing through the pipeline, the better is its efficiency
- In steady state of a pipeline, we have  $n \gg k$ , then efficiency should approach 1
- However, this ideal case may not hold all the time because of program branches and interrupts and data dependencies



# Important Terms

Throughput : The number of tasks that can be completed by a pipeline per unit time

$$H_k = \frac{n}{[k + (n - 1)]\tau} = \frac{nf}{[k + (n - 1)]} = \eta f$$

# Problem 1

Draw a space-time diagram for a six segment pipeline showing the time it takes to process eight tasks

Determine the number of clock cycles that it takes to process 200 tasks in a six segment pipeline

	1	2	3	4	5	6	7	8	9	10	11	12	13
S1	T1	T2	T3	T4	T5	T6	T7	T8					
S2		T1	T2	T3	T4	T5	T6	T7	T8				
S3			T1	T2	T3	T4	T5	T6	T7	T8			
S4				T1	T2	T3	T4	T5	T6	T7	T8		
S5					T1	T2	T3	T4	T5	T6	T7	T8	
S6						T1	T2	T3	T4	T5	T6	T7	18

## Problem 2

~~0.0909 Mega/sec~~



Assume each task is subdivided in to 6 subtasks and clock cycle is 10 microseconds.

- Determine the number of clock cycles that is taken to process 50 tasks.

$$T_k = (k + (n-1))\tau$$
$$\therefore (6 + 49) \times 10 \mu s \Rightarrow 550 \mu s$$

- Determine the number of clock cycles that is taken to process 50 tasks in non-pipeline processor

$$T_1 = k \times n \times \tau = 6 \times 50 \times 10 \mu s$$
$$= 3000 \mu s \Rightarrow 3 ms$$

- Compute speed up, efficiency and throughput

$$\text{Speedup} = \frac{T_1}{T_K} = \frac{3000}{550} = \frac{ms}{\mu s} \Rightarrow 5.454$$
$$\text{efficiency} = \frac{\text{Speedup}}{k} : \frac{5.454}{6} = 0.909 \Rightarrow 90.9\%$$
$$\text{Throughput} H_K = \frac{n \times f}{t}$$
$$= \frac{0.909}{10 \mu s} = f$$



# Pipeline Hazards / Conflicts

- **Resource Hazard**: access to same resource by two segments at the same time
- **Data Hazard** : an instruction depends on the result of a previous instruction, but this result is not yet available
- **Control Hazard**: arise from branch and other instructions that change the value of PC

# Resource Hazard

	Clock cycle								
	1	2	3	4	5	6	7	8	9
I1	FI	DI	FO	EI	WO				
I2		FI	DI	FO	EI	WO			
I3			FI	DI	FO	EI	WO		
I4				FI	DI	FO	EI	WO	

— —

# Problem



Consider the following code. Assume that initial contents of all the registers is zero.

```
START:    MOV R4, #1  
          MOV R3, #2  
          MOV R1, #2  
          MOV R2, #3  
          ADD R3, R1, R2  
          SUB R4, R3, R2
```

- Write timing of instruction pipeline with FIVE stages
- What is the content of various registers after the execution of program?

# Data Dependency Conflict



START:      MOV R4, #1  
                MOV R3, #2  
                MOV R1, #2  
                MOV R2, #3  
                ADD R3, R1, R2  
                SUB R4, R3, R2

time	1	2	3	4	5	6	7	8	9	10
I1	FI	DI	FO	EI	WO					
I2		FI	DI	FO	EI	WO				
I3			FI	DI	FO	EI	WO			
I4				FI	DI	FO	EI	WO		
I5					FI	DI	FO	EI	WO	
I6						FI	DI	FO	EI	WO

# Problem



START: ADD R3, R1, R2  
SUB R4, R1, R2  
JNZ NEXT  
MUL R5, R1, R2  
ADD R6, R3, R4  
:  
NEXT: STR3, LOCN1

HLT

- Write time – space diagram for instruction pipeline

# The Effect of a Conditional Branch on Instruction Pipeline Operation



Time →      Branch Penalty →

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction 1	FI	DI	CO	FO	EI	WO								
Instruction 2		FI	DI	CO	FO	EI	WO							
Instruction 3			FI	DI	CO	FO	EI	WO						
Instruction 4				FI	DI	CO	FO							
Instruction 5					FI	DI	CO							
Instruction 6						FI	DI							
Instruction 7							FI							
Instruction 15								FI	DI	CO	FO	EI	WO	
Instruction 16									FI	DI	CO	FO	EI	WO